




TECHNOLOGY IN ACTION™

Learn Raspberry Pi with Linux

*LEARN THE INS AND OUTS OF LINUX,
THE OPERATING SYSTEM THAT RUNS
RASPBERRY PI*



Peter Membrey and David Hows

For your convenience Apress has placed some of the front matter material after the index. Please use the Bookmarks and Contents at a Glance links to access them.



Contents at a Glance

About the Authors	xv
About the Technical Reviewer	xvii
Acknowledgments	xix
Introduction	xxi
■ Chapter 1: Your First Bite of Raspberry Pi.....	1
■ Chapter 2: Surveying the Landscape.....	31
■ Chapter 3: Getting Comfortable	53
■ Chapter 4: The File-Paths to Success	69
■ Chapter 5: Essential Commands.....	89
■ Chapter 6: Editing Files on the Command Line.....	109
■ Chapter 7: Managing Your Pi.....	129
■ Chapter 8: A LAMP of Your Own	149
■ Chapter 9: WiPi: Wireless Computing	181
■ Chapter 10: The Raspberry sPi	199
■ Chapter 11: Pi Media Center.....	223
Index.....	243

Introduction

Despite sounding like something Grandma would bake on Sunday afternoons or a noise that would make people glare and tut, the Raspberry Pi is in fact a computer. That much you probably knew (although, let's be honest, the name and logo don't really give much away) but the Raspberry Pi promises more than that. An awful lot more.

The venerable Commodore 64 was released in 1982, and with sales reaching upwards of 17 million, it is often considered the best-selling computer of all time. More importantly (at least from my perspective), it was also my first computer. For Christmas, just before my ninth birthday (when the C64 was nearly a decade old) I received the new model (C64C), which was identical to the classic machine in all but cosmetics. It arrived all set up and attached to a nice new 14-inch television (it even had a remote control!). I suspect my dad had hatched what he believed to be a most cunning plan; if he could sneak in and set everything up while I was asleep, come Christmas morning, I would be so busy playing with the computer that my parents might get an extra few minutes of sleep.

Sadly, things did not go quite according to plan. Although everything was set up, and even though the television was tuned to the computer's signal, one simple but key thing had been forgotten: it hadn't occurred to anyone to tell me how to actually load a game. Needless to say, a lie in was not forthcoming . . .

Games came on cassette tape. This was before CDs, and at least on the C64 had to be played in a special tape recorder called a datasette. (What do you mean what's a CD?) Sadly, the datasette spent more time in the shop than being attached to my computer and as it was the only way to load anything into it, I had no choice but to occupy myself with the manual. This I used to great effect and taught myself how to program good old BASIC (Beginners All Purpose Symbolic Instruction Code—can you believe I actually remembered that?).

While I'm sure this story is very gripping, you could be forgiven for wondering why I am boring you with it. Those events happened more than two decades ago, so what possible relevance could they have today? Well tinkering with that machine and then the Amiga that followed it (still my favorite machine of all time) gave me a real appreciation for what a computer could do. The Amiga was severely underpowered compared with PCs of the same era, yet it consistently beat them with better graphics, better reliability, and better sound. It was able to do all of this because the hardware was exquisitely designed. Amiga enthusiasts were some of the most resourceful people I've ever seen. Who'd have thought you could turn a real-time clock port into a connector for high-speed storage?

All of this was only possible because people really understood how all the parts fit together. They knew how to get the best out of the machine because they really knew how the machine worked. These days, I spend my working day trying to make fast things go faster. To have any hope of success, I too need to know how everything works. Companies need people like me to push things forward, but they're coming across a bit of a problem. People who really know computers inside out are getting much harder to find—we are a dying breed, and this is the situation that the Raspberry Pi Foundation is desperately trying to reverse.

So what happened? Well, things changed. Computers went from being the curiosity in the corner to being a basic part of everyone's lives. They evolved to the point where they just work and everyone knows how to use them. This is similar to the family car. Everyone has a rough idea how a car works, but few people are very interested. The car takes them from place to place, and that capability is what is interesting, not how the car achieves it. Computers are generally seen in the same light. People have a rough idea about turning them on, installing software and so forth, but how they actually work at a low level isn't really seen as relevant or interesting. This in turn means that not only are fewer people getting excited by computing itself but even fewer people think that there's more to it than double-clicking an icon.

This problem has drifted up through schools and into universities. Teaching programming is a relatively challenging task. It requires a certain way of thinking that for many people is tough to get a handle on. Traditionally, universities would start a computer science course by teaching about logic gates, how memory works, and how to program a CPU. Only once you understood what the bare metal was doing would you try to learn C because although C is a higher language, it reflects the hardware it runs on. Understanding the hardware makes understanding C that much easier.

But with larger class sizes, more limited teaching time, and students arriving with less and less knowledge of computing fundamentals, universities have had to adapt. Rather than teaching all that low-level stuff, now they teach Java and other scripting languages. Because these languages handle all of the “fiddly bits” for you, you can effectively pretend that they don’t exist (although this can cause some issues; see “The Infinite Memory Myth”).

This is simply fantastic from a productivity point of view, but when you do want to take it to the next level (maybe you’re processing data and your script is just too slow), you have no idea where to turn. Worse, when someone tells you the technique for improving that performance, you have no idea what they’re talking about.

Of course not all universities have taken this route. I’m studying at the Hong Kong Polytechnic University, and its course on Computer Architecture is very detailed and covers a lot of ground. If you want to get the top grade, you will need to implement a CPU cache for the CPU simulator program. Needless to say, there is a lot to learn for everyone on this course. That said, we need more than this. It’s too late to capture people’s interest when they’re starting graduate studies. I taught seven-year-olds how to program BBC BASIC when I was in my last year of primary school (they even got presented certificates by the school) and they loved it. Computing lets you create a virtual world with your mind (the children liked to create little text-based adventure games) and ultimate power rested in their hands. They got creative; they added spells, new roads, secret entrances, and much more. Okay, they needed a helping hand (they were only seven) but they had the desire to create and to build cool new things.

INFINITE MEMORY MYTH

Over the years, I’ve done a lot of consultancy work with large enterprise customers and that has inevitably meant I’ve come across Java on many occasions. One of the interesting things I have come across is what I’ve termed the Infinite Memory Myth. This seems to crop up more in Java applications than in other languages, but that’s probably because Java tends to be more widely used in those settings.

The short version of the myth is that developers seem to constantly create new objects, often to the point where the application consumes huge amounts of memory or crashes altogether. They tend to have no idea how much memory each object takes or, more worryingly, why they should care. As far as they are concerned, they request a new object and one is provided. When an object is no longer used (i.e., nothing points to it any more), Java will at some point get around to cleaning it up (called garbage collection). All of this is automatic; the developer doesn’t need to do anything.

The problem is that this leads people to forget (or in many cases were never taught at all) that memory is finite, and at some point it simply runs out. You can’t assume that you can read in every row in a table and that it will always work. You can’t assume that just because your test file is 50MB in size that the application will never be given a 5,000MB file to work on.

This lack of understanding stems from not being able to appreciate all the hard work Java is doing on the programmer’s behalf. It is running about and managing memory allocation and garbage collection, and the programmer remains blissfully unaware. A good understanding of computing fundamentals would give a developer keen insight into what Java is doing (both the how and the why) and thus appreciate that just because creating new objects is easy, memory itself is not free.

So this is what Eben Upton and the Raspberry Pi Foundation are trying to bring back to the world. They want to rekindle that lost art and make computers cool and interesting again. To do that, it has created a computer that even by today's standards is no slouch. Is it as powerful as your laptop? Well no, probably not, but can you buy your laptop for \$35, slip it into your pocket (possibly not a great idea), generate relatively little heat, and drink very little in the way of power? If you answered yes to those questions, I really really want to hear from you—that sounds like a laptop I need to buy!

However you choose to look at it, the main selling point (no pun intended) is the price. Anyone can pick up a Raspberry Pi without having to think too much about it. With a modest laptop clocking in at around \$500 + and a MacBook pro nearly four times that, it's not the sort of thing you can just splurge on without thinking, especially if it's going to be for experimentation and playing about. However at \$35, the Pi is cheaper than some monthly movie subscriptions, you could almost buy a new Pi every month!

Why Eat Raspberry Pi?

Whichever way you look at it, you will come back to the price. Whatever else the Raspberry Pi is and whatever other promises it has in store for us, all of them are interesting because of the price. There are two types of people who will be rushing out to get a Pi. The first group is already clued up on Linux and computing, and for them the Pi represents a server in a pocket and a cheap one at that. No longer do they need a full-size PC monster guzzling electricity and generating enough heat to rival a bar heater. Oh sure, you can get low-powered systems in nice shiny boxes, but they're still not all that cheap to buy, even if they're cheaper to run. However, a device like the Pi is cheap to run and cheap to buy, and it has just what you need to build a pretty respectful server.

If you're not one of those people, don't worry because this book is for you. You like the idea of a cool little computer for \$35 and you think you can do some pretty awesome stuff with it; you're just not really sure how. For us, the big benefit is that the Pi is at the sort of price where we can afford to buy it just for fun and use it for experimentation. It's not going to replace the family PC and you're not going to need to take out a mortgage to buy one. You can play around with a Pi completely guilt-free and try all manner of weird and interesting things without having to worry about cost or destroying your main computer (and thus incurring the wrath of everyone you live with).

Because the Pi is close enough to a normal PC (even though the architecture is a bit different) you can do PC-type things with it. In fact, that's the first thing we show you in this book. Thus you don't have to start from scratch; all that you know already you can apply to a Pi (it rhymes if you say that really fast) and so you can hit the ground running. No doubt you will want to do all the things that the first group of people wanted to do as well. Fear not, we have you covered—by the time you've finished this book you'll be able to do all that and more!

There are lots of reasons why everyone should rush out right now and get some Pi. Actually at the time of writing, there is still a three-week lead time for delivery, and when the Pi was initially released, one of the resellers took 100,000 preorders in just one day. That's a lot of Pi! Although the lead time will naturally keep changing, the short version is that the sooner you order, the sooner you will get your Pi!

So why all this interest? What is so special about the Raspberry Pi that it has achieved an almost cult-like following and has remained in heavy back order for months?

It Only Costs \$25

Okay, hands up all of you who are only interested in the Pi because it costs significantly less than a night on the town? If you put up your hand, you're not alone. The goals of the Raspberry Pi Foundation are laudable but they all center on getting this powerful machine into our mitts at a price point that won't break the bank.

What really has everyone drooling is not the fact that as far as computer hardware goes, the Pi is effectively free, but more that it is a full computer that can run Linux. That means servers, home automation, video streaming, and pretty much anything else you can imagine.

WAIT, IS IT \$25 OR \$35?

Throughout the book we do bring up the price of the Pi a fair bit; after all, it is one of its most distinctive features. However, we also mention two prices, \$25 and \$35, so which is it? Well there are two versions of the Pi, the Model A and the Model B (as uninspiring as those names sound, they're taken from the BBC computers, and from a geeky point of view, the names are quite inspired). There isn't a great deal of difference between the two models, with Model B having 10/100 Ethernet built in and an additional USB port. Model B also draws a bit more power. The Ethernet adapter is actually connected internally via USB, so there is no difference between the build in Ethernet and a USB device that you could plug into the USB port itself.

So which should you buy? Well if you think networking will be useful, the built-in Ethernet port is pretty much required. I love having built-in Ethernet connectivity because it just makes life so much easier. However if you aren't planning on using it much, there really is no need to get Model B. That said, for \$10, it might be worth spending more just in case you decide you want to play with networking later on...

MORE PI!

If one Pi is good, two must be better. Right? Well, \$50 would get you two Pis and while this doubles the cost, it also more than doubles the fun. Now you can experiment with networking and getting the Pis to talk to each other. After all, it's good to talk!

Experiment in Safety

When my computer is out of commission for even a short period of time, it is pretty inconvenient. I certainly don't want to fiddle about with something and accidentally erase my hard disk (been there, done that). You'd also be well advised not to try over-clocking your CPU on the brand-new computer that you were just given as it's not much fun to think you've totaled the machine within an hour of turning it on!

To kill hardware takes a fair bit of effort (such as taping over a pin on the CPU and removing the CPU speed multiplier lock), but it's fairly easy to remove your family photo album and the latest draft of the book you've been working on for the past six months (there is a reason why we dedicated a whole chapter and personal plea in the book to backups). If you have a Pi and you nuke it, the worst case is \$25 down the drain, which is a lot better than what would happen if you toasted your main machine.

Independence

I'm sure some people will point out (quite correctly) that most of the horrors I just described can be avoided if you play in a VM rather than on your main machine. Apart from it not being as much fun (real hardware just smells better), it doesn't give you all the benefits of a separate piece of kit. For a start, a virtual machine is only on when your main machine is running. If you happen to have a laptop that follows you everywhere, a virtual machine won't be a great option for a home web server. Also, if you ever reboot your main machine, your virtual machine will go down with it. If you were using it to stream movies to your TV, you could end up with some very displeased family members. By having a real piece of hardware, you can keep your experimentation completely separate from anything else that you might be doing.

Low Power

The Pi has a very modest power footprint. In fact, the Model A Pi only draws 300ma, which means you can power the whole thing from your USB port. According to Apple, my iMac draws 94 watts at idle and up to 241 watts when the CPU is maxed out. The Model A Pi draws at most 1.4 watts, and the Model B draws at most 3.5. That's an awful lot of power saved. Understandably, these figures only take into account the power requirements of the Pi itself, with a bit extra put aside to power modest USB devices. If you add lots of power-hungry devices to your Pi, these figures would increase accordingly.

APPLES TO ORANGES

I just know someone is going to cry foul about my power comparison—and for good reason. It is true that compared to an iMac, the Pi draws basically no power at all. However, it is also true that an iMac does a lot more. First, it has a big screen, hard disk, and CPU that would blow the ARM on the Pi into the middle of next week. So if I know that I'm comparing two totally different systems and I'm admitting that they're very different (who compares a golf cart to a Formula One car?) why am I wasting ink and paper with this description?

Well, although we are comparing two very different machines, we are comparing the same sort of tasks. If you want to have a little web server or stream video to your TV set, a Pi is more than powerful enough to do that for you. Bear in mind that the Pi is clocked in at 700Mhz; not too long ago, that was what you'd find on a very powerful desktop. It's not that long ago that you'd find this sort of performance on an enterprise grade server. Ten years ago, it was just a fantasy. In short, the Pi has more than enough juice to do most of the things you'd want from a server and it won't require a small nuclear power station to do it.

The Ingredients for a Raspberry Pi

After seeing all its benefits, you might be wondering why a \$35 computer stacks up so well to ones that cost many times that. If you're thinking that, it won't be long before you wonder why there is such a big difference in price. Surely if you could get something this cheap that does most of the things your main machine can do, something must be up with the price of the other machines. After all, if a powerful laptop could be made significantly cheaper, it would easily make more money in increased sales than it would lose in reduced profit margin.

Well that's true to a point, but the Pi, powerful as it is, will probably never be a direct replacement for your main computer. It's not any particular one thing that limits the Pi, but a combination of design decisions to balance features with cost that will ultimately prevent it from taking the crown. That said, it is still a fantastic platform, and we'll look at some of those highlights right now.

ARM CPU

The major and most obvious difference between the Pi and your desktop is that your desktop will almost certainly have a processor from either Intel or AMD at its core, whereas the Pi has an ARM-based CPU. The Central Processing Unit (CPU) is the part that actually runs the programs you provide. Before a CPU can run a program, it must be translated into a language that the CPU can understand. So all CPUs execute programs, but the program has to be in a language they can understand. An ARM CPU cannot understand instructions written for an Intel or AMD CPU, which effectively means that most off-the-shelf software (such as Microsoft Windows and games) cannot run on a Raspberry Pi.

Okay so you're not going to be running Windows XP on the Pi, but does the ARM offer any advantages? The first advantage is that ARM CPUs draw much less power. This isn't overly surprising as ARM really came from the embedded hardware industry where power usage and heat generation are a really big deal. In fact, you'll find some

form of ARM in almost every modern cellphone including the iPhone 4s and Samsung Galaxy S3. In fact they are one of the most widely used processors in the world and can be found in all manner of devices such as DVD players, appliances, and even cars.

ARM CPUs also generate very little in the way of heat. If you look at the Pi itself, you will notice the CPU doesn't have a heat sink. If you look at any Intel or AMD CPU, you can't fail to miss the huge cooling fan that it requires in order to prevent it burning out. Some people even user water cooling systems to keep their PC processors running at a reasonable temperature.

The last benefit is really a cost-to-performance ratio. For the vast majority of things, the real bottleneck is not CPU power but how fast data can be fed to it. CPUs have long been much faster than hard disks, and even the bus that links all the computer components together can't keep up with even a modest CPU. So what do you get out of this? A low-cost processor that almost certainly will do everything you need without the cost penalty.

WINDOWS 8 ON ARM

Early last year, Microsoft announced that Windows 8 would support ARM processors. This is not overly surprising as ARM is present in practically every modern mobile device; and with the performance of the high end chips gaining significant ground on Intel and AMD offerings, it probably won't be long before we start to see more powerful desktops being powered by these beasts.

Microsoft has long dominated the PC market and regardless of your thoughts on who should be leading the way, everyone agrees that Microsoft is far from done. Of course, where Microsoft leads, others will follow. If Microsoft starts supporting the ARM platform directly, it won't be long before hardware manufacturers will start providing the necessary drivers and support for their devices. Game developers will also need to pay attention to ARM-based platforms and that could create some very interesting challenges and compromises.

What does it mean for your Pi? Well, you probably won't be able to run Windows 8 on it; it just wasn't built with that in mind, but it does mean that the technology underpinning your Pi is not based on some old bits and pieces found at the back of someone's bottom desk drawer. Rather it is quite possibly based on the platform of the future!

512MB of RAM

The Model B Pi comes with 512MB of RAM (originally 256MB, but they changed to 512MB in October 2012), which is shared between the system and the Graphics Processing Unit (GPU). Now considering that the laptop I'm using to write this has 8GB of RAM, my desktop at home has 16GB (that is 32 and 64 times, respectively), and that up to half of this is shared with the GPU, it sounds a little weak. What can you really do in 512MB of RAM? Or even 256MB if you have a Model A or an older Model B?

Well the answer to that is a lot more than you might think. Remember, Windows 95 was able to operate with a couple of running programs in just 8MB of RAM (which conveniently enough is 32 times less than the Pi), and the good old Commodore 64 came with only 64KB (the Pi has something like 4,000 times that) and it was able to run thousands of games to entertain the masses.

So why do we have so much memory in new machines? Memory is cheap these days, and although we could do lots of great stuff in very little memory, it's a skill that has gone out of fashion. Why spend so much effort on optimization memory usage when it will most likely never matter? If 8GB of RAM costs \$50, is it worth hours if not days of a programmer's time to save a few megabytes here and there? Probably not.

Remember, the Pi is meant as an experimentation platform, not as a general PC replacement, and Linux (especially without a GUI) will run with plenty of memory to spare for all your programs. By keeping the memory to a reasonable minimum, the Pi is able to hit its price point without greatly hindering what you can do with it.

GPU

The GPU is really a specialized form of a CPU. A CPU is generic in what it can do, and it tries to be good at everything it does, the classic jack-of-all-trades. A GPU, on the other hand, does one thing and one thing well. It is specifically designed to handle the intense mathematical calculations needed to render complex displays. This started off predominantly in 3D graphics rendering, but has more recently gained traction in day-to-day computing in which rather than computing graphics for display, a GPU can be harnessed to execute similar types of instruction for the user. For example, when an application such as Photoshop enhances a photograph, it applies an algorithm to the image and historically the CPU would do this. Today, Photoshop can offload that processing work on to the GPU, which because of its very specialized design can do the work much faster—all without needing any assistance from the CPU itself.

The main reason why a GPU is important in a device like the Pi is that even with a modest processor, it can still handle high-quality displays and decode high-quality video streams. This makes the Pi useful as a media device as well as allowing for a full graphical display that still feels snappy even with a slower processor.

For the most part, the GPU is not something you will directly care about, but by knowing that the heavy duty graphics work can be offloaded somewhat from the main CPU, you can be more confident that the ARM that powers the Pi will be able to deliver enough brute force for your needs.

Ethernet Port (Model B only)

You might think that there's not really much to say about an Ethernet port. After all, it's pretty straightforward: you plug it in and you can access the network. If you don't have one, then you can't. Simple enough?

For me, the ability to connect to a wired network is essential. WiFi can often have issues and sometimes isn't even available (especially if you decide to turn your Pi into a WiFi access point). Although this feature costs an additional \$10, in my experience not having the network card when you really need it will cost you much more than \$10 in time, effort, and general hassle. It is true you can add a USB network card to the Model A (and technically the card attached to the Model B is actually connected via USB), but that would cost at least \$10 and then you'd have a USB device flapping in the wind. You might say to yourself "If it's USB, I can use it with my other devices, too!" but in reality, you probably never will and you're far more likely to lose the damned thing altogether (right when you most need it) than it is to come to your rescue in a time of need.

However, one reason why you might decide on a Model A is that because it doesn't have these extra components, it consumes significantly less power. This probably won't be a major concern for most users, but if you're planning on using the Pi in a battery-powered product, you would be very keen to lower the power requirements as much as possible.

Overall, my recommendation is to get the Model B because you get the network card. If you're absolutely 100% positive that you'll never ever need Ethernet (or you want the smallest possible power footprint), there's probably little point paying the extra cash for this version.

USB

This isn't so much a feature these days as a true requirement. Almost all peripherals connect this way and the Pi is no exception. It will work with all standard USB keyboards and mice, and assuming Linux has a driver for it, other USB devices as well (such as the Ethernet card in the Model B).

Both models come with USB 2.0 support, although Model A only has a single port, whereas Model B has two. If you want to be able to plug multiple USB devices into your Pi, you might want to consider Model B, but if you're setting it up as a basic desktop, you could use an external USB hub to get a similar effect.

GPIO Ports

General Purpose Input/Output (GPIO) ports are a very interesting and key addition to the Pi. They provide an easy way to connect hardware to your Pi that you can then control through software. If, for example, you wanted to add a thermometer or light sensor, you could build a device that connects to your Pi via the GPIO ports.

If hardware projects don't really interest you, then you can probably forget about the GPIO ports. But if you're looking to integrate your Pi with various bits of hardware or make your Pi the brain of some wacky invention, GPIO ports will give you an easy way to do that.

Baked to Perfection

By now, you'll have picked up a pretty good appreciation for what the Pi is and what makes it special. It was designed to be low cost so that everyone who wanted one could get their hands on one, but the designers have gone to great lengths to make sure that even though the Pi can't deliver everything a desktop can, it delivers more than enough to make it a fully functional computer laboratory not to mention a very nice server platform.

Whistle-Stop Tour

So, what do we have in store for you in the rest of the book? The book is broken down into three core parts. In the first part I'll show you how to get up and running with the GUI. If you've never heard of a GUI, this is definitely where you want to start. The Pi comes with everything you need, but out of the box it needs a bit of tweaking. Don't worry; we'll get you up and running quickly, and you'll soon be off exploring.

Part 2 takes you back to the command line and teaches you how to move around. Many people (myself included) spend the vast majority of time at the command line. It's fast, powerful, and always available (even when a GUI is not, such as when you want to make changes via the network), but it is a little bit different from the GUI than most people are used to. Fear not; we all have to start somewhere and in this section we'll make sure you're comfortable and at home on the command line.

In Part 3, we actually start to do more interesting things. By this stage you're happy on the command line (which puts you ahead of most of the crowd) and now you want to delve a bit deeper and actually make the Pi work for you. We spend this section covering some of the great things you can do. I'll give you the solid foundation you'll need to do all of the hardware and software projects in Brendan Horan's "Practical Pi Projects."

Your First Bite of Raspberry Pi

In Chapter 1 we look at what Linux is and why it's on your shiny new Raspberry Pi. We take a brief look at different "distros" and explain why not all "Linuxes" are the same. We take a closer look at the Raspberry Pi and what makes it just that little bit different, and then we get your graphical interface up and running. We round off Chapter 1 with a quick tutorial on how to move around the desktop.

Surveying the Landscape

Following on is a tour of the desktop and all the fun things that you can find. We start with the essentials and show you how to get online, check your e-mail, and get set up with instant messaging. We'll then expand on that a little and shoot for some music and video. By the end of this chapter, you should be able to do your day-to-day computing without having to think twice about using a Raspberry Pi.

Remodeling the Landscape

Once you're happily browsing the web and sending e-mail, we will look at how to go about customizing your new home and making it your own. We'll start with the obligatory "how to change your desktop background" and then move on to various other little tweaks and enhancements that you can do to make your time on the Pi as pleasant as possible.

Getting Comfortable on the Command Line

Now for something just a little bit different. You can use your Pi just like your other computer, but now we're going to take it up a notch. In order to get the full benefit from Linux and to get your hands dirty with some very interesting projects (automatic cat detection and prosecution, anyone?) you'll need to lift off the crust and start getting closer to the metal.

The first place to start is the command line where you'll be interacting with Linux on a very precise and powerful level. We'll be starting off slowly and easing you into what is likely to be a very alien environment. We'll discuss the shell, where it came from, and why it's important. We'll also look at a bit of the history involved because knowing where it came from (while fascinating in and of itself) will help you get the full benefit from the experience.

Storing Stuff and then (Hopefully) Finding It Again

Once you're up to speed on what the command line is, it's time for the most important subject. Most books start by telling you about all the shiny commands you can run, and that's all well and good, but first you need some context.

I'll start you off by looking at the Linux file system and explaining the "everything is a file" philosophy that sits at the core of every UNIX-based operating system. This sounds a bit scary and we won't go into too much depth, but with this under your belt you'll be able to easily pick up new commands and make full use of them.

Getting Down to Business

Now you know about the command line and how Linux lays out all your files, we'll actually showing you how to make Linux do things. I'll show you how to become root (does anyone remember the "power up" scene from He-Man? No?) and how to install new applications and tools. We will also briefly look at some useful command-line tricks that will make you a keyboard demon. You will soon find that you can do things far more quickly on the command line than by clicking your mouse . . .

Editing Files

Next on our tour de force is editing files. I will show you my two editors of choice. One (nano) is simple, easy, and great for general use. I use it a lot for making small changes to config files. For more serious heavy lifting, I prefer vim. It has all sorts of magic codes and key presses (it's like playing Dungeons and Dragons—not that I ever have), and if you can hold all the key-combos and commands in your head, well, you will never want to leave! For the rest of us, well, there's always nano...

From Beginner to Admin

Now we're starting to take a look at system services and what makes Linux tick. You'll learn that as smart as Linux is, it's not that different from all the tools and things you've been running so far. You'll be able to apply your newfound knowledge and be able to start and stop services to your heart's content.

Network services are the bread and butter of any server. Most people have used them from a user's perspective but now it's time to experience them from the pointy end. We'll look at the usual suspects, such as OpenSSH and the Apache web server. We'll explain the very real dangers of running your own servers on the Internet. There are people out there who would quite like to gain entry to your nice little server, and they do not much care who you are. Don't worry; we'll provide a simple list of do's and don'ts that will put you ahead of the game.

Magic Spells for the Command Line (aka Scripting)

A script is simply a collection of shell commands in a text file that Bash (or a shell of your choice) will execute. Even in this simple form, they are very powerful, but you can do so much more. In short, you can automate a lot of tasks with relatively little effort, and you can build some pretty impressive tools, too. This chapter won't make you an expert, but it will certainly give you a taste for the power at your fingertips. You might laugh at me calling them magic, but wait until you turn a tedious 3-hour task into a single command that completes in 4 seconds!

Pi Web Server

We start the project section with a classic: getting a full LAMP (that is Linux, Apache, MySQL and PHP) stack up and running on your Pi. We will show you why your Pi is an ideal little web server and how it not only provides the perfect environment for development dynamic websites but also makes a great little portable website demonstration tool!

WiPi: Making Your Pi Wireless

What's more fun than a tiny little network server? A tiny little network server that has no wires! We'll show you how to hook up a USB wireless network adapter so your Pi can talk to the world without any wires. If that isn't enough, we take it one step further and look at ways of cutting the cords altogether and getting your Pi to run off a battery.

The Raspberry sPi: Security Cam and Messaging Service

Want to know who is sneaking into your room and stealing all your pens? Your sPi has you (or your office) totally covered! In this chapter, we will show you how to hook up a webcam to your Pi and have it take snapshots of any movement in the room. Not only that; it will promptly e-mail you a copy of the intruder in real time! Even James Bond doesn't have one of these.

MC-Pi: A Pi for All Your Media Needs

Last, but by no means least, we show you how to turn your Pi into your personal media center. We show you how to stream video to your TV and how to act as an airport device for sharing music around the house. With high-resolution streaming video, we really start to push the Pi and show what it is truly capable of!

Onward

We've covered a huge amount already. We've looked at how computing has changed over the years and how this has not only had an impact on how we see and use computers in our daily lives but also it has affected the way computers are perceived in education and the level and depth of knowledge that students are taught. We briefly touched on the Raspberry Pi Foundation and how it hopes that the Pi will start to reverse this trend.

Next we looked at some of the more compelling reasons why you might want to get a Pi and why the Pi makes an ideal platform for development and experimentation.

We then looked at what the Raspberry Pi is at the hardware level and discussed the differences between your main computer and the Pi. We examined the trade-offs and looked at why some of the decisions were made and what impact they will have on you as a user.

Last but no means least, we rounded off the introduction with a whirlwind tour of all the things that we're going to cover in the rest of this book.

So let's head to the kitchen and make some Pi!

CHAPTER 1



Your First Bite of Raspberry Pi

This chapter is where we finally get our hands dirty. If you decided to skip the introduction in your eagerness to get up and running, that's fine; you won't have missed anything critical for this chapter. However, make a note to pay the introduction a visit in the not too distant future as it gives you a lot of background on the Pi and what makes it so special.

Now, back to dirty hands! We are going to start off by unpacking the Pi and going through the list of things that you need to actually get it up and running. Once we've hooked everything up, we'll then need to sort out something for it to run—in this case, Raspbian Linux (more on this a little later). Once we get Raspbian running, we still need to configure it and some of the options are a little technical and Linux-specific; but don't fret, we have you covered there, too. To round off the chapter, we will bask in the glow of the Raspbian desktop before heading off to Chapter 2 to see what we can actually do with it.

Your Freshly Baked Pi Arrives

Okay, so the postman has just dropped off your long-anticipated package, and after eagerly ripping away the padded envelope you're left with a little box, (*little* being the operative word). You simply can't help but look into the depths of the envelope to see if there is perhaps something lingering at the bottom. Failing to find anything with a visual inspection, you'll no doubt proceed to the old faithful approach of turning the envelope upside down and giving it a bit of a shake. Although you might shake loose the packing sheet (somehow they always seem to super glue themselves to the inside), you won't find anything else.

When you order a Raspberry Pi, this, dear reader, is all that you're going to get (see Figure 1-1).



Figure 1-1. What falls out of the envelope

WHY A PICTURE OF JUST A BOX?

The honest reason is because when this is all that turned up, our first thought was that something must have gone missing. Yes, the site tells you when you order that this is all you are going to get, and it really makes perfect sense.

We want to emphasize that this is all you get, and it's not enough to get your Pi hooked up and in business. You will need other bits and pieces (some or all which you might already have), and it's important to highlight it right at the beginning. In the next section, we will cover the kit you actually need to get started.

You might spot the big label at the bottom that (in big bold letters) says *GETTING STARTED*. It also provides the address of the Raspberry Pi website with a direct link to what you need to get started. Not too challenging, you might think, assuming that you actually looked at this side of the box rather than simply ripping it open to get at your new toy. We would love to say that the first thing we did was visit that website, but we were somewhat overexcited about seeing a Raspberry Pi with our own eyes. Admittedly, the complete lack of any writing on the other side of the box was probably a dead giveaway that we were missing something important, but we were somewhat preoccupied with showing everyone in the office just how small a Raspberry Pi actually is. Believe it or not, we have company in the *fail* corner; another technically savvy new Pi owner converted his box into a Pi stand and missed the label altogether (and his packing sheet was MIA for nearly three weeks).

Unsurprisingly, going to <http://raspberrypi.org/downloads/> is actually a very good place to begin your journey, and we will be heading there shortly. If you have yet to order a Pi (or are waiting for it to turn up in the mail), Figure 1-2 shows one sitting next to a mouse.



Figure 1-2. A Raspberry Pi next to an average-sized mouse

Even though the size of the Pi is well known and it comes in such a small box, it isn't until you actually hold it in your hands that you realize how small it truly is! After "aw-ing" to ourselves (and anyone else who sadly for them happened to be within earshot), we figured it was time to get the show on the road.

And that's where things came to a somewhat abrupt halt. We had the Pi, but it had only just dawned on us that we didn't have the faintest idea what we actually needed to make it go. Even though we're not exactly new to all this, we let ourselves get caught up in the moment and for whatever reason hadn't considered that the Pi would need anything special (or at least something that your average geeky type person wouldn't have hanging around the house).

If you're not equipped with a cupboard of odds and ends straight out of Dexter's lab, you might need to pick up a few things. Fortunately, they are all easy to find, and you should be able to pick up everything you need at your local computer shop. If you don't fancy going on a treasure hunt, many companies are offering starter kits that come with everything you need. Because these kits are rapidly changing, the best way to find what's available is to go online and search for "raspberrypi starter kit".

Note Although we only cover HDMI in this book (so it's the only display in the list), it is possible to use different types of display such as DVI, composite, or SCART with appropriate cables or adapters.

List of Ingredients

To bake your Pi to perfection, you're going to need the following ingredients:

- Raspberry Pi
- Micro USB lead (for power)
- USB power adapter (also for power)
- HDMI lead: Type A to Type A (to connect to your monitor or TV)

- HDMI display
- SD card: 2 GB to 64 GB (for storage); 8 GB minimum recommended
- SD card reader
- USB keyboard and mouse

Micro USB lead

If you happen to have an Android phone or a Kindle sitting about, chances are good that you've got a lead that you can reuse to power the Pi. This lead isn't actually used for transferring data, and although you can draw power from a USB port (you can plug it into your main PC or laptop), you can't use the USB connection for anything else. The different types of USB connectors are hard to describe if you haven't already seen one. Take a look at Figure 1-3 for some examples.




Figure 1-3. Different types of USB connectors

The connector you're interested in is the first one on the left, known as a *micro USB*. Be careful because on quick inspection, a micro USB plug can easily be mistaken for a mini USB plug (second on the left). The last thing you want to do is make a special trip to the store and then come back only to find out that you picked up the wrong one!

USB Power Adapter

You might actually be able to get away with this one. A normal USB port can power a Model A Pi (the one without built-in Ethernet), and for testing you can probably get away with doing the same with a Model B. That's all well and good as long as you want to have your Pi sitting next to your PC, but chances are you want it to be somewhat independent. Even if you do plan on using it in the same location, powering it from your laptop can prove to be

problematic when you want to take your laptop somewhere, but your Pi is busy doing something and you really don't want to unplug it.

Fortunately, the sheer number of devices that adopted USB as a means of charging means that you can get main adapters really cheaply and easily. As far as which adapter to get, that really comes down to personal choice. However, as the Model B requires 700 mA and you always want to have a bit of room for expansion, you probably should aim to get an adapter that can provide at least 1000 mA (or 1 amp). From our highly scientific tests (wandering around numerous shops and squinting at packaging labels), it seems that 1000 mA is actually the most common rating. We did come across some rated at 500 mA and although that's enough for most USB devices (and even the Model A Pi), it isn't really enough for your needs.

HDMI Lead

Over the past few years, HDMI has become the *de facto* standard in connecting a myriad of devices to both monitors and TVs. This is really handy because it means that if a device supports HDMI, it can be easily connected to any display that supports it. This might not sound all that impressive, but it wasn't too long ago that TVs and monitors were very separate things and usually there was no direct way to connect, say, a computer to a TV or a VCR to a monitor (although interestingly if you go back 20 years, all home computers connected directly to the TV (for example, the Commodore 64 or Spectrum). Of course, you could get special hardware, and some higher-end devices did offer a range of different connectors, but as a rule, the two worlds didn't really mix.

Fortunately, the Pi uses HDMI, so we can ignore the irritations of the past. To connect your Pi to a display, you will need a "Type A to Type A" lead. Type A is the size you will find in the back of your TV or monitor; all you have to do is find a lead for connecting some device via HDMI to a TV and where both ends are the same size. Most consumer electronics use Type A, so if you have an Xbox 360 or your laptop has HDMI, chances are you already have a lead of the correct type.

HDMI Capable Display

You probably saw this one coming, but you're going to need some sort of display that supports HDMI. As it has been widely adopted over recent years, practically every new TV comes with one or more HDMI ports. These days, it's fairly common to see three or four ports on a TV because you'll need that many to hook up all your new digital devices.

It's pretty easy to determine whether your TV or monitor supports HDMI. All you have to do is look for a physical HDMI port. You are much more likely to find it on your TV rather than your monitor, but many of even the more basic monitors these days seem to support it.

In our case, it turned out that the aging monitor didn't support HDMI, although the TV in the living room did. Of course, it being the main TV meant that people wanted to watch TV on it and admittedly we were not looking forward to sitting cross-legged on the floor in front of a big TV, trying to convince a Pi to boot. In the end, we decided to get a new monitor that supported it. Although we could have bought a DVI converter, we decided we couldn't pass up on the opportunity to get a new shiny toy.

If you can't reuse your TV or computer monitor, you should be able to get your hands on a basic TV or monitor that supports HDMI quite easily. As mentioned earlier, the Pi does have composite output as well as HDMI, so you could certainly use that if you prefer.

SD Card

Most computers use a hard drive of some sort as their primary form of storage. Even laptop-sized disks are larger than the Pi, and although the newer solid state disk (SSD) models draw very little in the way of power, they would certainly drastically increase the amount of power your Pi would need. Fortunately, we have an alternative. Rather than using something heavy duty like a hard disk, we can instead borrow technology that cameras have been using for many years: flash memory. Although the cards can't match a hard disk for space or for performance, they are exceptionally

good for power usage and despite being smaller than their erstwhile cousins; 64 GB (the maximum for your Pi) is still a reasonably impressive amount and probably more than you will need for your Pi.

SD cards are standardized so there isn't too much more to say about them. You do want to get a high-performance card if you can (often referred to as *Class 10*) but they ultimately all do the same job. That said, we have heard of some SD cards that haven't worked with the Pi, but if you stick to a well-known brand you should be fine.

As there are lots of different types of memory cards floating about these days, Figure 1-4 shows what an SD card actually looks like.



Figure 1-4. SD card with micro SD and adapter

The card on the far left is your typical SD card; on the far right is a micro SD card. They have become pretty popular for use in smartphones (particularly Android-based phones), and because most phones come with small cards (in terms of storage space), it's not unusual to either buy a larger card or have one thrown in as part of your phone package.

If you happen to have one of these floating about, you can use it with your Pi by using a micro SD adapter (shown in the middle of Figure 1-4). They tend to come with micro SD cards because often one would use the micro SD in a phone but the laptop would only have the full-sized SD card. We haven't seen the adapters for sale individually as they are basically just a physical adapter, but they do tend to end up in the strangest of places so you never know!

SD Card Reader

Now that you have your SD card, you need some way to actually use it. The Pi has you covered and has an SD card slot built right in. However it doesn't actually have any other storage on the device, so it's effectively a blank slate until you insert an SD card with something useful (like Linux) installed on it. The thing is, to put Linux on the card, you need a device that's already up and running that can also read and write to the card. This is a classic Catch-22 situation. Even if you could borrow someone else's card with Linux already installed, you couldn't simply swap out the cards because as soon as you do that, you won't have Linux anymore!

Again, thanks to the proliferation of digital cameras, many computers come with card readers built in. Many (often dubbed as *media PCs*) come with a whole range of slots for various different card types. So chances are you already have a way to read the card. If you don't have a reader already, you can pick up any cheap multicard reader from your local computer store. They're generally inexpensive and will support lots of different types of cards. Just make sure that it has an SD card before you hand over your hard-earned cash! For reference, the adapter we are using looks like the one in Figure 1-5.



Figure 1-5. Multicard reader

USB Keyboard and Mouse

Last but not least is the good old keyboard and mouse. Advanced as the Pi is, it doesn't yet have telepathy so you're going to need some way to control it. This is common sense, but with modern computers often coming with a wireless keyboard and mouse and a fair few of those using Bluetooth with no USB adapter (even some with an adapter have been known to cause problems), you might find that your current keyboard and mouse just won't work with your new toy (no Bluetooth on the Pi, I'm afraid). There's really no need to go into any detail here as any keyboard and mouse will do, and this is mostly a reminder to check what you actually have before you go out to buy your Pi-making ingredients to avoid untold frustration when you get home and find you're completely stuck.

Whew, We're Done!

Finally we have everything we need! Hooking all this stuff up to your Pi is very easy, especially as each item has a unique shape so it will fit into only one slot (see Figure 1-6 for the finished Pi).



Figure 1-6. A fully loaded Pi

If you put everything together, plug in the power, and turn on your TV you should see... a completely black screen.

Don't Panic!

When your PC or laptop boots, there is a piece of software called the BIOS that kicks everything off for you. It tests the memory, sets up a basic display, and allows all your devices to initialize. On some machines (notably those from Apple) the machine will have an EFI instead of a BIOS. For all intents and purposes (at least from our point of view), they're basically equivalent. Regardless of the technology used, it is this system that finally hands control over to your bootloader, the piece of software responsible for starting your operating system.

BIOSs are by nature very noisy, and if you get something wrong it will either bleep a lot (with some magic number code that you can only find in the BIOS manual that you threw in the bin 3 years ago) or display some helpful yet cryptic message on the screen. In short, it may not do what you want, but at least you know that your computer is still alive. Although it has a little red LED "on light," the Pi won't do a thing unless it has a bootable operating system on its SD card. If you were expecting some sort of splash screen or other sign of life (hands up; we know we were), you're out of luck (and probably thinking you have a dead Pi).

Now that you know what you need to restore the Pi to life (a bootable operating system) it's time to move on to Phase 2 of our master plan and get us some Linux!

Linux

As this is a book for beginners, we're going to take the time to talk a little about Linux, where it came from, what's special about it and some of the little gotchas that trip people up from time to time. If you already know this stuff (or simply don't care), feel free to skip to the next section "Downloading Raspbian." For those who want a quick refresher, just keep on reading.

What Is Linux?

Ah, this simple question opens a big can of worms that many people go out of their way to avoid. The reason is that in technical terms, Linux means one thing but in general speech, it tends to mean something else. When talking about operating systems in general, we see Microsoft Windows and Apple OSX to be discrete, whole things. If you say “I run Windows,” everyone knows what you’re talking about. With Linux, it’s a bit different.

Linux is just an operating system kernel, which means that it handles all the low-level bits and pieces such as handling device drivers, and providing easy access to networks and hard disks. What actually makes Linux usable is all the software that is wrapped around it. Not much trouble there, but it starts to get complicated when you realize that people have differing opinions about what software should be wrapped around it. There are no simple or minor opinions when it comes to computing!

As this software is open source, and anyone can put it together in pretty much any way they like, people have been able to build their own Linux distributions. This is an operating system with Linux at its core, but with the surrounding ecosystem set up to match the goals of the people who built it. For example, Red Hat Enterprise Linux (RHEL) is built to be robust, supportable, and stable over long periods of time. Fedora, on the other hand, is released every 6 months or so and has the latest and greatest of everything in each build. Gentoo requires that you build your software from source (so it can be completely optimized for your machine), and Debian goes to great lengths to remain stable and secure at the cost of introducing new features.

OPEN SOURCE

In the olden days, when computers cost millions of dollars, it was the machine itself that sold, not the applications that ran on it. It was common practice for a company to write software for you, in order to get you to buy their machine. It wasn’t until computers became commodity items that suddenly the value was in the software, and companies moved to start protecting the software they’d written. One way to do this is to provide the software ready to use without the source code (which is effectively the blueprint on how to make the software). This meant that you couldn’t change the software or make improvements to suit your changing needs.

Many people believe that software should always come with the source code so that people can make changes. Open source revolves around the idea of people being able to freely exchange, modify, and improve software collectively. This can only be done when the source code is provided freely and permission is given (through an open source license) to make those changes.

Entire books can and have been written on this topic, but you can find a presentation on this topic that Peter gave at his university last year: <http://www.youtube.com/watch?v=c-1LQIGH6cI>

So which one is best? Well, that depends on your needs! There’s no perfect distribution; just the best fit for a particular job. For the Pi, the official and supported platform is Raspbian, which is based on the Debian distribution. Because it is supported and because it’s the easiest to use (and quite likely the fastest to get updated and fixed when things go wrong), we’re going to stick to using Raspbian in this book. If you do fancy something a little different, Brendan Horan’s upcoming book (also from Apress), *Practical Raspberry Pi* will show you in great depth how to install Fedora and (if you’re feeling particularly brave) how to make a custom build of Gentoo!

When the Linux kernel first debuted in the early 1990s, no one really appreciated the huge impact it would have on the world of computing. Open-source software had already been around for a long time before this, and countless tools for the UNIX platform had already been released (such as the awesome GCC compiler). However, they were just tools and software packages. They still needed a proprietary operating system to run (the only kind available at the time). To be fully compatible with the open-source ethos, what was needed was an open-source kernel to power these systems, and this is what Linux delivered. While many will talk about how this brought freedom and hope to the world of computing, we will save you from that particular lecture (although admittedly there is quite a bit of truth to it) and simply say that it was the Linux kernel that really brought open source to the eyes of the general public.

Now that the world has Linux, what exactly can we do with it? Almost anything we like—such as installing it for free on our Pi.

Downloading Raspbian

At last we're off to <http://raspberrypi.org/downloads/> to get our hands on Raspbian. Not only does the page contain links for the download we're looking for but it also has links for various useful tools as well as links for other prebaked Linux distributions for the Raspberry Pi. We touched on this lightly in the previous section, but if you're thinking about getting one of the more exotic distributions, take a look at the "Why Raspbian?" sidebar.

WHY RASPBIAN?

As you can see from the download page, there are numerous distributions to choose from, and especially if you've used Linux elsewhere you might be tempted to pick something other than Raspbian. However, before you download that exotic distribution, you should know that the majority of people using a Pi run Raspbian—and as such most of the support, blog posts, and tutorials (not to mention this book) are all based on it. This means if you pick something else, you might find it much harder to get help if something doesn't go quite according to plan.

The other more subtle issue is that the Raspberry Pi is an ARM-based device. It doesn't quite work the same way as your PC does. For example, the Pi doesn't have a PCI bus, so any tools that expect there to be one (and there are more than a few) just won't work. There are also quite a few gotchas for these other distributions, and they might not be as well maintained as Raspbian. In short, you will be taking on much more work if you don't pick Raspbian to start out.

Traditionally, a Linux distribution will come on a set of CDs or DVDs. When you download your distribution of choice, you would download a set of images. These images are basically a clone of the real thing. The idea is that you take the image and once you write it to a CD, that CD will be identical to the one the image was originally cloned from. Raspbian (and the Pi in general) is a bit different. It doesn't have a built-in CDROM or even a floppy drive. Instead of downloading an image of a CD or DVD, we're going to download an image of an SD card that contains Raspbian all set up and ready to go.

You can download Raspbian either directly from the site (there are numerous mirrors) or via BitTorrent. BitTorrent can sometimes be faster, and it helps take some of the load off Source Forge. However, there are some places where you probably shouldn't be firing up a BitTorrent client (i.e., at work), and some connections perform very badly due to the way BitTorrent works (such as mobile 3G). If in any doubt, just select the direct download. We want the mainline version of Raspbian (Raspbian Wheezy), which is the first one in the list of available distributions for download.

Getting Raspbian onto Your SD Card

Getting Raspbian onto your SD card isn't as easy as simply copying across the file. Remember we downloaded an image that is a direct representation of what should be on the card. If you just copy the image onto the card, it will still just be a card with a copy of the Raspbian image on it. What we need to do is write the image directly to the card, and for that we're going to need some of those handy tools we briefly mentioned earlier.

Note The image file you downloaded is a compressed ZIP file. On Windows 7, double-clicking this file will open a Compressed Folder. You will need to copy the image file itself to another location such as My Documents. On the Mac, simply double-clicking the file will extract the image and place it in the Downloads folder.

Using Image Writer on Windows

The tool recommended for writing images on Windows is the somewhat predictably named Image Writer. You can find the link on the Raspberry Pi downloads page; you want the binary version highlighted in Figure 1-7.




Figure 1-7. Download the binary version of Image Writer

Depending on what software you have installed on your PC, you will either be able to open the download as a folder or open it in your ZIP program of choice. Either way, you need to extract all the contents before you can run the application. If you are just using the standard Windows tools, when you double-click the program (Win32DiskImager), you will be prompted to extract the contents to another location.

1. Select Extract All and make sure that Show Extracted Files When Complete is checked on the next window.
2. After you press Extract, you should see the same folder layout as in the ZIP file.
3. Double-click the program in this new folder and you should be able to open the tool. Because the program isn't digitally signed, you will get a warning like Figure 1-8.




Figure 1-8. Windows will warn you when you run Disk Imager

4. Because the tool needs low-level access to perform its main function (i.e., writing directly to disk-based devices), Windows will warn you that the program wants to make changes to your computer. Click Yes to continue.
5. We're not sure if you will see the message shown in Figure 1-9, but it happened on both the machines we tested it on. In both cases, it didn't seem to affect the program, but just in case you do see it, we thought you'd like to know you're not the only one!




Figure 1-9. Error occurred when opening Image Writer

Note Image Writer will look only for cards to write to when it starts up. That means you'll need to have your SD card in the slot before you start the application; otherwise, you will get an empty list, even if you later insert the card.

6. In our case, the device we want to write to is drive G:\, but you should choose the right one for your system. If the Device drop-down menu is empty, try closing Image Writer and reopening it as sometimes it seems to miss things on the first try. It goes without saying that you should make sure that you've actually picked the right drive—the last thing you want to do is accidentally erase the card from your camera and all those special family moments that you've captured over the years!
7. Once you've selected the right device (take a quick look in My Computer if you're not convinced) and you've picked the Raspbian image by browsing to it, click the Write button and sit back and relax. It does take a fair bit of time to write the image (it depends on how fast your card is, among other things). However, it gets faster as it goes along (it's mostly writing zeroes at this point), so it will probably be done within a few minutes. Figure 1-10 is an action shot of Image Writer during our tests. From the screenshot, you can see that the image we're writing to the SD card is the current (at least at the time of writing) Raspbian release.




Figure 1-10. Action shot of Image Writer writing the image

8. Once complete, take the card out of its slot and feed it to your Pi.

You're ready to rock and roll! To get started, flip through to the section, "First Boot," to get the show on the road!

Note If you look closely at the filename in Figure 1-10, you can see that it has an extension of .img. One of the more common mistakes when writing an image is to accidentally write the ZIP file rather than the decompressed image file (after all, they have almost identical names). When you come to boot up your Pi, if nothing happens, come back and double-check that you picked the right file.

Using dd on the Mac

If you are using a Mac, we want to show you how simple it is to write an image to your SD card. Alas, we can't do that because there isn't a simple, easy way to do it. Disk Utility, which excels at so many other things, simply does not know how to handle our standard image files. We tried converting it to a .dmg file (the native image format for a Mac), which didn't work, either. We also looked at online tools that you might have been able to download to help with the problem, but we didn't find anything that made the task significantly easier than doing it by hand in the first place.

So we're going to get our hands dirty on the command line a little bit earlier than planned. We're not going to go into too much depth about what it is we're actually doing because most of these commands you'll see a little later in their Linux incarnations. For now, you'll just have to trust us.

Finding the Terminal

First, we need to open the terminal. You can find it in the Utilities directory under Applications. To put the following into context, we end this section with a complete transcript of all the commands that we used. The process is reasonably straightforward; it just has a fair number of moving parts.

We opened the terminal and now we need to use sudo to become the root user, the equivalent of Administrator on a Windows machine.

Now, tinkering with low-level devices is dangerous for security reasons (you don't want a virus being able to write directly to a hard disk, for example), so a normal user is not allowed to issue commands at that level. Because we need to write our Raspbian image directly onto an SD card, we simply have no choice: we need direct access to the card. Fortunately, root has effectively unrestricted access to the machine, so as root we can write the image to the card. The command we need to use is this:

```
sudo -i
```

This command opens an interactive prompt as the root user and gives “super user” privileges. You can use sudo to execute commands directly, but because we are going to be executing a few of them, it's more convenient to do it this way. sudo will then prompt you for your password (the one you normally use for your user account). Once sudo can confirm that it really is you sitting at the keyboard, you will end up at a root prompt.

Caution We now have permission to do anything we like on the system. And although that means we can now make our image, it also means that we could really mess things up if we aren't careful. Under UNIX, it is assumed that root knows what it is doing and there is very little in the way of safeguards. Be very careful when running any command as root, double-check it (and then do it again), and only when you're completely sure should you hit the Enter key. As soon as you're done, you should also close the terminal window, just in case.

Using the Terminal to Write the Image

Okay, so what now? One of the challenges with writing to the SD card is to know which device on the system that card actually is. You don't want to accidentally mistake your main hard disk for it because that would have some rather unpleasant consequences. There are generally some assumptions you can make (for example, we know that /dev/disk0 is usually the system disk, so you never ever want to write to this), but in order to limit the risk, we're going to do this one by the numbers.

First, make sure your SD card is not attached to your Mac and then run the `mount` command like so:

```
mbp:~ root# mount
/dev/disk0s2 on / (hfs, local, journaled)
devfs on /dev (devfs, local, nobrowse)
map -hosts on /net (autofs, nosuid, automounted, nobrowse)
map auto_home on /home (autofs, automounted, nobrowse)
mbp:~ root#
```

Now, you will probably see something similar to this, but it might be a little different. We can ignore `devfs` and the `map` lines as they're really just parts of the operating system. What we're interested in here is the first line that identifies `/dev/disk0s2` on `/`. This is the system disk, and we want to make sure we don't touch this one. You might have additional entries in your list if you have any network drives mapped or disk images attached. That's okay because we're not really looking for anything in particular at this stage; we just want to establish a base line for your system.

Okay, now it's time to slot in your SD card. SD cards come preformatted for use in cameras (although these days many cameras will reformat them, anyway) and so you should see it show up in Finder under Devices. Our card looked like Figure 1-11.




Figure 1-11. What you'll see in Finder after inserting your SD card

As the card has been automatically mounted, we can now go back and rerun our `mount` command to see how things have changed. Here's what we got when we reran it:

```
mbp:~ root# mount
/dev/disk0s2 on / (hfs, local, journaled)
devfs on /dev (devfs, local, nobrowse)
map -hosts on /net (autofs, nosuid, automounted, nobrowse)
map auto_home on /home (autofs, automounted, nobrowse)
/dev/disk2s2 on /Volumes/UNTITLED (msdos, local, nodev, nosuid, noowners)
mbp:~ root#
```

What we care about here is the addition of a new disk, in this case `disk2`. We can see from where it is mounted (`/Volumes/UNTITLED`) that this is the disk we are interested in. So now we know where we want to write our image.

■ Info BSD systems split disks out in slices. In our example, the mounted file system was on `disk2s2`. This translates to `disk2`, slice 2. This allows you to reference a slice directly, but for our purposes we want to write to the device. So in our case we will be writing to `/dev/disk2`.

There is one little wrinkle, though. We can't write to the SD card directly while we have it mounted and usable. If we eject it from within Finder, it will not only unmount the filesystem but it will also eject the device as well. That doesn't help us because we still need the device to be present in order to write to it. To solve this issue, we will unmount the filesystem manually using a Mac-specific command called `diskutil`. The Mac does have the general

UNIX `umount` command (which we'll cover in later chapters), but if anything on your Mac happens to be looking at the SD card, `umount` will fail and claim that the device is busy. `diskutil` is aware of this and can usually unmount the device without any problems. We ran the following command:

```
mbp:~ root# diskutil unmount /dev/disk2s2
Volume UNTITLED on disk2s2 unmounted
mbp:~ root#
```

So we gained super-user privileges, we isolated the device that we want to write to, and we unmounted the filesystem so that we can write directly to the card itself. All we have to do now is actually go ahead and do the writing!

The tool of choice for this task is called `dd`. This tool is a bit old school in that it only cares about reading from a device and writing to a device. It has no interest in filesystems nor is it aware of any of the subtle differences between a hard disk and an SD card. It just reads and writes without caring where it reads from or writes to. This is a fundamental UNIX principle that we will be revisiting in later chapters. For now, we're going to take advantage of this to make an exact copy of our image file on the SD card.

`dd` only needs to know two things: where to read from (the `if=` argument) and where to write to (the `of=` argument). We figured out the writing bit (it's going to be `/dev/disk2`), but what about the source? If you followed the earlier instructions, you should have the image file sitting in the `Downloads` directory in your home area. The easiest way to get the information we need into terminal is to type this in:

```
dd if=
```

Then go to your `Downloads` directory in Finder (or wherever you extracted the image file to), click and hold on the image file, and then drag it on to your terminal window. This should then paste in the full path for you without you needing to type anything in. This should mean you end up with something pretty similar to this:

```
dd if=/Users/myuser/Downloads/2012-08-16-wheezy-raspbian.img
```

Now we have our input file specified, all we need to do is specify where exactly we want `dd` to write it. Thanks to our earlier experimentation we know that this will be `/dev/disk2`. We're also going to add `bs=512` which lets `dd` write the image in bigger chunks (meaning you won't have to wait four hours for it to complete). So the full command we need for `dd` looks like this:

```
dd if=/Users/miggyx/Downloads/2012-08-16-wheezy-raspbian.img of=/dev/disk2 bs=512
```

Running this command might take some time (hours, not seconds). The image is about 2 GB in size, and your average smart card reader is not particularly fast. With our USB adapter, our recent MacBook Pro took a good two hours to actually write the image. After coming back later, you should find that `dd` has completed its task and you're back at the command prompt.

Okay, admittedly this wasn't quite as simple or straightforward as writing the image on Windows, but if you look at everything we had to do, using the command line allowed us to express what we wanted in a very compact and precise way. This is something we come back to in Chapter 3, in which we look at the various benefits of learning to use the command line.

Here is the full transcript of the commands we entered in order to write the image to the disk. If something in our explanation didn't quite click, or you're not entirely sure what the end result should look like, you can use the following transcript as a reference to compare your version with ours:

```
Last login: Wed Sep 19 13:48:40 on ttys000
mbp:~ pmembrey$ sudo -i
Password:
mbp:~ root# mount
/dev/disk0s2 on / (hfs, local, journaled)
devfs on /dev (devfs, local, nobrowse)
```

```

map -hosts on /net (autoofs, nosuid, automounted, nobrowse)
map auto_home on /home (autoofs, automounted, nobrowse)

<<SD Card inserted>>
mbp:~ root# mount
/dev/disk0s2 on / (hfs, local, journaled)
devfs on /dev (devfs, local, nobrowse)
map -hosts on /net (autoofs, nosuid, automounted, nobrowse)
map auto_home on /home (autoofs, automounted, nobrowse)
/dev/disk2s1 on /Volumes/UNTITLED (msdos, local, nodev, nosuid, noowners)
mbp:~ root# diskutil unmount /dev/disk2s1
Volume UNTITLED on disk2s1 unmounted
mbp:~ root# dd if=/Users/pmembrey/Downloads/2012-08-16-wheezy-raspbian.img of=/dev/disk2 bs=512
mbp:~ root#

```

With our (painstakingly) prepared SD card, we can finally move on to the next step.

First Boot

Okay, that was a bit of a marathon, but hopefully it's something you will only need to do once. Once you get the hang of it, it's not so bad; it's just that there are quite a few moving parts.

Armed with our freshly imaged SD card, we can at long last boot up the Pi. Plug in your keyboard, mouse, and monitor, and slot your SD card into your Pi. All you have to do now is attach the power lead and you're good to go! All being well, you should see a screen like Figure 1-12.




Figure 1-12. The Pi's first boot—it lives!

If you can see something similar, congratulations, your Pi is up and running! If not, there are a few common things that might have gone wrong.

The first thing to check is whether your Pi has a nice red glowing power LED. Now we know this sounds obvious, and we know that it's probably the first thing you checked, but not checking it is also the most common mistake to make. If your desk looks anything like ours, you will have a fair collection of cables draped about. It's not too much of a stretch to see how you might have plugged in the wrong cable or that the cable you're using isn't actually plugged in to a live power source (some laptops, for example, turn off USB power when they go to sleep). Either way, this is what you should check first because if you did make a mistake, only you will know about it. This is much better than having an extended discussion on the forums trying to get help only to then realize that whoops! you forgot to plug it in.

The second thing to check is whether you have selected the right input on your TV or monitor. Most TVs have multiple inputs these days, and monitors often do not default to HDMI (ours doesn't), even if that is the only input signal they are receiving. Obviously, we can't tell you exactly how to fix or check this one as it all depends on the equipment you're using. However, it should be fairly easy to flick through the different inputs to see whether or not your Pi turns up. In one interesting case, even though on the TV the port clearly said HDMI2, the Pi was actually visible on HDMI3. We have no idea why, but had we thought to flick through the inputs before fiddling with the Pi, we could have saved ourselves a fair bit of grief.

The third thing to check is that the SD card is fully inserted into its slot. The Pi will quite literally do nothing unless it can find the card and boot from it, so if all you're getting is a black screen, this could well be the reason. If you are sure the card is inserted properly, but you're still not having any luck, it might be that something went wrong when you wrote the image to the card. In that case, flip back to the last section and have another go at writing the image to the card. If you still aren't getting anywhere and you've checked that everything is all plugged in nicely, you have the right input on your monitor, and you have a happily glowing power LED, then it's probably something more onerous. Unfortunately, without being there with you to look at your poorly Pi, we can't really give any more specific advice. Fortunately, you're not alone; there are lots of people on the Raspberry Pi discussion forums who are ready to help out. To visit the forum, go to www.raspberrypi.org and click the Forum link (just below and to the left of the Raspberry Pi logo). This will get you to the main forum page; the particular board you're looking for is Basic Usage and Setup under Using the Raspberry Pi.

Configuring Your Pi

After the torrent of incomprehensible text (though we're sure it means something to someone) you should end up at the configuration shown in Figure 1-13. This is where we'll configure your Pi for general use so that next time you turn it on, you will be greeted with a nice desktop rather than some chunky text.




Figure 1-13. First boot: time to configure your Pi!

We are only going to cover the basic configuration options here (and you can always go back and change them once you've got your Pi up and running) as we want to get you in front of a working computer as soon as possible. To start with, ensure that your Pi is configured to boot to the desktop (select the `boot_behavior` option to make sure that the desktop will appear next time you boot your Pi).

Expanding the Filesystem

The first option of interest is `expand_rootfs`, which, as the description suggests, will expand the filesystem to fill the entire card. The image we downloaded was around 2 GB after we decompressed it and this is what we wrote to our SD card. So although you might have an 8GB card or larger, you will be able to see only the first 2 GB of it. By expanding the filesystem, we can take advantage of all the space that you have available. This approach works pretty well because it means we don't need a different image file for each disk size; we can all use the same small image and simply expand it to fit our needs. Using the arrow keys, move the bar down to highlight `expand_rootfs` and press Enter. You should see a window that tells you that the partition has been resized, and the filesystem will be resized after you next reboot. Press the Enter key to get back to the main menu.

Configuring the Keyboard

The next option of interest is `configure_keyboard`. By default, Raspbian assumes a UK keymap, but chances are you will want something else. Be warned, though, that Raspbian knows a huge number of keyboards, and the sheer quantity can be confusing. For instance, do you know if you have a generic 104 key keyboard or is it actually a 105 key version? We weren't sure either, so we stuck with the default: the Generic 105 key (Intl) PC shown in Figure 1-14.




Figure 1-14. Choose your keyboard

If you can find your exact model, by all means feel free to select it. If you aren't sure, you're probably best off following our lead and sticking to the default. Once you have settled on your keyboard of choice, press Enter to move through to the next screen.

Now we get to choose which keymap we want to use. By default, it lists various UK keymaps (as shown in Figure 1-15), but if you aren't lucky enough to own a British keyboard, move the bar down to Other, press Enter, and then scroll through the list to find something more suitable for your needs. Once you've found the right keyboard, again just press Enter to move on to the next screen.




Figure 1-15. Choose your keyboard layout

If you're not sure what to select from the menu (shown in Figure 1-16), you're not alone. For most people, the keyboard defaults will be fine because they will include any special key combinations that your particular keymap already offers. This is one of those times when if you need to make a change, you probably already know why and what to choose. If you are in any doubt, just leave the default selected and press Enter. Now you will see another menu that looks suspiciously like its predecessor (shown in Figure 1-17), at least in terms of knowing what to pick. Again you probably already know if you need to change anything here; if you're not sure, you should just select the default by pressing Enter.




Figure 1-16. Customize your keyboard configuration




Figure 1-17. Customize your keyboard some more

The next menu (shown in Figure 1-18) only really means anything to old-school Linux users. Back in the day, if your GUI decided to crap out on you, you could effectively “reboot” the session by holding down Ctrl + Alt + backspace. As the GUI has become more stable, and Linux has become more mainstream, this feature is more often than not disabled. We haven’t had to make use of it for years, although we do happen to think it is quite useful. The default is No, and this is probably what you want. You can always go back and change this later if you find yourself craving this particular feature. To pick the default (and finally return to the main menu), just press Enter.




Figure 1-18. Do you want to enable the X server kill switch?

Changing the Password

Next we’re going to change the password for the pi user account. This is the account you’ll be using throughout the book, and for most people this will be the only account you’ll ever need on your Pi. At the moment, although the password is set to raspberry and it is the default, it means that anyone can log in to your Pi. While securing your Pi against unauthorized intruders is probably not top of your list of priorities, it is still good practice to change this just in case. Later on, you might make your Pi available via the Internet and in your excitement you might forget that you didn’t actually get around to changing the password. Bad things will happen. So let’s take the opportunity to make the switch now. Highlight `change_pass` and press Enter, and you’ll see something like Figure 1-19.



Figure 1-19. Time to set your password

The configuration tool has executed the standard `passwd` command, which is why you're looking at a text prompt rather than a prettier menu. Just enter your new password (remember that it is case-sensitive: *HeLLo* is not the same as *hello* or *Hello*) and press Enter. For security reasons, the password won't be shown on the display and it won't show any stars to let you know you've actually typed something. You will then get prompted to enter your password again; after you press Enter, you will be back at the configuration menu. (Getting a sense of *deja vu* yet?).

Configuring the Locale

Next is configuring your locale. By telling Linux where you are in the world, it can determine useful things such as the separator between numbers and which currency symbol to use. For example, in the United States we might write 1,000.00, but in many places in Europe it would be written as 1.000,00. Either way for maximum comfort and to ensure that everything just looks right, we need to set the locale. Highlight the `change_locale` option and press Enter.

You'll now be presented with an extensive list of different locales to choose from (see Figure 1-20). The first two letters in the name specify the language while the second pair of letters specifies a regional difference. For example, `en_US` would be English customized for US users but `en_GB` would have settings specific to English speakers in the UK. You should select the language and country pairs that most closely match your needs and choose them by pressing the spacebar. You probably want both the ISO-8859 and the UTF-8 versions. If you decide to hedge your bets and select all locales, be warned that you will be in for quite a wait as the Pi generates all the locale settings for you. In fact, we tried this option and after waiting ages for just two locales to be generated, we got fed up and pulled the plug. After you select the locales you want, press the Tab key once to highlight the OK button and then press Enter to move to the next screen. Here you'll need to pick the locale you want to use by default (see Figure 1-21). Once you have made your choice, highlight it and press Enter. Once again, we're back at the main configuration page.




Figure 1-20. Select which locales you want to use




Figure 1-21. Now pick which locale you want to use by default

Changing the Time Zone

The `change_timezone` option does pretty much what it says. After selecting this option, you'll be shown a list of regions (shown in Figure 1-22). In this case, we were configuring Pete's Pi, so we selected Asia. The next screen will give you a chance to fine-tune your selection and pick your country or city (see Figure 1-23). Once you've found the closest match to you, select the option and press Enter.




Figure 1-22. Select your region




Figure 1-23. Select your country or city

Allocating Memory

The `memory_split` option needs a bit of explanation. Your Pi has two main processors: the CPU (or central processing unit) and the GPU (or graphics processing unit). The CPU is responsible for the general running of your machine, and the GPU is responsible for handling the display and provides features such as 2D and 3D acceleration as well as hardware support for decoding high-quality video streams. A GPU is dedicated to handling graphics (for the most part) and cannot perform the role of a CPU. However, it does mean that if the CPU can offload all the complex graphics work to the GPU, the CPU doesn't have to be quite so powerful. By having both, the Pi keeps the costs and power requirements nice and low but without sacrificing much in the way of performance.

So far so good; but what does that have to do with splitting memory? The Pi comes with 256 MB of RAM in total and it needs to supply both the CPU and GPU. As the GPU is effectively an independent unit, it needs its own allocation of memory. Rather than having dedicated memory for both, it is possible to split the 256 MB between the two processors. By default, the Pi allocates 192 MB for the CPU and only 64 MB for the GPU. This is ideal if you want to use the Pi as a server or you never plan to attach a display to your Pi. If you're not using graphics, allocating additional memory to the GPU is somewhat pointless.

On the other hand, if you plan to use your Pi as a mini desktop computer or to display high-quality movies, you're going to want as much memory allocated to the GPU as you can get away with (which in the case of the Pi means splitting it down the middle and having 128 MB for each). It would be difficult to anticipate how everyone will want to use their Pi, and apart from adding more memory to the Pi itself (and raising the cost) there is no easy solution. Fortunately, we have the ability to choose how we want to allocate the memory ourselves.

For general usage and "getting started"—type tasks, we recommend you split the memory 192 MB/64 MB. To set this up, select the `memory_split` option from the menu, press Enter, and you'll have something similar to Figure 1-24. Type in the memory split that you want and then press Tab to highlight the OK button and then press Enter.




Figure 1-24. Choose your memory split

At Last! It's Configured!

Wow, we finally made it to the end! After all that effort, you can now reboot the Pi and (finally) start using it! Press the Tab key twice to highlight the Finish button and press Enter. Your Pi should shut down and then reboot. If you decided to expand the root filesystem earlier, you'll see this (as promised) during the boot process, as shown in Figure 1-25.



```
[info] Loading kernel module snd-bcm2835.  
[ 24.452557] bcm2835 ALSA card created!  
[ 24.468005] bcm2835 ALSA chip created!  
[ 24.479396] bcm2835 ALSA chip created!  
[ 24.490482] bcm2835 ALSA chip created!  
[ 24.498896] bcm2835 ALSA chip created!  
[ 24.507291] bcm2835 ALSA chip created!  
[ 24.515565] bcm2835 ALSA chip created!  
[ 24.523608] bcm2835 ALSA chip created!  
[ ok ] Activating lom and md swap...done.  
[....] Checking file systems...fsck from util-linux 2.20.1  
done.  
[ ok ] Mounting local filesystems...done.  
[ ok ] Activating swapfile swap...done.  
[ ok ] Cleaning up temporary files....  
[ ok ] Setting kernel variables ...done.  
[....] Starting resize2fs_once:resize2fs 1.42.5 (29-Jul-2012)  
Filesystem at /dev/mmcblk0p2 is mounted on /: on-line resizing required  
old_desc_blocks = 1, new_desc_blocks = 1  
Performing an on-line resize of /dev/mmcblk0p2 to 1925120 (4k) blocks.
```

Figure 1-25. Resizing the filesystem on first reboot

This process can take a while depending on how big your SD card actually is. As it is being resized online, it won't actually interfere much with the startup process, but because it's happening in the background, and it's doing fairly intensive work on your SD card, you might notice that things feel a bit sluggish. That's completely normal, and it will speed up considerably once this process has finished.

After all that effort (and a bit more waiting while the Pi boots up), you should finally be at a pretty graphical display that should look something like Figure 1-26.




Figure 1-26. At long last, we have a desktop!

Congratulations! You now have a fully functional Raspberry Pi! If you don't end up with a pretty desktop, you might need to tweak your configuration a little. At the login prompt, you can log in with the username `pi` and the password you set during the install. If you then run the following you will be able to change the settings (`boot_behavior` is the option you want to change):

```
$ sudo raspi-config
```

You can also try starting the graphical interface manually with this:

```
$ startx
```

Summary

This chapter has taken you from getting your hands on your Pi all the way through getting it up and running. We deliberately took it slow and covered all the basics as it's always the so-called simple things that trip people up (us included!). We looked at what comes in the box (and more importantly what doesn't) and provided a shopping list of all the things you'll need to bake your Pi.

We then touched briefly on Linux—what makes a distribution and why we decided to stick with Raspbian for this book. We then covered how to download Raspbian and get it on your SD card (which turned out to be much more involved than we first thought).

Armed with a fully loaded card, we fired up the Pi and walked through the initial setup and configuration process that ultimately ended with us looking at the very pretty Raspberry Pi logo in the middle of an equally impressive desktop.

In the next chapter, we'll start to explore your new desktop. We give you a quick tour of what comes with Raspbian out of the box, and we finish up by showing you how you can tweak and customize it to suit your needs.



Surveying the Landscape

Although the majority of this book focuses on getting you comfortable on the command line and really getting your hands dirty with Linux, we would be completely remiss if we didn't first show you how to use the graphical interface that comes with your Pi. It's not quite as flashy as what you might find on a Windows PC, a Mac, or even some Linux distributions, but it's fast, easy to use, and will let you do all the things you're used to doing. And, of course, it's all on a device that fits in your pocket and draws less power than the monitor or TV it's connected to.

We start off this chapter by looking around the desktop and seeing what interesting things we have to play with. Some are obviously useful applications in their own right and some are probably less useful at this stage (for example the Python development environment won't be of much interest to you unless you're programming in Python). We'll then trundle through the various menus and have a look at what we find.

Welcome to LXDE

Unlike Windows, Linux isn't integrated with a particular GUI; it uses a separate application (a client and a server, in this case) to provide one. This application is known as the *X-server* and it used to be by far the most complicated thing in the Linux universe to set up. Although it has gotten a lot better, not too long ago, great time and effort needed to be spent if you wanted to use the scroll wheel in that newfangled mouse of yours. Needless to say, the fun of that wears out pretty quickly. These days, the X-server is almost a black box and very rarely do you need to actually touch it. For that reason, we won't go into any more depth on that subject in this book, but if you are a bit curious to find out more, check out http://en.wikipedia.org/wiki/X_Window_System.

The X-server provides the framework for using a GUI, but it doesn't actually provide the pretty desktops that we're used to seeing. They are provided by desktop environments. Again, we won't cover how these all work, but suffice it to say, like everything else in the Linux world, you have a choice about which one you prefer. For most people, it comes down to which one comes with their Linux distribution of choice, although most distributions offer a selection. The more advanced ones have all sorts of features and special effects, and understandably consume a fair amount of resources while they're at it. On the latest and greatest desktop or laptop, this is hardly a problem, and most people jump at the chance to make their expensive machine actually do some work. However, it does become something of a problem when you're running a less-than-high-spec machine such as the Pi.

It's not just the Pi that can struggle—computers of just a few years ago could easily run Linux, but struggled with the new highly graphical interfaces. When you consider even older machines that would often get thrown out, a suitable display manager could easily manifest as a perfectly useful PC. There's still more to it, though; modern netbooks are low-powered by design, with the aim of high efficiency for battery consumption as well as cost. They also benefit from a minimalist design when it comes to the desktop.

When you need to stretch your processing budget, the best place to cut corners is in the feature set of the graphical interface. Do you really need all those fancy features? Are transparent windows essential to what you're doing? Do you really need that 3D-rendered clock? If not, and you want to cut some fat, a desktop environment such as LXDE will be right up your street.

LXDE (short for *Lightweight X11 Desktop Environment*) was designed for machines in which energy efficiency and the need to run on hardware at the lower end of the performance spectrum was of critical importance. It aims to be efficient in terms of CPU cycles and memory usage, and although some window managers eschew all unnecessary features, LXDE tries to remain full of features at the same time.

Although we now know that the GUI is made up of a lot of moving parts, LXDE is mostly transparent when used. Really there is no harm in thinking of LXDE as the main environment for your Pi and if you plan to spend most of your time using the GUI, we feel there's no harm in considering it a core part of the operating system. You can simply just use it and accept it as being available. With that in mind, let's move on to exploring what we actually have when we are using LXDE.

What Do We Have Here?

Once you've installed and set up your Pi, and you've powered it up for the first time, you should be left with a screen that looks like Figure 2-1.




Figure 2-1. First boot of your Pi

How the icons are laid out really depends on the screen you connected to your Pi. It should look pretty similar; the only real difference is how the desktop icons are laid out. And because Raspbian is constantly being improved by the developers, it's possible that you will see new icons or perhaps some replacements.

There are some key areas here that are hard to see in one large screenshot, so we're going to break each part down individually. For simplicity (not to mention the fact that they have become common idioms) we'll refer to various GUI elements by the name most people are familiar with. Thus we will discuss the start menu and taskbar. Let's start with the bottom left of the screen.

The Start Menu and Bottom Left of the Task Bar

The first button on the bottom left (the one that looks like an alien spider—see Figure 2-2) is actually the LXDE logo and acts as the start button. The second button (which is indistinct even on a high-quality monitor) will start up PCManFM, the LXDE file manager. This is the tool you reach for when you need to move some files around or do some general tidying up. The next button (a globe) starts the default web browser that happens to be Midori in this case. Like the other apps, you'll find that most of these applications are designed to be lightweight and efficient.



Figure 2-2. Start button and friends

So far, so good. None of these applications is really that different from what you're probably used to doing on your main computer. The next two, however, are a bit different: they control the environment rather than start individual programs. The first button (the two overlapping windows) acts like a "minimize all" button when clicked: all the windows currently being displayed will be immediately minimized. Clicking the button again doesn't restore the windows, though, so if you do have a lot of windows open, don't expect to be able to click this button again to get them back. This button has a second feature: if you click with the middle button (most scroll wheels can also be used as a button), all the windows will be "shaded." When you shade a window, it rolls up inside its title. We don't really use this feature (in all honesty, we never saw the point, but to each his own), but if it's something you like, feel free to use it.

Next up are the desktop buttons. Microsoft Windows by default has a single desktop, and most people are unaware that it can actually have more than one. This feature was added in recent versions of Windows, but Linux has had it from day one, mostly because other UNIX-based operating systems have had it for considerably longer.

The feature is quite powerful in that it lets you have multiple desktops, each with its own applications. This allows you to keep all your work applications on one desktop while keeping your chat and browser windows in another. By default, LXDE provides two desktops, but you can have more if you want.

The rest of this part of the bar is taken up with the various applications that you've opened. Like Windows, all running applications (at least those that provide a graphical interface) will show up here and you can use the bar to select which application you want to interact with.

And on the Right Side...

After the taskbar, there are more icons, as seen in Figure 2-3. The first one appears to be just a black box, but it actually acts as a CPU monitor. The more work you give your Pi, the greener this box will become. It acts like a chart, showing you how your CPU is being used over time. You can actually amuse yourself (at least for a few minutes) by rapidly moving your mouse pointer. The faster you move it, the more CPU is used, and the box becomes greener. Fun and games aside, though, this monitor is quite useful because if your Pi starts acting a bit sluggishly, a quick look at the monitor can tell you whether you are maxing out your CPU.



Figure 2-3. The right side of the task bar

Next to the CPU monitor is the clock, which doesn't really need much in the way of explanation. The monitor with the padlock in the middle is used to lock your screen (although we have heard that this doesn't work for everyone). This is very useful if you're stepping away from your computer and you have some friends around who think it might be funny if you made some "amusing" posts on Facebook. By locking your screen, their efforts will be thwarted—assuming you remember to change your password to something other than the default, of course!

Last up is the power button that gives you numerous options for shutting down or rebooting your machine, as shown in Figure 2-4 (although if you started manually with startx, then you will only see logout and cancel).



Figure 2-4. Clicking the shutdown button

Clicking Shutdown will do pretty much what you expect, as will Reboot and Logout. There is one thing that is worth remembering, though: the Pi runs Linux, which is a multiuser operating system. That means that things might be happening in the background that you're not aware of. For example, someone might be streaming a video from your Pi. If you shut down or reboot, you will disconnect anyone else who is connected to the Pi. This might not be an issue for you right now, but it is something that you probably should be aware of.

And that's pretty much it for what you have on your screen. We'll cover the apps on your desktop after the next section, in which we look at the mysteries of the start menu.

The Start Menu

The start menu concept has become almost ubiquitous across most operating systems (notably not the Mac) since its introduction with Windows 95. Although various people might scoff at that and point to places where it might have been seen before or insist that the start menus in Linux are nothing like those found in Windows, we try to move beyond that. The fact is, the start menu concept works very well regardless who thought of it, and there is a reason why it's been adopted so far and wide.

Enough of the soap box; let's see what our start menu has to offer (see Figure 2-5).




Figure 2-5. The start menu

As before, there's nothing really groundbreaking here. The first section lists all the applications that are available, all neatly grouped together: in this case, Accessories, Education, Internet, Programming, and System Tools. Apart from System Tools (and Accessories doesn't look that inspiring), it's hard to decide which of these options we should delve into first. Before we can begin, though, we need to wrap up the rest of the items.

Preferences and Logout are fairly straightforward. We'll look at some of the settings in Preferences that you can tweak later on; if you press the Logout button, you will, well, log out. No surprises there. The final item is the Run command, which lets us execute specific commands without the need to pick them from a menu, click an icon, or open a terminal. Just as with Windows, this option can be a useful convenience. It also has autocomplete functionality, as shown in Figure 2-6.




Figure 2-6. The run command and autocomplete

Here you can see we typed in **gno**, and it autocompleted to the app that matches the pattern. This can save you a lot of time as you will get to know how many characters you need to type before the autocomplete kicks in with the right answer. Quite often, command names (especially for GUI apps) can be pretty long and this can be real time saver!

Accessories

Now let's take a look at what we have under Accessories. Like all the other menus, the items here can change as you install more software, and on newer versions of Raspbian it might look a little different. On our Pi, the Accessories menu looks like Figure 2-7.




Figure 2-7. The Accessories menu

Debian Reference

First up is Debian Reference. As its name suggests, it's a guide to all things Debian and because Raspbian is based on Debian, it gets taken along for the ride. The manual is nicely structured and well written, and has had the additional benefit of evolving over an extended period of time. The manual is a set of HTML pages, and when you click the link it will open in the Midori browser. Simply click the first link as shown in Figure 2-8.

Debian Reference (version 2)

English

[HTML \(multi files\)](#) , [UTF-8 plain text \(gzipped\)](#)

Figure 2-8. You want the HTML version on the left

The second link contains the same content in plain text, but it is stored as a compressed archive and so it's not something you'd use for online viewing. As the manual is taken from the Web, this is a holdover from the days when you might have wanted to download a copy for browsing offline. All being well, you should end up with the page shown in Figure 2-9.




Figure 2-9. The Debian reference manual in all its glory

You can navigate the manual just like a web page. Although having a local copy is very useful, it's not as easy to search as asking Google, which will usually take you straight to what you want. That said, if you're trying to figure out why you can't get online, you'll be happy for a local copy.

File Manager

As touched on earlier, LXDE ships with the PCManFM file manager. Of course you can also access it from the task bar, as shown earlier. Taking its lead from the minimalist design movement, this file manager is lightweight, simple, and easy to use, but still has all the features you're going to need. It works in a very similar way to Microsoft's Windows Explorer, so if you're coming from a Windows background you'll feel right at home with PCManFM (see Figure 2-10). If not, you'll be able to pick it up easily with just a few moments of clicking about.




Figure 2-10. The PCManFM file manager

When it initially opens, PCManFM will drop you into your home area. It will show you how much space you have available on the bottom right (always useful in devices with relatively limited storage such as the Pi) and lists the most common places on the top left. For browsing, you can simply click your way through the folders and you can use the navigation icons on the toolbar to help you move quickly through your files. If you're looking to get a bit more hands on, we cover how to manipulate files in Chapter 4.

Image Viewer

Image Viewer is one of those applications that you'll probably never open directly; you'll almost always open it by double-clicking an image file that by default will then open Image Viewer for you. Simple though it may be, Image Viewer is not lacking in features and has all the basic tools you've come to expect, as shown in Figure 2-11.




Figure 2-11. Image Viewer's main interface

The first two buttons let you move back and forth through a collection of images. This is actually surprisingly useful because people are now accessing albums rather than specific photos. To help with exactly that, there is also a slideshow feature that you can access by pressing on the play button.

The next section has controls for how you want to actually view the image.

- The first two are the traditional zoom controls
- The third icon is for “fit to screen”
- The fourth is show at full size (which, if it is a photo from a digital camera, means you’ll have to scroll)
- Last, the full-screen option to really take advantage of that 60-inch flat screen you just got your hands on

In the third section are the controls for altering the image. The first two are rotate controls (rotate left and right, respectively), and the second pair of controls allows you to flip the image horizontally and vertically. This is especially handy when you took the photo at an odd angle and don’t fancy tilting your head 90 degrees in order to see your picture the right way up.

The next set contains file controls. They let you open, save, save as, and delete the file, respectively. The last two icons provide access to the preferences panel (nothing too interesting in there) and allow you to exit the program.

Leafpad

In keeping with the minimalist theme, the Leafpad text editor does the basics and that's about all. It's very similar to later versions of Microsoft's Notepad and it is similarly easy to use (see Figure 2-12). It's not much good for writing letters to Aunt Maude or to the council to complain about the next door neighbor's dog, but it's perfect for when you need to do some light text editing.




Figure 2-12. Leafpad has a similar look and feel to Microsoft's Notepad

Despite its simplicity, Leafpad still provides some of the most useful tools, namely being able to do a find-and-replace, and enable and disable word wrap. Armed with those tools, you can accomplish pretty much anything.

Note If you feel that Leafpad is a little underpowered for your tastes, you might want to flip ahead to Chapter 6, in which we show you some of the more powerful native text editors that you can get your hands on.

For the most part, Leafpad does exactly what it says on the box: it edits text, and if that's all you need to do, Leafpad is an excellent choice.

LXTerminal

LXTerminal is a terminal emulator for Linux that gives you access to a virtual console. We won't go into any depth on how to actually use this particular application because we'll discuss the details in the next chapter. Suffice it to say, this tool is used to give you direct command-line access to the operating system. For many people, this is a place where angels fear to tread, yet for seasoned administrators, it is somewhere they call home. To find out whether the command line is for you, though, you'll have to wait until Chapter 3.

LXTerminal is a simple emulator (see Figure 2-13), but it gets the job done. It doesn't have all the bells and whistles that some emulators provide, but it is perfectly fine for the majority of people.




Figure 2-13. LXTerminal up and running

One thing we will mention here is the prompt. When the font color is green, and the end of the prompt is a dollar symbol, you know you're running as a normal user and you don't have root privileges (you can't execute any commands that can really hurt the system). Again, we'll cover this topic in more depth in the next chapter; we mention this distinction here so that the next program makes sense.

Root Terminal

This is all but identical to LXTerminal, except that instead of running as the normal pi user, it runs as the root (system administrator) user. When you first run this application, you'll get a popup message, as shown in Figure 2-14.



Figure 2-14. Not a warning; you're running admin-level commands without needing a password

Although it looks like an error message, it just lets you know that the system is configured to let you run an administrative-level command (in this case, running LXTerminal with root privileges) without having to provide authentication. For the Pi, this is almost certainly what you want, but on a machine that's connected to sensitive resources and in which other people could feasibly access the machine, you might want to have it password-protected.

A root terminal is very similar to a standard LXTerminal, but (as you can see in Figure 2-15) the prompt looks a little different.




Figure 2-15. LXTerminal running as root

Here you can see that that prompt is now in gray text, and the end of the prompt is the pound or hash symbol. This means you are operating with elevated privileges and you could indeed run commands to effectively destroy the operating system. Linux doesn't try to second-guess you; if you tell it you want to erase everything on the disk, it will obey that command and do so precisely. There is no safety net when running as root, which is why most people recommend that you stay in your own user account as much as possible. We'll come back to this theme many times in this book.

XArchiver

XArchiver is less likely to be called directly than Image Viewer, but nonetheless is probably an application that will get a fair bit of use on the average system. This tool allows you to create and decompress archives of various formats, which is essential when you're downloading content from the Internet. XArchiver looks like Figure 2-16.




Figure 2-16. XArchiver when it's opened by hand

Here we downloaded the documentation to the Erlang platform (it was the first thing that popped into our heads when we needed something to demonstrate XArchiver). After double-clicking the file, XArchiver was opened automatically. The only button we have used so far is the extract all button, which is the second one from the right. Clicking this will extract all the contents of the archive, which is usually what you actually want to do.

Education

The Education menu has only two items, but as they comprise some of the more interesting applications on the Pi, it's worth checking them out. The menu on our Pi looks like Figure 2-17.



Figure 2-17. Education menu

Scratch

Scratch was developed at the MIT media labs by the Lifelong Kindergarten group led by Mitchel Resnik. The goal was to create a way to teach people how to program without having to teach all the fiddly bits first. This made the tactile approach in *Scratch* (in which you can drag and drop blocks of logic) much easier to pick up.

We can't claim to be experts in *Scratch* (in fact, we first came across it when we were playing with the Pi) but we are convinced that these sorts of applications will have a huge impact as the Pi spreads out across the waters. Getting people hands-on with technology and providing them with a powerful system that lets them get started without having to master all the underpinnings first is definitely a winner in our book.

If you're not convinced, check out Figure 2-18; it certainly made us take a second look.




Figure 2-18. Scratch programming environment

Squeak

Squeak is a programming language that is preinstalled on your Pi. For licensing reasons, however, the install is not quite complete, so you won't actually be able to use it out of the box. We are not sure if this is a temporary issue and one that might be resolved in future versions of Raspbian. At the moment, if you open Squeak you'll see a window that looks something like Figure 2-19.



Figure 2-19. Starting up Squeak for the first time

After doing a bit of digging, we discovered the reason and also a post on the Raspberry Pi forums by jstout that provided a handy shell script that would install the missing components for you. Thanks to jstout, we can help you get Squeak up and running. You'll need to open up LXTerminal and enter these commands (you could also come back to this section once you've read Chapter 3 and are comfortable on the command line):

```
$ wget -r -O SqueakV41.sources.gz http://ftp.squeak.org/4.1/SqueakV41.sources.gz
$ gunzip -f SqueakV41.sources.gz
$ wget -r -O Squeak-4.3-All-in-One.zip http://ftp.squeak.org/4.3/Squeak-4.3-All-in-One.zip
$ unzip -j -o Squeak-4.3-All-in-One.zip "Squeak-4.3-All-in-One.app/Contents/Resources/Squeak4.3.*"
$ sudo mv SqueakV41.sources /usr/share/squeak
$ sudo mv Squeak4.3.image /usr/share/squeak
$ sudo mv Squeak4.3.changes /usr/share/squeak
```

This isn't the time or place to discuss what those commands actually do, but the short version is this: they download the missing components and place them in the right location for Squeak to find them. Once you've entered these commands, you'll be able to run Squeak without seeing that popup box. Instead you should see Figure 2-20.




Figure 2-20. Squeak is fully operational

Squeak is a dialect of Smalltalk, which is an object-oriented, reflective language. This makes it a very powerful and expressive language (you can do more with less code) that has earned it quite a following over the years. It has inspired various new programming languages (including Ruby) and was also used to implement Scratch, which we discussed just a few pages ago.

We're not Smalltalk experts either, but we do know that it is a very interesting and powerful language, and it might make an interesting place for you to start your exploration of programming languages if you feel so inclined.

Internet

The Internet menu unsurprisingly contains applications that are specific for accessing the Internet. On our Pi, the Internet menu looks like Figure 2-21.



Figure 2-21. The Internet menu

Raspbian ships with three different browsers out of the box. They are all designed to be minimalistic and your choice will come down to personal preference. For us at least, the preferred browser has to be Midori, which is the one we'll cover in this section. At the end of the day, if you know how to browse the Web on any computer, nothing in these browsers will be new to you, and we're covering it here more for completeness than out of real necessity.

So rather than going through all the browsers, we're just going to touch on Midori because it offers both a standard browsing experience and a private option. Most browsers on other platforms look very distinct when they're in private mode (usually a different color scheme or icon set), which ensures that you don't accidentally type the wrong thing into the wrong window. Midori, on the other hand, simply changes the name of the title bar that isn't really all that distinctive. Take a look at Figure 2-22.




Figure 2-22. Midori in standard and private mode

As you can see, there's not a lot of difference here. The private browser doesn't allow you to add bookmarks, and the ability to open previously closed tabs is not available. Other than that, they look pretty much identical. If you do decide to take advantage of this feature, you should always double-check before you hit the Enter key.

Programming

The Programming menu consists of four items (see Figure 2-23), two of which we covered when we came across them under the Education menu.



Figure 2-23. The Programming menu

Because we've already covered Scratch and Squeak, all we have left is IDLE and IDLE 3. Actually, these two are pretty similar in that they are both integrated development environments (IDEs) for the Python programming language. IDEs provide a rich development environment in which all the tools you need are close at hand. They usually include tools to run and test code as well as nicely formatting your source code so it's easier to read. The reason why there are two different versions is because there are two current versions of Python that for various reasons aren't mutually compatible. In many ways, scripts that you might write for version 2.x of Python will work just fine in 3.x, but there are quite a few gotchas and changes in the language that make guaranteeing compatibility impossible. This is why you'll often see two versions of Python installed on any given machine.

Note At the time of writing, Python 2.x is still where most people are at, and many projects are still actively developed using it. In fact, when we write new Python scripts we also use 2.x because it has the most libraries, a history of documentation and blog and forum posts, and generally a large community. On the other hand, Python 3.x has had a lot of tweaks and language improvements that will never make it to Python 2.x, and all the new ideas and development are focused on the new series. Ultimately, Python 3.x is the way to go, and if you're just starting out in your programming career, we recommend you start out on Python 3.x. It certainly can't hurt, and you can always look at Python 2.x if for some reason you need to use it, and in reality going from one to the other is really not a big problem.

Because the two versions of IDLE are effectively identical in terms of how they're used, we're going to show you IDLE from the 2.x branch (rather than IDLE 3), as shown in Figure 2-24.




Figure 2-24. Firing up IDLE

When you first start up IDLE, you'll be dropped into a Python shell. There's nothing special about this prompt; in fact, you can get the same prompt by running Python at the command line. IDLE really shines because it provides an editor that supports syntax highlighting. If you're new to programming, syntax highlighting is where the text that you enter (the source code) is color-coded to make it easier to read the program. Keywords and structures are highlighted, so it's much easier to scan through a program (and check that you haven't made any obvious typos) than if you simply had black text on a white background. If you're not convinced, try it sometime; you might be surprised.

Figure 2-25 shows a very simple program that prints "Hello world!" 20 times and exits.

```
for x in range(1,20):
    print "Hello world!"
```

Figure 2-25. Using the IDE component of IDLE

As you can see, the structure is a lot easier to read when you add a splash of color (or if you're reading this in the black-and-white printed version, you might have to use your imagination). You can open an editing window by clicking File and selecting New Window.

Python is a very powerful language and is used in a huge range of places for all sorts of things. If you're interested in learning Python, check out www.python.org for more information on how to get started.

System Tools

This section won't even need a subheading as there's only one tool in this menu group, as shown in Figure 2-26.




Figure 2-26. The somewhat lonely System Tools menu

The Task Manager lists all the processes currently running on your machine along with how much of your system resources are being used. Figure 2-27 is a snapshot of what Task Manager reported on our Pi (of course, yours will likely be quite different).




Figure 2-27. I spy on my little Pi: Task Manager shows what your Pi is up to

As you can see, the CPU is practically idle at the moment. Considering that it is running a full operating system (and an advanced one at that), it might be surprising to see that the CPU isn't doing any heavy lifting. Although when compared with high-end Intel processors, the ARM chip in the Pi is far from impressive, this truly demonstrates that you can get a perfectly good desktop machine running with surprisingly few resources if you are careful.

Our RAM usage tells a similar story. Admittedly, we don't have much open at the moment (although you can see from the process list that we have a VNC remote desktop session running), but we are still well under 50% usage and with the new Model B's packing 512 MB of RAM, this will allow even more applications to run at the same time.

The process list gives you a good indication of what's running on the machine and how many resources are being consumed on an individual basis. This is mostly useful when your Pi is acting a little bit funky, and you need to see whether an application has suddenly decided to eat all your RAM or set your CPU on fire. If you do find an application that is getting a bit too greedy, you can right-click its name and then select Kill, which will forcibly shut it down and thus hopefully restore our Pi to full health.

Preferences

Now it's time for a whistle-stop tour of some of the things you can tweak to make your desktop a little more personalized. Because the tools in Raspbian are pretty much self-describing (not to mention fairly obvious), we're going to cover just the two main topics and let you explore the rest of the preferences for yourself. Because you really can't do much in the way of damage, feel free to explore to your heart's content.

Desktop Preferences

To get to the desktop preferences, right-click your desktop and click Desktop Preferences. You should see a window like Figure 2-28.




Figure 2-28. Desktop Preferences

We're not going to cover the Advanced tab (it only has one option, and it isn't one that most people will be interested in) and so we won't show a screenshot of it. Everything on the Appearance tab is fairly straightforward, and there are no unusual options or special tweaks that you need to do to get anything to work.

Customizing the Look and Feel

If you want to customize the look and feel of your entire desktop experience, you can tweak themes, icons, fonts, and more from the Customize Look and Feel tool. You can access it from the Preferences menu, as shown in Figure 2-29.




Figure 2-29. Finding Customize Look and Feel

This will open the Customize Look and Feel app, as shown in Figure 2-30.




Figure 2-30. The Customize Look and Feel tool

Because the tool is fairly straightforward, we'll cover only the first screen, which allows you to change the color scheme for the entire environment in a single mouse click. The default scheme is Mist, which we actually prefer, but there are a number of other predefined schemes that you can choose from. The sample window will immediately reflect your new choice, and so gives you a preview as to what your environment will look like once you hit the Apply button.

Summary

And that's pretty much it for the whirlwind tour of the desktop! We've really only scratched the surface, and the aim is just to get you comfortable enough with what's available to start exploring by yourself. Remember that Pi exists to allow people to experiment and learn more about computing, so getting your hands dirty (even if you break your Pi now and then) is really nothing to worry about.

This chapter covered the applications installed by default and accessible via the start menu, and we skimmed over some of the interface features that are somewhat unique to the Linux world. We then finished off by looking at the two most common configuration changes that people want to make.

In the next chapter we start to get really serious (sort of) and take you back in time to the command line, where we introduce you to an interface that once you've gotten used to, you'll wonder how you ever managed to live without it.

Getting Comfortable

Normally we don't start a chapter with a figure (and technically speaking we haven't in this one, either), but we're going for some good old-fashioned shock and awe. Take a look at Figure 3-1 (courtesy of Wikipedia).



Figure 3-1. *TeleVideo 925 computer terminal*

You might be wondering why we thought it necessary to show you some technology from the early '80s (late 1982 in this particular case). In the olden days, computers lacked all the niceties that we have today, such as colorful displays and mice. Back then, it was just text, often green on a black background. Mice were decades away, and the only way to communicate with a computer was either through text-based menus or the command line. Because the command line gave a great deal more flexibility than a set menu, the command line was by far the most common interface and the one you'd be using if you accidentally stepped 40 years into the past. You'd sit at one of these screens called *terminals* and tapping away endlessly into the night.

As such, the majority of the second section of this book is all about the command line and getting you comfortable using it. Although things have improved somewhat in the intervening decades, the technology is basically the same.

In this chapter, we take a really quick look at the history of computer interfaces and why the terminal has survived as well as it has. We'll then look at the various terminals available to you and how you go about getting to them. We'll then quickly explain the command prompt and let you in on a little-known secret about the different shells that are available.

Ye Olde Computer

Although you can now quite literally buy a computer for \$25, in the '70s you could expect to pay at least 100,000 times that amount for a computer with any serious amount of processing power. Such machines would take up entire rooms (if not entire floors), require incredible amounts of power, and generally look very impressive to passersby. A key problem that owners of these things had to deal with was one of scale.

When your computer is worth more than the building it is sitting in, you want to make sure that you get the most out of it. One approach to solving the challenge was to use time-sharing. The idea was that compared with a human interacting with the computer, the computer was incredibly fast. In fact, even though these computers were puny by today's standards, they could still run rings around the people who were actually trying to use them. So this begs the question: if the computer is spending most of its time waiting for the operator to enter commands and data, couldn't it be doing other things while it waits?

It wasn't long before someone figured out that if the computer sits waiting for input, it can be better utilized with more than one person entering data. That way, multiple people can enter data at the same time, and we're no longer limited by the speed of a single person.

Great idea, but how to carry it out? To make it work, we'd really want to put a screen in front of everyone who needs to talk to the computer and with the available computer technology, even a very basic computer would be impractical. Then again, we didn't really need a computer on their desks; we just needed a way for them to communicate with the big computer downstairs. In effect, we just want a keyboard and a screen with very long cables.

Say Hello to the Dumb Terminal

This is where the dumb terminal came in. It was called a *dumb terminal* because it was quite literally just a screen and keyboard. Using instructions relayed from the computer itself, it would draw the text onto the screen. Whenever a key was pressed, it would be sent to the server that would then reply with more updates for the screen. The dumb terminal didn't have any real processing capability; it just worked as an interface and that made it relatively cheap and easy to install and maintain.

The technology did improve over the years, and the dumb terminal became "less dumb." Support for different colors and brightness was added, as were fancier features for redrawing and managing the screen. However basic the concept, that of the dumb terminal just being an interface to a remote machine never really changed all that much. In fact, although dumb terminals started to be replaced when more modern networks became available (not to mention computing power became much cheaper), the approach continued to be used in the newer systems. This wasn't just because there was a lot of legacy systems that people had to keep using (ever see the old green-on-black screens in some big department stores?), but it actually had some very useful properties that continue to be useful even today—properties that we'll go into once we've seen what a modern terminal looks like.

GUIs AND TUIs

You'll often hear people referring to GUIs (pronounced "gooeys") and TUIs (pronounced "tooey"), and we'll be referring to them throughout the book as well. Fortunately, they are pretty straightforward, and you'll be using them in no time (assuming you actually want to, of course).

GUI, which is short for *graphical user interface*, is used by all modern computer systems in some shape or another. It used to be really easy to spot a GUI because it was pretty much any device that used a mouse. The lines have blurred somewhat in recent years, thanks to devices such as the iPad and Android smart phones, you can have a GUI that just uses your fingers. It's now much easier to know what a TUI is as generally anything that isn't a TUI is considered a GUI.


TUI, which is short for *text-based user interface*, is often used interchangeably with the command-line interface (CLI). A TUI can present basic menu systems, but those systems are navigated via keyboard control rather than a mouse.

In short, if you drive it from the keyboard, and there's no mouse pointer, it's a TUI. If you can click, point at it with a finger, or see pretty pictures, it's a GUI.

Modern Terminals

There's quite a bit of evolutionary history between the dumb terminals we just talked about and what we use on modern Linux systems today. Often you will find that you have a computer (for example, your Pi) and you'll want to connect to it over the network. When you do this, you are running a *client* on your desktop machine and you are connecting to a *server* that's running on the machine you want to connect to; in this case, it is your Pi (which is acting as a server). Once connected, you are effectively "remote controlling" the Pi, so this type of connection is known as a *remote terminal* (a software version of the dumb terminals of the past). Remember that as far as Linux is concerned, whatever you type on your local client machine is as if you were typing on the Pi directly. This is just like a TV, on which changing the channel, either on the TV itself or via the remote control, causes the same result.

In the interests of preserving your sanity (not to mention crossing the boredom threshold), we'll leap ahead to the modern day and show you what a remote terminal looks like on your Pi (see Figure 3-2).



A screenshot of a terminal window titled "pi@raspberrypi: ~". The window displays a text-based login screen for the user "pi". The text includes the following information:

```
pi@raspberrypi: ~
login as: pi
pi@10.0.0.20's password:
Linux raspberrypi 3.1.9+ #272 PREEMPT Tue Aug 7 22:51:44 BST 2012 armv6l

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.

Type 'startx' to launch a graphical session

Last login: Sun Sep 16 15:48:56 2012
pi@raspberrypi ~ $ █
```

Figure 3-2. SSH connection to a Raspberry Pi

As you can see, it is still a text-based interface: the programs that run will all happen on the remote server. There are some nicer features (color, for a start) and we can make the terminal as big as our screen can hold (traditional terminals tended to be fixed size). In this particular case, we are connected to the Pi via Secure Shell (SSH) (more on this later) over TCP/IP.

Why Do We Still Care About These Things?

Good question! In the modern era of iPads, iPhones, and Androids, why would we ever want to use a text-based interface? For most people, especially if they have never used one before, a text-based interface really sounds like something of a backward step, not a skill that they'd be especially interested in acquiring. Needless to say, we are most fond of the venerable command line and here are some of the reasons why we think you should invest your time and effort in getting to grips with it.

Simple

One of the key reasons is simplicity. A text-based interface is pretty hard to beat when it comes to being simple. Literally all it does is send and receive text. You won't be using a terminal to look at your holiday snaps, but it does give you a clean crisp interface to issue commands to your Pi.

Fast

You probably won't believe us right now when we say this, but using the command-line interface will allow you to complete many of your tasks much faster because you can be very precise and express a lot of work in a single line of text—the sort of thing that would require multiple clicks or a strange convoluted process to express using a mouse.

Lightweight

Anyone who has had the misfortune of trying to remote control a desktop PC over a slow Internet connection will appreciate exactly what we mean here. Sending screen updates across the Internet (especially for high-resolution screens with lots of colors) by its nature simply takes up a lot bandwidth. It can also take a while for the data to arrive at its destination, which is what gives remote sessions that unresponsive and blocky feel to them.

Text-based sessions don't suffer from this problem because they only send text. They were designed for systems in which the maximum speeds were a mere fraction of what is available to us today. So even on a very slow dial-up link, a remote terminal session tends to perform extremely well.

It also puts less strain on the server. Maintaining a graphical interface means that the computer has to do quite a bit of work. Having a mouse pointer means that the computer has to track that pointer and make sure that it reacts properly when it touches an applications window or when the user clicks a button. A surprising amount of work goes on behind the scenes, and this is why Linux servers rarely have a graphical interface installed, much less running.

Powerful

The command line gives you an awful lot of power. You can express complex thoughts and issue commands in a way that just wouldn't be possible via a graphical interface. For example, you can string commands together to make more complex chains. You can take the output of one command and use that as input for another. This technique uses something called *pipes* (we cover it more thoroughly in Chapter 5).

Always Available

Although things have gotten better in recent years, one of the most fickle pieces of software on a Linux box tends to be the graphical interface. It always seems to be the first thing to go, and when your server stops booting and you need a rescue environment, it's unlikely that a graphical interface will be available to you.

A terminal, however, is the native way to talk to Linux (and most other UNIX-like operating systems) and as such, it's the one tool that will always be available to you in some form or another. Knowing how to use the terminal will let you get out of scrapes and deal with situations that would completely stump you if you weren't comfortable with the command line.

Convinced Yet?

Hopefully, these highlights have at least convinced you that there might be something to the command line after all. If you're not totally convinced, that's okay; it is hard to really get a feel for the benefits until you've actually spent some time experiencing them for yourself. In the next section, we'll cover how to actually get yourself to a command line, explain what the stuff you're looking at means, and cover some basic commands to get you started on your journey.

TERMINAL OR VIRTUAL TERMINAL?

For the most part, you won't have to worry about the differences between the two. You will almost always use a virtual terminal that is either a terminal window inside the Linux GUI or connecting via SSH. The only real difference between a terminal and a virtual terminal is that a *terminal* is a single physical device such as a serial port or sitting at the screen directly connected to the computer (without a GUI). A *virtual terminal* has a device assigned to it, but there is no corresponding physical device. In fact, under Linux a new device is created and assigned to each new virtual terminal.

We brought up the distinction here because some documentation (generally speaking, old documentation) might refer specifically to one or the other. These days, though, most of your time will be spent on a virtual terminal, and you won't need to worry about the difference.

Three Terminals

There are three main ways to get to the command line. The first is to sit physically in front of the machine and use what is often referred to as the console, the second is to open a virtual terminal inside the GUI, and the third is to connect over the network using SSH.

Console

Although the *console* is simply just another terminal, it is a term that has gained a somewhat mythical status over the years. When you're working on the console, you are directly working on the machine because it is literally a keyboard and monitor plugged directly into it.

It used to be very common for administrators to restrict certain users (such as the root user, more on that later) to be able to log in only when sitting directly in front of the machine. Because the computer was likely in a locked server room, this provided an additional level of security.

Most distributions do not come with this sort of security enabled, and now that remote access over the network is the norm (and sometimes an administrator doesn't even have physical access or the machine is on the other side of the planet), this feature is far less useful.

Still, if you are sitting in front of your Pi, you can access the console by holding down the Ctrl and Alt keys and then pressing F1 through to F6. F7 by default is where the GUI lives, so once you have finished with the console, you can flip back to your GUI.

When would you need to do this? Well, if you can't connect to your Pi via the network and you're having trouble with the GUI, the console will be your new best friend. Generally speaking, you won't use the console very often because there are usually more convenient ways to get the same effect (we cover them next). However, when setting up the system or if you've decided not to use the graphical interface, the console will be your only real choice for accessing the machine locally.

Getting a command line through the other two methods is easier and more flexible because as often at least one of those options will be available to you.

Opening a Virtual Terminal in the GUI

Raspbian uses the *Lightweight Desktop Environment (LXDE)*. As its name suggests, it focuses on being lightweight, which is just what you want when your computer sits squarely in the featherweight category!

The native terminal for LXDE is called LXTerminal, and fortunately for us it not only comes preinstalled but also features on the Raspbian desktop, as shown in Figure 3-3.




Figure 3-3. Finding LXTerminal

We took the liberty of double-clicking the icon, and you can see that we have what can only be described as a working terminal. Of course, there's really not too much to show in this case because a terminal is by definition pretty simple. You will probably find yourself using the terminal a fair bit if you're planning to sit in front of the Pi and use it as a more traditional computer. If you're thinking more about working remotely, then you need to read the next section, which is quite honestly a lot more interesting.

Connecting via SSH

SSH is a great way to connect securely to your Pi over the network (see the sidebar “What’s So Special About SSH?”). It gives you all the benefits of a virtual terminal, but you can access it from any machine on the network. Potentially this means that you could connect to your Pi at home from your work PC and fiddle about with it in complete security. We find that SSH is by far the most common way we interact with servers in general, and because many people will use the Pi as a headless device (i.e., no keyboard or mouse attached to it—maybe not even a monitor), we think SSH will be the primary candidate for you as well.

There are two parts to the SSH equation. You need an SSH server and you will need an SSH client. Fortunately, both are freely available and easy to set up.

WHAT'S SO SPECIAL ABOUT SSH?

SSH is the *de facto* standard for accessing any UNIX-like machine remotely. It is even supported by many network-enabled devices such as printers and high-end switches. The reason why it is so popular is because it encrypts all the data between your computer and your server. If someone is listening to your connection, all they will get is garbage. It also allows you to authenticate that the server is the server you think it is.

Before we had SSH, everyone used telnet, which is a very simple protocol and doesn't have any sort of encryption. This meant that if anyone could see your traffic (think public machines or WiFi access points), it would be possible to simply read your password off the network in clear text. Telnet also had no way to confirm the identity of the machine you were connecting to, which meant you could never be sure that you weren't accidentally providing your username and password to somebody else.


It is still possible to install a telnet server on Linux, but this is highly discouraged. There really is no good reason to prefer telnet over SSH, so you should always use SSH when working on your server remotely.

Setting Up an SSH Server on the Pi

This one is actually very straightforward because unless you specifically disabled it when you installed Raspbian, the SSH server will already be running (Raspbian enables it by default). In fact, all you actually need at this stage is the IP address of your Pi.

One of the nice things about covering the Pi's virtual terminal first is that we can now use that to figure out what our IP address is. The Pi takes advantage of Dynamic Host Configuration Protocol (DHCP), which means that on the vast majority of networks it will be able to sort out an IP address for itself. This not only saves us time but also means you don't need to figure out how to choose an appropriate address and then configure it. The marvels of modern technology!

To find out the address, we're use the `sudo ifconfig` command. This command (short for *interface configuration*) will show you the configuration of your network interfaces. Figure 3-4 shows what it looks like on our Pi.



```
pi@raspberrypi ~ $ sudo ifconfig
eth0      Link encap:Ethernet HWaddr b8:27:eb:9f:75:52
          inet addr:10.2.42.132 Bcast:10.2.42.255 Mask:255.255.255.0
                  UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
                  RX packets:2642 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:1439 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:1000
                  RX bytes:1003078 (979.5 KiB) TX bytes:1125901 (1.0 MiB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
                  UP LOOPBACK RUNNING MTU:16436 Metric:1
                  RX packets:0 errors:0 dropped:0 overruns:0 frame:0
                  TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
                  collisions:0 txqueuelen:0
                  RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

pi@raspberrypi ~ $
```

Figure 3-4. Running ifconfig on the Pi

Here you can see that we have two network interfaces. The local loopback device (or `lo` to its friends) is a virtual interface that network applications can use to talk to each other on the same machine. Because no hardware is involved, it's more efficient to use the `lo` interface as well as it being always available and configured in the same way. Under Linux, most system services (such as printing and even the GUI) operate as client/server applications and make extensive use of this interface.

Right now, though, we're not all that interested in the `lo` device; we're far more interested in the network interface that connects the Pi to our physical network. `eth0` (short for Ethernet device number 0; computers tend to start counting at zero) represents our link to the real world. At this stage, we won't bore you with what all the information means, and it's not really something you would look at on a regular basis. The most common reason for running `ifconfig` is to find out the IP address and that is, of course, why we are here. In Figure 3-4 you can see the following line:

```
inet addr: 10.2.42.132 Bcast: 10.2.42.255 Mask: 255.255.255.0
```

We're really only interested in the IP or `inet` address, which in this case is 10.2.42.132. Armed with this piece of information, we should be able to connect to the Pi from anywhere on your network.

Sorting Out an SSH Client

We are making the assumption that you're either connecting from a Windows PC or from a Mac. If you are using a different operating system, don't panic: there is probably a suitable client for you to download. You can still follow along with the other two tutorials regardless of which client you use (or on which platform)—the general process is the same.

Putty for Windows

Putty is pretty much the standard that all other SSH clients are compared against and for good reason. Not only is it full of features but it also happens to be free. You can download Putty from this address:
<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>.

Because it is a little long for typing, you can get a similar result by searching Google for "Download Putty." Once you get to the download page (see Figure 3-5), choose the first link on the page: `putty.exe`.

Binaries

The latest release version (beta 0.62). This will generally be a version I think is reasonably likely to work well. If you have a previous version, you can download it from the same place.

For Windows on Intel x86

PuTTY:	putty.exe	(or by FTP)	(RSA sig)	(DSA sig)
PuTTYtel:	puttytel.exe	(or by FTP)	(RSA sig)	(DSA sig)
PSCP:	pscp.exe	(or by FTP)	(RSA sig)	(DSA sig)
PSFTP:	psftp.exe	(or by FTP)	(RSA sig)	(DSA sig)
Plink:	plink.exe	(or by FTP)	(RSA sig)	(DSA sig)
Pageant:	pageant.exe	(or by FTP)	(RSA sig)	(DSA sig)
PuTTYgen:	puttygen.exe	(or by FTP)	(RSA sig)	(DSA sig)
A .ZIP file containing all the binaries (except PuTTYtel), and also the help files				
Zip file:	putty.zip	(or by FTP)	(RSA sig)	(DSA sig)
A Windows installer for everything except PuTTYtel				
Installer:	putty-0.62-installer.exe	(or by FTP)	(RSA sig)	(DSA sig)
Checksums for all the above files				
MD5:	md5sums	(or by FTP)	(RSA sig)	(DSA sig)
SHA-1:	sha1sums	(or by FTP)	(RSA sig)	(DSA sig)
SHA-256:	sha256sums	(or by FTP)	(RSA sig)	(DSA sig)
SHA-512:	sha512sums	(or by FTP)	(RSA sig)	(DSA sig)

The latest development snapshot. This will be built every day, automatically, from the current development code - in whatever state it is at.

(The filename of the development snapshot installer contains the snapshot date, so it will change every night. It is not offered by default.)

Figure 3-5. Downloading Putty

Putty is a stand-alone tool, so you can run it as soon as you've downloaded it; you don't have to install anything on your machine. This is generally pretty useful because it means you can stick it on a USB stick and carry it around with you, and you can run it on other machines without needing Administrator privileges. If you keep losing track of Putty (it's amazing how easy it is to do), you might want to download the "Installer" version, which will install it on your machine and set up the Start menu and Desktop shortcut icons for you.

Now that we have Putty, it's time to fire it up. Double-click the icon, and you will probably see something that looks a lot like Figure 3-6.




Figure 3-6. Windows getting excited

Windows is just warning you that the program hasn't been digitally signed. You've probably come across this one before, but when you're installing some security-enhancing software, the last thing you want to do is install something unpleasant instead. In this case, the official version of Putty isn't signed, and you can ignore this error. If you plan to use Putty a lot, you should probably uncheck "Always ask before opening this file." Trust us, if you don't do it now, you will after the tenth or eleventh time you run it.

After you've convinced Windows to stand down, you'll see a connection box that looks like Figure 3-7.




Figure 3-7. Starting up Putty

The default settings are fine for connecting to your Pi; you just need to supply the IP address you discovered earlier. Punch this into the Host Name (or IP address) box and then click Open. Because it will be the first time this particular machine has connected to your Pi, you will see a warning similar to the one in Figure 3-8.



Figure 3-8. The initial warning from Putty

Don't worry; this is part of the SSH security system. It is letting you know that it hasn't seen this particular server before (well, it wouldn't have, would it?) and is showing you the key for confirmation. We don't know anyone who has memorized the SSH fingerprint (although remembering the first four and last four digits goes a long way to making sure you have the right one); and in any case, we were expecting to see this warning. Once Putty has remembered the key, it will check it against the IP address and hostname each time you connect in future. If the remote server ever sends back a different fingerprint, Putty will warn you that there could be a problem. More often than not, this occurs when you rebuild or buy a new server, so you know why the key has changed. If you see a warning and you can't think of a good reason for it, think very carefully before logging in to that machine. In this particular case, you know the key will be new, so you can simply select Yes.

Now you will be back to a somewhat familiar black screen with some text. So far it will only be this:

login as:

Because we haven't gotten around to setting up different user accounts you're going to log in as the 'pi' user so just enter **pi** and press Enter. So far, so good. Now we just need to provide it with the password you set when you did the initial install. We know that you changed the password as we suggested (you did, right?), but just in case you didn't, the default password is *raspberry*. Punch the password in, press Enter, and at long last you should see the command prompt:

```
pi@raspberrypi ~ $
```

If you don't have a Mac (or aren't particularly interested in using SSH from one), jump ahead to the next section: the command prompt.

LOOK BEFORE YOU LEAP

This note falls under the “obvious-until-proven-otherwise” category. When you’re working with remote terminals, it is disturbingly easy to get disoriented and start running commands on the wrong machine. If you are trying to reboot the Pi on your desk, you don’t want to accidentally reboot something else such as a corporate server (you never know where your new Linux skills will take you!).

SSH on the Mac

If you’re using a Mac rather than a PC, you won’t be able to use Putty. That’s not a problem because thanks to its UNIX roots, the Mac comes with a good SSH client built in. All you need to do is open a terminal window and run the SSH command.

First, get Terminal up and running. You can find it in the Utilities directory inside Applications. You can also find it by using Spotlight (the magnifying glass at the top right of the menu bar and typing in **Terminal**).

You’ll find that it isn’t too different from the Terminal you were using on the Pi, and there’s a good reason for that. As mentioned earlier, the Mac is actually based on UNIX, just as Linux (and hence Raspbian) is. Although the Mac does a good job of keeping its UNIX side well hidden (you can use a Mac for years and never come across it), you might want to dig into it a bit more if you find Linux and the command line to your liking.

Now that Terminal is open, connect it to your Pi. You’ll use the `ssh` command; because there is no GUI, you have to specify the relevant options on the command line. Based on the previous example, your terminal should look something like Figure 3-9.

```
0a022aaa:~ miggysX$ ssh pi@10.2.42.132
The authenticity of host '10.2.42.132 (10.2.42.132)' can't be established.
RSA key fingerprint is 31:04:39:be:64:a6:99:45:bc:5b:d1:c1:2f:71:94:b5.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '10.2.42.132' (RSA) to the list of known hosts.
pi@10.2.42.132's password:
Warning: No xauth data; using fake authentication data for X11 forwarding.
Linux raspberrypi 3.1.9+ #272 PREEMPT Tue Aug 7 22:51:44 BST 2012 armv6l

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.

Type 'startx' to launch a graphical session

Last login: Mon Sep 17 08:24:15 2012 from 0a022aaa
pi@raspberrypi ~ $ 
```

Figure 3-9. Connecting to the Pi from a Mac

To save space, Figure 3-9 shows the whole process rather than just an empty terminal window. The bit we're interested in right now is at the very top:

```
$ ssh pi@10.2.42.132
```

The command is `ssh`, but we have provided some additional information: the user and IP address that we want to connect to. You can read the options as "user `pi` on server `10.2.32.132`." A bit like an e-mail address, it is a nice succinct way to show who we want to connect as and where we want to connect to. You should, of course, replace the IP address with the one of your Pi.

The first block of text is something you'll see whenever you connect to a new machine:

```
The authenticity of host '10.2.42.132 (10.2.42.132)' can't be established.  
RSA key fingerprint is 31:a4:39:be:64:a6:99:45:bc:5b:d1:c1:2f:71:94:b5.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added '10.2.42.132' (RSA) to the list of known hosts.  
pi@10.2.42.132's password
```

This is just the `ssh` way of warning you that it can't determine whether you're connecting to the server you think you're connecting to and shows you the key that the server used to identify itself. In this case, you can ignore the warning and say Yes. If you want to know a bit more about what it all means, we covered it briefly in the "Putty for Windows" section. After you enter your password (`raspberry` is the default) you should end up at the following prompt:

```
pi@raspberrypi ~ $
```

Now that you have your command prompt, you'll look at what it actually means.

Welcome to the Command Line

Whichever route you take to get to a terminal, the end result will ultimately always be the same: you'll end up sitting at a command prompt (or simply a prompt) that will look something like this:

```
pi@raspberrypi ~ $
```

There is actually quite a bit of information here, but what does it all mean? You can probably figure most of it out from what we've been doing so far in the book, but for completeness, here is how it breaks down.

The first part (`pi` in the example) is the username. This isn't especially useful (or awe inspiring) when you have only a single Pi and a single user. It is, however, quite useful when you end up having to manage many users or when you have a variety of accounts of different machines. This is where the second part comes in (`raspberrypi` in the example) because it tells you the hostname of the machine you're connected to. Together, this lets us know who we are and where we are. Again, not terribly useful with one Pi, but this information will come in handy some day.

Caution We've highlighted this elsewhere, but it's so important that we're going to say it again. When you have more than one terminal open, it is all too easy to run a command on the wrong server. Sometimes you end up running more than one command before you realize your mistake. Often no harm is done, but if you accidentally rebooted a critical server at work when you actually meant to reboot your Pi, you're going right to the top of your IT department's revenge pending list.

The next item on the prompt (the ~ or tilde) actually tells us which directory we're currently in. The tilde is a special case because it refers to the home area of the current user. You can think of it as shorthand because really that's all it is. Any time you refer to your home directory (which in this case is /home/pi), we can simply replace it with the tilde. Time for a quick example:

```
pi@raspberrypi ~ $ pwd
/home/pi
pi@raspberrypi ~ $ cd ~
pi@raspberrypi ~ $ pwd
/home/pi
pi@raspberrypi ~ $
```

Here we used the `pwd` (short for *print working directory*) command, which prints out the full path of the directory we're currently in. In our example, it printed out `/home/pi` because that's the pi user's home directory. We then used the `cd` (short for *change directory*) command to go to the ~ directory. We used `pwd` again and were able to verify that after changing the directory we were back where we started. We'll be explaining these commands more thoroughly in the next chapter, but isn't this command-line stuff at least a little bit interesting?

When referring to the ~ in general speech, you refer to it as *home directory*. So if you sat at the prompt in our example and someone asked you what directory you were in, you would say you were in your home directory or if you wanted to be explicit you would say that you were in the pi user's home directory. If you wanted to get really specific you could even mention the hostname as well. You wouldn't say that you are in the tilde or "wiggly line" directory.

However, when you are giving someone a path to enter, you would say *tilde*. For example, if you wanted the user to go to `~/test/`, you would say "type cd space tilde slash test slash and then press Enter." This is because tilde in this case refers to something specific that you are using as a convenience to get to a location rather than actually telling someone what that location is. It also helps prevent confusion with relative paths, but again, that's something for the next chapter.

Note We've heard tilde pronounced as both *tilled* and as *till-duh*, so you should expect to hear both of them in the wild. Interestingly, though, you might find that whichever one you use, people might have trouble understanding you, and this applies even to experienced administrators. It's not because you're pronouncing it wrong; it's just they don't know that the "wiggly line to the left of the 1 key" is actually called *tilde*. If you find yourself in that situation, referring to it as "the wiggly line one" seems to be a fairly safe bet.

The last part of the command prompt shows you what your access level is. Linux systems are pretty basic when it comes to system privileges. It can generally be summed up this: you are either the root user or you're not. Now, there are exceptions to this rule and being a member of certain groups can allow you to do certain things that would normally require root access. That said, this is really the exception rather than the rule. If you are root, then you can do pretty much anything you please without hindrance. If you are any other users, expect to find your abilities drastically curtailed. This isn't as true today as it once was because there are tools such as `sudo` that allow an administrator to grant the ability to run certain commands with root privileges without that user having access to the root account. Of course, ultimately `sudo` runs as root and it just decides whether it will run the command that you asked it to, so it still obeys the all or nothing rule.

From the prompt's perspective, if you are a normal user, you get a dollar sign as a prompt. We're logged in as the pi user (which obviously isn't the root user), so we too have the \$. If we were logged in as root, we would be a privileged user, so we would get the hash symbol (or the pound sign if you're not from the UK). A root prompt looks like this:

```
root@raspberrypi ~ #
```

The hash symbol is a sign (and a warning) that you are running with elevated privileges and that you should therefore be very careful about what you type. The prompt also shows the username as root, which provides an added indication. However, these subtle hints are easy to miss (people don't tend to study the prompt before each command), so always check who and where you are before you press the Enter key.

That's really all there is to the prompt. Most shells allow you to customize it to show different information, and occasionally some administrators do change this. In our experience, most people don't bother, and we rarely feel the need to change it ourselves. The reason why we're bringing it up here is because if you ever do log in to another Linux machine, and the prompt is wildly different, it's not a cause for concern. The terminal works in exactly the same way except the prompt has been customized.

Different Shells

Okay, maybe we oversimplified things when we said that all terminals work alike. In truth, there are some minor differences, especially in what features they support, but more often than not you won't notice. In fact, what will likely catch you out is when the remote machine is using a different shell. A shell is a program like any other but rather than being a tool for browsing the web or reading e-mail, it provides a command line for you to interact with the operating system. It operates like a wrapper around the kernel and that's how it got its name: by being the shell around the kernel just like a peanut.

So far, we've been making use of the BASH or Bourne Again Shell. It is by far the most popular shell in use today and it's available on every Linux distribution. It's also usually the default, which is why you'll never have to worry too much about which shell your system is running.

The reason we're bringing this up here is because, like everything else in the open source world, opinions vary on what makes the best shell. Some take a bare-bones approach and only support the most basic features while shells like BASH come with batteries included and have more features than you can shake a stick at. Despite BASH's popularity, it is conceivable that at some point in your illustrious Linux career, you might come across something else, and forewarned is forearmed.

We're not going to spend any more time looking at the different shells because quite honestly you'll probably never need to worry about it. Even if a machine is using a different shell, you can probably start BASH to replace it (most systems have BASH installed, even if you get a different shell when you log in). Because system commands are the same regardless of which shell you use, even if you can't get to BASH you'll probably find you can do everything you would normally do. In fact, most shells are interchangeable for the most part, and chances are even if you were using another shell, you probably won't even know it!

Summary

This chapter introduced you to the terminal and its rich (if not exactly fascinating) history. We then discussed how the terminal is still used today and how it has been adapted to suit our modern needs. We then looked at the three main ways that you can get to a terminal on your Pi: the console, a virtual terminal, and over the network via SSH. Next we gave a quick description of the command prompt and touched on some interesting commands to highlight a point (which we absolutely promise will be explained when you turn the page), and we wrapped up the chapter by giving you a heads up about the different shells out there and how they may (but probably won't) trip you up.

Coming up in the next chapter is all you need to know to happily navigate around your Pi, create files, modify files, delete files, and much more. This is where we start to get really serious about the command line and where you'll start picking up some real hands-on skills.

CHAPTER 4



The File-Paths to Success

This chapter takes a look at what a filesystem is and how different operating systems try to make sense of them. We explain the difference between separate root filesystems and unified root filesystems, and why knowing this up front might save you a lot of headache.

We then move on to looking at the filesystem presented to you on your Pi. We explore the standard layout and explain what goes where and what it actually does. We will also show off some simple commands that highlight some benefits of the UNIX “everything-is-a-file” approach to system design.

Next we will show you how to move around your system, and you’ll learn about fully qualified and relative paths. We’ll show you how to create directories and files (and then copy, move, and delete them), and we’ll explore Linux file permissions and show you how Linux decides who can do what to your files. We’ll also touch on users and groups, and show you how to set ownership and file permissions on your own files.

We’ll then round off the chapter by showing you how to create the Linux equivalent of shortcuts.

What Is a Filing System?

We actually planned to start this section with a dictionary definition of what a filing system is, followed by our interpretation. Unfortunately, most of the definitions we could find basically said “a filing system is a system for filing,” we’ll skip that part and go straight to our own definition. It does suggest, however, that although a filing system is easy to talk about, it is a lot harder to define. We think a filing system is this:

A way of sorting and categorizing data that makes finding something easier

This definition covers everything from having a simple in-tray (at least you know where the document is, even if you can’t put your hand on it right away); to the way you view your contacts on your phone (that’s an alphabetical system right there); all the way to some of the more exotic systems such as the Dewey Decimal System, used to categorize nonfiction content at your local library.

Actually, libraries are an interesting use-case for filing systems because most of them have at least two going on at the same time. Generally libraries tend to arrange fictional work by the surname (and then forename) of the author, whereas the nonfiction section is sorted based on Dewey Decimal numbers. By using two systems at once, they solve a fundamental problem. Storing books by the author’s name is great if you have a favorite author and want to find more books. When it comes to nonfiction, though, you probably want to browse a selection of books on a given topic, without a particular author in mind. If those books were sorted by author name, you’d probably never be able to find anything.

Some libraries take this a step further by sorting fictional books first by genre (say Sci-Fi or Fantasy) and then within that category they sort by author name. This has some pretty neat benefits. You can still go straight to your favorite author because you’ll know the genre, but you can also browse for similar books with the same genre. If you are looking for the next best thing in the fantasy genre, you can simply stand in that section and flick through.

So far, so good. Filing systems make finding things easier, and we can see that it’s important to use the correct filing system for the task; otherwise, you’d probably be better off with no filing system at all.

But what does this have to do with computers? Well, one of the major tasks for any computer is to process information, which means that it has to be able to store and retrieve information easily. Older computer systems did use the alphabetical system and it worked very well, but as soon as there was a buildup of a number of files, it started to get very messy; and finding information became a challenge.

This issue was addressed by having the ability to create a directory (often called a *folder* these days), in which you could put related files. This helped a lot, but the same old problem started to creep up on people. Yes, you could have a directory for accounts and another for tax returns, but what if you were an accounting firm and you had lots of accounts and lots of tax returns? Soon you were back to where you started because you either had a small number of directories with large amounts of files or a large number of directories with only one or two files in them.

Around the same time, the term *filing system* (used predominantly for paper-based systems) became shortened to *filesystem*, which is almost exclusively used to refer to the way a computer stores data on a disk. Because this is the term you will come across most often, we'll switch to using the *filesystem* term as well from now on.

The last change to filesystem structure that remains familiar to this day was to allow directories to contain other directories. This would allow you to have a great deal of flexibility to store content in the most appropriate ways. As always, there's more than one way to arrange a given set of files, but by and large the system has served us well and with modern search technology, finding what we're looking for is easier than ever.

More than One Filesystem

Here is where we hit a bit of a snag with this approach. A filesystem sits on top of some form of storage. That can be a hard disk, a USB stick, a DVD, or any number of storage media. Each device is effectively independent from any other. This makes sense because you can take a USB stick from your laptop and then use it on your desktop machine. Clearly there is no link between the filesystem on the USB stick and your laptop because otherwise your desktop wouldn't have all the information it would need to find the files.

Each filesystem has a *root directory*, which is the main directory on a device and holds all the other files and directories. Like the root of a tree, all the files and directories branch off from this central location. This does pose a bit of a problem, though. If each device has its own filesystem and root directory, how do you easily present this structure to the end user? It turns out that there are two approaches to solving this problem. The first is separate roots; the second is a unified filesystem.

Separate Roots

This approach was adopted by Microsoft all the way back in the days of MS-DOS and it's still with us today. The approach is very simple. Each device has its own root entry and the user can then use this entry to locate the device of interest and then can simply navigate the filesystem as normal to find the file. Under Windows, each root is assigned its own letter. For historical reasons, the system disk on Windows is referred to as C:, which you'll often hear people refer to as *the C drive*. Floppy drives are assigned A: and B: (PCs with hard disks tended to have two floppy drives) and because few people had more than one hard disk, the CD-ROM drive tended to be assigned to D:.

This system is really easy to use, and there is never any confusion about where a particular file is because you can tell simply by looking at the drive letter. However, there are a couple of issues with this design:

- First, because root devices are assigned to letters, and there are only 26 letters in the alphabet, you are limited to 26 devices. Hardly an issue to home users even today, but back when businesses were using large mainframes, they could have hundreds of such devices.
- The second issue is that the user has to be aware of the physical location of the file, which can actually add complexity because it ties the location of the file to the physical device that the file is on. If a bigger disk is added to the system, and the files were moved to that new location, their root device would change. Any person (or software) that depends on a file's location would need to be updated, and that's never fun.

In short, although the system employed by Windows is simple and effective, it can become an administrative headache on very large and distributed systems. These days there are lots of ways around that, and new techniques have emerged to hide this structure beneath the surface, such as the Microsoft Distributed File System (DFS) for network shares and libraries for grouping content in one place on your desktop such as documents and music. So it really isn't a huge issue for Windows systems today. However, back when UNIX reigned supreme, these technologies weren't available, and they went with the other approach: a unified filesystem.

Unified Filesystem

Linux, being based on UNIX, has a unified filesystem. That means that unlike Windows, which has multiple root devices, a Linux system has only one and it is always mounted on `/`, which is UNIX-speak for the root directory. If there is only one root directory, then how does Linux handle additional devices, all of which we know contain their own filesystems?

The solution is to take the root directory of the new device and then attach it on an existing directory in the tree. This is known as *mounting a filesystem*, which allows Linux to have an almost unlimited number of devices because it can attach them anywhere in the existing tree. This means that you could mount one device and then mount another device on a directory inside the first device. From the user's point of view, they can move through all these different directories as if they were all on one big device. The physical structure (that is all the disks and the raw data on them) is completely hidden from the end users. They have no way to tell they've crossed from one device to another.

This solves the issues highlighted with the separate root approach in that the unified filesystem is consistent regardless of what the underlying mechanism is. In fact, you can mount network file shares and even virtual filesystems into the tree. Of course, the benefits of the separate root approach have now become the downside for the unified approach. It's no longer easy to see where things are going, and this again can make things more confusing. Like Windows, Linux systems have evolved over time so that these problems aren't quite as pronounced as they would otherwise be. However, these benefits are mostly seen in the GUI and tend to be far less pronounced when using the command line.

The Mac has to be Different

As an interesting comparison, OS X on the Mac makes use of both approaches. Under the covers, OS X is a UNIX-based system and so it does have a unified filesystem. However, unless you go to the command line (and most Mac users don't), you will never see evidence of this. When you attach a USB stick to the Mac, it will mount it in a similar way as Linux, but it will show it to the user as if it were its own root device in a similar approach to Windows. However the Mac doesn't assign a drive letter, it just sets up a unique name.

Bring it All Together

Admittedly that's a fair bit of theory and you're probably wondering where all the fun commands are that we promised to show you. You might also be wondering why we just bored you with all that, and when (if ever) such information might be useful. The reason we've put emphasis on this theory up front is because when we started out with Linux, a lot of the things that caught us out were preconceptions from using other operating systems. When one of us first installed Linux some 15 years ago, it took several hours of swearing before he figured out what "you need a root / device" meant. This isn't just a Linux thing, a dabbling with the BSD family of operating systems and its alternative way of handling disks caused us to wipe the wrong disk because we were thinking Linux and not BSD.

So hopefully this last section will save you from some of the pain that we went through when we first started out. You don't need to memorize all the theory, but if you come away from this section with an appreciation that despite many similarities, Linux is not Windows or a Mac, then you will be ahead of the game.

Everything as a File

Now normally when we discuss this topic, we start by talking about hard disk partitions because most people are familiar with them, and there's a nice easy mapping between the way Linux presents them and the physical device. Because the Pi doesn't technically have a hard disk (it pulls some strings to make the SD card look and feel like one), this easy example isn't actually available on the Pi. The thing is, it's too good an example to pass up, so we're going to stick with the classic. So without further ado, here are the files that describe the disk set up on one of our servers:

```
/dev/sda
/dev/sda1
/dev/sda2
/dev/sdb
/dev/sdb1
```

As you can see, these files all live in the `/dev` directory, which is where Linux keeps all the device files. There's quite a bit more to say on this, but we'll come back to it a little later because, trust us, it will make much more sense if we cover it last. So for now, let's ignore the directory and focus on our collection of files. You'll notice that in this example, all the files start with the `sd` prefix because they are all "SCSI disks" and share the same driver. For machines that have IDE-based disks, the prefix is `hd` for *hard disk*.

So now we know that we have some SCSI-based disks (SATA disks also show up as SCSI), but what else can we tell from those files? Well, you'll notice that we have `sda` and we have `sdb`. The first SCSI disk on the system is assigned to `sda`, and the second is assigned to `sdb`. No surprise that the third disk will end up being assigned to `sdc`. Because we don't see an `sdc` in our example, we can assume that there are only two disks on this system.

Now we're really getting somewhere. There are two SCSI disks on our system, and we just have one more thing to discuss: the numbers stuck on the end of the names. In this case, the *number* refers to the partition number on the drive. The first partition is 1, and the second partition is 2. Nice and simple with no real surprises, we're sure. With this additional information, we not only specify which disk but also which partition on that disk.

This does, of course, beg the question of why you should care. Surely there must be an easier way to find out which disks are attached to the system? Well you'd be right, but we're not looking at this to discover what's connected to the system. Instead, we're looking at how we can specify a particular device that we want to access. This is because everything in Linux is a file, so even physical devices are represented in that way. When we want to access a hard disk or we want to access a particular partition, we do so by accessing the relevant file. Not only physical devices follow this rule; directories are files as well. In day-to-day use, you don't have to think about this, but it's a very powerful feature. By treating everything as a file, Linux provides a standard way for all programs to talk to all things. This greatly simplifies development and allows us to do some impressive things, such as redirecting input and output or using the output of one program as the input for another. We cover this in more depth in Chapter 5.

Some real-world examples are useful here. When you want to make changes to the partition table on a hard disk, you use the `fdisk` tool. The `fdisk` tool will need to know which disk you want to work. Assuming that you want to partition your second disk, the command would look like this:

```
$ fdisk /dev/sdb
```

Because we're partitioning the drive, we want to refer to the whole device, and `/dev/sdb` lets us do that. But let's say we have done our partitioning, and we have one big partition that we want to format for Linux to use. The command for that is `mkfs.ext4` (on modern Linux distributions), and as before, you're going to need to tell the command what you want to format. In this case, we want the first partition on our second disk so the command would look like this:

```
$ mkfs.ext4 /dev/sdb1
```

Again, this makes sense because we're referring to a specific partition on a specific device. No big leaps of faith here, but we can now share a little secret with you. The `mkfs.ext4` command doesn't really care what filename

you give it. If you forgot to put the 1 on the end, it would have happily formatted the entire device. This is a case of “everything as a file” at work because everything is a file, and files are all accessed in the same way. Our tools simply read and write to files; they neither know nor care what it is they’re writing to!

This concept takes a bit more effort to really understand. Under the covers, Linux manages all these different devices for us, each with its own specific requirements and drivers. To keep things simple for application developers (not to mention the users), Linux hides all this complexity and instead presents each device as a specific file. This brings us nicely back to the /dev directory (we told you it would come up). The files in that directory are special in that they don’t really exist, at least not as physical files that are stored anywhere. Linux creates a virtual filesystem that it then populates with files for the devices that are available on a given computer. Because our Pi doesn’t have any SATA or SCSI disks, you won’t find any /dev/sd files on your Pi.

Filesystem Layout

Let’s look at how the filesystem is laid out on your Pi. You’ll find that the structure is very similar to what you’d find on any Linux system, so not only will this section help you master your Pi but you’ll also be able to navigate servers as well.

/ (Root Directory)

This is the root directory (not to be confused with the root user’s home directory) that represents the top of the file hierarchy. Everything goes somewhere in or under this directory, no exceptions. As we discussed earlier in the chapter, Linux has a unified hierarchy, and this is where it all starts.

/root

This is the home directory for the root user and spends a large amount of its time being confused with the root directory (/). It can get more confusing because if you are in the root directory and someone asks you to “go into the root directory,” which one do they actually mean? Often you can easily tell from the context, but if not, don’t be afraid to ask for clarification; you’d be surprised how often this one comes up.

/etc

/etc is arguably one of the most important directories on your system. It contains most of the configuration files for not only your system but also the applications that you might have installed (such as the Apache web server). Many users new to Linux take care to back up their applications, but often forget about the configurations stored in /etc. At the end of the day, the application can usually be replaced quite easily, but getting that config right from scratch can be a real headache. This is a directory you want to take great care of.

/proc

/proc is a virtual filesystem that the kernel uses to provide easy access from userland tools (that is, software that is run outside of the kernel). Everything you ever needed to know about the system state or running processes can be found in /proc. Two common examples are the CPU configuration (stored in /proc/cpuinfo) and memory usage (stored in /proc/meminfo). Most of this information is read-only, which makes sense because it is just a virtual representation. However, some files do allow communication to go both ways and you can potentially use it to tweak kernel and system settings while the machine is running.

/var

/var is generally where you'll find files created by your applications and the system itself. For example, most applications will store their logs in /var/log/, and many will store the lock files in /var/run/. Lock files don't have any meaningful content; an application just uses them to work out whether it is already running so that you can't accidentally have two copies running at once. On Raspbian, the Apache web server uses /var/www/ for storing a website's files. In other distributions, this is not the case—such files tended to be located in the /srv/ directory. That said, more mature applications may still be in /var, even on modern distributions.

/boot

Traditionally, the /boot directory actually lived on its own small partition on the first hard disk. At the time, the majority of computers were unable to boot from a single big partition, so it was very common to see these split out. On modern machines (although the Pi still requires it), this is no longer a problem, so this directory is often included in the root directory directly. As its name suggests, it holds the key files needed to boot a system, including the bootloader and the Linux kernel.

/bin and /sbin

These locations store user and administrative programs, respectively. When you try to run a program from the command line, there are specific places that the system will look for the application. This list of locations is known as *the path*. Usually a normal user only has /bin in their path, so they effectively can't see applications in /sbin. There are some applications that the user can access even though they're generally used only by administrators, but the user needs to know where they are. This usually doesn't pose a problem, and you won't ever need to go looking for anything.

/dev

We already touched on this directory in the “Everything as a File” section. It contains a file for each device or subdevice on the current system and provides a way for system tools (and, of course, users) to easily access the hardware on a particular machine. Besides the disk devices discussed earlier, there are also devices for graphics cards, sound cards, virtual terminals, and more.

Note You won't find network cards in the /dev directory because they're considered a special case. To find information on network interfaces, you need to use the `sudo ifconfig -a` command. This will list all the network devices that Linux knows about. For more information, check out Chapters 3 and 9.

/home

Traditionally, all user home directories would be stored under /home and this was often its own partition, disk, or network share. The idea was to keep user data separate from system data and applications. Most UNIX-based systems still follow this rule, but there are occasions when you will find some home directories living elsewhere (i.e., the root user's home directory). The Mac, for example, stores home directories in /Users, and some businesses will put different users in different locations based upon their needs. Although /home generally contains only home directories, there's no requirement that home directories actually reside here.

/lib

This directory contains library files that are needed by various applications. Libraries allow functionality to be packaged up and then shared by other applications. A good example might be a database driver that allows applications to access a database. So that these applications can find those libraries, they need to be installed in a known location and with a known file name. It's rare that you'd need to poke around in this directory, but if you do, you should be careful because breaking things in here could affect the stability of your system.

/lost+found

We mention this one for completeness, but it's not really part of the filesystem structure *per se*. Rather it's where files get placed when the filesystem loses track of them. For example, if a disk were to be damaged, you would need to run a disk repair. Some files may be recovered, but for various reasons, it may be impossible to determine where that file came from. If that happens, Linux will place those lost souls in this directory. We've never had cause to look into this directory, and most likely you'll never need to look in there either.

/mnt

Short for mount, the /mnt directory was traditionally where you would mount additional filesystems. If you wanted to attach a network share or an external hard disk, you would create a directory in /mnt and mount in there. Floppy disks and CDROMS also need to be mounted before use and were also generally mounted here as /mnt/floppy and /mnt/cdrom, respectively. However, in recent years, this directory has mostly fallen by the wayside because more people have removable media rather than always attached storage (see the next section on /media).

/media

This directory is relatively new to Linux and was added to make a clear distinction between mounted external devices (such as those in /mnt) and removable media such as USB sticks, cameras, and media players. These are usually handled automatically under Linux, and so usually you won't manually add or remove anything in this directory.

/usr

This is where most of the software on the machine ends up and so it is often the largest directory on a server (at least if you don't count user home areas). Although it's useful to know where the software lives on your machine, because everything is handled automatically for you, everything in here should "just work".

/opt

This directory is a jack of all trades. On some systems, it is packed full of applications; on others, it remains completely empty. It's usually used for third-party software and applications. For example the Oracle database server installs in /opt/ by default. You probably won't find much cause to use this directory, but be careful if you do because it is easy to forget things that are sitting in /opt/ when it comes to taking backups and so forth.

/srv

Another relative newcomer, the /srv/ directory is the designated location for storing data for services that serve files. Although this directory seems to be present, some applications still don't make use of it either out of custom or simply

because everyone is used to the contents being somewhere else. If you’re looking for things that used to be in `/var/`, this is probably a good place to look next.

`/sys`

This directory contains system information, and unlike `/proc/` (which is only in memory), the contents of this directory are written and stored on disk. It doesn’t seem to get an awful lot of use, and we have never even had to look inside this directory. However, it needs to be useful only once for you to be glad it was there.

`/tmp`

This directory is the computer’s scratch pad, and all sorts of applications create files in this directory. It is used whenever temporary file storage is used (such as during processing). This directory is emptied by default on a reboot, but it can still potentially get quite large. In theory, all applications are supposed to treat `/tmp/` as transitory; that is, they should never expect a file to be there after a restart.

Wrapping it Up

That pretty much sums up the key areas on the Linux filesystem. No doubt there are other little nooks and crannies that you might come across as you look around your system, but they will probably be a subset of these. Remember that although this structure is followed by most systems, most doesn’t necessarily mean all—and you might come across some differences or where directories are used in other ways.

Putting it to Work

We’ve covered a huge amount of information so far in this chapter. We’ve looked at how filesystems work and how UNIX systems (and Linux in particular) have adopted a unified approach. We looked at “everything as a file” and we showed how these special files and filesystems fit into the grand scheme of things. We wrapped all that up with an overview of what goes where on the filesystem and what they do. What we haven’t done yet is actually used the filesystem.

So far, we have been pretty much “hands-free,” but that’s all going to change in this section. We’ll start by showing you how to create directories and move about on the filesystem. Once we’ve given you the power to create, we’ll then give you the power to destroy, and you’ll be able to remove files and directories at will (not always a good idea). We’ll wrap this section up with a quick overview on Linux file permissions and how to read and set them. Let us begin!

Where Are We? Using `pwd`

The first thing we need to show you is how to figure out where you are on the system. The easiest way to get your bearings is to look at the command prompt. We talked about this in the last chapter—as a quick refresher, here is what the command prompt looks like while we’re sitting in our home directory:

```
pi@raspberrypi: ~$
```

The tilde (the squiggly line) is shorthand for the current user's home directory. We highlighted this feature already when we first introduced it, so we won't go through it again, but what if we want to know which directory we are actually in? What we need is the `pwd` command. We touched on this useful tool in the last chapter as well, but in case you missed it, here it is again:

```
pi@raspberrypi: ~$ pwd
/home/pi
pi@raspberrypi ~$
```

This tool is very useful for telling us where we are, but it doesn't tell us anything about what's in the same location with us. It's like being blindfolded and told you're standing in the kitchen; it's a great start, but you're still effectively blind. You'd certainly like to know who and what is in the room with you; for that, you need the `ls` command. This command does have a large number of options, though, and so we will only cover the most common ones, those that we use every day. Actually you can just remember them as recipes because you'll often just pass the same options time after time (or at least we do).

What's in Here with Us? Using `ls`

Now let's see what's in here with us:

```
pi@raspberrypi: ~$ ls
Desktop    python_games
pi@raspberrypi: ~$
```

Although we can't easily show it in the book, `Desktop` and `python_games` are both colored deep blue. This tells us they are directories. At present, we don't have any files in this directory, or do we? Actually we do, but by then they are considered hidden files. Under Linux, any file that begins with a period (full stop) is considered hidden. There's nothing special about the files; they are various config or temporary files that various applications have created. We generally don't want these files cluttering up our display, so `ls` and friends don't show them. We can, however, force `ls` to show us those files with the `-a` flag like so (although you will most likely see slightly different results on your Pi):

```
pi@raspberrypi: ~$ ls -a
.
..          .cache
.bash_history .config
.bash_logout   Desktop
.bashrc        python_games
 dbus          .ssh
pi@raspberrypi: ~$
```

Creating Files to Play with: Using `Touch`

For now, we'll leave the `-a` flag alone and create our own files to play with. Because we haven't covered how to create and edit text files yet (we will show you how to do that in Chapter 6), we'll introduce you to another little tool called `touch`. Under Linux, files have two timestamps:

- The creation timestamp
- The last modified timestamp

These timestamps allow you to see when a file was created and when it was last updated. This is useful from an administration point of view because you can see which files are actively being used, but various tools (such as backup scripts) use this timestamp to figure out whether a file has changed since they last looked at them. Sometimes it is useful to be able to update that timestamp without changing the contents of the file, and this is where `touch` comes in. It touches the file that updates the timestamp, but if the file doesn't currently exist, `touch` will create it for you. In other words, it's a great tool for creating empty files. So, let's start by creating a couple of originally named files:

```
pi@raspberrypi: ~$ touch raspberry
pi@raspberrypi: ~$ touch pi
pi@raspberrypi: ~$ ls
Desktop pi python_games raspberry
pi@raspberrypi: ~$
```

And that's all there is to it. As you can see from the `ls` that followed it, we now have two additional files. This time they are gray in color, which tells us that they're normal files. By coloring the entries for us, `ls` makes it much easier to see what we're doing. For example, any file that is executable will be colored green, but we'll come back to file permissions later in this chapter.

So far we only have two files, but if we had 10 or 20, it would start to get a bit crowded in here. The way to handle that is, of course, to create directories to store our files (and potentially other directories) in, and that's what we'll look at next.

Note Directories and folders are basically the same thing. Originally called *directories*, Microsoft started referring to them as *folders*, which they felt was a better description. Although Linux used to use the term *directories* because it has become more of a desktop operating system with people moving over from Windows, the folder terminology has become increasingly common.

Somewhere to Store our Files: Using `mkdir`

To create a directory, we use the `mkdir` (or *make directory*) command. Unsurprisingly, it will create a new directory. However, if there is a file with the same name, or if the directory already exists, you will get an error message. Let's create a directory called `pifun` to store our files:

```
pi@raspberrypi: ~$ mkdir pifun
pi@raspberrypi: ~$ mkdir pifun
mkdir: cannot create directory 'pifun': File exists
pi@raspberrypi: ~$
```

As you can see, trying to create the directory twice will cause an error. Don't be distracted by `File exists`; this could actually refer to either a directory or a file. Another quick `ls`, and we'll see that things are moving along quite nicely:

```
pi@raspberrypi: ~$ ls
Desktop pi pifun python_games raspberry
pi@raspberrypi: ~$
```

Making Use of a New Directory: Using the mv Command

So now that we have our directory, let's tidy up the mess we've been making. Because we want to move the files into our new directory (rather than just create a copy of them), we'll need to use the `mv` (*move*) command. This command is a bit more complicated than the ones we've covered before because it takes two arguments (that is, the things you want the command to act on) rather than one. This makes sense, though, because not only do we need to tell `mv` what we want to move but we also need to tell it where we want it to move the file to. As with most file commands under Linux, the first argument is the source, and the second argument is the destination. Let's move those files now:

```
pi@raspberrypi: ~$ mv pi pifun
pi@raspberrypi: ~$ mv raspberry pifun
pi@raspberrypi: ~$ ls
Desktop pifun python_games
pi@raspberrypi: ~$
```

So far, so good. Now we want to make sure that our files arrived in one piece. There are two ways to do that. We could go into the `pifun` directory with `cd` and then run the `ls` command or alternatively, we could simply give `ls` the path to the directory that we want to look into. We've already used the first approach, so let's try give the second approach a try:

```
pi@raspberrypi: ~$ ls pifun
pi raspberry
pi@raspberrypi: ~$
```

Time for Some Cloning: How to Use the cp Command

Admittedly this isn't very exciting, but as a newly minted administrator, you'll be spending a lot of time moving files about and checking where things are. So far, you've learned how to move a file, but what if you just want to copy a file instead? When making a backup or getting a selection of files ready to send to a friend, you want to actually keep the originals. For this, we use the `cp` command (short for *copy*). Let's move into our new directory and copy one of our files:

```
pi@raspberrypi: ~$ cd pifun
[pi@raspberrypi pifun]$ cp pi pi2
[pi@raspberrypi pifun]$ ls
pi pi2 raspberry
[pi@raspberrypi pifun]$
```

That worked well, we now have `pi` and `pi2`, just as expected. Let's try the same thing again, only this time we'll copy a directory:

```
[pi@raspberrypi pifun]$ mkdir moarpi
[pi@raspberrypi pifun]$ cp moarpi moarpi2
cp: omitting directory 'moarpi'
[pi@raspberrypi pifun]$
```

This time it didn't quite go according to plan. The reason our copy attempt failed was because by default `cp` will only copy individual files. It won't copy entire directories. Even though directories are technically treated as files, the way we use them via various commands mimics a real-world folder. The reasoning behind this is that when you copy a directory, you copy everything within it, including all its files and directories. This could be a lot of data and could

take up a lot of space as well as a lot of time to complete. Forcing us to be explicit about our intentions (which soon becomes second nature) means that when we mean to copy a single file but accidentally pick a directory, we will get stopped before any copying takes place.

That's all well and good, but what if you really did want to copy that directory? When we first looked at the `cp` command, we mentioned tasks such as taking backups, and (let's be honest) you're much more likely to want to back up a directory than a list of specific files. We can get the behavior we're looking for by telling the `cp` that we want to copy *recursively*. This will then copy the directory and anything within that directory to the destination. We specify this by using the `-r` flag like so:

```
[pi@raspberrypi pifun]$ cp -r moarpi moarpi2
[pi@raspberrypi pifun]$ ls
moarpi moarpi2 pi pi2 raspberry
[pi@raspberrypi pifun]$
```

Unlike the `copy` command, when you move a directory, there is no need to specify that you want to do so recursively because moving a directory without its contents wouldn't make an awful lot of sense.

The Power to Destroy: Using the `rm` Command

So far, we've shown you how to create files and directories, and how to copy and move them about. Now we're going to show you how to destroy those files with the `remove` or `rm` command. It goes without saying that the `rm` command is one of the most dangerous in your arsenal. It can quite easily destroy an entire server if you're not careful, and we know people who have accidentally done just that.

For a change of pace, let's look at how we can delete an empty directory. Before we look at how to do it with `rm`, we can use the `rmdir` command for this (which is short for *remove directory*). Now the catch with this command is that it will delete only directories that are completely empty. If there is even a single file inside, this command will fail. This makes it very safe to use, but not all that practical because generally when you delete a directory, you also want to remove all its contents. Let's kill two birds with one stone:

```
[pi@raspberrypi pifun]$ rmdir moarpi2
[pi@raspberrypi pifun]$ rm moarpi
rm: cannot remove 'moarpi': Is a directory
[pi@raspberrypi pifun]$
```

We could delete `moarpi2` with `rmdir` because the directory was empty, but when we tried using the `rm` command, it refused to cooperate. This is because the `rm` command was written with similar reasoning to the `cp` command. Because removing a directory is far more dangerous than simply copying it, this is probably a good thing. We can use the same `-r` flag to tell `rm` to delete recursively:

```
[pi@raspberrypi pifun]$ rm -r moarpi
[pi@raspberrypi pifun]$
```

Success! Now sometimes when you try this, especially on a large directory with lots of files and subdirectories, you can end up with lots of issues that cause `rm` to give up. For example, some files might be write-protected. You can suppress these errors by using the `-f` flag, which means *force* and is akin to saying, "Damn the torpedoes! Full speed ahead!" That sounds like a great idea until you stop to think what would happen if you ran this command (which you should *never* do):

```
[pi@raspberrypi pifun]$ rm -rf /
```

If you accidentally run that command, `rm` will proceed to delete absolutely everything on your system. If you happen to have a USB hard disk attached or you've mounted some network shares, you're in real trouble because `rm` won't confine itself to your internal disks; it will crawl through the entire tree deleting everything in its wake. This is one of the reasons why you should use a normal user account for day-to-day tasks. Your own user will not have sufficient privileges to delete anything critical to the system, but even then, chances are high that you can still damage all your attached media. You need to be very, very careful whenever you use the `rm` command and you'd better double- and triple-check it because Linux will assume that you know what you're doing and won't ask for confirmation!

The `rm` command can also remove files simply by providing the path to the file. You don't need to use the `-r` flag for this operation, so you can simply do this:

```
[pi@raspberrypi pifun]$ rm pi2
[pi@raspberrypi pifun]$ ls
pi raspberry
[pi@raspberrypi pifun]$
```

And that in a nutshell is how you move about and manipulate the filesystem.

Fully Qualified and Relative Paths

In Linux, there are two ways to specify a path to a file or directory:

- You can give a fully qualified path that starts with a forward slash
- You can give a relative path that starts with a filename, directory name, a dot, or two dots

Strange though these might sound, they are both just ways of providing a specific location to your programs.

A path is considered fully qualified when it starts from a fixed reference point (i.e., the root directory). Regardless of where you are on a system, a fully qualified path will always point to the same location. It's like the old bell tower in the middle of town; if you give anyone directions using that as your reference point, you have a common anchor that both you and your friend know how to reach.

On the other hand, a relative path depends on your current location to make sense. You can specify paths using `./` to mean the current directory or `../` to mean the parent directory (the next one up from the current location). If you had a path that looks like `.../test.txt`, this would only work from a few specific locations (where the file is two levels above your current directory). However, it's nice and short as well as being easy to type. Of course, the same file (if it were created in your home directory) would be accessible with `/home/pi/test.txt`, which unlike the relative path, can be used from anywhere on the filesystem pretty easily. Relative paths can be tricky when you're several levels deep and you're not sure where you are, but you do know the full path of the file you want. If you are in `/home/pi/Documents/Work/Projects_2012/Spreadsheets/` and you want to access `test.txt`, you could either use the full path or alternatively use `../../../../text.txt`. As you can see, in this case, the relative path is more confusing than the fully qualified path.

So when should you use one or the other? You should use whichever option is the most convenient or makes the most sense for the task. Sometimes it's faster or easier to use a fully qualified path. Other times, you are buried deep down in the tree, and writing fully qualified paths would be tedious at best and totally confusing at worst.

Users and Groups

We're only going to touch on the basics of users and groups here so that you know enough to understand the file permissions section that is coming up next. Users and groups are key to the way Linux secures your files, and you'll need to know about them before we move on to the next section.

In the UNIX way of thinking, all people have their own usernames. A *username* identifies a particular person or entity (for example, a web server might have its own username) on a particular system. So far, we've spent most of our time as the pi user, but we've also seen that we can become the root user. Your username is the key that Linux uses to identify you as you.

Groups are similarly straightforward. Each user belongs to one primary group, but may actually be a member of any number of groups on the system. On a university system, a student's username might have its own private group (standard practice on Red Hat and Debian systems these days) but they might also belong to a student group and a research group. They might also belong to a group specific to their department. Groups are useful to administrators because we can group a selection of users together and treat them as a single entity. This makes things such as file permissions much easier to manage.

When you create a new user on your Pi, you will automatically create a group with the same name. On some systems, users would by default join a user's group, but as you'll see in the next section, this could lead to accidentally giving people access to files that they shouldn't have. Because a private group is by definition private, no one else will be a member, and no one can gain access to your files just because they happen to be in the same group as you. This is why on modern systems you'll usually see that the owner and group of a file happen to be the same.

File Permissions

File permissions allow you to determine who you want to be able to access your files and what exactly they're allowed to do with them. There are three different permissions that you can set. The first is whether or not someone can read your file. The second is whether they can write to your file. And third is whether they can execute it (i.e., run it like an application).

Of course, just being able to set these permissions on a file isn't particularly flexible. You might want to give access to only a certain group of people and restrict everyone else. This is where users and groups come into play. On Linux, there are three roles that a given user might fall into:

- *User*: Refers to the owner of the file
- *Group*: Refers to the group that owns the file
- *World*: Also often referred to as *other*

Each role gets its own combination of permissions; that is, you can define whether any of those three roles can read, write, or execute your files. We're going to show you how to do just that but before we do that, we need to show you how to see what permissions are actually in effect and so now is a good time to show you how to use the -lh option for the ls command. Let's try running it now:

```
[pi@raspberrypi pifun]$ ls -lh
-rw-r--r-- 1 pi pi 0 Oct 7 16:29 pi
-rw-r--r-- 1 pi pi 0 Oct 7 16:58 pi2
-rw-r--r-- 1 pi pi 0 Oct 7 16:29 raspberry
[pi@raspberrypi pifun]$
```

The -lh argument specifies that we want ls to show us a list of files (-l) and that we want file sizes to be in human readable format (-h). Without the human readable flag, ls will show us all sizes in bytes, which when you're dealing with large files is not very easy to read. It doesn't really matter in this example because our files are empty, anyway.

There are two things that we're really interested in as far as file permissions go. The first describes the permissions currently in force, and the second shows us which user and group owns the file. Let's break it down for the Raspberry file:

```
-rw-r--r-- 1 pi pi 0 Oct 7 16:29 raspberry
```

The file permissions part is:

```
-r--r--r--
```

There are ten possible slots in that list. If a particular permission is missing, `ls` will show a hyphen. However, a normal file always has a hyphen in the first slot. If it were referencing a directory, the first slot would be a `d` to highlight that it's not a file. This slot can also be an `l` if the file is a link (or shortcut) and we'll show you how to use these in the next section. For now though we can ignore the first slot and focus on the final nine slots.

The remaining nine slots are grouped in threes to give us three groups. These correspond to user, group, and world roles, respectively. Each of three slots in each group represents a specific permission: read, write, and execute. Where the permission is set, you will see a letter, but when the permission is not set, you will get the hyphen. If our `Raspberry` file had all permissions set, it would look like this:

```
-rwxrwxrwx
```

Let's split that out a bit so it's a little easier to read:

```
- rwx rwx rwx
```

So if we look at the first three permissions, we can see that the owner has read, write, and execute permissions. We can also see that group and world also have full permissions. To interpret what these permissions mean, though, we really need to know who actually owns the file. Let's take a look at the part of the line that shows who owns the file:

```
pi pi
```

Well, that wasn't too painful. Remember that on modern Linux machines, users have their own private groups that are named after the user. That's what we're seeing here. The first `pi` refers to the owner, which is, of course, the `pi` user. By default, when a file is created, the group ownership is set to the user's default group. In this case, that would be our private group which is also called `pi` (the second `pi`). So look at our original file entry:

```
-rw-r--r-- 1 pi pi 0 Oct 7 16:29 raspberry
```

We can read this as "The `pi` user has read and write privileges. The group has read privileges and world also has read privileges." Linux applies these permissions in a specific order based on who you are.

1. If your username matches the owner of the file, the user permissions will apply when you try to access it.
2. If you're not the owner, but you are in the same group as the file, Linux will apply the group permissions to you.
3. If you're neither the owner or in the same group, Linux will apply the permissions from the world role.

In our example, though, the permissions for world and group are identical, so if you're not the owner, you will get the same level of access. However only the owner can actually save changes to the file. There is one exception to this rule, however: the root user. The root user is effectively immune to file permissions and can change permissions and file ownership for any file on the system, regardless of who the owner actually is.

The execute permission allows you to execute a file as a program. This is a security feature so you can effectively stop people executing commands that you don't want them to. However, you have to be careful because if users can read your file, there is nothing stopping them from copying it to their own file and making that file executable.

The execute bit has another purpose when it comes to directories. Obviously, you can't execute a directory so instead this flag means that the user (or group or world) is allowed to browse the directory—they can do an `ls` on it. They might not be able to access anything in the directory, but they can still have a peek and see what's hiding in there. If you give users permission to read a directory but not to execute, they will be able to read a file inside, but they wouldn't be able to browse for it. They would have to know the filename in advance. This also applies to access files in GUI applications, but it can be annoying typing in a specific path because most applications expect you to browse to the file you want.

That's really all there is to it. There is a feature called "extended file attributes," but we're not going to cover those in this book. They provide a great deal more flexibility than the standard model, but are similarly more complicated. If you're used to the way Windows handles permissions, you'll find that extended file attributes are a bit more in line with what you're used to.

Setting File Permissions

First we're going to look at how we can set file permissions (see Table 4-1).

Table 4-1. Available arguments for the `chmod` command

Part of Permission Setting	Possible Arguments
Role	u - user g - group o - other /world a - all
How to apply	+ - add - - remove = - explicitly set
What to apply	r - read w - write x - execute

We will be using the `chmod` command, which changes file permissions. You can specify permissions as a combination of the preceding values. These can be combined in three different ways.

- Add permissions
- Take permissions away
- Explicitly set permissions

The difference is that the first two will leave all the other permissions intact after they've done their thing. If you explicitly set permissions, any unspecified permissions will be revoked.

Let's start out by removing all permissions from everyone for our `pi` file:

```
[pi@raspberrypi pifun]$ chmod a= pi
[pi@raspberrypi pifun]$ ls -lh
----- 1 pi pi 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

All file permissions have been removed from the file, but how does the command actually work? Permissions are specified with three parts (corresponding to the three rows in Table 4-1):

- Who you want the change to apply to
- How you want the change applied
- What you want the change to be

In this case we applied the changes to a, which is basically shorthand for ugo (it applies the changes to everyone). We used the equals sign, which means we want to explicitly set the permissions and then we didn't actually supply any permissions. If a permission is absent, it is assumed not to be set, so in our example by not supplying any permissions, we effectively revoked them all, regardless of what they were previously.

Seeing as it is our file, we want to give ourselves full permissions. Admittedly, the execute bit is not much use in this case (you'll find it invaluable when you start scripting—see Chapter 7), but we're going to give it to ourselves anyway. We can do that with this command:

```
[pi@raspberrypi pifun]$ chmod u+rwx pi
[pi@raspberrypi pifun]$ ls -lh
-rwx----- 1 pi pi 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

Let's pick this command apart. We specified that we wanted to change only the user's permissions, that we wanted to add them (not that it mattered in this case because we'd removed all permissions beforehand so an equals sign would have done the same job) and that we wanted read, write, and execute privileges. To wrap up this example, let's restore read access to the group and world roles:

```
[pi@raspberrypi pifun]$ chmod go+r pi
[pi@raspberrypi pifun]$ ls -lh
-rwxr--r-- 1 pi pi 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

Just for completeness, let's step through this last example. We want to apply the permissions to the group and other roles, we want to add the permissions to what is already there, and we want to grant read privileges. And that is pretty much it for setting file permissions. There is, however, an alternative style that uses numbers rather than letters to specify what permissions you want to set. To get the same effect as what we have already (i.e., it doesn't have any effect), we would use this:

```
[pi@raspberrypi pifun]$ chmod 744 pi
[pi@raspberrypi pifun]$ ls -lh
-rwxr--r-- 1 pi pi 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

In this system, each permission has its own value:

- Read is 4
- Write is 2
- Execute is 1

To set the permissions, you add up the numbers to get the total for each set of permissions (user, group, world). For example, to set all permissions you'd add 4 and 2 and 1 to get 7. For read, you would simply do 4 plus 0 plus 0, which of course gives you 4. Put them all together for the user, the group and world we get 744. This syntax is the original syntax used on most UNIX systems. Using the letters is a relatively new idea, but at the end of the day they

both achieve the same results. The main benefit of the new syntax is that it's a lot clearer and easier to pick up. Personally, we tend to use the number style, but that's only because we've been doing it for so long and it's become second nature to us. You should feel free to use whatever system you feel the most comfortable with.

So now you can manipulate permissions like a guru, but we are still missing the second part of the puzzle; we haven't shown you how to change ownership of the file. This is actually a lot less common than you might think; far less common than tweaking the occasional file permission, that's for sure. There's also another little wrinkle. A normal user (anyone other than root) cannot actually change which user owns the file because if you accidentally assign the file to another user by accident, you have no way to actually get that file back again. Of course, being all-knowing and all-seeing, the root user can change the ownership for any file on the system.

We can simulate "rootness" by using sudo. As discussed earlier, this little command acts as a filter of sorts. It always runs as root, regardless of who executes it, and executes commands as root on their behalf. To prevent any shenanigans, sudo will check the user and the command they are trying to run against an approved list. If you're on that list (and the pi user is), you can execute all sorts of magic without every technically becoming root yourself.

To use sudo, all we have to do is prefix the command we want to run with the sudo command. That's pretty much it. When you first run sudo, it will ask you for a password. This is the password for your particular user, not the password for the root user. The aim is that you can prove that you are the pi user, and then sudo will check to see what the pi user is allowed to do. This means that if you have lots of users on your computer, and you want to let them do some more powerful commands but don't want to give them root access, you can set up sudo to allow them to execute a specific command without having to hand over the keys to the mansion.

Let's start off by trying to give the file to the root user using the chown (or change ownership) command:

```
[pi@raspberrypi pifun]$ chown root pi
chown: changing ownership of 'pi': Operation not permitted
[pi@raspberrypi pifun]$
```

The operation not permitted is Linux's way of telling us to get stuffed. As mentioned earlier, only the root user is allowed to change the owner of a file. To pull this off, we'll need root privileges, so let's put sudo to work for us and run the command again:

```
[pi@raspberrypi pifun]$ sudo chown root pi
[pi@raspberrypi pifun]$ ls -lh
-rwxr--r-- 1 root pi 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

Success! We were able to change the owner to the root user. This will work with any valid user and any file or directory that you wish to change. There is another command called chgrp which you won't be surprised to know allows you to change which group owns a particular file. Now there is a bit of an issue with this command as well. Although normal users are allowed to change the group, they are only allowed to change it to a group of which they're a member. If your user is only a member of its private group, then you won't be able to do an awful lot with this command either.

Once again, it's root and sudo to the rescue. Because root can do whatever it pleases, it can change the group accordingly. As it happens, it looks an awful lot like our last command:

```
[pi@raspberrypi pifun]$ sudo chgrp root pi
[pi@raspberrypi pifun]$ ls -lh
-rwxr--r-- 1 root root 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

And there we go—the file now belongs to the root group and the root user. When you do have to change file ownership, it's much more common to need to change both the user and the group that owns the file. It's relatively rare to change just the group (we can't remember when we last used the chgrp command). The chown command

provides a shortcut that allows us to set both a new owner and a new group at the same time. Let's use this shortcut now to return the ownership of the file to our pi user. We'll still need to use sudo, of course:

```
[pi@raspberrypi pifun]$ sudo chown pi:pi pi
[pi@raspberrypi pifun]$ ls -lh
-rwxr--r-- 1 pi pi 0 Oct  8 03:52 pi
[pi@raspberrypi pifun]$
```

With the shortcut, you just specify the user and the group separated by a colon. One last thing we need to cover with these commands is that they only operate specifically on the file you provide. If you provide a directory rather than a file, it will set the permissions on the directory, but those changes won't filter down through to all the files. Sometimes that's what you want, but more often you want the changes to propagate. Unlike the cp and rm commands that use -r, these two commands use -R (they use the capital letter rather than the lowercase letter). Be careful when you use this because often file permissions are precisely set, and if you waltz through obliterating them with your new version, there's no way to undo the damage. As always, double-check what you've typed before you press the enter key.

Shortcuts and Links

Linux allows you to create links (or shortcuts) by using the ln command (short for link). There are two types of links:

- A *soft link* is like what you might see on a Windows system after using the Create Shortcut feature. It creates a file that is just a pointer to the real location of the file elsewhere on disk. You create a soft link by specifying -s when using the ln command.
- The *hard link* is more interesting. When you use a hard link you have effectively created two names for the same file. That might sound like semantics, and with most modern applications being able to follow a soft link, there's rarely a need to use a hard link. Hard links are also restricted to a single filesystem that has to support them (most Linux filesystems do). The main benefit of a hard link is that the hard link is completely indistinguishable from the original file; they are simply two names pointing to the same location. You don't need to specify anything to create a hard link as ln will do so by default.

Note To avoid confusion and to allow links to work across filesystems, you should use a soft link.

Let's do a quick example to show this in action:

```
[pi@raspberrypi pifun]$ ln pi pi1
[pi@raspberrypi pifun]$ ln -s pi pi2
[pi@raspberrypi pifun]$ ls -lh
-rw-rw-r-- 2 miggyx miggyx    0 Oct  8 08:14 pi
-rw-rw-r-- 2 miggyx miggyx    0 Oct  8 08:14 pi1
lrwxrwxrwx 1 miggyx miggyx    3 Oct  8 08:33 pi2 -> pi
[pi@raspberrypi pifun]$
```

Let's have a look at what we've got here. pi and pi1 are identical in every way, but that's not really a surprise because apart from the name, they are the same file. Notice that the number after the file permissions block now shows 2 for pi and pi1. This tells us that there are currently two filenames pointing at this particular file. Also not much of a surprise because we're the ones that created the second entry.

Much more interesting is `pi2`, which we created with a soft link. First we can see that the file permissions have all been set. This isn't a problem because when Linux follows the soft link to the real file, it's the real file's permissions that will be used to define who can access the file. The soft link really just points out the location. We can also see that the filename itself is a bit different. It shows the filename that we originally give the soft link, but it also shows the file that the soft link points to. In this case, the file happens to be in the same directory, but it could just as easily have been anywhere on the system.

That's really all there is to it for creating links. They can be useful when you want to make one directory or file appear to be in a new location. For example, a program might write to a data directory, and you want to move that directory on to a bigger disk. No problem; you can move it to the bigger disk and then create a soft link to it with the same name. The application probably won't even notice. This can really save you a lot of headache, especially when time is something of a premium. (And let's be honest: when is it not?)

Summary

This chapter has given you the inside scoop on all things filesystem. We've looked at the history and shown why our filesystems look the way they do. We then touched on how they hang together and how the Linux filesystem is structured. We put that to good use and brought you up to speed on all the basics for creating, copying, moving, and deleting your files.

We discussed file permissions and how they are enforced, and how we can set them to match our needs. We also looked at the more traditional way of setting file permissions (in case it is ever needed). We then showed you how Linux applies these permissions and how you can change which user and group owns a specific file. We rounded everything off by discussing how you can create links and the differences between both soft and hard varieties.

In the next chapter, we're going to expose you to all the most common commands that you'll find on your Pi. These are the commands that will become part of your toolbox that you'll regularly dip into. In fact, we use all these commands in our daily work. So, onward to Chapter 5!

CHAPTER 5



Essential Commands

You now have this marvelous new Raspberry Pi that you have installed and configured to accept and run a version of Raspbian Linux. You have logged in and can comfortably move around within the system. You can do the basics on your Pi, but now feel that you are ready to take on more.

Well that's what this chapter is dedicated to: teaching you a whole host of commands so that you can empower yourself as a Linux user. These commands will cover a whole host of functionality ranging from becoming the root super user, examining disk space, terminating processes, and configuring your user space and monitoring environment variables. All these tasks are common ones for the systems administrator and can have a big impact on how your system functions. So without further ado, let's get started.

Become the Boss

Root is a word that describes the base of a tree, something from which all the rest grows. In Linux and UNIX, the term is used to denote the start point of a filesystem (as we saw in Chapter 4) and to denote the master administrative user on the system (covered briefly in Chapter 3). To perform any system or administrative function, you need to be identified as root in order to have permission to change core functionality. Because of this, being root is very much a two-edged sword: unlimited power to change your system and unlimited power to break system functionality. Probably one of the truest things I can say about working in a Linux environment is that eventually you will issue a command as root that will break something (if not everything) and this will cause you no end of hassle or may even require you to reinstall your operating system. For this reason (and related security reasons) its advised that a user never actually login to the shell as the root user.

To solve this problem, we use `sudo`, which we've seen a little of in Chapters 3 and 4. As you know, this tool is a wrapper for any shell command you can think of. This means that if you prefix a command with `sudo` (e.g., `sudo ls`), you will no longer be executing this command as your own user but rather as the root super user. This does make working as the super user slightly more complex, but also far less risky. In fact, many distributions of Linux no longer set a root password, but rather assign a new user the power of `sudo` on creation. This means that all administration must be done via an empowered user with the `sudo` command to issue all the requisite admin functions.

So that you can execute as root, `sudo` requests your user password in order to validate that you are who you say you are, not just someone with access who left their keyboard unattended. In addition to asking all users for their passwords, `sudo` also has a steering file that governs which users and groups have access to the `sudo` command and even the capability to limit which programs they can execute via `sudo`. There is a special command, `visudo`, to edit this steering file (called the `sudoers` file) that must be run as root (so prefix with `sudo`):

```
$ sudo visudo
```

This code fires up a special instance of one of the system's text editors (normally `vi` or `nano`) with the `sudoers` file already loaded to be edited. For a quick runthrough of how to use `vi` and `nano`, please check out Chapter 6.

Inside the sudoers file are a number of things that control how the sudo environment works, including a PATH variable (which we will cover later). The important thing, however, is the line down in the userauth section that looks like this:

```
root    ALL=(ALL:ALL) ALL
```

This line states that the root user on all hosts can access all commands on all other hosts (a *host* is a server on the network). This is a bit confusing, but it basically grants full access to any user that can identify itself as root. There is also another line:

```
%sudo  ALL=(ALL:ALL) ALL
```

This line says that members of the sudo group can also do the same things as the root user: accessing from all hosts on the network all commands that can be executed on all hosts on the network. You will notice the leading % that is used to signify that this is a group name. The basic format of one of these lines is this:

```
<user> <From which hosts>=<(On which hosts):<which commands>> < command options and commands >
```

This syntax is a bit confusing at first so let's have a look at creating a particular set of permissions for your pi user as an example. First we will need to create a new basic set of user permissions that will look just like the root users, but will instead refer to our pi user. This line will look like:

```
pi    ALL=ALL    ALL
```

I have removed the <on which hosts: which commands> section and replaced it with one generic ALL to make things a bit easier to understand. Now let's say that we want the pi user not to be asked for a password. To do this, we add a NOPASSWD: before the final ALL that represents commands. This gives us a new line:

```
pi    ALL=ALL    NOPASSWD: ALL
```

If you add this now, you can run all commands as the pi user without needing to even provide a password. This is a very dangerous command, so let's modify it slightly to make it only the visudo command that can be executed. To do this, we replace the final ALL for commands with the command we wish to use, which in this case is visudo. But we can't just write visudo, we need to provide the whole path to the application, which is /usr/sbin/visudo. So we get the following:

```
pi    ALL=ALL    NOPASSWD: /usr/sbin/visudo
```

Now you should be getting the idea of what you can do within the sudoers file to control how the sudo command is accessed. You can have multiple entries that govern how these commands can be run, by whom and under which conditions. There is also a large amount of functionality for adding aliases, which can be large combinations of commands, hosts, or users. Much of the syntax for these commands is available from the manual for sudoers.

Finally, one last command for you to be aware of is the su command, which is short for switch user. This command does exactly what it implies and switches which user you are, and if you aren't root it will prompt you to enter the password of the user you wish to switch to. One of the niftiest cheats for gaining root access is to combine

sudo and su to switch into the root user, with only the current user's password instead of the root user's. This command is the way that many people get round the "no known root password" limitation placed by many Linux distributions:

```
$ sudo su -
```

The minus sign denotes that you wish to gain a login to the environment, and because no other arguments are provided, it will attempt to log in as the super user. You can also use this to switch users by adding the username for a given user prior to the minus. You can also use the sudo command with the `-i` option to gain an interactive shell that will achieve the same thing:

```
$ sudo -i
```

RTFM

One of my favorite refrains from the tech world is *RTFM*, which is short for *Read The Flaming Manual*. While I have found that asking people for assistance is good, it's always comforting to know that there is a manual you can read, and Linux is no exception. Most applications come with a manual document that can be accessed by using the `man` command! This command takes whatever you wish to get the manual page for as its first argument. So, if you wanted to look at the manual page for the `sudoers` command in order to get a better idea of how to write your `sudoers` file, you would type this:

```
$ man sudoers
```

And then the manual page would appear. Navigating a man page is done in the same way one would navigate a document opened with `less` (which is a command we will cover further on in this chapter), so navigation is done with the arrows and with page up and page down keys. You can also search for a string by pressing `/` and then typing what you wish to search for. The `n` key will help you navigate your way through the file. So, if you ever find yourself in need of some more information on a command, remember to Read The Flaming Man Page!

System Resource Monitoring

One of the key things that a systems administrator will need to do is monitor the usage of resources on his system. While some of these commands may be of less use to your average Pi user, there are some times when being able to see what your system is doing will come in handy.

The first system command to be aware of is `top`, which is a command for displaying the current usage of resources within your system and displaying which processes are consuming what level of resources. The `top` command is incredibly powerful and displays a whole host of information about your systems resource usage. The output of `top` normally looks like Figure 5-1.

top - 22:22:11 up 6 days, 20 min, 1 user, load average: 0.11, 0.08, 0.06											
Tasks: 66 total, 1 running, 64 sleeping, 1 stopped, 0 zombie											
%Cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 99.9 id, 0.0 wa, 0.0 hi, 0.1 si, 0.0 st											
KiB Mem: 123156 total, 109784 used, 13372 free, 30340 buffers											
KiB Swap: 0 total, 0 used, 0 free, 41608 cached											
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
7445	raspbian	20	0	4608	1312	960	R	16.0	1.1	0:00.08	top
1	root	20	0	1832	680	592	S	0.0	0.6	0:22.11	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.45	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	ksoftirqd/0
5	root	20	0	0	0	0	S	0.0	0.0	0:59.85	kworker/u:0
6	root	-2	0	0	0	0	S	0.0	0.0	0:03.20	rcu_kthread
7	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	khelper
8	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kdevtmpfs
9	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	netns
10	root	20	0	0	0	0	S	0.0	0.0	0:04.19	sync_supers
11	root	20	0	0	0	0	S	0.0	0.0	0:00.18	bdi-default
12	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kblockd
13	root	20	0	0	0	0	S	0.0	0.0	0:00.28	khubd
14	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	rpciod
16	root	20	0	0	0	0	S	0.0	0.0	0:00.42	khungtaskd
17	root	20	0	0	0	0	S	0.0	0.0	0:00.51	kswapd0
18	root	20	0	0	0	0	S	0.0	0.0	0:00.00	fsnotify_mark
19	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	nfsiod
20	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	crypto
27	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kthrotld
28	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	VCHIQ-0
29	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	VCHIQr-0
30	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	dwc_otg
31	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	DWC Notificatio
32	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kworker/u:1
33	root	20	0	0	0	0	S	0.0	0.0	2:11.51	mmcqd/0
34	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	ext4-dio-unwrit
125	root	20	0	2904	1324	728	S	0.0	1.1	0:00.65	udevd
190	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	bcm2708_spi.0
951	root	20	0	0	0	0	S	0.0	0.0	1:53.66	RTKTHREAD
956	root	20	0	5588	1204	848	S	0.0	1.0	0:19.18	wpa_supplicant
983	root	20	0	6332	2148	396	S	0.0	1.7	0:00.01	dhclient
1102	root	20	0	2304	796	584	S	0.0	0.6	0:02.77	rpcbind
1133	statd	20	0	4112	1728	1148	S	0.0	1.4	0:00.06	rpc.statd
1145	root	20	0	2736	444	264	S	0.0	0.4	0:00.00	rpc.idmapd
1452	root	20	0	27928	1452	1008	S	0.0	1.2	0:03.53	rsyslogd
1482	daemon	20	0	2084	324	200	S	0.0	0.3	0:00.02	atd
1524	root	20	0	3744	784	628	S	0.0	0.6	0:02.29	cron

Figure 5-1. Top in action!

The output from the top command is huge and if you watch, all the values there are updated regularly. Let's go over the top output line by line so you can understand what you are looking at.

Uptime and Load Average

The first line gives information about the current time and then uptime of the system. It then shows the number of users on the system and finally the load average. *Load average* is one of the most complex things to understand; it is a “meta” variable because it refers to what is going on within your system relative to your system’s capability to perform.

It contains 3 values that give the average of the current minute, the last 5 minutes and the last 15 minutes. Load average represents the average of actual computation usage and demand for your system's computation resources.

Probably the best way to think of load average I have heard is as a freeway. Normally the percentage value of CPU usage represents how much of the current available computation resources are being used. This is akin to measuring how much traffic is currently on the freeway. The load average compares both how much computation is being used and how much demand there is for CPU usage. This is like examining both the freeway and its onramps. The difference between the two comes when looking at it from the perspective of a "busy lunchtime," meaning that there are lots of cars on the freeway, but just because "it's a busy lunchtime," everyone can get onto the freeway and use it, nobody is stuck waiting on the ramp, and traffic is flowing nicely.

This is why load average is such a powerful metric; it allows you to see demand for your system's computing power in real time so that you can see when your demand begins to spike and can thus take action to lower demand.

Wow, that's a lot on just one line!

Tasks

The second line of the `top` output lists the tasks that the system is currently performing:

- The total number of processes
- The number of actively running processes
- The number of processes that are sleeping in the background
- The number of stopped processes
- The number of zombie processes

Most of these processes are general purpose for being aware of what is going on within your system. But zombie processes can be a real problem. These processes are processes that have finished executing, but are still resident in system memory—often because they need to send a return value to the process that spawned them, but are yet to do so and thus must wait there as the living dead before they can send their final message and depart to the afterlife.

Having a long running zombie process can indicate a problem with the application that is used to spawn it, which is called its parent. If you see zombie processes on your system have a look to see which process spawned them and check over its running state and log files (log files are output from applications that are written to disk, so you as the systems administrator can see what they are doing), as this may be indicative of a larger issue.

CPU Utilization Percentage

The third line of the `top` command is the CPU utilization percentage value. This shows in terms of percentage how much of your current computation resource is allocated to which domain of your system. The values are the following:

- us (for user applications)
- sy (for system applications)
- ni (for processes that have been "niced" in order to give them more or less CPU priority)
- id (representing idle)
- wa (for processes that are waiting for I/O completion)
- hi (for processes waiting for hardware interrupts)
- si (for processes waiting for software interrupts)
- st (which indicates time stolen by a hypervisor—a *hypervisor* is software that runs a virtualization platform like Xen or VMWare ESX)

Of these values, the one that will normally have the most computation resource allocated to it is `id`, which represents free available computation that can be used elsewhere. Seeing high values of `wa`, `hi`, `si`, or `st` are never good signs as they suggest that your system is waiting on a particular piece of hardware or another function to finish processing. If these numbers stay high, you should look at investigating which processes cause them and see if you have any hardware issues.

Memory Usage

Lines 4 and 5 represent the memory usage of your system. Line 4 is for the actual memory usage, and line 5 is for the swap space usage. Both of these lines will show the total, used, and free memory in their given space. This final value represents the buffers and cached value. These two are interrelated and are the cause of much concern for users new to Linux. This has to do with the way that RAM is allocated. Linux can sometimes take up all available RAM and place it into a buffers, which means that it is taken by the buffers but can still be used later for your applications. Linux will treat this RAM as being used, but it is just sitting there in the buffers waiting. The buffers value can be a little hard to understand; this leads to another command showing how these resources are used (I will cover it next).

Process Table

After the fifth line there is a large section that contains a process table, which lists all running processes and, in near-real-time updates, which processes are consuming how much of which resource. The following shows a number of different columns and the pieces of information they represent:

- *PID*: The processes' ID number
- *USER*: The user who owns the process
- *PR*: The priority of the process
- *NI*: The nice value for the process
- *VIRT*: The amount of virtual memory consumed by the process
- *RES*: The size of actual resident virtual memory
- *SHR*: The amount of shared memory the process is using
- *S*: The process status (e.g., sleeping, running, zombie)
- *%CPU*: The percentage of CPU consumed
- *%MEM*: The percentage of RAM consumed
- *TIME+*: The total CPU time the task has used since it started
- *COMMAND*: The actual name of the command

`top` also has a host of keypress actions that can manipulate the order in which these processes are displayed; you can check them out in the man page for `top`.

So `top` is a fantastic command for showing all the ins and outs of your system's resource usage and one of the best tools for diving into the exact nature of what your system is doing within itself.

Using free to View Memory Allocation

As mentioned earlier, there is an easier way than using top to see what your memory allocation is like. The command to do this is the free command, which you can invoke with the unit you wish to have it display in (-k for kilobytes, -m for megabytes, -g for gigabytes, and so on). For our purposes, megabytes is best, so invoke it as follows:

```
$ free -m
```

Your output should look like Figure 5-2.

	total	used	free	shared	buffers	cached
Mem:	231	157	74	0	18	79
-/+ buffers/cache:		59	172			
Swap:	99	0	99			

Figure 5-2. Free as in free

From this, you can see the +/- buffers and cache line that change the used and free values. This is a great way to question any time your system shows spikes of memory up to 100%, as it just may be lots of buffers and cache in action. From these two utilities you should be able to gain great insight into exactly what is going on within your system and how its resources are being used.

Disk Usage

Now that we have looked at CPU and RAM usage, which are fairly dynamic, it's time to have a look at something of a more static nature, disk utilization. In most operating systems you can quickly and readily view how much of the available storage space is in use, and Linux is no exception. The command to show the current usage of the whole system's disks is df, which is short for *disk free* and can be used on the command line with no arguments to show everything you need. If you do, your output will resemble Figure 5-3.

Filesystem	1K-blocks	Used	Available	Use%	Mounted on
rootfs	1828696	1757188	0	100%	/
/dev/root	1828696	1757188	0	100%	/
tmpfs	23688	504	23184	3%	/run
tmpfs	5120	0	5120	0%	/run/lock
tmpfs	47376	4	47372	1%	/tmp
tmpfs	10240	0	10240	0%	/dev
tmpfs	47376	0	47376	0%	/run/shm
/dev/nmcblk0p1	57288	34784	22504	61%	/boot

Figure 5-3. Disk free, ouch!

Figure 5-3 shows disk free in action and shows that the main rootfs mounted on / is basically 100% full—ouch. You can see the number of 1 k blocks being used by the system, but the values don't stand out very easily. You will also note that there is a permission denied error at the top.

Fortunately, there are solutions for both these issues. First is to run this command with sudo so that it has root access to the whole system. The second is to add the -h argument. This argument is fairly universal in being one of two things: either -h for *help* or -h for *human*, which is to say human readable units. So we run this command:

```
$ sudo df -h
```

The output should be much more understandable to the average human being (see Figure 5-4).

Filesystem	Size	Used	Avail	Use%	Mounted on
rootfs	1.8G	1.7G	0	100%	/
/dev/root	1.8G	1.7G	0	100%	/
tmpfs	24M	504K	23M	3%	/run
tmpfs	5.0M	0	5.0M	0%	/run/lock
tmpfs	47M	4.0K	47M	1%	/tmp
tmpfs	10M	0	10M	0%	/dev
tmpfs	47M	0	47M	0%	/run/shm
/dev/mmcblk0p1	56M	34M	22M	61%	/boot

Figure 5-4. *df* for humans

As you can see, the human readable flag has changed those Size and Used values over to a human readable number with a nice little flag on the end denoting unit. Experimenting with display flags is a great way to help you make sense of some of the complex data coming out of Linux commands. Normally you will find a huge section detailing (in way too much depth) all the flags for a command in its man page.

So, we have seen the full disk utilization of the system, but let's say we want to find the detailed utilization of all the files in a folder. There is another command we can use to do just this: du, which is short for *disk usage* and is used to summarize the disk usage of each file, recursively for directories. If you run du, it will tell you what the estimated usage of each file within this folder is and that of each subfolder and each item within it and its subfolders, and so on. du will also take advantage of the -h for human flag that was in df. So if we run it, we get an output like this:

```
$ du -h
4.0K ./cache/dconf
12K ./cache/menus
576K ./cache/midori/web
4.0K ./cache/midori/thumbnails
24K ./cache/midori/icons
608K ./cache/midori
4.0K ./cache/openbox/sessions
12K ./cache/openbox
640K ./cache
12K ./dbus/session-bus
16K ./dbus
5.0M ./tor
12K ./ssh
1.8M ./python_games
8.0K ./dillo
20K ./vnc
4.0K ./gvfs
124K ./gnupg
```

```

8.0K ./fltk/fltk.org
12K ./fltk
48K ./Desktop
8.0K ./config/lxpanel/LXDE/panels
16K ./config/lxpanel/LXDE
20K ./config/lxpanel
8.0K ./config/pcmanfm/LXDE
12K ./config/pcmanfm
8.0K ./config/lxterminal
28K ./config/midori
4.0K ./config/enchant
28K ./config/openbox
104K ./config
76K ./fontconfig
16K ./netsurf
40K ./local/share/gvfs-metadata
36K ./local/share/webkit/icondatabase
40K ./local/share/webkit
88K ./local/share
92K ./local
9.0M .

```

Wow, that's a big output for such a simple command. It shows the utilization of every folder within the current folder, and right down the bottom it shows the total usage. This is where we can take advantage of another flag in order to shorten the output: the `-s` flag that stands for *summary*. Doing this will only print the total allocation for this directory and all within it, rather than everything as we have just seen. So go ahead and execute the command; your output should look much neater:

```
$ du -sh
9.0M .
```

That's much better. You will notice that the output is using the `.` value to denote the current directory because `du` will work on the current working directory if not told where else to look. You can specify a directory to use by simply adding one to the end. The next example will look at `/var/log`, which is where all system logs are kept, so we will need to run this command as root using `sudo` as some of these logfiles are protected:

```
$ sudo du -sh /var/log
9.8M /var/log
```

And there you have it! With these two utilities, you can find out the current disk usage of your system and even work out which parts of your system are consuming the most space.

Managing Processes

Probably one of the most important tasks when working within a Linux environment is knowing how to manage running processes. After all, a computer is just a processing system, and if you can't manage what it is processing, how can you make truly good use of your machine? The first command you should become aware of is the `ps` command. The `ps` command is short for *processes snapshot* and can be used to list all the current processes being run by your user; when invoked, you should get an output like the one shown in Figure 5-5.

PID	TTY	TIME	CMD
1951	pts/0	00:00:01	bash
2162	pts/0	00:00:00	ps

Figure 5-5. Current processes

Yes, it's a little lackluster right now as I am only running two processes: the bash application that is running my user shell and the copy of ps. Since this is a snapshot, the application will always capture itself. This is useful if I want to control those processes run by my current user, but if I am to do more, I will need to add some arguments to my ps command. The argument to show every process on the system in the standard syntax is `-ef`, so go ahead and run that and see how your output changes (see Figure 5-6).

```

root 1756 1 0 17:05 ? 00:00:05 Xtightvnc :1 -desktop X -auth /root/.Xauthority -geometry 1280x800 -depth 24 -rfbwait 120000
root 1760 1 0 17:05 ? 00:00:00 /bin/sh /root/.vnc/xstartup
root 1771 1760 0 17:05 ? 00:00:00 /usr/bin/ck-launch-session /usr/bin/dbus-launch --exit-with-session x-session-manager
root 1805 1 0 17:05 ttys1 00:00:00 /sbin/getty --noclear 38400 ttys1
root 1806 1 0 17:05 ttys2 00:00:00 /sbin/getty 38400 ttys2
root 1807 1 0 17:05 ttys3 00:00:00 /sbin/getty 38400 ttys3
root 1808 1 0 17:05 ttys4 00:00:00 /sbin/getty 38400 ttys4
root 1809 1 0 17:05 ttys5 00:00:00 /sbin/getty 38400 ttys5
root 1810 1 0 17:05 ttys6 00:00:00 /sbin/getty 38400 ttys6
root 1811 1 0 17:05 ? 00:00:00 /sbin/getty -L ttyAMA0 115200 vt100
root 1815 1771 0 17:05 ? 00:00:00 /usr/bin/sshd-agent /usr/bin/ck-launch-session /usr/bin/dbus-launch --exit-with-session x-session-manager
root 1817 1 0 17:05 ? 00:00:00 /usr/sbin/console-kit-daemon --no-daemon
root 1887 1771 0 17:05 ? 00:00:00 /usr/bin/lxsession -s LXDE -e LXDE
root 1890 1 0 17:05 ? 00:00:00 /usr/bin/dbus-launch --exit-with-session x-session-manager
root 1891 1 0 17:05 ? 00:00:00 /usr/bin/dbus-daemon --fork --print-pid 5 --print-address 7 --session
root 1897 1887 0 17:05 ? 00:00:00 openbox --config-file /root/.config/openbox/lxde-rc.xml
root 1899 1887 0 17:05 ? 00:00:12 lxpanel --profile LXDE
root 1900 1887 0 17:05 ? 00:00:01 pcmmanfm --desktop --profile LXDE
root 1906 1 0 17:05 ? 00:00:04 /usr/lib/notification-daemon/notification-daemon
root 1908 1 0 17:05 ? 00:00:01 nm-applet
root 1910 1 0 17:05 ? 00:00:00 /usr/lib/arm-linux-gnueabihf/lxpolkit
root 1914 1 0 17:05 ? 00:00:00 /usr/lib/arm-linux-gnueabihf/libmenu-cache1/libexec/menu-cached
root 1916 1 0 17:05 ? 00:00:00 /usr/lib/gvfs/gvfsd
root 1918 1 0 17:05 ? 00:00:00 /usr/lib/gvfs/gvfs-fuse-daemon -f /root/.gvfs
root 1919 1724 0 17:05 ? 00:00:01 sshd: pi [priv]
root 1933 1 0 17:05 ? 00:00:00 /usr/lib/arm-linux-gnueabihf/gconf/gconfd-2
root 1937 1 0 17:05 ? 00:00:00 /usr/lib/gvfs/gvfs-gdu-volume-monitor
root 1939 1 0 17:05 ? 00:00:00 /usr/lib/udisks/udisks-daemon
root 1940 1939 0 17:05 ? 00:00:00 udisks-daemon: not polling any devices
root 1943 1 0 17:05 ? 00:00:00 /usr/lib/gvfs/gvfs-ghotoc2-volume-monitor
root 1945 1 0 17:05 ? 00:00:00 /usr/lib/gvfs/gvfs-afc-volume-monitor
pi 1950 1919 0 17:05 ? 00:00:03 sshd: pi@pts/0
pi 1951 1950 0 17:05 pts/0 00:00:01 -bash
root 1980 1724 0 17:05 ? 00:00:00 sshd: pi [priv]
pi 1987 1980 0 17:05 ? 00:00:00 sshd: pi@pts/1
pi 1988 1987 0 17:05 pts/1 00:00:00 -bash
root 2094 1724 0 17:48 ? 00:00:00 sshd: pi [priv]
pi 2102 2094 0 17:48 ? 00:00:00 sshd: pi@pts/2
pi 2103 2102 0 17:48 pts/2 00:00:00 -bash
root 2137 2 0 18:02 ? 00:00:00 [flush-179:0]
root 2138 2 0 18:08 ? 00:00:00 [kworker/0:1]
root 2189 2 0 18:18 ? 00:00:00 [kworker/0:0]
root 2208 2 0 18:25 ? 00:00:00 [kworker/0:2]
pi 2209 1951 0 18:26 pts/0 00:00:00 ps -ef

```

Figure 5-6. ps -ef

Wow, that's a lot of output, so much that it went off the top of the screen! There are always a large number of processes running on your system when you run `ps -ef` as this lists all the systems background processes along with all user processes. If you want to see them in a list, you can simply add a `| less` to the end of your command, which will let you go up and down and manipulate the list (I will explain what the `less` command actually does a little later in this chapter):

```
$ ps -ef | less
```

`ps` is a lot like `top`, but does require fewer resources and has a few other advantages that we will cover later in this chapter. For now, you should just be aware of what each of the columns in `ps` means:

- *UID*: Represents the user who owns this process
- *PID*: Represents the process's PID number
- *PPID*: Represents the PID number of the process that created this one
- *C*: Represents CPU utilization of this process
- *STIME*: Represents the start time of this process
- *TTY*: Represents the terminal controlling this process
- *TIME*: The cumulative CPU time consumed by this process
- *CMD*: The command-line argument of this process

So you might be saying, “Well fine. You have shown all the processes in the system, but what can I do with them now?” This is where the magic of `ps` comes in; it makes it very easy for us to pick out the PID of a particular process. The PID number is very important as it is a unique number for each process and is the way that the system references processes. This will come of use in our next command: `kill`.

Killing a Process

`kill` is used to do exactly what its name implies: kill processes. But it also provides a number of other useful functions via the mechanism through which it kills processes. In Linux and UNIX, a *signal* is a low-level system message that is passed into an application from the operating system. Signals have no actual function beyond being a signal, and it is up to the application to respond to it in the correct manner. This is not true for all signals, as some cannot be handled within the application, but it is true for most. The function of the `kill` command is to send signals to applications. There are 64 signals that the version of `kill` on Raspbian can send to an application. To list them all, issue this command:

```
$ kill -l
```

I won't list all the functions of the signals now, but you will need to be aware of these, the most common signals:

- *Signal 1 (SIGHUP)*: Short for *hang-up*. Tells the application to hang up its current connection. Normally used to make an application reinitialize itself.
- *Signal 3 (SIGQUIT)*: Signals that this application should gracefully close itself down and quit.
- *Signal 6 (SIGABRT)*: Signals that the program is aborting; it is to close immediately.
- *Signal 9 (SIGKILL)*: Signals to forcibly “pull the plug” on an application.

These signals are the most common as they are used to perform application shutdowns, with nine being the most prevalent. The format of the `kill` command is this:

```
kill <signal> <PID>
```

So to kill a process with the highest force, execute this:

```
$ kill -9 <PID>
```

This is the point where those PIDs out of the `ps` command come into play. Those PID numbers are the ones that you will feed into `kill` in order to direct the signals from `kill`.

Reading Information in /proc

Since we are playing with PIDs, it's probably good to revisit the `/proc` filesystem, first discussed in Chapter 4. This is part of the OS in which all the process information lives. If you use `ls /proc`, you will see a number of files and a set of numbered directories. These directories correspond to each of the process PIDs. If you look within one of these directories, you should see a whole host of files that relate to a given process. While you should never edit any of these willy-nilly, there are some files within these directories that are very useful.

Probably the most commonly used is a system-wide file that contains information on all the system's processors. This is the `/proc/cpuinfo` file. You can output the contents of a file to the screen using the `cat` command, which is short for *concatenate* and is normally used to get the contents of a file into a buffer (in our case, it is the console screen). Go ahead and issue this:

```
$ cat /proc/cpuinfo
```

Your output should look like this:

```
Processor      : ARMv6-compatible processor rev 7 (v6l)
BogoMIPS       : 697.95
Features       : swp half thumb fastmult vfp edsp java tls
CPU implementer: 0x41
CPU architecture: 7
CPU variant    : 0x0
CPU part       : 0xb76
CPU revision   : 7

Hardware       : BCM2708
Revision       : 0003
Serial         : 00000007a8a46ba
```

This is the Linux internal information on the Raspberry Pi's processor. If you are ever in doubt on some of the particulars of your processor, this is the place to look.

In addition to `cpuinfo`, there is another file within `/proc` that you can take advantage of when diagnosing system configuration information. It is the `/proc/<pid>/cmdline` file. This file lists the full command-line argument used to invoke a process. If you need to get your hands on the commands used here, executing `cat` against this file is the best way to do it.

File Commands

As mentioned earlier, you can use the `cat` command to output the contents of a file. There are a number of other tools you can use to work with the contents of files within your system. The first is one that was mentioned earlier: the `less` command. It is a command to display the contents of a file to the screen in a way that can be moved around and searched as if it were a text editor. This is a great command for working with logfiles or the output of long commands such as `ps`. `less` allows you to move up and down the file with the arrow keys or with the page up and page down keys. You can also search using simple word searches by typing `/` and then entering your search string. `less` will then highlight the found instances of your expressions, and you can tab between them with the `n` key to go down and the `N` key to go up.

In addition to being able to search within a file for things, `less` has the ability to watch the file and constantly display any new content on the last line. This is great for watching growing logfiles as you can pick up new lines

without needing to reopen the file, to get less to perform this function, simply press the F key. Along with this, you can also automatically navigate to the end of a file by pressing the G key. This process will also grab any new lines that have been added to the file and display them. You can also use the lowercase g to navigate to a specific line number, so if you wish to go to the first line, would press 1g. Finally, less is able to respond to all of its command-line arguments within itself. This is useful for enabling line numbers as you can simply give less the command -N<enter> and it will display numbers next to each line.

After less, one of the commands that provide a startling amount of flexibility is the `find` command, which is used to find files. It does this by examining every file within the given directory and all of its subdirectories. `find` is useful on its own, but extraordinary when paired with other commands as you will see later. It also has a healthy cache of options that are interchangeable and can also be combined. I will cover a few that I have found incredibly useful here:

- The first is `-L`, which says to follow symlinks, which `find` does not do by default.
- Second is `-maxdepth <number>`, which represents how many directories deep this command will search through. Just the one? Just this and all its subdirectories?
- Third is `-newer <file>`, which shows anything that was changed more recently than a given file.
- Fourth is `-empty` and it, as its name implies, finds empty files.
- Fifth is `-atime <number>`, where the number represents the number of days that have passed since a file was accessed.
- Sixth is `-name <filename>`, which searches for files named exactly like the name passed in.
- Last is `exec <command>`, which tells `find` to execute a given command on each of the files and is where the real power of `find` comes out. You can use this command to find specified files and then execute commands on them. In this way, using `find` is a fantastic way of executing a cleanup of older unwanted files for instance. I'd also advise that when using `find` in this manner you should test first and test thoroughly because one incorrect key stroke could spell doom for your system or your data.

Combining these commands, you can achieve something that looks like this:

```
$ find /mnt/Volume1 -empty -name fooBar -exec rm
```

This command will find all empty files within `/mnt/Volume1` that are empty and named `fooBar` and then remove them! There is so much power in the `find` command that you can harness when administering your Pi.

Another command that has a significant if limited use is the `file` command. The `file` command is used to go through the contents of a given file and determine what “kind” of file it is. *Kind* is a bit of a loose term here because `file` knows only a moderate number of different file types. However, it is fantastically useful when you want to know an answer to this question: “Can I safely open that file in a text editor, or is it a binary application?” This particular command can save you much heartache. You can also use `file` to determine which archival format an application is written in, which is a great help when you need to extract from an archive but don’t know what kind! As an example, if you run `file` against a binary executable such as `/bin/bash`, you should see the following:

```
$ file /bin/bash
/bin/bash: ELF 32-bit LSB executable, ARM, version 1 (SYSV), dynamically linked (uses shared libs),
for GNU/Linux 2.6.26, BuildID[sha1]=0xe370a0be978f49569f115c9f41bc54688be9fe32, stripped
```

The code shows that the file is an ELF 32-bit LSB executable compiled for an ARM system, so the file should be compatible with your Pi. This is a good way to diagnose any issues you may have with applications downloaded from the Internet because they may be compiled for x86!

The last file command I have for you in this section is one of those commands that many people hate and many more love: grep. grep is short for *Generalized Regular Expression Parser* and is best known as being a pattern matcher. Like less and find, grep is both useful on its own and useful when combined with other commands. You can use grep to search a given file for any lines that contain a given text string. The basic format of grep is this:

```
grep <pattern> <file>
```

Both these values will take the wildcard format, so you can use grep to search through every file in an entire directory for a string. As an example, let's look for the word *find* in the */etc./init.d* folder:

```
$ grep find /etc./init.d/*
```

The output would be this:

```
/etc./init.d/checkfs.sh:# Include /usr/bin in path to find on_ac_power if /usr/ is on the root
/etc./init.d/checkroot.sh:# Include /usr/bin in path to find on_ac_power if /usr/ is on the root
/etc./init.d/kbd:           LIST_CONSOLES='cd /etc./init; find -name 'tty*.conf' -printf '%f '
| sed -e 's/[^0-9 ]//g'
/etc./init.d/kbd:           # eventually find an associated SFM
/etc./init.d/nfs-common:# particularly heavy daemon, so we auto-enable it if we find an
/etc./exports
/etc./init.d/rc:# Now find out what the current and what the previous runlevel are.
/etc./init.d/sudo:          find /var/lib/sudo -exec touch -t 198501010000 '{}' \;
```

grep also has some wonderful command-line options, including *-i*, which makes the search case insensitive. grep also has the *-v* option, which turns a search on its head and returns everything *not* matching the given search string! This should give you an idea of how powerful a tool grep can be when you wish to find a given piece of information and thus why it is an awesome part for a systems administrator's arsenal of commands.

Combining Commands

Now we get to the fun part! The ability to combine commands so that you can use the output of one command and give it to a second command to work with. There are also a number of commands that I will detail here as they are fantastic when combined with others. But first you need to understand how to combine commands and thankfully it is a rather easy thing to accomplish. To combine commands, we use the pipe symbol (*|*).

Yep, it's that simple. Let's say we wanted to grep some things, but wanted the output in a more manageable way. We would simply write the grep, add a pipe, and then add a less. By doing this you will find yourself inside less looking at the output for your grep! This is where the fun truly begins. Want to search all the filenames in a given directory? Use *ls* to list the files and pipe the output to grep to search. In this instance, you don't even need to include a filename with the grep as the pipe will provide grep with what is known as a *buffer* that grep will treat the same as file contents! So the command is this:

```
$ ls <directory> | grep <search string>
```

Combine ps and grep

Now, if you want to play more with things, you can combine grep with *ps -ef!* In this way, you can search for a specific application; for example, to search for the SSH application *sshd*, we would run this:

```
$ ps -ef | grep sshd
```

This will generate the following output:

```
root      1722      1  0 19:24 ?          00:00:00 /usr/sbin/sshd
root      1956  1722  0 19:43 ?          00:00:00 sshd: pi [priv]
pi        1963  1956  0 19:43 ?          00:00:01 sshd: pi@pts/0
pi        2030  1964  0 20:33 pts/0        00:00:00 grep --color=auto sshd
```

Add an Inverse grep

Ah! There's a problem! We seemed to have caught our own grep in the command. But that's no problem we can simply pipe again to grep and do a `-v` for grep, which will return us every line without a grep. This gives us the command:

```
$ ps -ef | grep sshd | grep -v grep
```

This gives an output of this:

```
root      1722      1  0 19:24 ?          00:00:00 /usr/sbin/sshd
root      1956  1722  0 19:43 ?          00:00:00 sshd: pi [priv]
pi        1963  1956  0 19:43 ?          00:00:01 sshd: pi@pts/0
```

Cut Down the Results

Much better! Now you can begin to see what we can accomplish with pipe, but I am far from done with the things you can use this for! Let's introduce the next command: awk. awk is a tool for doing data extraction from a line of text. In this way, you can pull all kinds of information out of whatever commands you're running. By default, awk will pull each element out by splitting on *whitespace*, which is any tab or space between characters. You can set the character to split against using the `-F <delimiter>` argument. Once you have told awk how to break your text, you will need to output it. awk has an entire syntax for outputting commands, and it is an incredibly complex art form.

However, for most of your needs, simply printing the correct field will do. All awk outputs are wrapped in curly braces { and }, and you will normally need to wrap them in single quotes so the command line will ignore any spaces. This means the basic syntax looks like this:

```
$ awk '{ <command> }'
```

The most useful command for anyone working with awk has to be the `print` command because it is used to display the delimited variables. Each of these variables is treated as the number it is in the sequence, so \$1 for the first, \$2 for the second, and so on. Using this we can combine the awk with our `ps -ef` from earlier to print out the second space delimited field, the PID! This gives us the following command:

```
$ ps -ef | grep ssh | grep -v grep | awk '{print $2}'
```

Wow, these commands are getting *long!* But examine the output of this command:

```
1722
1956
1963
```

Aha! The quick among you will realize what is coming next! We can pipe this to kill! In this fashion, we can use a `ps`, `grep`, and `awk` to grab only the PIDs of processes matching a given pattern and then pass them off to be killed!

xargs to Work with Each Result

This is one of those pieces of magic that never gets old; however we can't pass a list of PIDs directly to kill because we want them passed one at a time. To do this, we need to take advantage of another command, xargs, which works the same way that the -exec does in find. It allows you to treat each line of output from the commands before as an individual argument for another command. To use xargs, we simply have them wrap the command in the same manner sudo does. So if we wanted to pass each of those PIDs to xargs via kill, we would simply run the following and that would be it:

```
$ ps -ef | grep ssh | grep -v grep | awk '{print $2}' | xargs kill -1
```

In this case, I will *not* be killing off my sshd process as I don't want to reboot my Pi. But you should get the idea of how you can begin taking output from one command and passing it through to another command to run further executions on these commands.

There are actually a number of other commands that come in really handy when doing these kinds of "shell-fu." The first is wc, which is short for *word count*. By default, wc will give you the newline, word, and byte counts for a given block of text or file, but it also has some great arguments, such as -l, which will give you the number of lines; -w, which will give you the number of characters; or -m, which will give you the number of words.

Next is sort, which sorts each line of the output into order. You often see sort paired with uniq, which is short for *unique*. With these two, you can pipe an output to sort then to uniq, and you will have only the unique values in your given output!

The final thing to be aware of when working with a number of commands is the output operator >. This operator is from the same family as pipe, but instead of being used to direct your command output to another command, it is used to direct your output to a file. This is a great way to save the output of commands to a file so that you can revisit it later or to have it useable as arguments for another command. By now, you should be getting the impression that the shell is incredibly powerful and you can do all kinds of tricky manipulations by combining commands with pipe!

User Environment

We are going to take a small break from examining commands and have a look at some of the things that make up the environment in which you work. Your environment is launched every time you log in, and each time is the same. This is due to a number of files and commands that provide directives on how your simple shell environment is to be made up and how it is to function. One of the key files involved in performing these tasks is your systems master controlling environment script /etc./profile. This file provides the basis of the shell and environment variables that are the base of each user's session.

The profile file sets a number of variables, including the PATH variable. This variable is one of the key variables in setting up how your shell environment works. So we will have a look at it in depth.

The PATH is one of (if not the most) important variables in your whole environment. It represents the systems search path that basically means where it will look for applications. As an example, I am sure that by now you are aware that all the commands we have run are just words entered into the shell, but in fact they are the names of fully functional applications that are located within the OS. The PATH contains a list of folders that your shell will search through and execute any command from within automatically. This means that you aren't always searching for the exact path of a command; you can just run it as needed by its name. The problem here is that you cannot be certain where the program is being run from, or if you have multiple versions, which one is being run. There is a command that solves this problem: the which command. It outputs the path of the application that your shell would use to execute a given command. For example, let's find the path of the ssh application, we type which ssh and we get this:

```
$ which ssh
/usr/bin/ssh
```

Yep, it's that easy. The PATH is used by a number of programs to find information that they need. Your path variable is initially set within /etc./profile but this can be modified. But before we go out and start playing with the path variable we should find out what it is currently set too. There are two ways you can do this, first is to use the echo command to print the variable. echo will interpret whatever you give it and then print to the command line, so if we execute echo \$PATH the echo command will output the contents of the PATH variable like so:

```
$ echo $PATH
/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/local/games:/usr/games
```

Okay, that's what the PATH variable is. The *path* is a list of folders separated by colons. This brings us to the second way of viewing the PATH variable: the env command displays all the currently set environment variables. This is a fantastically useful command to check the current state of your whole system environment. Running env should give you an output like this:

```
$ env
TERM=xterm
SHELL=/bin/bash
XDG_SESSION_COOKIE=eb95d80869be1ad62af36ec5502c41a1-1349772229.897557-1575821964
SSH_CLIENT=10.0.0.104 3643 22
SSH_TTY=/dev/pts/0
USER=pi
LS_COLORS=rs=0:di=01;34:ln=01;36:mh=00:pi=40;33:so=01;35:do=01;35:bd=40;33:01:cd=40;33:01:or=40;
31:01:su=37;41:sg=30;43:ca=30;41:tw=30;42:ow=34;42:st=37;44:ex=01;32:*.tar=01;31:*.tgz=01;31:*.arj=01;31:*.taz=01;31:*.lzh=01;31:*.lzma=01;31:*.tlz=01;31:*.txz=01;31:*.zip=01;31:*.z=01;31:*.Z=01;31:*.dz=01;31:*.gz=01;31:*.lz=01;31:*.xz=01;31:*.bz2=01;31:*.bz=01;31:*.tbz=01;31:*.tbz2=01;31:*.tz=01;31:*.deb=01;31:*.rpm=01;31:*.jar=01;31:*.war=01;31:*.ear=01;31:*.sar=01;31:*.rar=01;31:*.ace=01;31:*.zoo=01;31:*.cpio=01;31:*.7z=01;31:*.rz=01;31:*.jpg=01;35:*.jpeg=01;35:*.gif=01;35:*.bmp=01;35:*.pbm=01;35:*.pgm=01;35:*.ppm=01;35:*.tga=01;35:*.xbm=01;35:*.xpm=01;35:*.tif=01;35:*.tiff=01;35:*.png=01;35:*.svg=01;35:*.svgz=01;35:*.mng=01;35:*.pcx=01;35:*.mov=01;35:*.mpg=01;35:*.mpeg=01;35:*.m2v=01;35:*.mkv=01;35:*.webm=01;35:*.ogm=01;35:*.mp4=01;35:*.m4v=01;35:*.mp4v=01;35:*.vob=01;35:*.qt=01;35:*.nuv=01;35:*.wmv=01;35:*.ASF=01;35:*.rm=01;35:*.rmvb=01;35:*.flc=01;35:*.avi=01;35:*.fli=01;35:*.flv=01;35:*.gl=01;35:*.dl=01;35:*.xcf=01;35:*.xwd=01;35:*.yuv=01;35:*.cgm=01;35:*.emf=01;35:*.axv=01;35:*.anx=01;35:*.ogv=01;35:*.ogx=01;35:*.aac=00;36:*.au=00;36:*.flac=00;36:*.mid=00;36:*.midi=00;36:*.mka=00;36:*.mp3=00;36:*.mpc=00;36:*.ogg=00;36:*.ra=00;36:*.wav=00;36:*.axa=00;36:*.oga=00;36:*.spx=00;36:*.xspf=00;36:
ALL_PROXY=SOCKS://localhost:9050
MAIL=/var/mail/pi
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/local/games:/usr/games
PWD=/home/pi
LANG=en_US.UTF-8
SHLVL=1
HOME=/home/pi
LOGNAME=pi
SSH_CONNECTION=10.0.0.104 3643 10.0.0.56 22
_=~/usr/bin/env
```

There are a lot of variables here that cover a range of things, including the shell we are using (SHELL), the variables governing the colors displayed when we execute ls (LS_COLORS), the path, my current location (PWD), my home directory (HOME), and my username (USER) in addition to the PATH. All these pieces of information can be incredibly useful when trying to diagnose issues with running commands.

Now that we have been able to view the PATH variable, we should look at modifying it. To modify a variable, we need to use the `export` command in order to assign it into our environment. The syntax for setting a simple variable is this:

```
$ export VARIABLE="SOMETHING"
```

If you run this and run `env` again, you should see this new variable show up within its output like so:

```
VARIABLE=SOMETHING
```

If we were to do this to our PATH variable, we would reset it to whatever we set it to, which is not what we want. We could also simply rewrite the whole path, but that's quite arduous for a simple change. The easiest way to append something to the PATH variable is to include the PATH variable within the assignment. Let's work with an example: say we want to include the `/opt` directory into our path. We would start with `export PATH=` and then we would want to use the current PATH variable (as we would use to display it with `echo`), so we write `"$PATH"`. To the `$PATH` we would add a colon as a separator and the new directory `/opt`, which gives us this command:

```
$ export PATH="$PATH:/opt"
```

When run, this will change your PATH variable to this:

```
/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/local/games:/usr/games:/opt
```

Awesome; the `/opt` was appended to the end. This brings us to a key point; PATH has an order, left to right. Those on the left side of the PATH variable will be used before those on the right. In our example, `/opt` would be the last place we search for any given file.

Now that you understand your user environment and how PATH works, you should learn about a special file within your home directory. This is your local shell profile override file. In your home directory there are a number of hidden files that end with `rc`. These files are normally executed by your shell or whichever application upon its execution to load your personal preferences. The file for the standard Linux BASH shell is the `.bashrc` file. This file is responsible for a number of the environment variables we saw before; it is also the reason that some of your commands have color. For example, the `ls` command, which is normally located at `/bin/ls` (you can check with `which ls`). If you were to execute this instead of the normal `ls`, you would not see the colors! These colors were added specially by the `.bashrc` file using the `alias` command.

`alias`, as its name implies, is used to alias a command or a group of commands under a single shell command, just like a program. In the case of your `.bashrc` file, it has aliased the `ls` command to execute this:

```
$ ls --color=auto
```

This addition means that every time you issue an `ls` command, you are automatically adding the color argument to `ls`. The way the `.bashrc` achieves this is by adding this line:

```
alias ls='ls --color=auto'
```

The great thing about files like `.bashrc` is they are designed to be user modified—after all, they are your local system variables. In this way, your local environment is made from the base system values given by `/etc./profile` and the additional values given in `.bashrc`.

The cron Command

The last command is probably one of the most powerful and useful tools in a systems administrator toolkit: the `cron` command. `cron` is a timing application that executes commands at a given time based on rules passed to it by a steering file. `cron` has a steering file per user called a `crontab`, which it will check each minute to see if another command needs to be executed. To view your current `crontab`, you would simply need to execute this:

```
$ crontab -l
```

Unfortunately, you won't have one to begin with, so all you will get from this command is this:

```
no crontab for pi
```

Fear not because we will now go over the basics of creating a `crontab` timing file. The command to enter the special `crontab` editor is as follows:

```
$ crontab -e
```

This will open a copy of your system's standard editor and create an initial file with a nice block of comments that outline how a `crontab` entry is laid out. It provides the following example:

```
0 5 * * 1 tar -zcf /var/backups/home.tgz /home/
```

This command will perform a `tar` and `zip` of all users' home directories every Monday at 5 a.m. Recall that `tar` creates a single archive file of all the documents, and the `zip` will compress the file—the upshot being a single compressed file containing all the files. The five numbers at the start of the line govern the timing; these numbers represent the following things, in the following order:

- Minute of hour
- Hour of day (in 24 hour format)
- Day of month
- Month of year
- Day of week (Sun, 0 to Sat, 6)

As you can see from the example, this command runs on the zero-th minute of the fifth hour of any day of the month, in any month of the year as long as that day is a Monday. With this combination, you can set up any number of timing jobs in order to perform systems administrative functions. The best part is that you can use all the commands we have discussed earlier to perform these tasks for you! Want to kill a process at a certain time? Write the `cron` timing out for when you want things killed and then use the `ps`, `grep`, `awk`, `xargs`, and the `kill` command we worked up earlier! With the combination of `cron` and all the commands we covered earlier, there is no limit to the tasks you can do!

Summary

Congratulations, you should now be well on your way to being a fully fledged systems administrator. You have at your disposal more than two dozen individual commands and the ability to combine them in new ways to perform almost any task your heart desires. In addition, you should now have the tools to automate the execution of these commands using `cron` by editing your `crontab` file. All these tools form the basis of how a systems administrator will manage a Linux system as they provide the flexibility to write robust commands and have them executed automatically.

CHAPTER 6



Editing Files on the Command Line

Believe it or not, one of the things you'll end up spending most of your time doing on Linux is editing text files. They get absolutely everywhere, and you'll find them used for content, source code for applications, configuration files, and startup scripts. The reason is that despite their somewhat basic and boring nature, they are very useful for storing human readable information. When you combine this with ease of use, you know you're on to a winner. Linux doesn't have a registry like Windows, and it's extremely rare for applications not to use text files for their configuration. Even Sendmail, which has a historically evil config file (so evil in fact you need to write a config file for yet another application that will then create the Sendmail config file for you), stores its config file as plain old text.

This chapter gives you a brief overview of text files and how they're used today. We'll then introduce the two editors that we'll cover in this chapter (nano and vim) and then we'll get into the real meat of the chapter and show you how to actually get things done.

What is a Text File?

When a computer stores data on disk, it can store data in one of two formats: as a text file or as a binary file. So because computers store everything as binary data (those good old 1s and 0s), what exactly makes a text file different from a binary file? After all, surely the text file must also be stored in a binary format. If you were thinking along those lines, you're quite correct. It's not so much how the data itself is stored, but more to do with how that data can be read.

The best way to appreciate the difference is to think about the page you're reading right now. The reason why you can understand what I'm writing is because we have implicitly agreed to use English to communicate. The content of this page has structure that we both understand. We start sentences with capital letters; we use punctuation to show when sentences end or when to stop for a breath. In short, we have quite a complicated protocol here, but it works because we are both capable of reading and writing English. We can understand both the *structure* (how the data is to be read and written) as well as the *semantics* (what the words actually mean). So, what about the following sentence?

Goedemorgen. Hoe gaat het?

Straight away we can tell that it's not English; after all, if it were, we would be able to understand it. Unless you happen to speak Dutch (we don't, but our good friend does), you probably wouldn't know which language it was written in. Even though the content doesn't make sense, the structure itself does. We can see we have two sentences: one is a statement and the other is a question. In this particular case, we can guess what the first word means. If we pronounce it as we would English, it sounds very similar to good morning (which is actually what the word translates to). From that, we can probably guess that the question is "How are you?"

The reason we can do that is because although we don't understand the content, Dutch and English share a common structure or format. We share an alphabet (although Dutch does have some characters not present in English) and punctuation (full stops, commas, and question marks). Now you might not get the pronunciation right (it is, after all, a foreign language) but because we share a common structure, we can at least handle the content, even if we can't understand it.

Let's take this one step further with Figure 6-1.

早晨好。你怎么样？

Figure 6-1. Some Chinese text

At least for the average reader of this book, Figure 6-1 will be unintelligible. It's actually the same two sentences that appeared in our Dutch example, but this time written in Chinese using simplified characters. To the uninitiated, there is no way to tell from this collection of characters how they should be pronounced. In fact, depending on the dialect spoken, these characters could be pronounced in many different ways. This is because each character represents an idea rather than a sound, and once you recognize the idea, you know which word to say. This is why in China it is possible for two people who cannot speak the same language to communicate perfectly well by reading and writing characters.

So why did we sidetrack into foreign languages? Well, text files are like reading and writing English or Dutch. Whether the content makes sense or not depends on whether the reader can speak English or Dutch. This is similar to two applications and their config files. If you give an Apache web server configuration file to the Sendmail e-mail server, it will spit it back in your face (Sendmail doesn't speak Apache), but it can open and read in the file.

Chinese, on the other hand, is our example of a binary format. It is very alien for most Western speakers (although arguably it is far more efficient, elegant, and sophisticated than phonetic writing used in Western European languages). Unless you know Chinese, you can't understand the content or the implied structure. In computing terms, what you have is a blob of raw data that only makes sense to the application that wrote it.

Now you might be thinking that storing data in an arbitrary format is no big deal. After all, the application that wrote it is the application that will read it, and as long as it understands the file, surely that's enough. Well yes and no. What happens if the program breaks? What if you need to read the file on another machine? Perhaps the vendor has provided some special tools to do just that, but therein lies the problem: you need special tools.

Not so for text files; they can be opened in any text editor, and no special tools are required. Because the vast majority of Linux applications all use text files (we can't think of any off the top of our heads that don't), all you need to read or make changes to a configuration file is your favorite text editor, and Linux has a stunning variety of editors to choose from. Thus you can become familiar with a single tool and then use that tool to reconfigure all your applications, write source code when you're doing scripting, and even use it for writing books (which is precisely what we're doing right now).

The Contenders

So the text editor is your Linux Swiss army knife and something you never want to leave home without. Of course, just like Swiss army knives, there are many different text editors to choose from, and some have more features than others. Sometimes you want a basic text editor that just lets you open files, edit some text, and then save the results. Other times you might want to do a complex search and replace or delete to the end of the line. Of course, you can quite easily use one of the more feature-full editors for all your tasks; after all, you don't have to use the features if you don't need them. That said, it's been our experience that people are generally familiar with at least two text editors: one basic and one fully featured. In this book, that means nano and vim, respectively.

The first text editor that we'll cover is nano. It's a lightweight and easy-to-use text editor and it tends to be installed on most systems. Some later versions of nano have more features (such as syntax highlighting when writing source code), but all effectively look and work the same way. If you can use nano on your Pi, you'll be able to use it equally well on any platform.

The nano editor was actually written to be a drop in replacement for the pico text editor that used to ship with Pine, a text-based e-mail client. Pico can still be found on older systems (usually much older), and chances are if you can't find nano, you might still be able to find pico. Because it's a drop in replacement, the same commands and basic features used in nano will also be available in pico. Combined, this gives you coverage of most other UNIX platforms as well as Linux.

Nano's key selling point is its ease of use, but it does have a number of useful features. These features are accessed from menus, and while nano has more than enough features for simply editing a config file, sometimes you need something with a little more power. This is where the vim editor comes in. It can do all the things that nano can do, but it can do so much more so that's pretty much where the comparison ends. Vim enables the following: syntax highlighting, different color schemes, cut and paste, delete blocks of text, indent chunks of text, open multiple files and show both at the same time, and much more.

Now you might be thinking that if vim is this great, you might save yourself a lot of headaches by simply just learning how to use that text editor rather than any others. Like nano, vim is available almost everywhere in some form or another, and has the same basic features and accepts the same commands. The reason we're covering nano is that you can be up and running quite literally in just a few seconds. Its features are easy to explain, and getting it to do what you want is really easy. Although vim users may disagree, in our experience, vim doesn't exactly hold any ease-of-use awards. That said, if you really don't want to learn two editors and feel that you'd rather just learn vim, that's fine, you won't really be missing out on anything; it just might take you a bit longer to get productive.

Starting out with nano

We're starting out with nano for the simple reason that it's really easy to use and doesn't take much explaining. You can quite literally do everything you need to in nano, and while you might be able to do it faster in another editor, nano will help you get the job done just fine.

Nano can be run by using the nano command. Okay, no surprises there, but running nano by itself isn't all that useful. After all, you've got a text editor now and you want to do some editing. Unless making a few notes, no one ever just writes text without wanting to save it afterward. When you open nano, you probably want to open a particular file for editing. That file may or may not exist, but you probably know what it is you want to call it when it comes to pressing the save button. For the sake of this introduction, we're going to open nano and tell it that we want to edit test.txt. We can do that with this:

```
$ nano test.txt
```

All being well, you should get something like Figure 6-2.



Figure 6-2. Freshly opened nano with a new file

We've connected to our Pi using SSH but because nano is a terminal based application, it will look the same regardless of how you run it. One interesting thing about these sorts of terminal applications is that they are sensitive to how big your terminal is. A fixed terminal (such as the physical console) has a set screen size. When you connect via SSH or use a virtual terminal, you can adjust the size of the screen and this will cause nano to resize along with it. This makes it very flexible and can fit in with your needs regardless of the shape or size of your screen.

So What Does This All Mean?

Nano's screen can be a bit confusing at first (especially the bit at the bottom), but most of the time you won't actually be paying attention to anything other than your text. In fact, all the other stuff melts away in the background. For example, on the top left of your screen you can see that we're running "GNU nano 2.0.9." Until we came to write this book, we had no idea what version of nano we were running and quite honestly until we had to describe it here, if someone had asked us to draw the nano layout from memory, we'd have had a pretty hard time doing it. Actually, we wouldn't have been able to do it at all!

There are two other pieces of information that you can glean from nano's title bar. Smack bang in the middle is the name of the current file that you're editing. In our case, it's `test.txt`. Not exactly earth-shattering, but it's quite common to be working on multiple windows at once, and being able to quickly get your bearings can come in handy.

The third piece of useful information is whether the file has been modified; that is, are there any unsaved changes? This can be very handy when you want to check whether you've saved your latest work. It's also useful for simply jogging your memory and reminding you to save your work. To see what this looks like, just press the spacebar to make a change, and you'll see that "Modified" appears in the top right of the screen (see Figure 6-3).




Figure 6-3. Modified file in nano

That's pretty much it for the title bar. At the bottom of the screen there are actually two sections. The first one is the status line, which shows you important information (right now, it's telling you that it's a new file), and the second one is the shortcut or menu bar. We're not going to say much more about the status line because it's pretty much self-explanatory. Anything of importance that you need to know, nano will put it there. Of much greater interest (and in fact the key to using nano) is the shortcut bar. Let's take a closer look at it in Figure 6-4.



Figure 6-4. The shortcut bar

When you've been working in Windows or on the Mac, you've probably reached the point where rather than using the mouse to copy and paste things you instead use the keyboard combinations. On Windows, you would use the Ctrl and C keys, and on the Mac it would be the Command key and C. When you use more than one key in concert with another, you end up with a key combination. Nano works in a very similar way, but rather than just copy and paste, all its important functions are driven with key combinations. Because it would be somewhat difficult to remember all the possible options available (although the ones you'll use on a frequent basis will soon be committed to your long-term memory), nano displays them in the shortcut bar. Actually, it only shows the most common commands, but we have never needed to stray away from them.

To make use of any command, you need to use the correct key combination. You can find which combo you need by looking at the shortcut bar. A shorthand way of writing *control key* is to use the carat (^) symbol. Looking at the shortcut bar, you can see that if you want to get help, you need to use ^G; in other words, hold down the Control key and press the G key.

Saving Your File

Saving a blank text file is pretty boring, and it's hard to tell if it saved correctly if all you have is a space character. We're going to go for an old favorite here and punch in "Mary had a little lamb."

Unlike some editors (most notably vim, which you'll come across later), nano doesn't have the concept of different modes. That is, when you type into nano, it will be interpreted as text to put into the file. Vim, on the other hand, stands in command mode, and if you just start typing, you will get some interesting results. As you might expect, pressing Enter after each line will take you to the beginning of the next line. After typing in the nursery rhyme, you should get something like Figure 6-5.




Figure 6-5. *Mary had a little lamb, nano style*

We shrank the nano window down a bit so that we don't take up vast quantities of paper with blank whitespace. It also acts as a pretty good example of why you might want to resize a window and how it looks when you actually do. In this case, everything looks just fine, although admittedly for writing anything of a reasonable length, you'd need something a bit bigger than this.

You may have noticed that the status line seems to have disappeared. This is because the file stops being a new file once any content has been added to it. Our modified warning is still up in the top right hand corner though. It's time to get rid of that and get Mary safely on to disk.

Looking at the shortcut bar you can see the options we have available to use. The most common ones you'll use are WriteOut (which really just means *Save File*) and Exit. Exit actually doubles as WriteOut because if you try to exit nano with unsaved changes, it will ask you if you want to save your changes.

Let's start with the WriteOut command. Hold down the Control key and press the O key. Nano should now look something like Figure 6-6.




Figure 6-6. *Using the WriteOut command*

We made two changes here:

- The status line has made a comeback. This time it's asking for information, or at least it wants us to confirm it. Now that it's time to save the file, nano is asking us to confirm what the name should be. Because we passed it a name when we originally started nano, this is the name that nano will provide by default. If we press the Enter key now, it will save the file; but before we do that, let's take a look at the shortcut bar.
- The shortcut bar has changed to a whole new set of options because the shortcut bar is actually contextual; that is, it will show you the most relevant things based on what it is you're doing at the time. These options are only useful (or even available) when saving a file, so this is when you get to see them. We've never actually had to use any of these options apart from Cancel (when you suddenly decide that saving is not what you want to do at all).

Okay, back to saving the file. If you like, you can change the name that the file will be saved as, and nano will not only save it under your new name but will also continue using that name when you go back to editing, which is the same file and you should get a message like the one shown in Figure 6-7.




Figure 6-7. File has been written to disk

You can see that the shortcut bar has gone back to its usual look and we have a new message in our status line. You'll also have spotted that the Modified status has vanished. So far so good.

Let's quickly look at what happens if you try to exit nano while you have an edited file (see Figure 6-8):




Figure 6-8. Exiting before saving

All we did here was add a signature to the file (although admittedly it probably wasn't Mary who wrote it) and then tried to exit nano by using Control and X. This time the status line gives us a scary warning, and the shortcut bar gives us a somewhat reduced number of choices. We can say yes, in which case nano will save the changes before it exits; or we can say no, in which case nano will still exit, but we will lose all the changes we made since the last time we saved the file. We can also think better of exiting and change our minds. In this case, we can use Control plus C to cancel our exit request, which will just take us back to nano.

Moving Around in nano

You can move around your text just by using the cursor keys because you're probably used to from your word processor of choice. In fact the only reason we mention this here is because some editors don't use those keys at all because they take your hands too far off the home row (that is, they're slow to use because you have to move your hands a long distance to reach them). So rather than explaining cursor keys, we'll show you the two features we use most in nano, which are both coincidentally related to finding things.

The first thing you will probably want to use is the search feature that is somewhat oddly named “Where Is” under nano. We can issue that command by pressing the Control and W keys to get to the next step, as shown in Figure 6-9.




Figure 6-9. Searching for some text

Once again we have some new options available to us. The most common task, though, is to simply find a word. So let's search for Mary and see what happens (see Figure 6-10).




Figure 6-10. Nano found Mary

After pressing the Enter key, nano will try to find the text you searched for. In our example, we were already at the end of the file; that's why you can see “Search Wrapped” in the status bar—nano reached the end of the file and wrapped around to continue the search from the beginning. This is a really handy feature!

There will often be more than one occurrence of the word that you're looking for. If you want the second *Mary*, you repeat your search. You can do this by just pressing Control and W and then immediately pressing the Enter key. Nano remembers what you searched for last time, and if you don't supply a new word, it will simply search again with the old word. This means you can easily cycle through the document to find the specific entry that you're looking for.

The last feature that we'll cover is a slightly more specific form of search. Suppose that rather than searching for a word, you want to go to a specific line number. In day-to-day text editing this is not something that you'll actually do all that often. However, when you start writing scripts and they generate errors, you will generally be told which line that error occurred on. For short programs it often doesn't matter, but if you're building something fairly epic, you'll want to be able to get to the line quickly. This is true even in medium-length programs because if you have a particularly tricky problem, you'll have to keep revisiting a specific piece of code and you certainly don't want to have to use the cursor keys to find it each time.

To go to a specific line number, you need to start the search as you did before by using Control plus W; but rather than typing any text, you should then do Control plus T. This will change the status line as shown in Figure 6-11.




Figure 6-11. Go to a specific line number

All you have to do here is enter the line number that you want to go to and press Enter. You can also see that there are some other useful options here that allow you to go to the start of the file and the end of the file, respectively. You can also press command and T to go back to searching for text, or you can press command and C and cancel the search.

Wrapping Up nano

That's all we're going to cover for nano but as you can see, even the limited amount of functionality we've covered has already given you the ability to edit configuration files or even compose your own book. Nano is great for these kinds of tasks where you don't want distractions and you just want to get the job done. As it happens, a good deal of this book was written in nano, so it just goes to show that this tool is certainly not just a toy! That being said, let's move on now to look at what the vim editor can offer you instead.

Getting Started with vim

Vim is a bit more advanced than nano. Whereas nano was specifically designed to be simple, vim is designed to be a fully featured environment for doing all sorts of tasks. You can even integrate other applications (such as the shell and source code control) directly with vim, so you can do huge amounts of work without ever leaving your text editor.

The downside to this raw power is that you have to learn how to use it. Unlike nano, which is basically ready to use right away, vim has the concept of modes. Generally speaking you're either in the normal or insert modes, and they are the ones that we'll focus on in this section.

However, we do want to point out that even though we use vim on a regular basis, we're not what you would call vim experts. We know how to use the functionality that helps us with our work, but there are many features that we don't use or even know about. Because this is just a primer on vim, we don't feel that this is an issue, but if you think vim is the tool for you, there is a wealth of books dedicated solely to vim as well as numerous video tutorials on some of its more advanced and powerful features.

Rather than having to type in yet another poem, let's give Mary an encore and bring her back to the stage with vim. First though, we need to install it:

```
$ sudo apt-get install vim
```

Now that we have vim, let's see about that encore:

```
$ vim test.txt
```

Needless to say, starting vim is as straightforward as starting nano, but it does look a little different once it has started (see Figure 6-12).




Figure 6-12. Starting up vim

Gone are the nice headers, status lines, and shortcut bars. In the land of vim, like Linux itself, we're expected to know what we're doing. In the bottom left of the screen, you can see the name of the file that we're editing. Next to that is 6L, which tells us how many lines are in the file, followed by 115C which tells us the number of characters. On the other side of the screen we have another useful piece of information: 1,1, which tells us the line that we're currently on and the column that we're in. At present it is telling us that we're on line 1 and in the first column. That makes sense because we only just opened the file; if you look at the cursor, sure enough we've sitting on top of the first M.

Note Depending on the version of vim and how it is configured, it will usually take you to the location you were at when you last stopped editing that file. This is really useful when programming because it gets you immediately to where the action is.

You can also see several lines of tildes (~) beneath the last line in the file. They are here to remind us that although the screen is a certain size there is no content below Mary. This is especially useful when doing certain types of programming and you want to make sure there's no additional whitespace. It's really just a visual clue and won't get in your way.

The vim Modes

Vim comes across as being a bit quirky because (unlike word processors and more basic text editors) with vim you can't start typing right away because vim has different modes or personalities. When you first start vim, it is in command mode. Unlike nano, in which you give commands by using key combos, vim has a wide range of commands that are single or double letters. You might think that this would be hellishly confusing and, well, you'd be right. Once you know the magic key presses to make vim do your bidding, you'll love this way of controlling your text editor. Until then, however, it's likely to drive you to distraction.

So let's get out of command mode and slip into something more comfortable. We can replicate a nano-esque environment in vim by pressing the i key. This will take us out of command mode and place us in insert mode. You can tell when you're in this mode because you see --INSERT-- in the bottom-left corner, as shown in Figure 6-13.

A screenshot of a terminal window titled "miggyx — root@plexus:~ — ssh — 80x12". The terminal displays a poem in black text:

```
Mary had a little lamb
its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go.

-- Mary
```

The text "Mary had a little lamb" is selected with a blue highlight. Below the poem, there are several blank lines starting with a tilde (~). At the bottom left, the text "-- INSERT --" is displayed in red. In the bottom right corner, the coordinates "1,1" and the word "All" are shown.

Figure 6-13. Vim in insert mode

At this point you can move around your text file just as you did with nano. The cursor keys work as expected and you can add and delete text as you did under nano. The differences start when you want to do something other than writing text, though. How do you save your changes?

Saving Your Changes

To save your changes in vim the first thing you need to do is exit from insert mode by pressing the escape (ESC) key. This takes you back to normal mode, in which you can issue commands and consequently save your document. First, though, let's change the signatory to say John instead. This will give us the much-needed change for the document. To actually save the document we use this command:


```
:w
```

The command really is a colon followed by a lowercase w. Vim is case-sensitive, so if you tried to do: W, you would get an error message that it isn't a valid command, as shown in Figure 6-14.

A screenshot of a terminal window titled "miggyx — root@plexus:~ — ssh — 80x12". The terminal displays the same poem and signature as Figure 6-13. At the bottom left, the text "E492: Not an editor command: W" is displayed in red. In the bottom right corner, the coordinates "6,7" and the word "All" are shown.

Figure 6-14. Whoops, W isn't a command

There's no need to panic; just try the command again and you'll be fine, as shown in Figure 6-15.



A screenshot of a terminal window titled "miggyx — root@plexus:~ — ssh — 80x12". The window displays the poem "Mary had a little lamb" followed by a signature "— John". Below the text, the status bar shows the command ":w" and the message "'test.txt' 6L, 115C written". The status bar also indicates line numbers 6, 7 and the word "All".

```
Mary had a little lamb
its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
^
^
^
^
^
^
^
"test.txt" 6L, 115C written
6,7
All
```

Figure 6-15. :w works like a charm

If you want to change the file name while you were saving, you can simply type in the filename you want to use after the command like so:

```
:w test2.txt
```


This will then save the changes to the new filename and will then continue to edit the new file.

Getting Out of vim

This leaves us with just one more feature to bring you to the same basic level that you had with nano (we'll cover searching in just a moment). Once you've done all your editing and you've saved your files, you'll want to exit. The command for this is as follows:

```
:q
```

As long as you haven't made any changes, vim will immediately exit, and you'll be back at the command line. However, if you had made some changes and you try to do this, you'll once again catch the wrath of vim, as shown in Figure 6-16.



A screenshot of a terminal window titled "miggyx — root@plexus:~ — ssh — 80x12". The window displays the poem "Mary had a little lamb" followed by a signature "— John". Below the text, the status bar shows the error message "E37: No write since last change (add ! to override)". The status bar also indicates line numbers 6, 8 and the word "All".

```
Mary had a little lamb
its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
^
^
^
^
^
^
^
E37: No write since last change (add ! to override)
6,8
All
```

Figure 6-16. To quit without saving you'll need an override

Fortunately, vim has already told us how to fix the problem, but basically if you want to quit without having your changes saved, you'll need to explicitly tell vim that's what you want to do:

```
:q!
```

And that's all there is to leaving vim!

Searching in vim

In this section we'll cover only the basics and show you how to do the same sorts of searches that you did in nano. In actual fact, vim has very advanced search capabilities (you can use regular expressions for a start; see the following sidebar) but although you might come to depend on these features as you become more advanced in your usage, you'll still probably find that you spend most of your time doing the simple things.

To search for a word in vim, all you need to do is prefix it with a forward slash. To give this a try, while in command mode (tap the Escape key twice to ensure that you're in the right mode) type:

```
/Mary
```

This should then find all instances of Mary and highlight them for you, as shown in Figure 6-17.

```
miggyx — root@plexus:~ — ssh — 80x12
Mary had a little lamb
its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
~
~
~
~
~
search hit BOTTOM, continuing at TOP      1,1      All
```

Figure 6-17. Vim found Mary

You can see the “search hit BOTTOM, continuing at TOP” message, which is equivalent to the nano “Search Wrapped” message. One thing to watch out for with vim is that because it uses regular expressions for searching, it's also case sensitive by default. That means if you search for *mary* rather than *Mary*, expect to receive a complaint as shown in Figure 6-18.

```
miggyx — root@plexus:~ — ssh — 80x12
Mary had a little lamb
its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
^
^
^
^
^
^
E486: Pattern not found: mary
3,21
All
```

Figure 6-18. Vim can't find *mary*

One advantage that vim has over nano is repeated searches. To move to the next search result, all you need to do is press the **n** key. In our preceding example text, if you search for Mary with `/Mary` and then keep pressing the **n** key, you will see vim will cycle through the search results for you. We can't really demonstrate that effect in print, but it's a nice distraction for a few seconds if you'd rather not be doing something else. If you'd like to search in the opposite direction (i.e., toward the top of the file rather than the bottom), simply use a capital **N** instead.

REGULAR EXPRESSIONS

Regular expressions in the right hands are like magic. They slice through complex text-processing tasks like a hot chainsaw through warm butter. They are also frequently used to give system administrators bed-wetting nightmares and to scare small children. Regular expressions are very powerful once you know how to use them, but boy does that power come with a cost! Be that as it may, a basic knowledge of regular expressions will be very useful to you, and you can use them directly in vim for searching, replacing, and otherwise manipulating your text.

We don't have enough space to go into regular expressions here (often referred to as *regex*), but there are some fantastic resources available on the Internet to get you started, not to mention complete books on the topic. A great starting place is Jan Goyvaerts' website: www.regular-expressions.info/. Not only is this site full of useful and helpful information (it is very often our first port of call) but Jan has also written some very powerful software. We bought RegexBuddy many years ago and it's still our first port of call when we're trying to debug some complicated regex. If you want a simpler tool for testing out your regular expressions, check out <http://rubular.com/>, which does real time regex parsing in Ruby.

Moving to a Specific Line Number

The last thing we covered in nano was moving to a specific line number. This is actually much easier in vim, probably because it's something that programmers have to do all the time. To jump to a specific line number, all you need to do is place the line number after the colon. For example to jump to the 5th line all you'd need to do is this:

`:5`

Once you hit the Enter key, you'll be taken directly to the fifth line. As before, this feature isn't particularly useful for normal editing, but it comes in handy when you're fiddling about with a program that simply refuses to work the way you know it should!

Moving About in vim

Although we've already shown you how to move around your text file, we haven't been leveraging the power of vim. Vim provides additional commands that can allow us to navigate through text more quickly.

In command mode, you can still move around your text using the cursor keys, which is convenient. In effect, you can do the same thing using two other commands: the h and l keys. The h key will move the cursor back one character, and the l key will move the cursor forward one character. Admittedly that's not very exciting, but we can now introduce the b and w commands. They work just like the previous commands, but instead of stepping back just one character, they jump to the beginning of words. Suddenly that's more useful because now if you have a long line of text, rather than pressing and holding down the cursor key for extended periods of time (tempting, though, isn't it?), you can greatly speed up the process by hopping a word at a time.

We can take this one step farther, though. If you want to jump to the beginning of a line you can press the 0 (zero) key. To get to the end of the line, you simply use \$ (in our case, that means holding down the Shift key and pressing the 4 key). This is surprisingly useful because there will be lots of times when you want to get to either the beginning or the end of the line, and now you can do so with a single key press.

Table 6-1 makes it a bit clearer.

Table 6-1. Common Navigation Shortcuts in vim

Movement	Key press
Start of line	0
Back one word	b
Back one letter	H
Forward one letter	I
Forward one word	w
End of line	\$

So what we have now is a progressive way to move our cursor from one end of the line to the other. No longer do we have to fiddle about with the Enter or backspace keys, or fall asleep waiting for the cursor to trundle across the screen. That's just the beginning, though, because these keys can be used for a lot more than just moving around your file; they can also be used to specify the scope of the commands you feed to vim.

Deleting in vim

In command mode, to delete a single character all you need to do is press the x key. This will remove whichever letter happened to be under the cursor at the time. This is sometimes useful in its own right (usually when tweaking a typo), but it doesn't really offer you any great advantages over dropping into insert mode. The d key, however, provides a very powerful delete command that you can now combine with your newly acquired knowledge of moving the cursor to do some pretty impressive things.

The first delete command we're going to try is dd. Let's start out by deleting the first line in our file. To make sure we're all in the right place, use the following command to get back to the first line:

:1

Now we're going to delete the current line by typing the following:

dd

Your poem should now be somewhat shorter and look like Figure 6-19.



Figure 6-19. Deleting the first line

Now that we've ruined our poem, now would be a great time to introduce another very useful feature in vim: the "undo" command. To undo the change we just made, simply press the u key, and Mary will be restored to her former glory. This not only prevents panic attacks when accidentally deleting the wrong line but it also means we can delete parts of this line in numerous fascinating ways.

If you have more than one line that you want to remove, you can tell vim how many times you want it to execute a particular command by prefixing the command with a number. For example, if we want to remove three lines from our poem, we could do this:

3dd

This is equivalent to pressing the **d** key six times and gives you the happy result shown in Figure 6-20.



Figure 6-20. Deleting three lines in one go

You can see in the bottom-left corner that there are now “3 fewer lines” just as you hoped. Let’s restore Mary with the u key and try something else. With our cursor at the beginning of the first line, try executing this command:

```
dw
```

This command combines the delete command with movement commands that you saw earlier. In this case, we combined it with w, which means “one word to the right.” You can see the effect of that command in Figure 6-21.

```
miggyx — root@plexus:~ — ssh — 93x11
Had a little lamb
its fleece as white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
~
~
~
~

"test.txt" 6L, 115C
1,1 All
```

Figure 6-21. Deleting a single word

Rather than moving the cursor, we have effectively selected text that we wanted to delete. In this case, d will start deleting from the cursor’s current location, up to the first letter of the next word. We could have achieved this effect easily enough simply by pressing the x key multiple times, but using dw is clearly faster and more precise. Let’s take this to the next level and combine what we know. We know that we can make vim execute a command multiple times by prefixing the command with a number. We also know that we can combine the d command with w to delete a complete word. Let’s revive Mary (another press of the u key) and try this new improved command (the results of which are shown in Figure 6-22):

```
2dw
```

```
miggyx — root@plexus:~ — ssh — 93x11
@ little lamb
its fleece as white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
~
~
~

"test.txt" 6L, 115C
1,1 All
```

Figure 6-22. Deleting two words at once

Well isn’t that interesting? We’ve been able to combine a series of commands to express some reasonably complicated tasks. You can, of course, combine them in any way. For example, if you want to delete everything to the end of the line, you can use d\$. If you want to delete everything from the start of the line, you can use d0. It’s being able to combine commands in this way that makes vim such a powerful tool for working with text. Although these features are of fairly limited use when working with English prose, when you start working with programming languages, they become very useful indeed.

Miscellaneous Little Commands

Before we move on to look at some of the features of Vim's visual mode, we're going to provide a very quick overview of some other simple commands that will come in handy. The i enters into insert mode, but sometimes you want to write immediately above or below the current line. You can get this effect by using O and o respectively. These will both insert a new line and then switch to insert mode. Sometimes you don't need to actually edit the content of a file; you just need to change a single letter. This could be because of a typo or maybe you just want to increment a number. You can do this with the r key. This key allows you to replace the character under the cursor with any other character. Simply press r and then type the character you want to replace it with. Simple but very effective!

Visual Mode

Visual mode is not something we use very often, but it can save us a huge amount of time and effort. This mode enables you to select or highlight blocks of text for further processing. This is akin to highlighting text with your mouse in a word processor. There are two ways you can enter visual mode from command mode: by pressing either the V key or the v key. If you use the uppercase V, you will be able to select a block of text based on lines. With the lowercase v, you gain extra precision as you can select based on single letters rather than entire lines. The overall result is the same, and the option you choose depends only on which one would be most convenient to you.

So what exactly can you do with visual mode? Well it really doesn't do anything in and of itself; it just provides an easy way for you to tell vim what text the following command should be applied to. In keeping with our previous examples, let's select a chunk of Mary's poem, as shown in Figure 6-23. (You don't need to highlight the exact text that we have, but try to get a fairly random chunk.)

```
miggyx — root@plexus:~ — ssh — 93x11
Mary had a little lamb
its fleece as white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
^
^
^
^
^
-- VISUAL --
3,10          All
```

Figure 6-23. Highlighting text in visual mode

As you can see from the highlight, we used the lowercase v option as we wanted to highlight within a particular word. There's no real reason for that apart from a sense of the dramatic, so it could easily have been a bunch of lines selected with V. Now let's apply our favorite command: the delete command. By pressing d, you will find that the highlight text magically disappears. In this case, you don't need to provide additional information to vim to tell it what to delete because in effect you've already told it by highlighting the text in the first place.

Indenting Text

Of course, you can do other interesting things apart from deleting text (though deleting text is fun). For example, if you want to indent the first two lines of the poem, you can highlight the first two lines with V and then use the greater-than and less-than signs (> and <, respectively) as appropriate. We can't show you how that would look in print (honestly moving huge chunks of text back and forth is surprisingly therapeutic), but we can show you the end result of this command, as shown in Figure 6-24.

```
Mary had a little lamb
  its fleece was white as snow
and everywhere that Mary went
the lamb was sure to go.

-- John
~  
~  
~  
~
```

Figure 6-24. Indenting highlighted text

There is one more terribly useful feature that we haven't covered yet that works very nicely with visual mode: copy and paste.

Copy and Paste vim Style

Copy and paste in a terminal is something of an iffy proposition. The reason for that is an application running in the terminal effectively lives in its own little world. If you copy text from the terminal on your laptop, for example, that text is copied into your laptop's clipboard. The remote application (in this case, vim) has no idea that you did that. Likewise, if you paste content in to vim from your laptop, vim will simply receive it as a set of key strokes, it won't have a clue that it came from a clipboard. Generally speaking, this is okay. If you just want to move a bit of text about, using the clipboard in this way is probably fine.

There are times when this doesn't do what you want. For example, if you have a split screen and you have two files side by side and you want to select the first paragraph on the left, you're now a bit stuck. The terminal just shows you a representation of what's on the server; it's just a display and as such the terminal has no real understanding of what it's displaying. If you try to copy and paste the left side, you'll find the right side comes along for the ride. That's far from ideal (not to mention highly irritating).

You might want vim to do something with the text. If you just want to paste it once, no big deal, but what if you want to paste exactly 10,004 times? Vim can't help you there because it has no concept of your local clipboard. You'll find that when these issues do crop up, they'll be fairly niche and you can work around them. However it is likely to really irritate you because you will think "if only I could just copy this text!" while you sit there trying to work around the problem.

What we need then is a remote clipboard—somewhere for the remote application to store data for processing. There's no standard for this, but many terminal-based applications offer a solution; for vim, it's the yank and paste feature. To yank some content, you need to use the `y` key, which takes the same modifiers that the `d` command uses, so you can quite happily copy an entire line with `yy` (remember that the letter just doubles up) and a single word with `yw`. Also like the delete command, we can also use this with the visual mode.

First, simply highlight the text that you want to copy. For simplicity, let's go for the top two lines again. Once you've highlighted that with the `V` key, press the `y` key to yank it into the clipboard. Phase one is complete; now all we have to do is paste it back into our file. There are two ways you can use the paste command: with either `P` or `p`. As with the `V` and `v` commands, the upper- and lowercase versions are usually related in some way. In this case, the `P` key inserts the content before the current line, whereas the `p` key inserts it after. Feel free to move your cursor to any location in the file before hitting one of those keys. The effect is really what you'd expect, and vim will just paste content in. You can, of course, combine these commands with the numeric prefixes to repeat the number of times something is inserted. Figure 6-25 shows what happens if you end `40p`.

```
Mary had a little lamb  
Mary had a little lamb  
its fleece as white as snow  
Mary had a little lamb  
its fleece as white as snow  
Mary had a little lamb  
its fleece as white as snow  
Mary had a little lamb  
its fleece as white as snow  
Mary had a little lamb  
80 more lines
```

Figure 6-25. Mass pasting

Although you don't know it yet, you actually learned to cut as well as copy text. Whenever you use the d key, you delete the text from the file, but a copy is placed in the clipboard. If you select a line and delete it with dd, you can restore it with P. Bear in mind that if you use the p key instead, you'll actually be inserting the original line after the line you're currently on because you've removed the line. If you want to restore it to its original location, you must use the P key to paste before the current line.

Summary

This chapter looked at what separates a text file from a binary file and why text files are the *lingua franca* when it comes to configuring software. We covered two of the key text editors that you are likely to come across in your travels, and we've given you a good grounding in both of them. Certainly we have not been able to show you everything that is possible, but now you should be in a good position to explore and try new things yourself.

In the next chapter, we'll cover the BASH shell and show you some of the more important things that a system administrator needs to know.

CHAPTER 7



Managing Your Pi

So now you have this magical Pi, installed Raspbian, logged in and got your Pi the way you like it, and even have a few things in mind to do with it. Although you are becoming familiar with Linux, the problem is that every time your Pi is switched off or loses power, you need to go over to it, log in and physically set things up before you can use it again. Well, no more!

This chapter is about systems administration and some of the basic application functions that will allow you to manage the way your system functions at boot. To help, we will also give you a crash course in coding with the BASH language so that you can create start scripts of your own to initiate applications (we'll cover just what a script is, too). Finally, we will cover some basic security for your system, including user management, so you can change your password and add new users to your system.

Remotely Accessing the Pi

Probably one of the hardest things to get used to when working in a Linux environment when coming from a Windows or an OSX environment is working in the shell. For starters, it's not pretty, but more importantly humans are creatures who are very much used to manipulating things to achieve an outcome. This means that a GUI-based environment is much more intuitive—at first. Having come from a Windows environment, I can say for certain that the command line is such a great environment to work in. The amount you can do with a few deft keystrokes—compared with keyboard and mouse work, clicking, dragging, and performing context actions is amazing. Being able to work within a system purely from the command line is a great asset indeed. But you need to become familiar with it, and the only way to do this is practice!

Why am I talking about the command line when we need to talk about managing your Pi? Because most Linux management is achieved through the command line, with most if not all of the system functions you want to use being available from the command line. Moreover, it is far easier and less resource-intensive to access your system via the command line than it is to access it via the GUI. To access your system remotely via the command line, you simply need to use Secure Shell (or SSH for short), which you first saw in Chapter 3. You also need to ensure that your system's network access is available on boot as well. These two functions serve the basics of being able to manage your system from anywhere as you have network enabled and then can use SSH to connect to your system. So, let's examine these two functions a little more closely so that you understand how to ensure you can always access your Pi to manage it.

Networking

If you are like most people and have a router to provide multiple devices on your system with network access, you will probably be aware that your router automatically provides network IP addresses to all your devices. It doesn't do this through any sort of magic, but rather through a special protocol called DHCP, which stands for Dynamic Host Control Protocol. The purpose of DHCP is to automatically assign an IP address to a machine, which means there is no need

to manually assign it one. DHCP is invoked automatically by the network service when it is configured to do so. We cover how to configure networking in much more detail within Chapter 9 WiPi.

However, if for some reason you need to manually invoke DHCP to assign you an IP address and get you connected to the network, you need to use the `dhclient` command. You will need to run it as root, as you will need to do with most of these commands, because you are changing core system functionality. If you just run `dhclient` as is, you will try and get a new IP address for every interface on the system, which may not be desirable. You can actually specify a specific interface for `dhclient` to work with by simply adding it as the first argument of the command (i.e., `dhclient eth0`).

DNS

Probably one of the most important things to keep in mind when working with your systems networking is DNS, which is short for Domain Name System. DNS is the way that every system on the Internet is able to turn URLs (e.g., www.apress.com) into an IP address (e.g., 173.203.147.183). DNS on your Linux machine works by knowing where it should direct its queries. This is governed by the file `/etc./resolv.conf`, which contains a name server directive that tells your system where the name server it should query is located. Normally this is your router, but it can be an Internet DNS server such as Google's 8.8.8.8. These will appear in `resolv.conf` like this:

```
nameserver 10.0.0.1
```

SSH

SSH is a way that you can remotely gain a secure and encrypted connection to your Pi's shell without needing to physically do anything to it. Getting SSH installed and running, along with basic commands to access your system, are covered in Chapter 3 so we won't go into much detail and will instead be diving deep into the function of SSH. SSH is provided by the `sshd` daemon that is run at boot time on your Pi. You can start and stop this command in the same way that your system does, which is to use the `sshd` init script. All system init scripts are normally located in `/etc./init.d/` and they are normally run as root. To use them, add the action you want to take (start, stop, and restart) to the first argument of the script. So to restart SSH, run `/etc./init.d/ssh restart`. Additionally, these scripts are the ones your system uses to run programs or commands on boot. So let's have a look at how we can leverage this to write your own init scripts.

BASH: Basic Coding

Besides "Hey can you fix my computer?" the most common computing question I am asked is "Can you teach me how to write this application in C?" Not a month goes by without someone I am working with, a friend, or a relative comes and asks me to teach them how to write software. Normally, they ask for a quick rundown of how to write applications in C. This is not an unreasonable question for the most part, but often when I ask them what they want to do, they want to do something simple. And while I would happily advocate learning C, there is a much simpler way for them to get introduced to the wonderful world of software design, without taking on the task of learning C (which for what they are normally after is a lot of overkill).

That is the goal of this part of the chapter. We aim to introduce you to coding, to basic logic structures, and to some of the basic software design principles so that you can get the job done. We will achieve this by introducing you to BASH. While many people will scoff, BASH is a perfect way to introduce the core fundamentals of what programming is about. BASH is also one of the most widely used languages for software in the world because most applications on Linux systems will have some level of BASH software performing some kind of intermediary functions somewhere along the route.

What Is BASH?

As covered in Chapter 3, BASH is short for Bourne Again Shell and is the default shell environment for most Linux and UNIX environments. When you log in to your system via the command line and are presented with your prompt, this prompt is provided to you by the BASH shell. The shell will take commands and instructions and process them to perform system functions. Most of the commands you normally issue when working with or administering your system via the command line are ones to invoke other applications so that you can generate output.

While this seems like a very obvious and simple way to go about things, it is also a fantastic way of demonstrating the core precept of how any software application works. You take a given input, perform some kind of computation upon it, a response (called the *output*) is generated. This process is, in fact, the basic process of so many functions of life and work that many people take for granted. It is this basic model of input, compute, output that will form the basic model of how we go about understanding how programming functions.

So back to the first question: what is BASH? BASH is a shell, and a shell is a method for taking inputs and computing against them to generate an output. BASH has a number of tools that allow you to use the same basic logic constructs of just about all programming languages. This is how BASH works as a programming language as well as a shell—you can write a number of these logical operations and commands and then use them to form a computation. Now you have a rudimentary understanding of what BASH is, what a shell is, what the broad purpose of any computer program is, and how all these functions are available to perform the basic action of input, compute, and output. Now that you have a very broad overview of what we want to achieve, we can use this knowledge as the starting block for learning how to write a program with BASH.

Starting in BASH

So you know what BASH is and what a shell is and how the rudimentary logic of a program works that's often considered by many to be head and shoulders above the pack. In fact, several interviewers commented that while their employers had rigorous questioning of potential software engineers who could all espouse what we have covered, they complained that many employees and graduates could not form the logic for a basic program on a piece of paper. They went forward to propose a coding test, a very simple test that any software engineer should be able to solve quickly. This test, which is called the *FizzBuzz test*, is based on a children's game of the same name. The basic principle is that you must count from 1 to 100 and say *Fizz* for every number that is a multiple of 3, you must say *Buzz* for every number that is a multiple of 5, and finally you must say *FizzBuzz* for every number that is a multiple of both 3 and 5. It sounds simple; that's because it is. After all, most coding is simply building small blocks of logic and attaching them to other blocks of logic; mashing together lots and lots of little input, compute, and output nodes to form a larger computing system that does a much larger input, compute, and output process.

To this end, we will start by writing our own solution to the FizzBuzz problem in BASH so that you all can pass the basic software engineer competency test! That's right; we will be programming in BASH. The BASH interpreter provides commands that have the same function as common programmatic logic elements. These can be combined with normal shell commands to form whole programs. So, let's start writing (we covered text editors in Chapter 6, so you can use any you like). All you need is to be able to edit a file and execute it on the command line afterward. Go ahead and open up your favorite text editor and write the following:

```
#!/bin/bash
```

This is the first thing that anyone wanting to write a program in BASH should write. This is called the *shebang*, and it is a special symbol that when placed on the first line of any file means that it is the path to the script interpreting engine we are to use. In this case, we will be using the interpreter for BASH that is located in `/bin/bash` on your Pi and on just about every other Linux and UNIX operating system. This will tell the outside command prompt shell that when it is running this program, it should use the `/bin/bash` interpreter to execute it. You should always include some form of shebang at the start of your given program for any interpreted language.

Interpreted versus Compiled

This brings me nicely to the next point of order: interpreted vs. compiled languages. For these purposes, you won't need to have a deep knowledge of compiled languages, but it is always good to know that they exist and to be aware of some of the differences between an interpreted language and a compiled language. So while both sets of languages allow you to write your own programs for the computer to execute, they go about this in fundamentally different ways.

Interpreted languages (sometimes called *scripting languages*) write down a series of programmatic commands that are fed to an interpreting engine (as given by the shebang, for example) that will extrapolate their computational meaning and perform the functions that the program intended. This is the pen-and-paper equivalent of a recipe for a cake, whereby you write down the exact method for how to make the cake, and the interpreter cooks it for you and outputs tasty delicious cake.

Note This recipe nature is also much like a script for a play in that as long as you have the script and can read it you will perform the same actions each time. This is why interpreted languages are often known as scripting languages: their ultimate output is a collection of text that can be used to create the programmatic output the same way each time.

Compiled languages take a different approach. They also have a recipe (called *source code*), but instead of the recipe running, it will be taken to a specialized piece of software called a *compiler* that will take the recipe and create an executable package (called a *binary*) from it. This binary package is written in the computer's own machine language so when executed it performs all the actions itself. This is the equivalent of writing the plans for a specialized cake-making system and then installing it in your system and having it make the cake. In both cases you get cake, but the interpreted one may take a little longer to make from the moment you say go. The flip side of this is that creating the specialized cake-making machinery takes much more effort. Now you should understand why I said languages like C were overkill for small problems—C is a compiled language.

Output in BASH

Back to our application: if you were to save and run your program, it would not do anything for you. So the first order of business should be to make your program output something. This is accomplished with the echo statement, which is also available from your command prompt shell; you use echo to make the system output a given set of text. So let's add an output to your script to have it say hello to everyone out there in the big wide world. Modify your script so it looks like this:

```
#!/bin/bash
echo "Hello World!"
```

Now go ahead and execute your script; to do this, you will need to change the script's file permissions to allow it to be executed. My script is called `fizzbuzz.sh`, so the commands needed to make it executable and then to execute it are these:

```
$ chmod +x fizzbuzz.sh
$ ./fizzbuzz.sh
```

This code will generate the following output:

Hello World!

Congratulations! You have just written the world's most simple program, the *Hello World program*, which is the first program most software engineering students are given to create. It shows how to begin building a small piece of software and accomplishes the basic task of any software program: it generates output. In this case, the chain was that we gave it an input of "Please output 'Hello World!'" and the system computed that and then gave us the desired output: Hello World! I wrapped Hello World! in a pair of quotation marks. (This is done so that the line of Hello World! is treated as a string of characters instead of other commands to execute.) All languages need to make the distinction between what is a piece of data and what is an actual piece of programmatic logic. The quotation marks in this instance show that everything within them is considered to be a string of text.

You may ask, "So what if I want to output a quotation mark? Suppose I want to output "Hello World!" like that in quotation marks." Well, not a problem, we have you covered! All systems have the concept of what is called an *escape*. An escape is a special character that tells the interpreter or compiler to ignore the special properties of the next character and just treat it as part of a given string. In most systems, the escape character is a backslash (\). So if we wanted to make our Hello World! have quotation marks, we would modify it to be as follows:

```
#!/bin/bash
echo "\"Hello World!\""
```

This shows that we want to escape those inner pair of quotation marks, so if we execute, we should see the following:

```
$ ./fizzbuzz.sh
"Hello World!"
```

Fantastic! You can even escape an escape if you want to print a backslash as the output of something in your code. Remember this as you are writing software because this is one of the most common mistakes people make—having unescaped characters in their code, which causes the rest of their software to have issues because the strings are mixed up.

Note Probably the most commonly escaped character is \n, which is short for *newline*. This can be used to output text on the next line.

As an aside, this is where text editors that have syntax highlighting come into a realm of their own because they are an immense help in displaying which characters are part of a string and which aren't. Vim has syntax highlighting, as does nano (but not over SSH); unfortunately, the default GUI editor LeafPad doesn't, but there are a number of editors out there that do you simply need to find one that suits you. They also have other functions such as bracket matching and brace matching, but they won't help us much at this stage.

Recap

So what have we learned so far? We have covered how to start off a BASH application using the shebang so we can pick out the right interpreter. You learned how to print basic output to the screen and also how to use the escape character within a string. This brings us nicely to the next point on our list: variables.

Variables

Variables are an abstract representation of a given piece of data. Come again? A variable is a way that a computer can store a given piece of information so that it can be retrieved again. This is the way that most information within a program is stored and manipulated because most programs do not work like our simple Hello World! They take

various forms of information and turn it into other information. So let's depart from our first example and return to the problem of writing a FizzBuzz application. We need to start by counting from 1 to 100.

We could simply write each of the numbers, one per line, and then perform the FizzBuzz calculations on each, but that's no really a good and logical way of going about writing this program. What we need is an abstract variable that can contain the "number"; we can then run the FizzBuzz tests against this "number" and then we will simply increase the "number" by one and rinse and repeat until the number 100 has been tested. Then we are done!

So now you should have some understanding of what a variable is. It's a way that we can represent any single piece of information. We don't have to have a specific piece of information because the variable is just that—variable in nature? However, there are a few limitations of variables to consider. First, there are many different types of variables defined in the wider world of programming, such as Boolean, integer, double precision integer, long integer, character, string, floating point integer, vector, array, and so on. For BASH, there are only a few variable types, but we should be aware of the most important of the other types of variables because this can come into play with other programming languages:

- First is the *integer*, which is a numeric variable that is designed to represent only a number (and in most cases only a whole number). This type is defined to allow easy use of mathematical and other numeric functions such as addition, subtraction, multiplication, and so on.
- Second is the *string*, which is a type of written data and represents a long stream of written characters. Strings are normally related to user input and output because they are far too cumbersome to work with inside an application due to the highly variant nature of what can be included in a string of text.
- Third is the *array*, which is a special type of variable because it is a metavariable and is basically a carrier for a number of other variables. The subvariables within an array are called *elements* and are referenced by adding a pair of square brackets to the end of the array variable and a number that corresponds to the element of the array in question. Arrays do not start at 1; they instead start at 0 because 0 is logically the first number in the counting sequence. So if you want to retrieve the fourth element from an array you would be adding [3] to the end of your array variable.

So now that you know what a variable is, we can start to form our basic program logic for solving the FizzBuzz problem.

Logical Operation: if

Now we have both the ability to output something and to store something, which covers the first and last portions of what one needs to do with a computer program input and output. What we need to do now is to actually perform some calculations on our variables so that we can have our program do something meaningful. This is where we break out into logical operations. In programming there are two basic logical operations that you will need to become intimate with. These two logical operations form the basis of most of the programming that will be done anywhere. They do this because they allow people to perform tests and take divergent actions depending on the outcome of the test.

The first of these statements is the *if* statement. An *if* statement performs just as its name implies—if something is true, an action is performed. It is the test nature of an *if* statement that provides it with all its power because when combined with the use of variables we can test just about anything we want. So looking at the FizzBuzz example, we know we have our initial variable "number" that will represent the counting numbers between 1 and 100. We will need to perform tests on "number" to see if it is divisible by 3, divisible by 5, or divisible by both 3 and 5. If it matches one of these conditions, we need to output the correct word.

So now you need to understand how an *if* statement is written into a program. In BASH, the basic syntax is this:

```
if [ <TEST> ]; then
    DO SOMETHING
fi
```

The preceding is a simple BASH `if` statement and it really isn't that complex. You will notice that along with the `if` and the `test` and doing something, there is also a `fi`. The `fi` denotes the end of things that are to be performed, as you can perform multiple actions within an `if` statement. So, now we have our `if` statement, so we need to make it do things. So let's have a quick look at some mathematical operations.

Note Remember to leave a space on either side of the open square bracket, a space before the close square bracket, and a space after the semicolon; otherwise, you will get a syntax error.

Test Based Arithmetic

We know that we need to have a test for equality, so how does someone confirm if any given number is divisible by another given number? The simple answer is to check that when divided by a given number, the result is a "whole" number without any "remainder." This is, in fact, as hard a concept to program as it sounds. Whereas multiplication, addition, subtraction, and even division will give finite answers that can be tested upon, checking whether something is divisible by another number requires a different operation. Thankfully, there is a specialized mathematical operation designed for just this issue: the *modulus*. The modulus returns the remainder of a division of any two numbers and is denoted by the symbol `%`. So the modulus gives you results like these:

```
12 % 3 = 0
12 % 5 = 2
```

And suddenly our test has appeared: if number modulus 3 is equal to 0 we output Fizz, if number modulus 5 is equal to 0, we output Buzz, and if number modulus 3 and 5 is equal to 0, we output FizzBuzz. So, if we combine all of these, we get the following:

```
if [ number % 3 = 0 ]; then
    echo "Fizz"
fi
if [ number % 5 = 0 ]; then
    echo "Buzz"
fi
if [ number % 3 = 0 ]; then
    if [ number % 5 = 0 ]; then
        echo "FizzBuzz"
    fi
fi
```

Okay, that's a good-looking attempt, but there are a number of problems that relate to things about BASH we still need to cover in order to make this into valid executable BASH.

Let's start with the simple variables. In BASH, there are two modes that variables operate in. First is *assignment mode*, which occurs when we are creating a new variable and giving it a value. In these instances, the variables work just like the previous number, so we could create our `number` variable as 1 this way:

```
number=1
```

This is fine and looks like our example, but when we want to use the current value of the `number` variable, we need to access the variable, which is done in *dereferencing mode*. In BASH, this means adding a dollar sign (\$) to the front of every variable for which we want to use the value.

To recap, there are two modes for a variable: one for loading the variable (called *assigning*) and one for pulling the values out of the variable (called *dereferencing*). In the first mode, we use the variable as is; in the second, we use a \$ in front of the variable. Okay, not a problem. Let's update the code:

```
if [ $number % 3 = 0 ]; then
    echo "Fizz"
fi
if [ $number % 5 = 0 ]; then
    echo "Buzz"
fi
if [ $number % 3 = 0 ]; then
    if [ $number % 5 = 0 ]; then
        echo "FizzBuzz"
    fi
fi
```

So our code looks a bit better, but you may have noticed the next problem that occurred when we assigned a number the value of 1. The use of the equal sign (=) is used to assign values to variables, so in this context it looks like we are trying to assign the variable 5 with the number 0, which will cause all kinds of problems with the test in the `if`. Most languages handle this using a pair of equal signs (==) to denote a test for equality, but in BASH there are a number of special arithmetic testing operators we are given to test this:

- `-eq` for equality
- `-ne` for not equal
- `-gt` for greater than
- `-lt` for less than
- `-ge` for greater than or equal to
- `-le` for less than or equal to

We can take advantage of these operators in our code to do the comparison of the result of our modulus and 0. So make these changes to your code again and you should have the following:

```
if [ $number % 3 -eq 0 ]; then
    echo "Fizz"
fi
if [ $number % 5 -eq 0 ]; then
    echo "Buzz"
fi
if [ $number % 3 -eq 0 ]; then
    if [ $number % 5 -eq 0 ]; then
        echo "FizzBuzz"
    fi
fi
```

Finally, when doing an arithmetic operation (such as modulus), we need to tell BASH that this is meant as an arithmetic operation. To do this, we use a \$ and a pair of parentheses around the actual bit of arithmetic so the operation of $5 \% 3$ means that we need $\$((5 \% 4))$, and we can make the change:

```
if [ $($number % 3) -eq 0 ]; then
    echo "Fizz"
fi
if [ $($number % 5) -eq 0 ]; then
    echo "Buzz"
fi
if [ $($number % 3) -eq 0 ]; then
    if [ $($number % 5) -eq 0 ]; then
        echo "FizzBuzz"
    fi
fi
```

There you have it; you have written three simple tests for the FizzBuzz application; and used if statements, a variable, the print, and the special arithmetic operator. That's a lot of code, so let's test it and make sure what we have got now works!

We can do this simple test by creating the number variable with a single value for now. Let's assign it to 15 so we can guarantee some output. With the number variable added, your script should look like this:

```
#!/bin/bash
number=15
if [ $($number % 3) -eq 0 ]; then
    echo "Fizz"
fi
if [ $($number % 5) -eq 0 ]; then
    echo "Buzz"
fi
if [ $($number % 3) -eq 0 ]; then
    if [ $($number % 5) -eq 0 ]; then
        echo "FizzBuzz"
    fi
fi
```

When executed, the output will be this:

```
$ ./fizzbuzz.sh
Fizz
Buzz
FizzBuzz
```

Troubleshooting

Oh dear. That output is a problem. We said all three words at once, not just the FizzBuzz when the number was divisible by both 3 and 5. We just failed! Argh! Okay, take a few deep breaths; this isn't the end of the world. In fact, it's a great time to introduce a new feature. The if statement does something when the test is true, but there is also

another part we can add: the `else`. The `else` is the flip side of the `if`, and the logic flows like this: *if* the test is true, do something, *else* do a different thing. We can use a combination of these two to perform our tests for FizzBuzz. The syntax of an `if, else` statement is as follows:

```
if [ TEST ]; then
    DO SOMETHING
else
    DO A DIFFERENT THING
fi
```

The syntax is almost identical to the original `if` statement, so it is easy to make a few changes to make your `if` statements into `if, else` statements. But now we run into the much bigger problem—and one that is the cause of most people's problems with programming; it's the logical order flow. We need to create a sequence of tests so that we can determine if something is divisible by 3, divisible by 5, or divisible by both; and perform a wholly different action in each case. This kind of thinking is what makes programming so hard for so many people, and there is no trick to it: you simply have to work it out. So let's have a go at working this one out together.

If we find something that is divisible by 3 or 5, and output something before checking whether it is divisible by both, we run into the potential for outputting a Fizz, a Buzz, and a FizzBuzz, which is not what we want. So let's start with a test for both:

```
if [ $($number % 3) -eq 0 ]; then
    if [ $($number % 5) -eq 0 ]; then
        echo "FizzBuzz"
    fi
fi
```

Okay, now we have FizzBuzz, but what if the answer *is* divisible by 3 but not by 5 (which is the case for Fizz)? Then we can just output the Fizz. Thus, if we add an `else` to the divisible by 5 tests, we will be guaranteed a number that is divisible by 3 and not by 5. This means we avoid the issue when we output both Fizz and FizzBuzz at once. We can use the new `else` statement here to accomplish this, so when we add the `else` to the divisible by 5 test, we get this:

```
if [ $($number % 3) -eq 0 ]; then
    if [ $($number % 5) -eq 0 ]; then
        echo "FizzBuzz"
    else
        echo "Fizz"
    fi
fi
```

Fantastic. Now we only have one case left: when a number is not divisible by 3 but is divisible by 5. We can do this test in another `else` statement off the divisible by 3 test. By doing this as part of an `else` statement and adding a second divisible by 5 test, we can ensure that a number is not already divisible by 3 when we output when it is divisible by 5. This means we avoid running into the issue from before, so add this test in and your code becomes the following:

```
if [ $($number % 3) -eq 0 ]; then
    if [ $($number % 5) -eq 0 ]; then
        echo "FizzBuzz"
    else
        echo "Fizz"
    fi
else
```

```

if [ $($number % 5) -eq 0 ]; then
    echo "Buzz"
fi
fi

```

And that should do it! We have ourselves a new series of tests to run. Go ahead and replace the old statement in your script and run this new one:

```
$ ./fizzbuzz.sh
FizzBuzz
```

Perfect! That's the correct result for 15! You can change the value of number and execute a few more times to test, but this test should stand up well because we have thought through our programming. Now we need to have the numbers from 1 to 100 counted. To do this, we will use the second logical operation!

Logical Operation: Loop

A loop is a special type of logical operation that functions at its core like an `if` statement, but instead of running the code if a statement is true, it will run the statement over and over again as long as the statement is true. This is how we will be able to increase the value of number by 1 over and over again to count to 100. A loop in BASH uses the following syntax:

```

while [ TEST ]; do
    DO SOMETHING
done

```

A loop is very similar to an `if` statement by design. Both are testing that some condition is true and both will then execute a section of code. The difference is that an `if` statement will execute only once and a `while` loop will execute until something is no longer true. This brings be to the first and biggest warning about loops: if you forget to have a leaving condition for your loop, then you won't EVER leave it. This is called an *infinite loop*, in which your program will get stuck, do nothing, and run forever. There are a few situations when this is desirable but not many. Luckily for us we have the `Ctrl+C` shortcut that will send a terminate signal to whichever program we are running and we can use it to snap programs out of their infinite loops.

You might be thinking, “Oh, so that's why so many programs get stuck!” This is 100% true. Sometimes cases that you may not have considered arise, and your program can wind up running in a loop forever, so consider yourself warned.

For our FizzBuzz program, we want to count the numbers 1 to 100, so we can do a simple test of numbers less than or equal to 100. So for all situations in which our number is less than or equal to 100, we will execute the loop. So let's go ahead and write the loop around our existing block of tests. When done, it should look like the following. *Don't execute it yet* because we have no way to increase the number from 1 to 100 (which means we'll get an infinite loop because we can never reach 100 and leave the loop).

```

#!/bin/bash
number=1
while [ $number -le 100 ]; do
if [ $($number % 3) -eq 0 ]; then
    if [ $($number % 5) -eq 0 ]; then
        echo "FizzBuzz"
    else
        echo "Fizz"
    fi
fi

```

```

else
    if [ $($number % 5) -eq 0 ]; then
        echo "Buzz"
    fi
fi
done

```

Okay, we have our code block, now we need to add the increasing number. We simply need to increase the value of `number` by 1 each time it goes around the loop. But we also want to perform all the calculations prior to changing the number. We can do this by using the arithmetic operations and assignment operation we have already. And we can add this block to the bottom of the code before the `done` statement:

```
number=$((number +1))
```

Once you have made this change, execute your script! It will run and give you an output like this:

```

Fizz
Buzz
Fizz
Fizz
Buzz
Fizz
FizzBuzz
...

```

This is great, and it looks like it worked, but we should go ahead and print the current value of `number` for each trip round the loop. We could output the value at the start before we Fizz or Buzz or FizzBuzz, but that means we will have numbers on each line. Better if we can integrate it into our `echo` statements. To do this, we simply add the `$number` variable within the string of text to print. You learned earlier that the quotation marks mean everything is treated like a string of text, but there is one marker that supersedes it: the `$` sign operator. This operator is used to make your life easier when you want to print out data from a variable in a block of text.

We should also add one final `else` statement, which will take care of the output of all the “other” numbers. Try and work it out for yourself; then check against this code:

```

#!/bin/bash
number=1
while [ $number -le 100 ]; do
    if [ $($number % 3) -eq 0 ]; then
        if [ $($number % 5) -eq 0 ]; then
            echo "$number - FizzBuzz"
        else
            echo "$number - Fizz"
        fi
    else
        if [ $($number % 5) -eq 0 ]; then
            echo "$number - Buzz"
        else
            echo $number
        fi
    fi
number=$((number +1))
done

```

And there you have it! A working FizzBuzz solution. Congratulations!

Troubleshooting

Having trouble with your code can be a bit of a nightmare. Probably the best tool at your disposal is the echo command. You can output anything you want, so if you are unsure of where your application is within your code, write an echo to output stuff and then check. Are you not entering the loop correctly? Output the variables that go into the test and compare by hand. Not sure why the if statement didn't work? Add an else, output the full test case, and see if you went in correctly. Using these outputs to trace where you are in your code is the best way of doing diagnostics. Otherwise, pay attention to any errors; most should provide a line number and will tell you (in a manner of speaking) what's wrong with what line. If you are in doubt, feel free to Google around because it's almost certain someone, somewhere has had the same error before you and asked for a solution.

Practical BASH: An Init Script

You should be aware by now that most Linux applications are started by a special script called an init script. These scripts aren't magic; they are simply scripts in BASH that take a given value and perform a series of actions based on what they are told to do—just as any good program does. The init scripts do have a small amount of special information in them, but none of this is actually program logic.

Let's go through one together and examine exactly how one of these init scripts works. Following is the init script for XBMC, which you will see again later in this book:

```
#!/bin/bash
### BEGIN INIT INFO
# Provides:          xbmc
# Required-Start:    $all
# Required-Stop:     $all
# Default-Start:    2 3 4 5
# Default-Stop:      0 1 6
# Short-Description: Start XBMC
# Description:       Start XBMC
### END INIT INFO
DAEMON=/usr/bin/xinit
DAEMON_OPTS="/usr/lib/xbmc/xbmc.bin"
NAME=xbmc
DESC=XBMC
RUN_AS=root
PID_FILE=/var/run/xbmc.pid
test -x $DAEMON || exit 0
set -e
case "$1" in
  start)
    echo "Starting $DESC"
    start-stop-daemon --start -c $RUN_AS --background --pidfile $PID_FILE --make-pidfile --exec
$DAEMON -- $DAEMON_OPTS
    ;;
  stop)
    echo "Stopping $DESC"
    start-stop-daemon --stop --pidfile $PID_FILE
    ;;
  *)
```

```

restart|force-reload)
    echo "Restarting $DESC"
    start-stop-daemon --stop --pidfile $PID_FILE
    sleep 5
    start-stop-daemon --start -c $RUN_AS --background --pidfile $PID_FILE --make-pidfile --exec
$DAEMON -- $DAEMON_OPTS
;;
*)
    echo "Usage: /etc./init.d/$NAME{start|stop|restart|force-reload}" >&2
    exit 1
;;
esac
exit 0

```

The first thing you will notice is there are a lot of lines that start with #. These lines are comments; because # is the comment symbol in BASH, any line that starts with a # will not be executed as part of the application. The top lines of INIT_INFO are in fact very important for an init script. These are special comments that can be processed to show how the application controlled by the script is to be run. The blocks look like this:

```

### BEGIN INIT INFO
# Provides:          xbmc
# Required-Start:    $all
# Required-Stop:     $all
# Default-Start:    2 3 4 5
# Default-Stop:      0 1 6
# Short-Description: Start XBMC
# Description:       Start XBMC
### END INIT INFO

```

These blocks provide a few bits of description and say for which application they provide a function; in this case, the application provided is XBMC. They also mention what applications are needed to be running before this application can start or stop. The \$all symbol means that this application will be started last, so that it is guaranteed that all other applications on which it can depend are started before it. Probably the more important set of operators are Default-Start and Default-Stop. These operators correlate with the Linux system's *run levels*, which govern the various stages of the Linux system boot process.

The Linux specification defines the following run levels:

- *Level 0*: Halt; shut down and power off
- *Level 1*: Single User Mode; only basic system functionality (used for repairs)
- *Levels 2 % 3*: Multi User; adds networking and multiple user functionality
- *Level 4*: User Defined
- *Level 5*: Normal System running state
- *Level 6*: Reboot

The numbers next to the start and stop switches correlate to these levels, so XBMC will start in run levels 2–5 and will shut down in levels 0, 1, and 6.

Pick and Match with the case Statement

Once all the initialization is done and the script has set some initial variables, it performs a few quick tests and then moves on to the `case` statement. A `case` statement is very much like a series of running `if` statements against the same variable. Insofar as you take a given variable and for each of a series of given values for said variable, you perform a function. In most `init` scripts, this is used to confirm which action it is to perform.

The `case` statement in this `start` script is working on the special `$1` variable. This variable represents the first thing passed to the script from the command line; so when trying to start the `xbmc` application, you will execute with this:

```
/etc./init.d/xbmc start
```

This gives the `$1` variable the value of `start`. The `case` statement offers a number of different options of the potential value of the `case` statement:

- `start`: Self-explanatory.
- `stop`: Self-explanatory.
- `restart|force-reload`: The values of `restart` or `force-reload`, which perform the stop application and then the start application in one process. The pipe (`|`) between them is used to denote an OR operation, which means that if either of the values here match, we treat this as having been a match.
- `*`: Anything else.

After each `case` value, there is a small block of code that is executed in that case, followed by a pair of semicolons. These semicolons are like a `fi` and are used to denote the end of a block of `case` statement code.

Application within Application: Forking

Once within the `case` statement, you can see that each line will perform some action with the `start-stop-daemon`. But this isn't a variable or anything special; it's another program. The most powerful part of BASH is the bit we haven't gotten to yet. BASH has the capability to execute any command-line process from within your programs just by using their names. Because BASH is the shell we use most of the rest of the time, this actually makes a lot of sense.

The last function of an `init` script: it will use the `start-stop-daemon` to create a new process for you. This action is called *forking* because you are forking something off of the currently running application. The `start-stop-daemon` is a small application that is used to fork off processes and will manage their running for you. The `start` script corrals the arguments for it, works out what action needs to be performed, and then tells `start-stop-daemon` which action needs to be taken against which application.

Now that you are armed with this knowledge, you should be able to write your own simple `init` scripts to start and stop any process you want by providing the starting comment block to describe how the script should be loaded, having a `case` statement determine which action to take, and finally executing the `start-stop-daemon` process to launch and manage your application!

Update the Run Files

Now you have a fully working script ready to go. You need to use `chmod` to set the executable flag on your file and then you can test it with the `start` and `stop` commands. It does as intended—fantastic! Now we need to add it into the system's boot logic. Historically, this involved linking the file in a special way into the various run levels. However, this has been made much easier with the `update-rc.d` application. The `update-rc.d` application takes advantage of those special comment headers we added to our BASH scripts in order to know which run levels should have which shortcuts.

All we need to do to add the `init` script to the boot sequence is to issue the `update-rc.d` command with the name of the `init` script we want, which is `xbmc` in the example. Then we need to add the argument to say that we should use the defaults from the comments in the script; the argument unsurprisingly `defaults`. This gives us the command `update-rc.d xbmc defaults`, which will need to be executed as root. When run, the outputs should be as follows:

```
$ sudo update-rc.d xbmc defaults
update-rc.d: using dependency based boot sequencing
```

And that's it; if you reboot, your application should run on boot!

Creating Your Own `init` Script

We have covered how `init` scripts are created; now we can create one. The process is relatively simple because we simply need to re-create the structure and input our own code in the correct place to perform the functions we need. You can have these `init` scripts do quite literally whatever you want, but for this example I will be creating a file with the `touch` command and then deleting it with `rm` when done. The `touch` command is used to simply create a file on disk with nothing in it; if it runs on an existing file, it will change the last modified time of the file. As mentioned earlier, you can have your programs managed by `start-stop-daemon` or you can simply perform system functions, as I will do in my example.

As you will recall, the first thing we need in this script is the shebang, which is the starting point for all BASH scripts and is thus the best place to start.

```
#!/bin/bash
```

Now that we have our shebang, recall the XBMC start script we worked with earlier: the first thing in that script is the opening comment block that gives us details about which `init` levels will run this script, and so on. So let's add that next. You should fill in all your own details here. My script will be called `touchfile.sh` and will provide the Touch File service. My opening block of script will look like this:

```
### BEGIN INIT INFO
# Provides:          touchfile
# Required-Start:    $all
# Required-Stop:     $all
# Default-Start:    2 3 4 5
# Default-Stop:      0 1 6
# Short-Description: Run TouchFile
# Description:       Run TouchFile
### END INIT INFO
```

Now that I have defined the comment block, I can define my variables. They will be the variables that will cover what my application does. Anything you can think of needing in order to start this that can be configured should be defined as a variable. In the case of my `touchfile` command the thing I need to have defined as a variable is the filename of my `touchfile`. Remember that this is your code, and thus you can define literally anything you desire as a variable. My variable is this:

```
TOUCHFILE="/var/tmp/touch.file"
```

Now that my touchfile is defined, I need to add the case statement to cover the available actions. In this case, I want to have the following actions:

- Start: Touch the file
- Stop: Remove the touched file
- Restart: Remove and then re-touch the file
- Reload: Touch the file
- Default Case: Tell people about how to use this file

This means I will have four case statements within my case. I will define the case statements (sans code to perform actions) as such:

```
case "$1" in
  start)
;;
  stop)
;;
  restart)
;;
  reload)
;;
  *)
;;
esac
exit 0
```

Now that we have the case entries, we just need to make them do something. I mentioned earlier what I want to do in each of the cases, so let's write out the commands into each block and combine them into the script we have already. This will give us the following script:

```
#!/bin/bash
### BEGIN INIT INFO
# Provides:          touchfile
# Required-Start:    $all
# Required-Stop:     $all
# Default-Start:    2 3 4 5
# Default-Stop:      0 1 6
# Short-Description: Run TouchFile
# Description:       Run TouchFile
### END INIT INFO
TOUCHFILE="/var/tmp/touch.file"
case "$1" in
  start)
    echo "Creating $TOUCHFILE"
    touch $TOUCHFILE
;;
  stop)
    echo "Removing $TOUCHFILE"
    rm $TOUCHFILE
;;
esac
exit 0
```

```

restart)
    echo "Recreating $TOUCHFILE"
    rm $TOUCHFILE
    touch $TOUCHFILE
;;
reload)
    echo "Re-Touching $TOUCHFILE"
    touch $TOUCHFILE
;;
*)
    echo " Usage: touchfile.sh <start|stop|restart|reload>"
;;
esac
exit 0

```

And that's it. The code is a simple init script, which accepts the `start`, `stop`, `restart`, and `reload` commands and outputs what it is doing and performs actions based on what arguments we provide it. This can be placed in `/etc./init.d` and can be set to be executable with `chmod +x`. Once they are done, you can test the script to ensure that it works. Finally, you can use `update-rc.d` to add this to your system boot process as we did earlier. It is in this fashion that you can write your own start scripts to run applications as you design.

Security and User Management

Security is one of the most neglected areas of systems administration. Because it is perceived as a black art to many, it often gets neglected.

The Rules of Raspbian Security

There are some very simple ways to make your system secure; just follow a few rules:

- Don't log in as root, but if you must, log out when done.
- Use `sudo` whenever possible for admin tasks.
- Choose a nontrivial password—one that is long enough and complex enough that it cannot be readily guessed or worked out by hand without taking a long time.
- Change your password regularly.
- Review your system logs regularly, specifically `/var/log/auth.log` because it lists all user authentications to your system.
- If you don't need an application running, don't run it.
- Expose as little of your system to the Internet as possible.
- Restrict file permissions whenever possible.

Most of these rules will seem like common sense and that is because they are. Yet most people seem to fall into the trap of thinking that they won't be attacked and ignore most of this advice. Now that you know these rules, you are probably saying, "I will change my password if only you would show me how!" Changing your password from the Linux command line is relatively easy: you simply use the `passwd` command and you will be prompted to enter a new password (twice, so that you can't misspell it by accident the first time).

You can also change the password for other users as root. To do this, you simply append the name of the user whose name you are changing as the first argument of the command. This is incredibly useful when configuring new user details or resetting passwords.

Adding a New User

Because you aren't supposed to use root at all times, it's important to be able to create new users and understand how user creation takes place. In Linux operating systems, the users are governed by the file `/etc./passwd`. You can have a look for yourself and you will see a number of lines like this:

```
root:x:0:0:root:/bin/bash
daemon:x:1:1:daemon:/usr/sbin:/bin/sh
```

These lines show the user, X for the password, the user's ID number, the user's group ID number, the user's group identifier, the path of the user's home directory, and the user's shell. These are the basics of what is needed for a user to function within the Linux environment.

While this file has a password field, it is no longer used and normally contains a placeholder (in this case, X). The passwords are stored in the shadow file, `/etc./shadow`. The shadow file was separated so that only root could access password data and all the more common user data could be read by other more generic system applications.

Now that you understand how Linux stores users, add a new one called `raspberry`. To add this user, we will use the `useradd` command and a few arguments. There is another command, `adduser`, which performs user additions, but instead of being supplied with arguments will prompt you for details. `adduser` is a great command to use in a pinch when you forget a necessary argument for `useradd`.

First we want to specify the user's home directory, which is done with the `-d` flag and the full path of the directory we want to use. Most home directories for users are `/home/<username>`, and we will do this for `raspberry`. Currently, `/home/raspberry` doesn't exist, but before you rush out and add it, we can use the `-m` flag to tell `useradd` to do it for us! This gives us this command:

```
$ sudo useradd -d /home/raspberry -m raspberry
```

We could go a step further and even specify the group with `-g <groupname>` to set the primary group for the user. Or even specify the shell that the user would use with `-s <shell path>` or finally a password with `-p <password>`. Once you have executed your command, you can go and inspect the changes to `/etc./passwd`, where you should see something like this at the bottom of your file:

```
raspberry:x:1001:1001::/home/raspberry:/bin/sh
```

That's it; you have now successfully added a new user! If you didn't specify a password, you can use the `passwd` command we covered earlier to change your new user's password and then you're set to go. However, let's say that you forgot to set the shell path with `-s` earlier and you want to change that so you start with `/bin/bash` rather than `/bin/sh`. Well, first you could edit the `passwd` file to change the shell or you can use the `usermod` command. The `usermod` command functions exactly the same as the `useradd` command, right down to the arguments; it simply adjusts the values. So run the following:

```
$ sudo usermod -s /bin/bash raspberry
```

We can expect the shadow file to be changed and, yes it is:

```
raspberry:x:1001:1001::/home/raspberry:/bin/bash
```

Fantastic! With these tools, you should be capable of creating new user accounts to allow people to access your Pi.

Summary

This chapter covered quite a number of things. We introduced SSH, DHCP and DNS. We covered some of the basics of how these systems function and how daemons are launched via `init` scripts. We then dove deep into learning BASH so that we could write our own `init` scripts. Finally, we covered some common security do's and don'ts; then went on to cover how the basics of Linux systems manage users and their passwords. We even went so far as to create a whole new user ready for use.

These skills should enable you to manage the startup and networking of your Pi and should have given you an introduction into the wonderful world of software development!

CHAPTER 8



A LAMP of Your Own

One of the key things that many people who come to Linux want to know is how to establish themselves by setting up a small website of their own. And many have difficulty finding a place to start. Which OS, which applications, what hardware, and how exactly to do all this? All these are legitimate questions as there are so many facets to the greater Linux environment that it is very easy to get lost.

Our goal is to provide the first few stepping stones on your way to being able to navigate the Linux environment. One of the key stepping stones on this path is to show you how to install, configure, and maintain a web server. Although you will be able to connect this server to the Internet, in these examples we will only cover using your Pi as a local network server (in other words, for running an intranet). In addition, there is also the extra challenge of learning how to make a basic interactive website. While this may seem a lot to take on, and others devoted entire books and lives to finding the absolute optimum solution this problem, we aim to cover it for the Raspberry Pi in one chapter.

To this end we will be using a LAMP. No, not one of those bedroom luminescence devices, but a Linux, Apache, MySQL, and PHP system designed to make delivering web content easy. This is one of the easiest ways to create your own web server and there are a number of variations on the original LAMP, which include, but are not limited to, the following:

- WAMP (Windows, Apache, MySQL, PHP)
- MAMP (Mac OS, Apache, MySQL, PHP)
- Replacement Ps with Python or Perl
- Replacement M with MariaDB

So, as you can gather from the amount of mixing, matching, and improving that people have done, it is a great way to get a powerful web server up and going. Best of all, it is simple to set up and can be run on just about any hardware.

One of the most regularly cited potential uses of the Pi is to enable the ability to build, run, and maintain a website. What better way to inspire a generation of young developers than to teach them how to make their own websites with a minimum of cost!

First Steps

Now it is time to get real about what we can accomplish here. Within this chapter, we will provide you the skills to set up a full website stack and show you how to integrate the functions of a MySQL database into a website using PHP to dynamically generate website content. However, there are a few things we won't be able to do for you here:

- Set up and configure a DNS or domain naming
- Advanced SQL and database administration
- Full coverage of HTML and HTML development methods

- Full coverage of PHP development and PHP development methods
- JavaScript
- CSS

There are a number of resources that you can take advantage of to get help with these other areas, including other great and weighty texts by Apress. We still have our work cut out for us, so let's get started.

L is for Linux

There are only a few prerequisites for setting up a LAMP stack. You will need to

- Have the Raspbian OS (the *L* of LAMP) installed and configured
- Have the memory split set to 240/16 (run `sudo raspi-config`)
- Be familiar with working within the Raspbian OS using the command line
- Have a basic understanding of how to write simple software
- Have your Pi connected to a network and the Internet (how else can we serve content if we aren't?)

There are only a few other provisos for Raspberry LAMP. We will be doing all of this from a shell. Given that the website will be viewed from outside the Pi, there is no need to use any extra resources to run the GUI. Now that we know what we need to get going, let's do it.

A is for Apache

Probably the most important part of any web server is the web server application. A web server is an application that will take requests for a website and then return the requested web page content. Most web servers can also provide a large number of other functions that can help enhance their capability to provide web content to end users. For our LAMP stack, we will be using the Apache web server.

Apache is considered by many to be the world's premiere web server. It was first released in 1995 and has become renowned for playing a role in the growth of the World Wide Web by providing a simple, powerful, and free web server to the masses.

As reinforcement of this fact, it is estimated that Apache delivers between an estimated 54 to 58 percent of the world's websites.

There are two stories that are told about how Apache came to be named:

- The first is that it was named for it being "A patchy web server" when first developed all those years ago.
- The second (and much better from a storytelling perspective) is that Apache is said to be named for the Apache tribes of America.

Being named for a tribe is not even a misnomer as acting like a tribe is one of the key features of Apache. Apache's first task when it starts is to create its own small tribe of "workers" that are tasked with the actual serving up of your website's content. Now that you have a clearer understanding of what Apache does, the first step becomes to get Apache installed, up and running.

For our installs, we will rely (as ever) on the ever-faithful `apt-get` tool:

```
$ sudo apt-get install apache2
Reading package lists... Done
Building dependency tree
```

```

Reading state information... Done
The following extra packages will be installed:
  apache2-mpm-worker apache2-utils apache2.2-bin apache2.2-common libapr1
    libaprutil1 libaprutil1-dbd-sqlite3 libaprutil1-ldap ssl-cert
Suggested packages:
  apache2-doc apache2-suexec apache2-suexec-custom openssl-blacklist
The following NEW packages will be installed:
  apache2 apache2-mpm-worker apache2-utils apache2.2-bin apache2.2-common
    libapr1 libaprutil1 libaprutil1-dbd-sqlite3 libaprutil1-ldap ssl-cert
0 upgraded, 10 newly installed, 0 to remove and 71 not upgraded.
Need to get 1,348 kB of archives.
After this operation, 4,990 kB of additional disk space will be used.
Do you want to continue [Y/n]?

```

Now that you have installed Apache, we should go and validate that it is all up and running. There are a few ways we can do this. First, run this:

```
$ ps -ef | grep apache
root      2306      1  0 Sep17 ?          00:00:09 /usr/sbin/apache2 -k start
www-data  2309  2306  0 Sep17 ?          00:00:00 /usr/sbin/apache2 -k start
www-data  2311  2306  0 Sep17 ?          00:00:00 /usr/sbin/apache2 -k start
www-data  2315  2306  0 Sep17 ?          00:00:00 /usr/sbin/apache2 -k start
```

From this output, you can see that Apache is indeed up and running. From this point, there are another few important things to note. The first process listed with the following line is different from the others:

```
root 2306 1 0 Sep17 ? 00:00:09 /usr/sbin/apache2 -k start
```

This is the chief of the Apache tribe mentioned earlier. It is owned by root (as indicated in the first column), and its process ID (PID) is 2306. The other three are owned by a user called www-data, which is a user specifically for Apache. While each worker has a different PID, they all have a PPID (parent process ID) of 2306, meaning that the Apache chief process made them (its PID is 2306).


Now for the second and by far more fun test; on this test, we will actually get Apache to display its default starting web page! For this, you can either log in to the GUI or you can get the system's IP address using ifconfig. Your output should look like the one following with the relevant IP addresses highlighted:

```
$ ifconfig
eth0      Link encap:Ethernet  HWaddr b8:27:eb:8a:46:ba
          inet addr:10.0.0.20 Bcast:10.0.0.255  Mask:255.255.255.0
            UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
            RX packets:213812 errors:0 dropped:0 overruns:0 frame:0
            TX packets:5119 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:1000
            RX bytes:19226371 (18.3 MiB)  TX bytes:495394 (483.7 KiB)

lo       Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
            UP LOOPBACK RUNNING MTU:16436  Metric:1
            RX packets:8 errors:0 dropped:0 overruns:0 frame:0
            TX packets:8 errors:0 dropped:0 overruns:0 carrier:0
            collisions:0 txqueuelen:0
            RX bytes:1104 (1.0 KiB)  TX bytes:1104 (1.0 KiB)
```

Assuming your system is (like mine) attached to a network, you can use the first IP address (10.0.0.20). This address in the `eth0` block represents its network port. The address in `lo` block is the loopback address, which is used for internal self-reference and is always 127.0.0.1.

The next test is to access the Apache default web page; and to do so, you simply need to input one of those IP addresses into a browser (assuming that you are on the same network as the Pi, of course!). If you want to test using the Pi's onboard browser, that's completely fine, too, and you can choose between either address (10.0.0.20 or 127.0.0.1). So what are you waiting for? Open a browser window and have a look at the default web page, as shown in Figure 8-1.



It works!

This is the default web page for this server.

The web server software is running but no content has been added, yet.

Figure 8-1. It works!

As you should by now be aware, it works! This is the default web page for the Apache web server and as it suggests there is no content added yet, we will be getting to that later. For now, let's have a look at how that one default page is generated, as this will give you some insight into how Apache truly works.

Apache Configuration

Getting your head around the Apache config is, to be frank, a bit of a nightmare. If you're game, go ahead and have a look in `/etc/apache2` for an idea of what I mean. There are a number of folders in there with different meanings and how those folders interrelate changes how your Apache server instance works. Thankfully, though, once you understand where you need to make changes, everything becomes much easier.

The first file of note in the Apache configuration is the `apache2.conf` file. This is the governing config file that says where each of the other separate pieces are located, and how the application should load them. The second things to be aware of are the `mods-available` and `mods-enabled` folders. If you have a look at the `mods-available` file, you will see that there are a large number of `.load` and `.conf` files. These files represent where Apache should load a module from and any configuration information this module needs, respectively.

Next is `mods-enabled`, an initial look over this folder will make it seem just like `mods-available`. However, if you run an `ls -l` on this directory, you will see this:

```
$ ls -l /etc/apache2/mods-enabled/
total 0
lrwxrwxrwx 1 root root 28 Sep 17 21:44 alias.conf -> ../mods-available/alias.conf
lrwxrwxrwx 1 root root 28 Sep 17 21:44 alias.load -> ../mods-available/alias.load
lrwxrwxrwx 1 root root 33 Sep 17 21:44 auth_basic.load -> ../mods-available/auth_basic.load
lrwxrwxrwx 1 root root 33 Sep 17 21:44 authn_file.load -> ../mods-available/authn_file.load
...

```

This shows that, in fact, each of the files in `mods-enabled` is actually a link to a file in `mods-available`! This is the first intricacy of the Apache configuration, while there are a number of available modules at any given time only a

certain number are actually loaded. Using symlinks like this means that only one copy of the configuration needs to be kept and maintained at a given time. And more to the point, any changes made to one will be made to both.

If you look inside the `sites-available` and `sites-enabled` folders, you will see the same layout, except there might be fewer sites. This brings us to the true object of our search: the first default site that is available within the Apache system. So open up either `/etc/apache2/sites-available/default` or `/etc/apache2/sites-enabled/000-default` because it doesn't matter; they are the same thing! So, now that you're inside, go over what everything in the config means. Many of them will be important for you later on if you wish to add additional websites or modify how websites are displayed.

```
<VirtualHost *:80>
    ServerAdmin webmaster@localhost
    DocumentRoot /var/www
    <Directory />
        Options FollowSymLinks
        AllowOverride None
    </Directory>
    <Directory /var/www/>
        Options Indexes FollowSymLinks MultiViews
        AllowOverride None
        Order allow,deny
        allow from all
    </Directory>
    ScriptAlias /cgi-bin/ /usr/lib/cgi-bin/
    <Directory "/usr/lib/cgi-bin">
        AllowOverride None
        Options +ExecCGI -MultiViews +SymLinksIfOwnerMatch
        Order allow,deny
        Allow from all
    </Directory>
    ErrorLog ${APACHE_LOG_DIR}/error.log
    # Possible values include: debug, info, notice, warn, error, crit,
    # alert, emerg.
    LogLevel warn

    CustomLog ${APACHE_LOG_DIR}/access.log combined
</VirtualHost>
```

The first line is probably one of the most important; it specifies that this web server is available to any incoming address on port 80. Changing this means you can run multiple web servers on the same site by using different ports to reference each.

Note You can specify a port when connecting to just about any application by adding `:<port number>` at the end of the URL.

Following the `VirtualHost` address, the next relevant thing is the `DocumentRoot`. This is critical as it says where Apache will load all the web server content from. This means that anything that is in the `/var/www` folder will become

available via the basic website on port 80. For example, if you go into the `/var/www` folder, you will see one file: `index.html`. If you examine this file, you will see some very familiar content:

```
<html><body><h1>It works!</h1>
<p>This is the default web page for this server.</p>
<p>The web server software is running but no content has been added, yet.</p>
</body></html>
```

Yep, that's the website content that is used to generate the basic start page that Apache is currently showing. I try to think of it this way: the web server is providing you access to that folder (and all subfolders and documents within it). As proof of this, try executing the following command as root:

```
$ echo "Hello World" > /var/www/foo
```

Note You can gain root by executing `sudo su -`.

Once you have successfully executed this command, direct your browser to `http://<IP Address>/foo` and you should see `Hello World` displayed there. You may be wondering why we didn't need to use `index.html` in the first request to our website (i.e., use <http://10.0.0.2/index.html>). This is because files named "index" are special ones. They are the default page displayed if no other content is available.

The next portions of the config you should be aware of are the `Directory` config sections. These are a list of directories that can provide special rules as to how they are accessed and even who can access them. In this instance, there are two directories that Apache has configuration settings for: the `/root` directory and `/var/www`. While there are a number of very interesting things you can do with the directory values, there isn't much we need to do with them here. For more information, you should head over to the Apache website and read through the far more detailed documentation available there.

After the directory details comes `ScriptAlias`, which governs how CGI scripts are accessed by the application. The `ScriptAlias` is a special directory in which these CGI programs live and are executed from. As PHP is a scripting system much like BASH (covered in Chapter 7), you may think this has a direct bearing, but it is not the case. PHP is run by a special interpreter module (`modphp`) rather than by a single CGI program.

The `directory` value below the `ScriptAlias` is related to it. The quick give away is that both reference `/usr/lib/cgi-bin/`. This directory contains the `+ExecCGI` directive. Again, this relates to how scripts and things such as PHP are executed. Finally, the configuration file comes to a close with a few very useful config lines. These values are the logfile directives for the Apache server; they say where the logfiles will go and which files they will write to. Both access logs (so you can spy on everyone who visits you) and the much more standard error logs are covered here. The `LogLevel` directive in the middle governs at which level logs regarding the running of the server should be output, and right now it is set at `WARN`, which is sensible.

Given that you will probably need to take advantage of these logs later when working with PHP, you should become familiar with them now. But first let's make a small change to the log that we will make Apache incorporate. Let's change that logging level to `notice`, so go ahead and change it then save the newly edited file. Once the file has been changed, we need to tell Apache this so it will reload the config. Unlike making a content change to your website, we need to inform Apache of config changes. Thankfully, Apache makes this easy by including a reload function that we can issue to it. Go ahead and issue this:

```
$ sudo /etc/init.d/apache2 reload
[....] Reloading web server config: apache2: Could not reliably determine the server's fully
qualified domain name, using 127.0.1.1 for ServerName
. ok
```

This has told Apache to reload its config and get going again. Now have a look at the logfiles to see the changes we made. The Apache log directory is /var/log/apache2, and the file we are interested in is the `error.log`. Once you are looking inside, you should see something like this:

```
[Tue Sep 18 22:46:56 2012] [notice] SIGUSR1 received. Doing graceful restart
apache2: Could not reliably determine the server's fully qualified domain name, using 127.0.1.1 for
ServerName
[Tue Sep 18 22:46:56 2012] [notice] Apache/2.2.22 (Debian) configured -- resuming normal operations
```

This says that my Apache instance was given a signal and it performed a graceful restart; then it resumed normal operation within the same second. This is the configuration reload we performed just now. We can also look at the access log and see ourselves accessing content off the web server, go ahead and open up the `access.log` file in /var/log/apache2 and have a look. You should see something like this:

```
10.0.0.104 - - [18/Sep/2012:20:31:17 +1000] "GET / HTTP/1.1" 200 482 "-" "Mozilla/5.0 (Windows NT
6.1; WOW64) AppleWebKit/537.10 (KHTML, like Gecko) Chrome/23.0.1262.0 Safari/537.10"
10.0.0.104 - - [18/Sep/2012:22:21:22 +1000] "GET /favicon.ico HTTP/1.1" 404 498 "-" "Mozilla/5.0
(Windows NT 6.1; WOW64) AppleWebKit/537.10 (KHTML, like Gecko) Chrome/23.0.1262.0 Safari/537.10"
10.0.0.104 - - [18/Sep/2012:22:28:36 +1000] "GET /foo HTTP/1.1" 200 274 "-" "Mozilla/5.0 (Windows NT
6.1; WOW64) AppleWebKit/537.10 (KHTML, like Gecko) Chrome/23.0.1262.0 Safari/537.10"
```

These lines show that someone on IP 10.0.0.104 made a GET request on the / of the web server, which is the /var/www folder and the index.html file within (because index is the default file). It also shows someone accessing the foo file within this directory. They were both me accessing the website in sequence. You will also notice that there is a request in the middle, which is for a favicon.ico. This represents my browser trying to request the favorite's icon of this website; there is not one, unfortunately.

Troubleshooting

For the most part, setting up an Apache server is a breeze. It's only when you start making a lot of changes that you might run into trouble. However, if you did find yourself running into issues, try the following. Check that you can see the Apache process running with `ps -ef`. If Apache is not running, try and start it. If that fails, try examining the error logfile as this will contain information on any errors that Apache has suffered. Most of the entries in the error log are fairly self-explanatory. However, if you find yourself at a loss, head over to the Apache website where there is a large amount of help available there for those in need.

M is for MySQL

Now that you have your web server up and running, you are probably asking, “I see that you can make a website by just writing content into a web server and it will be displayed. What more software could I possibly need?” And it’s true; right now, you could write a web page in basic HyperText Markup Language (HTML) and be done with it. The problem with writing in HTML alone is that HTML is static, meaning that any time you wish to change what is displayed, you need to open up and edit the file to change what is said. This is not going to make for very good management when you want to display different things to different people. You would need to be manually sitting there and editing a whole bunch of things—that’s pointless.

This is where we begin to see the power of the full LAMP stack. We have Apache to serve content, PHP will handle the dynamic portion, and the MySQL database will actually contain the various different bits of information you want to display. It is only when you have all these components in and working together that you begin to see the true power of a LAMP system.

Installing MySQL

A database (DB) is a system for the easy storage and retrieval of data. They take large volumes of input and then structure the data so that information can be easily retrieved with a query. All SQL databases are named for the Structured Query Language, which they use to process requests for information retrieval. MySQL is the world's most widely deployed open source database. It is what is called a *relational database system* that is used because it is easy to set up and to work with, thus its wide adoption. To install MySQL, we will again turn to our old friend apt-get; this time, we will be installing the mysql-server package. Execute the following:

```
$ sudo apt-get install mysql-server
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
heirloom-mailx libaio1 libdbd-mysql-perl libdbi-perl libhtml-template-perl
libmysqlclient16 libnet-daemon-perl libplrpc-perl mysql-client-5.5
mysql-common mysql-server-5.5 mysql-server-core-5.5
Suggested packages:
libipc-sharedcache-perl libterm-readkey-perl tinyca
The following NEW packages will be installed:
heirloom-mailx libaio1 libdbd-mysql-perl libdbi-perl libhtml-template-perl
libmysqlclient16 libnet-daemon-perl libplrpc-perl mysql-client-5.5
mysql-common mysql-server mysql-server-5.5 mysql-server-core-5.5
0 upgraded, 13 newly installed, 0 to remove and 84 not upgraded.
Need to get 9,770 kB of archives.
After this operation, 91.5 MB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Once the actual MySQL install has finished, you will be asked to provide a root password for your DB; this is the master admin password for your database. Make sure you remember it; you will need it later. Now that the installation has finished, confirm that the process has started with this:

```
$ ps -ef | grep mysql
```

The output should show a few different things, including the MySQL daemon, which is highlighted here:

```
root      5254      1  0 20:23 ?        00:00:00 /bin/sh /usr/bin/mysqld_safe
mysql    5592  5254  3 20:23 ?        00:00:03 /usr/sbin/mysqld --basedir=/usr --datadir=/var/
lib/mysql --plugin-dir=/usr/lib/mysql/plugin --user=mysql --pid-file=/var/run/mysqld/mysqld.pid
--socket=/var/run/mysqld/mysqld.sock --port=3306
root      5593  5254  0 20:23 ?        00:00:00 logger -t mysqld -p daemon.error
```

That's one big process command, but it does show that MySQL is up and running. We can go over the MySQL config, but thankfully the system will work for all our basic needs without the need to modify the configuration. Given that, let's start working on some basic ways to both insert data and retrieve it.

Structured Query Language

Structured Query Language (SQL) is a language for inserting and retrieving information from an SQL database. In SQL, there are a number of different queries you can do, ranging from the simple to the extremely complex. Given that there are (as with everything) a number of other resources that spend much more time and go into far more detail

about syntax and origin and whatever, we won't cover those here. Besides, for our purposes there should only be five types of queries we need:

- Basic admin commands
- Insert new data
- Find data
- Update data
- Delete data

While it sounds incredibly simple, it is very easy to get lost in the syntax if you wish to do something more complex. If that is the case, and you wish to do something slightly more complex than what we have covered here please head over to the MySQL web page as they have full documentation on how exactly to perform each query and what the large number of options on each command entail. If you want something a little more focused on web development with all the Apress niceties, have a look at *Beginning PHP and MySQL 5*, by Jason W. Gilmore.

Anyway, enough product placement; let's start with the basic administration stuff and make a database. The syntax for this one is dead simple:

```
create database <database name>
```

To perform commands on your MySQL instance, you need to connect to it and get into the `mysql` shell. This is the command to do this:

```
mysql -uroot -p
```

So jump into the shell, provide your root password, and execute `create database pi`; which should generate the following output:

```
mysql> create database pi;
Query OK, 1 row affected (0.00 sec)
```

Okay, awesome. We have a database called `pi`. In an SQL database system, a database is the highest-level construct. Underneath that is a *table*, which is the object that holds your data (a database holds multiple tables). Each of these items of data in a table is called a *row*. So our next step is to create a table; the problem is that we need to give our table some structure. We need to tell it what kind of data it should hold and how we will lay out that data, which means we need a project. The most obvious for now is a simple to-do list, so let's make one of those.

Note All MySQL commands end with a semicolon (;). It signifies that this is the end of a statement. This is done as you can put lines in for one command over multiple lines to make it easier to follow.

Creating Tables

Just stepping aside quickly, I have departed from the normal project planning method a little here so that we can run through a basic database setup together. When undertaking a big project, it is far better to plan ahead and fully understand what will be going into a database long before any data is ever added to it.

So, for our to-do list, we want to keep the following information for each “thing to do”:

- Description
- Person doing the task (owner)
- Date to be done by (date)
- Location
- Importance
- Who set the task (creator)

Okay, so with this information, we have our rough table structure; so now let’s go ahead and create it. The syntax for a table creation is this:

```
create table <tablename> (
<column name> <column type>,
<column name> <column type>,
...
<column name> <column type>
);
```

That’s the basic syntax, but there is one final thing we need to make sure we have when creating our table: a key. Given that we need to ensure we are getting the right value back each time, we need to have something unique about each individual piece of data. Many of these tasks could wind up almost identical, excluding one or two minor variations. For this reason, we will need what is called a *primary key*, which is a unique identifier for each row. In this case, it’s best just to use a simple count that will automatically go up for each row we add. Now we know what we want; let’s flesh it out into the syntax from before and see what we have.

```
create table todolist (
idnumber <column type>,
description <column type>,
owner <column type>,
date <column type>,
location <column type>,
importance <column type>,
creator <column type>
);
```

Well, looking better, but we still don’t have those <column types>, nor do we have anything that will say that our idnumber is the primary key, or should be automatically updated for each new row. This brings us nicely to what the column types are about. MySQL will need to know what kind of values each of our data elements will be so it can know how to store them and what kind of queries they can be involved in. There are at least 30 different data types in MySQL, which can perform all kinds of operations, but for our simple purposes there are only three we need to be concerned with: one is text or a string of text that is known as a VARCHAR and is given a maximum number of characters, the second is a number or integer known as an INT, and the last is a date value that is known as a DATE. Given that we now know some data types, we can set them out into our create_table command. Go ahead and do that now and remember that text strings (VARCHAR) need to be given a maximum number of characters long:

```
create table todolist (
idnumber INT,
description VARCHAR(200),
```

```
owner VARCHAR(40),
date DATE,
location VARCHAR(40),
importance VARCHAR(10),
creator VARCHAR(40)
) ;
```

That looks much better and will probably run, but remember we wanted to have that number be our unique identifier and be automatically updated for us to make our lives easier. The syntax for the `idnumber` being automatically incremented (updated to be +1) is `PRIMARY KEY NOT NULL AUTO_INCREMENT`. This gives us a final command of this:

```
create table todolist (
idnumber INT PRIMARY KEY NOT NULL AUTO_INCREMENT,
description VARCHAR(200),
owner VARCHAR(40),
date DATE,
location VARCHAR(40),
importance VARCHAR(10),
creator VARCHAR(40)
);
```

So, go ahead and run it:

```
ERROR 1046 (3D000): No database selected
```

Oh, whoops. We need to tell MySQL which database we are using, and thus under which database this table will reside. To change which database we are using, we need to use the `USE` command. To change to using the `Pi` database we simply execute `USE pi;` and MySQL will tell us that the database has changed. If you have forgotten your database name, you can use the `SHOW` command to see things, `SHOW DATABASES;` will show you all the databases on your system. Now go ahead and execute the table creation again. This time you should see this:

```
Query OK, 0 rows affected (0.91 sec)
```

Fantastic, we have a table. Let's have a look and see if we can see it. Issue the `SHOW` command, but this time for tables. You should see the following output, which lists off all the tables within the `Pi` database:

```
mysql> SHOW TABLES;
+-----+
| Tables_in_pi |
+-----+
| todolist    |
+-----+
1 row in set (0.00 sec)
```

Yahoo! Now, say for example that you have forgotten what exactly the table looks like; you will want MySQL to describe how the table is made up. You can use the DESCRIBE command to do just this, so let's try describing our new todolist table so we can see how MySQL understands it.

```
mysql> DESCRIBE todolist;
+-----+-----+-----+-----+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+
| idnumber | int(11) | NO | PRI | NULL | auto_increment |
| description | varchar(200) | YES | | NULL |
| owner | varchar(40) | YES | | NULL |
| date | date | YES | | NULL |
| location | varchar(40) | YES | | NULL |
| importance | varchar(10) | YES | | NULL |
| creator | varchar(40) | YES | | NULL |
+-----+-----+-----+-----+-----+
7 rows in set (0.00 sec)
```

Awesome; you can even see that our `idnumber` is the `PRI` key and has automatic incrementing on the end! Everything registered successfully. We have successfully created a database and a table. We have given structure to our table. Now it is time to put it to good use, but before we start playing with data there's one last admin command I want to cover: creating a user other than root. This way, we don't have to constantly use the root user, which will cut down the potential security risk. The syntax for this command (called a GRANT) is this:

```
GRANT ALL ON <database>. <table> TO '<username>'@'<user location>' IDENTIFIED BY '<password>';
```

So, let's say we want to grant the default pi user on our system access to our todolist table with the password `raspberry`. The command will become this:

```
GRANT ALL ON pi.todolist TO 'pi'@'localhost' IDENTIFIED BY 'raspberry';
```

Run this command and then quit MySQL shell by typing `quit`. Now try logging back in with your new username and password. Remember, the syntax for this:

```
$ mysql -u<username> -p
```

Note Quotation marks are needed when putting in any text data in MySQL that is not for something that MySQL already “understands” (i.e., table names and column names).

Inserting Data

Now you're logged back in as the `pi` user, let's start learning some actual commands to insert data into MySQL. The first command to cover is one for inserting data. How else can we be expected to work with MySQL data if there is no data available! The command for inserting data is the aptly named `INSERT` command. The basic syntax for `INSERT` is this:

```
INSERT INTO <TABLE> (<FIELD1>, <FIELD2>, ... <FIELDX>) VALUES ('<VAL1>', '<VAL2>', ... '<VALX>');
```

So, now that we know how an insert is supposed to look, let's go ahead and make one into the database. Let's insert a pair of tasks into our to-do list. The first one will be my writing of this chapter. So the command will be this (after ensuring we typed `USE pi`):

```
INSERT INTO todolist (description, owner, date, location, importance, creator) VALUES ('Finish LAMP Chapter', 'David', '2012-09-22', 'Australia', 'HIGH', 'David');
Query OK, 1 row affected (0.43 sec)
```

That worked. Let's add another, just for good measure. Let's make Peter do something now:

```
INSERT INTO todolist (description, owner, date, location, importance, creator) VALUES ('Finish GUI Chapter', 'Peter', '2012-09-22', 'Hong Kong', 'HIGH', 'David');
Query OK, 1 row affected (0.48 sec)
```

Now that we have a pair of records to play with. But how are we certain that they are correct? How can we check that the idnumbers are incremented, given that we didn't even add them into our `insert` statement? For that we will need to issue a query to our SQL database!

Querying a Database

Unlike most of the other statements, a data query is not done with a `QUERY` command. This is because all of the commands we have been executing are considered to be queries themselves. The command for retrieving data is called a `SELECT` and its syntax is as follows:

```
SELECT <Fields1>, <Field2>... <FieldX> FROM <TABLENAME> WHERE <INFORMATION QUERY>
```

Yes, I know that syntax is a little strange, but once we fill it out you will get more of an idea of why it is the way it is. For starters let's just grab everything. Normally you put field names that you wish to get, which limits the amount of excess data being transmitted, but in this case we can use the special wildcard `*`. Thus to query everything from our `todolist` table, we will execute this:

```
mysql> SELECT * FROM todolist;
+-----+-----+-----+-----+-----+-----+
| idnumber | description | owner | date | location | importance | creator |
+-----+-----+-----+-----+-----+-----+
| 1 | Finish LAMP Chapter | David | 2012-09-22 | Australia | HIGH | David |
| 2 | Finish GUI Chapter | Peter | 2012-09-22 | Hong Kong | HIGH | David |
+-----+-----+-----+-----+-----+-----+
2 rows in set (0.00 sec)
```

You will notice that I left off the `WHERE`, this is done, as there is no actual limitation of what we want, we want everything. After this, you can see that everything else we entered is in the format we gave it and more importantly the ID numbers are incrementing! This validates everything that we inserted before. Now that you understand the very basic syntax we can do some slightly more advanced queries. Let's say I want to find out whom to blame for assigning me (David) a task. The relevant information in this case is the creator as that's all we want, whom to blame. This will give us the first half of a query:

```
SELECT creator FROM todolist
```

Now the next part we need is the portion in which we specify that the owner must be David. The syntax for this is a very simple WHERE owner = "David" giving us the final query:

```
SELECT creator FROM todolist WHERE owner = "David";
+-----+
| creator |
+-----+
| David   |
+-----+
1 row in set (0.00 sec)
```

Beautiful. I now know that I have only myself to blame for this mess. In the same vein, let's try another. I want to know the description and priority of all tasks that are due after today (which, in this example, is the 20th of September). So again we start building our query. We want the description and importance fields this time, so in they go. We also want a date greater than 2012-09-20. Thankfully, MySQL understands date data, so all we need to do is give it our date and the greater-than symbol, which is >. This will give you the query:

```
mysql> SELECT description, importance FROM todolist WHERE date > "2012-09-20";
+-----+-----+
| description      | importance |
+-----+-----+
| Finish LAMP Chapter | HIGH      |
| Finish GUI Chapter | HIGH      |
+-----+-----+
2 rows in set (0.01 sec)
```

There is one last simple part of a SELECT query that will likely be relevant later. This is the ORDER BY statement that can be added to the end of a query. Take our last one for instance; let's say we want them ordered by the order in which they were added (by idnumber). This will make our query into this:

```
mysql> SELECT description, importance, idnumber FROM todolist WHERE date > "2012-09-20" ORDER BY
idnumber;
+-----+-----+-----+
| description      | importance | idnumber |
+-----+-----+-----+
| Finish LAMP Chapter | HIGH      |      1 |
| Finish GUI Chapter | HIGH      |      2 |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

And there it is, sorted. Okay, so we can't see it as it is in ascending order and has been previously displayed as such. To reverse (or force) an order, we can add an ASC or DESC to the end of the statement, so let's try that with the same query as before and watch the change:

```
mysql> SELECT description, importance, idnumber FROM todolist WHERE date > "2012-09-20" ORDER BY
idnumber DESC;
+-----+-----+-----+
| description      | importance | idnumber |
+-----+-----+-----+
| Finish GUI Chapter | HIGH      |      2 |
| Finish LAMP Chapter | HIGH      |      1 |
+-----+-----+-----+
2 rows in set (0.00 sec)
```

That's better; you can see the change in order now! Having things returned in order is great as this means the database system is doing the sort for you—something that it is programmed to do very easily. If we were to attempt to write this sort, it would likely take us much longer than the database to perform. Now we have covered our second basic command, and we can insert and retrieve data. The next command we will look into is one to update our data.

Updating a Database

Now that we are past the slight silliness of the SELECT command to query, we are back into named-as-they-mean territory with the UPDATE command. An update is somewhat of a hybrid of both the INSERT and SELECT queries, and rightfully so, given that we need to both find something and then update that something. The basic syntax of an UPDATE is this:

```
UPDATE <table name> SET <column name1> = "<value1>,<column name2> = "<value2>"... <column nameX> = "<valueX>" WHERE <Information Query>
```

Let's work our way through an example. Let's say I need a deadline extension on this chapter because I'm working way too hard and want to take it easy for a night. This means I will want to increase my deadline to the 23rd instead of the 22nd. So we know what we want to update, which gives us this command:

```
UPDATE todolist SET date="2012-09-23" WHERE
```

Now we just need a query; we can't pick something on the 22nd as there are two things using that date; the same with severity and creator. We could use owner or description or idnumber. In this case I would choose ID number, as we have set this as the primary key and thus totally unique identifier. This makes our UPDATE query into

```
mysql> UPDATE todolist SET date="2012-09-23" WHERE idnumber=1;
Query OK, 1 row affected, 1 warning (0.48 sec)
Rows matched: 1  Changed: 1  Warnings: 0
```

You will notice that I did not add quotation marks around number 1 on the end. This is because MySQL treats numbers as INT differently to numbers as a VARCHAR; while this seems semantic, it is an important distinction as mathematical operations cannot be carried out against a VARCHAR, but can be against an INT. MySQL has not displayed the actual output of the data; to see it we will need to issue another select, so let's go ahead and examine the change:

```
mysql> SELECT * FROM todolist WHERE idnumber=1;
+-----+-----+-----+-----+-----+-----+
| idnumber | description | owner | date | location | importance | creator |
+-----+-----+-----+-----+-----+-----+
| 1 | Finish LAMP Chapter | David | 2012-09-23 | Australia | HIGH | David |
+-----+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

And there you have it; I've just given myself a night off. But because of this, Peter and I will both have to work harder on the chapters we have remaining. This means that we will need to set the importance of both our chapters up to highest! The first part is simple:

```
UPDATE todolist SET importance="HIGHEST" WHERE
```

This is where we can look at changing how we do queries again. Do we wish to do everything? And set everything to HIGHEST by removing the WHERE? Let's search for HIGH and make it HIGHEST, which makes our final query:

```
mysql> UPDATE todolist SET importance="HIGHEST" WHERE importance="HIGH";
Query OK, 2 rows affected (0.49 sec)
Rows matched: 2  Changed: 2  Warnings: 0
```

Again there is no output, so we must retrieve the data again:

```
mysql> SELECT * FROM todolist;
+-----+-----+-----+-----+-----+-----+
| idnumber | description | owner | date | location | importance | creator |
+-----+-----+-----+-----+-----+-----+
| 1 | Finish LAMP Chapter | David | 2012-09-23 | Australia | HIGHEST | David |
| 2 | Finish GUI Chapter | Peter | 2012-09-22 | Hong Kong | HIGHEST | David |
+-----+-----+-----+-----+-----+-----+
2 rows in set (0.00 sec)
```

We were able to update both those HIGH values to HIGHEST. Now for the one last remaining task: delete.

Deleting Data

We want to be able to remove tasks as they are finished. The syntax of a DELETE command is constructed almost exactly like a SELECT command:

```
DELETE FROM <table name> WHERE <information query>;
```

Here's one final example. Let's say my reason for a night off wasn't that I was tired; it was because I was finished with this chapter early! Woohoo! So let's build the delete. Again let's go by the idnumber to be certain. The command will be this:

```
mysql> DELETE FROM todolist WHERE idnumber=1;
Query OK, 1 row affected (0.42 sec)
```

As with UPDATE, there is no information returned other than the short output saying that we deleted one row. Let's issue a SELECT again and see what has changed:

```
mysql> SELECT * FROM todolist;
+-----+-----+-----+-----+-----+
| idnumber | description | owner | date | location | importance | creator |
+-----+-----+-----+-----+-----+
| 2 | Finish GUI Chapter | Peter | 2012-09-22 | Hong Kong | HIGHEST | David |
+-----+-----+-----+-----+-----+
1 row in set (0.01 sec)
```

Only Peter's chapter remains, and he has to work the hardest since it's on highest. Okay, let's be kind—one last command. This one gets rid of a table; it's called the DROP command and is very simple:

```
DROP TABLE <table name>;
```

Let's drop our `todolist` table since we are done with these MySQL examples and we don't want it taking up room:

```
mysql> DROP TABLE todolist;
Query OK, 0 rows affected (0.49 sec)
```

Okay, now we only have our `pi` database left. If you wish to delete it, you simply need to replace the word `TABLE` with the word `DATABASE` in the `DROP` command. It's up to you; we will be reusing this one later. For now, let's head on to the last part of our LAMP stack: PHP.

P is for PHP

PHP is a system for allowing you to add dynamic functions into your web page. Originally, all web pages were made from HTML alone, which doesn't allow for much flexibility based on inputs and actions. HTML just displays a static page of content that needs to be modified to show something different. To solve this initial issue, the Common Gateway Interface (CGI) was developed. This provided a way for web servers to take requests that would allow them to return content.

Originally, most CGIs were full-blown applications that output various pieces of HTML depending on their input, and whole portions of these applications were dedicated to outputting large chunks of the same HTML over and over. This is where the development of PHP came in. PHP was designed to be a language in which one could add snippets of actual dynamic code into the static HTML and presto: a dynamic web page would be formed.

Since its creation, PHP has become widespread with installation figures of 20 million being cited. This success is in no small part due to the ease in which PHP can be used as a CGI compared with the original application form CGIs. Now that you have an understanding of what PHP is, let's install it on the Raspberry Pi. As ever, we will rely on the services of `apt-get` to fetch and install PHP for us. The command to run is this:

```
$ sudo apt-get install php5
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  apache2-mpm-prefork libapache2-mod-php5 libonig2 libqdbm14 php5-cli php5-common
Suggested packages:
  php-pear
The following packages will be REMOVED:
  apache2-mpm-worker
The following NEW packages will be installed:
  apache2-mpm-prefork libapache2-mod-php5 libonig2 libqdbm14 php5 php5-cli php5-common
0 upgraded, 7 newly installed, 1 to remove and 84 not upgraded.
Need to get 5,707 kB of archives.
After this operation, 16.3 MB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Notice that PHP is making a modification to our Apache installation. This is to replace the Apache tribesman modules with ones that are preferred for utilizing PHP. Once the PHP has finished installing, we want to go ahead and test that it is working. There is a simple PHP page that we can write, which will show all our installation settings and confirm that PHP is up and running. To test this first page, we need to replace `index.html` in `/var/www` with a new `index.php` that has the following line inside it:

```
<?php phpinfo(); ?>
```

Once you have removed `index.html` and added `index.php`, direct your browser to your Apache server. Figure 8-2 will be your reward.

System	Linux raspberrypi 3.1.9+ #272 PREEMPT Tue Aug 7 22:51:44 BST 2012 armv6l
Build Date	Sep 13 2012 01:19:39
Server API	Apache 2.0 Handler
Virtual Directory Support	disabled
Configuration File (php.ini) Path	/etc/php5/apache2
Loaded Configuration File	/etc/php5/apache2/php.ini
Scan this dir for additional .ini files	/etc/php5/apache2/conf.d
Additional .ini files parsed	/etc/php5/apache2/conf.d/10-pdo.ini
PHP API	20100412
PHP Extension	20100525

Figure 8-2. PHP info page

This is the PHP info page, which describes the entire current configuration of your newly installed PHP instance. Since we are viewing it, it also proves that the PHP interpreter is working correctly. It also shows that `index.php` has been picked up as a default page by Apache. The code that we added to generate this was the following:

```
<?php phpinfo(); ?>
```

This is a great example of PHP code; first we have the opening markers that show that this is PHP and to be interpreted as such (`<?php` and `?>` to close) and we have a single function call inside to `phpinfo();` that displays all our data.

A function is a reference to a predefined piece of code that we run by invoking the function name. Functions always end with a pair of parentheses, which can contain data (variables that will be passed into the function for use there). In addition to functions, PHP uses arrays, much like the arrays you saw in Chapter 7. To recap, arrays are variables that contain a number of values rather than one. An array can be manipulated as a whole array or by each of its individual elements. To access the array elements, you need to add a pair of square brackets to the end of the array variable and enter the number of the element you wish to access. These element numbers start from 0 as the first element and increment.

The other important things one needs to know for PHP are these:

- All PHP statements end with a semicolon
- All PHP variables start with a \$ sign (recall variables from BASH in Chapter 7; PHP variables are used for the same thing)
- All PHP array variables start with an @ sign
- All PHP function's internal code are surrounded with curly braces { and }

Next we need to have a look through the PHP Info page; start by searching for MySQL. Oh dear, there's no mention of it. PHP will need to understand how to communicate with MySQL in order for us to be able to display database content on our web page. Thankfully, apt-get comes to the rescue again. This time, we need to install the php5-mysql package:

```
$ sudo apt-get install php5-mysql
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  libmysqlclient18
The following NEW packages will be installed:
  libmysqlclient18 php5-mysql
0 upgraded, 2 newly installed, 0 to remove and 84 not upgraded.
Need to get 711 kB of archives.
After this operation, 3,547 kB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Again, you will see that this install adds a new MySQL client library for communications and the new PHP MySQL interleave. It will also automatically restart your Apache web server and change several PHP configuration files for you. Once again, open up and direct your browser to the Apache web page and search for MySQL, and there we go! (See Figure 8-3.)

MySQL Support	
Active Persistent Links	0
Active Links	0
Client API version	5.5.24
MYSQI_MODULE_TYPE	external
MYSQI_SOCKET	/var/run/mysqld/mysqld.sock
MYSQI_INCLUDE	-l/usr/include/mysql
MYSQI_LIBS	-L/usr/lib -lmysqlclient_r

Directive	Local Value	Master Value
mysql.allow_local_infile	On	On
mysql.allow_persistent	On	On
mysql.connect_timeout	60	60
mysql.default_host	<i>no value</i>	<i>no value</i>
mysql.default_password	<i>no value</i>	<i>no value</i>
mysql.default_port	<i>no value</i>	<i>no value</i>
mysql.default_socket	/var/run/mysqld/mysqld.sock	/var/run/mysqld/mysqld.sock

Figure 8-3. PHP MySQL config info

We have confirmed the following:

- We have a working PHP system
- Apache can render the PHP
- MySQL and PHP know how to communicate

This means that we have just accomplished creating a working LAMP setup! Now let's put it to use.

Simple Web App

Now that we have built our new LAMP stack and it is up and running on your Pi, it is time to make good use of it. Good use can be literally anything you want as you now have a web server of your own that you can write web applications on. For many of you, however, this may seem a daunting task, as you are unfamiliar with PHP or even HTML. Fear not; like Linux in general this seems daunting at first and there are intricacies, but nothing insurmountable and nothing that you can't get over with a little help. I've always found that the best way to learn is to actually do it, which means we can undertake a project together, which will illustrate how you can use your new LAMP stack to create a web application.

In order to build a web application, we need to understand both HTML and PHP. One final thing: this will be a bare-bones web page and may look a little rough. There are so many ways to improve the visual aspect of your web page, but there are whole books dedicated to this, which will cover a larger amount of content and interactions. So with this in mind, let's have a look at understanding HTML.

HTML

Hyper Text Markup Language (HTML) is the language that makes the Internet. HTML provides the basis for all transmitted web application content, it is made up as a series of tags and they are sent as a HTML document to your browser. Your browser then uses these tags to assemble its version of your web page. This means that web content can be very small and can be viewed by any system that has the capacity to decode HTML tags. The downside here is that different software can display websites with slight differences in them.

HTML tags are words surrounded by angle brackets. These tags are written in what is called a markup language. HTML tags come in two varieties, on their own or in pairs, with pairs being the much more common. Here are two sets of HTML tags

- `<html></html>`: An open and a close HTML tag (these tags tell your browser where the HTML page starts and finishes)
- `<hr>`: An hr tag on its own (`<hr>` will draw a line across your web page because it stands for headed rule)

The slashes at the front of the second HTML tag show that this is the close of this tag. The slash at the end of the hr tag shows it is meant to be on its own. Have a look at a very basic website:

```
<html>
  <head>
    <title>Pi Brand - Todo List App</title>
  </head>
  <body>
    <h1>Pi Todo List App</h1>
  </body>
</html>
```

If you place this into `index.php` on your Pi, open up your browser you should see a basic white page with the phrase Pi Todo List App in large black text. Additionally up the top of your browser should say the same. This website is wholly generated by the HTML tags we just entered. Let's go over these as an explanation of how the website works:

- First you will see that all the content is wrapped within a pair of HTML tags, which say that all the content within them is HTML.
- At the next level are the head and body tags, which represent the next level of tags, the heading information of the page and the actual body of the page.
- The `title` tags are used to set the title bar at the top of the page.
- And within the body is one pair of `h1` tags that are the big block of text. The `h1` is short for *header 1*, which represents the largest size of header tags. There are an almost unlimited number of `hx` tags you can use; as long as they are defined, most browsers will only go as far as `h4` or `h5` without one being defined separately.

So we have now covered five basic HTML tags but there are many, many more. For a reference of a host of the common HTML tags, some examples and great tutorials, head over to [w3schools.com](http://www.w3schools.com) or have a look at "Beginning HTML with CSS and XHTML," by Craig Cook and David Schultz.

Here are a few more HTML tags and how they are used:

- `<p>`: Paragraph tag, used to create text into a paragraph block
- `
`: Break tag, used to insert a break into the page
- `<hr>`: Horizontal rule, creates a horizontal line across the screen
- `<a>`: Hyperlink tag, for a link to another page
- ``: Image tag, for adding images to a page
- Table tags
 - `<table>`: Top-level tag of a table
 - `<th>`: Element of the tables heading row
 - `<tr>`: Row for the table
 - `<td>`: Cell of a table within a row
- Form tags
 - `<form>`: Top-level element of a form
 - `<input>`: Input element of a form

These are some of the basic HTML tags we will use in our project to create output and drive input to the to-do list. In addition to the raw tags, there are a number of different options you can use to change the various aspects of how these elements are displayed; unfortunately, it would take more than the rest of this book to cover each element, so we will have to discuss them as they come along. Now that you're familiar with what HTML can do and what some of these tags look like, we can move on to trying to add some PHP into the mix here. But first, let's have a look at actually designing our page.

Starting Our Page

The basic concept for our page will be a website in two parts; first will be a page that contains a to-do list, listing off all the elements in the database. Next will be a form for the submission of new elements to the database down the

bottom. We will also add a check box that will allow you to delete elements you no longer need. Now that we have the basic layout, there are some functional decisions that need to be made regarding how the processing of the page should occur. Since the design and output will remain constant with just some processing to connect to the database to load the elements or remove elements, we can just go ahead and have one page, which will attempt to perform any of the actions it needs (the additions and deletions before displaying the page).

Now we can start actually writing the page out. To start, add all the content to the page and get the output going. However, first we need to have our database set up so we have data to display. We have already created the Pi database and we have the to-do list table structure from earlier, so let's use them. The create statement for this table was as follows:

```
create table todolist (
idnumber INT PRIMARY KEY NOT NULL AUTO_INCREMENT,
description VARCHAR(200),
owner VARCHAR(40),
date DATE,
location VARCHAR(40),
importance VARCHAR(10),
creator VARCHAR(40)
);
```

Now that we have a dataset to work with and a layout for our data, we can begin to add display elements. Since our data is tabular, we can use the table elements to display it. This will make arranging the data much easier. So, we have our basic page layout from earlier; next we should set up our data connection. A MySQL PHP connection uses the `mysqli` interface. Thus to establish a connection to the database we must create a new `mysqli` that has the details of our database and connection:

```
<?php
$mysqli = new mysqli('localhost', 'pi', 'raspberry', 'pi');
if ($mysqli->connect_error) {
    die('Connect Error (' . $mysqli->connect_errno . ') '
        . $mysqli->connect_error);
}

$mysqli->close();
?>
```

The preceding is a small segment of PHP code, which does the following:

- Creates a new `mysqli` object to connect to the system, which is available on the local machine using the username `pi` and password `raspberry`.
- Checks that the connection to the database was successful; and if not, displays an error message.
- Closes the connection as we don't want to leave a potentially open connection there unused.

Now that we have this block of code for database connectivity, add this to the original block of code in `index.php`:

```
<html>
<head>
<title>Pi Brand - Todo List App</title>
</head>
<body>
```

```
<h1>Pi Todo List App</h1>
<?php
$mysqli = new mysqli('localhost', 'pi', 'raspberry', 'pi');
if ($mysqli->connect_error) {
    die('Connect Error (' . $mysqli->connect_errno . ') '
        . $mysqli->connect_error);
}
$mysqli->close();
?>
</body>
</html>
```

Display the Database Content

We now have that original piece of HTML from before, with the addition of one little `mysqli` connection object. This won't do much unless there is an error because nothing else will be displayed. Now it's time to add another block of code to display the to-do list. This will take two forms: the table block and the table head. Then we will need to display everything within the MySQL table. Displaying them means creating a loop that will display the content of each row, and placing it within the correct set of HTML tags. Let's start with the basic HTML table layout:

```
<table>
<tr>
<th>Description</th>
<th>Owner</th>
<th>Due Date</th>
<th>Location</th>
<th>Importance</th>
<th>Creator</th>
</tr>
...
</table>
```

This code sets out the table and its first row, which are all tagged as being table header elements. Now add the PHP to display all the content from the MySQL table:

```
<?php
$result = $mysqli->query("SELECT * FROM todolist");
while($row = $result->fetch_assoc()){
    print "<tr>";
    print "<td>".$row["description"].".</td>";
    print "<td>".$row["owner"].".</td>";
    print "<td>".$row["date"].".</td>";
    print "<td>".$row["location"].".</td>";
    print "<td>".$row["importance"].".</td>";
    print "<td>".$row["creator"].".</td>";
    print "</tr>";
}
?>
```

The PHP section performs this nicely. First, it creates a new variable, \$result, which contains the output of executing a query on mysqli. The query is, of course, SELECT * FROM todolist. Each line of the result is then passed out one by one in the while loop by the \$result->fetch_assoc() call and assigned to the row variable. For each row that comes out, we print the row value for each field we request. You can see the resemblance between the output section of the PHP and the raw HTML. This is deliberate as we want the output of the PHP to marry up with the table header.

Now we should add this back to the original code block, but when doing so we need to make one more change, which is to move the \$mysqli->close(); to the bottom of the new block, just before the close of the PHP segment. This move needs to occur because otherwise we will have closed off the database connection before we have actually pulled the data out of the database.

```
...
<body>
<h1>Pi Todo List App</h1>
<?php
$mysqli = new mysqli('localhost', 'pi', 'raspberry', 'pi');
if ($mysqli->connect_error) {
    die('Connect Error (' . $mysqli->connect_errno . ') '
        . $mysqli->connect_error);
}
?>
<table>
<tr>
<th>Description</th>
<th>Owner</th>
<th>Due Date</th>
<th>Location</th>
<th>Importance</th>
<th>Creator</th>
</tr>
<?php
$result = $mysqli->query("SELECT * FROM todolist");
while($row = $result->fetch_assoc()){
    print "<tr>";
    print "<td>".$row["description"].".</td>";
    print "<td>".$row["owner"].".</td>";
    print "<td>".$row["date"].".</td>";
    print "<td>".$row["location"].".</td>";
    print "<td>".$row["importance"].".</td>";
    print "<td>".$row["creator"].".</td>";
    print "</tr>";
}
$mysqli->close();
?>
</table>
</body>
</html>
```

Website Data Insertion

This HTML is now working in concert with the PHP to generate a full page of content. The static HTML provides a framework and then we have two pieces of PHP: one piece to establish a connection and a second to pull the results out of the to-do list table and add them to the page. Now that we have the basic display working, we need to add a form to give ourselves a form to submit new content. This basic form should have an input stream for each of the elements we will insert into the table. We will also need one more element, a special hidden one that will tell our processor what kind of things it is to do with the data. In this case, I like to use a variable called "action" and assign it as needed. The final element we will need is a submission one, which allows us to push the data to the server to be processed. In addition to all these elements, we need to give the form an action and method variable, which says how it is to invoke CGI on our web server and what method it should use. This block when it is put together looks like this:

```
...
</table>
<form action="index.php" method="POST">
<input type="hidden" name="action" value="insert" />
Description: <input name="description" /><br/>
Owner: <input name="owner" /><br/>
Date: <input name="date" /><br/>
Location: <input name="location" /><br/>
Importance: <input name="importance" /><br/>
Creator: <input name="creator" /><br/>
<input type="submit" />
</form>
</body>
</html>
```

This code block can then be added below the table and even separated with an `<hr />` tag, which means that you will have a table of display content and a section below it that will give you the ability to add new content. When put together and installed into the web server, the content that it generates will look like Figure 8-4.

Description	Owner	Due Date	Location	Importance	Creator
Finish LAMP Chapter	David	2012-09-22	Australia	HIGH	David

Description:

Owner:

Date:

Location:

Importance:

Creator:

Figure 8-4. To-do list app with insert table

If you press the submit button, you will wind up right back on the page, but you can see that the page is referenced in the URL. This is the action block sending the CGI command to run on this page. So now we will need to add some CGI processing to our PHP in order to process this data.

Recall that we set a method of POST earlier, which is one of two methods that we can use to pass data from our page into the CGI system to be processed. The other method is GET, and the difference between the two is limited as both pass data in basically the same manner. For all practical purposes, the only difference is that GET will display the content data on the web address, whereas POST will keep it hidden. You can check this out for yourself by changing the POST to GET and then pressing the submit button. Your URL will look something like <http://10.0.0.20/index.php?action=insert&description=&owner=&date=&location=&importance=&creator=>.

There is actually a large amount of data being transmitted that will need to be processed. Thankfully, PHP has methods to make this much easier; it has special variables that are automatically populated with the data from CGI requests. There are three special variables that you can access (as we did with the SQL assoc variables): `_POST`, `_GET`, and `_REQUEST`.

To process our CGI, we need to do quite a few things. First, check if the `action` variable is set and contains data; that it will when we press the submit button as we have it set to insert (and later have an action of remove). Once we have ensured that `action` is set, we can check what it is set to. Once we know which action we are performing, we can simply pull apart the remainder of the output and then perform the required actions on the database. Finally, if we do this before the main page load, we will actually have the latest version of the data displayed automatically!

Our CGI should start looking like this:

```
if(isset($_REQUEST["action"])){
switch($_REQUEST["action"]){
    case "insert":
        $SQL="INSERT INTO todolist (description, owner, date, location, importance, creator)
VALUES (";
        $SQL=$SQL."".$_REQUEST["description"]."";
        $SQL=$SQL."".$_REQUEST["owner"]."";
        $SQL=$SQL."".$_REQUEST["date"]."";
        $SQL=$SQL."".$_REQUEST["location"]."";
        $SQL=$SQL."".$_REQUEST["importance"]."";
        $SQL=$SQL."".$_REQUEST["creator"]."";
        $SQL=$SQL.");";
    if ($mysqli->query($SQL)=== FALSE) {
        printf("Error - Unable to insert data to table " . $mysqli->error);
    }
    break;
    case "delete":
        print "Delete function yet to be added!";
        break;
}
}
```

The first thing is to see whether the `action` variable specified in the hidden field within the form is set. If it is, we can be reliably certain that an action needs to be performed. Next, if `action` is set, we go into a `switch` (or `case`) statement to work out which function we are performing. I've added a case for both `insert` and `delete`, but the `delete` function is just printing that we will add it later. Besides, how would we access it without a `submit` method for it and an `action` field in a form?

Inserting to the Database

In the `insert` case, we simply need to create and the build the SQL command. The command we are trying to build is exactly like the previous SQL inserts. First, I laid out the static content: the frame of the `INSERT` statement, the table name, and which fields we will be writing into. Then I have started adding variable by variable to that statement. You can see that the additions have a lot of funny stuff around them because each of these variables will need to be surrounded by single quotation marks and have a comma on the end. Each of these little pieces of text needs to be treated as such, which means wrapping them in a pair of double quotation marks like so:

```
$SQL=$SQL."\"".$_REQUEST["description"]."\",";
```

This assigns the `SQL` variable the value of the current `SQL` value (given by `$SQL=$SQL`) and then adds to it (marked by the period) a single quote (given by `"\ "`) and the value of the `REQUEST` variable `description`, then a single quote, and then a comma. I know it seems a lot of work, but we have been able to pull out each of the variables we need and have the whole SQL statement in one tiny variable.

Caution It's also important to note at this point that the above is not to be considered secure. Those variables can quite literally be anything and there are special functions that you can use to perform some sanity checking on those values. The recommended function for performing a good selection of these sanity checks is the `mysqli_real_escape_string`.

Once we have created the `SQL` variable, we need to insert it by simply invoking the `mysqli` query function again. In this example, I've also added a check to see that the execution of the query failed. And that's it; this block of code will process the CGI insert into our database from our PHP. All that remains is to insert the block correctly into the right place in our code. I have chosen to append this to the first block of PHP we added instead of giving it its own. This means the flow of the page becomes first to draw the header of the page, then to create the database connection and perform any CGI actions, then to display the actual content of our to-do list, and then to display the final form. Once the code is in, you can test that it works by placing values into the form and pressing submit. This should go ahead and add an extra row that will show up in your table, as shown in Figure 8-5.

Description	Owner	Due Date	Location	Importance	Creator
Finish LAMP Chapter	David	2012-09-22	Australia	HIGH	David
Test LAMP Chpater	David	2012-09-22	Australia	HIGHEST	David

Description:
 Owner:
 Date:
 Location:
 Importance:
 Creator:

Figure 8-5. All work and no play...

Remove Entries

Now that the add functionality is working, we need to create a delete functionality. As with all things in life, there are a number of ways to can approach it. The easiest is to simply add a form and delete option for each individual row, which means that if you want to delete, you need to issue them one at a time. The second option is to have a series of check boxes that will delete all the ticked elements. I am a fan of the second option as it gives a lot more flexibility with the deletions.

We need to make two changes for this process to work. The first change is to add a form around the table that will be the delete form, a hidden field with the action delete, a submission button below the form, and check boxes to each element. To add the check boxes, we will also need to add an empty element to the start of the table header row. Figure 8-6 is what your page should look like with the added check boxes and submit button.

Description	Owner	Due Date	Location	Importance	Creator
<input type="checkbox"/> Finish LAMP Chapter	David	2012-09-22	Australia	HIGH	David
<input type="checkbox"/> Test LAMP Chpater	David	2012-09-22	Australia	HIGHEST	David

Description:

Owner:

Date:

Location:

Importance:

Creator:

Figure 8-6. Now with extra check boxes

The code section for the form content generation is shown following. You can see that there is another form wrapping the table that has the same method and action. Next is the new hidden input value that sets the action to delete, allowing us to access the correct section in our case statement. Probably the most important change is the addition of the empty td pair and the additional input of type check box. This check box is probably the most complex set of elements we have added—it contains the type of check box, the value that is specially configured to be accessible as an array that is done by adding the square brackets before and after. The final element is the value, which I have set to be the value of the check box. This means that when we go into CGI mode, we will take the array value of check boxes and it will contain a list of the idnumber values for the elements we wish to delete.

Here is a short summary of all the changes we have made to the form to add the check boxes:

```
...
<h1>Pi Todo List App</h1>
<form action="index.php" method="POST">
<input type="hidden" name="action" value="delete" />
<table>
<tr>
<td></td>
...
</tr>
<?php
$result = $mysqli->query("SELECT * FROM todolist");
while($row= $result->fetch_assoc()){
    print "<tr>";
    print "<td><input type='checkbox' name='checkboxes[]' value='".$row["idnumber"]."'" /></td>";
    print "<td>".$row["description"]."</td>";
    print "<td>".$row["owner"]."</td>";
    print "<td>".$row["date"]."</td>";
    print "<td>".$row["location"]."</td>";
```

```

        print "<td>".$row["importance"]."</td>";
        print "<td>".$row["creator"]."</td>";
        print "</tr>";
    }
$mysqli->close();
?>
</table>
<input type="submit"/>
</form>
...

```

Remove Data from DB

We have one final change to make, which is adding the delete processing in the `case` statement. We can use a similar logic to the `insert` statement, whereby we create a SQL variable and then iteratively add the remove for the `idnumbers` one after another. In this case, we will need to have a `for` loop that will run through each element of the `$_REQUEST['checkboxes']` variable that we can get from the `count` function being run on the variable. This gives us a full `for` loop as follows:

```
for($i=0; $i < count($_REQUEST['checkboxes']); $i++){
```

Once we are within the loop, we simply need to pull out the `$i`'th element of the `$_REQUEST['checkboxes']` array each time in the loop. In addition, we will need to add an `or` to the end of each of them so that we will delete each of the elements by `idnumber`; we use `or` because we want it to delete the variable if the `idnumber` is the first number, the second number, or the third, and so on.

However, we will have one issue with the `or`: we will add it to the end of the last element that we don't want. We will need to use the `rtrim` function to trim off the `or` from the end. The `rtrim` function will trim off a given value from the right side of a string, which will give us the SQL we want. Finally, we perform the same query action as in the `insert` and will again check to see whether there is an error. This will give us a code block like this:

```

$SQL="DELETE FROM todolist WHERE";
for($i=0; $i < count($_REQUEST['checkboxes']); $i++){
    $SQL=$SQL . " idnumber=" . $_REQUEST['checkboxes'][$i] . " or";
}
$SQL= rtrim($SQL, "or");
if ($mysqli->query($SQL)== FALSE) {
printf("Error Unable to delete value " . $mysqli->error);
}

```

Figure 8-7 is the final result. I have used the `delete` to remove the original entry so we are just left with the one. Remember that all examples from this book are available online via the Apress website. If you are unsure about how all this code is assembled, please download a copy and have a look!

Description	Owner	Due Date	Location	Importance	Creator
<input checked="" type="checkbox"/> Test LAMP Chpater	David	2012-09-22	Australia	HIGHEST	David

Description:

Owner:

Date:

Location:

Importance:

Creator:

Figure 8-7. Lowering my workload

Troubleshooting

Now let's go over some ways to troubleshoot some issues. First, try making each code change in sequence and then reload the page. You can use the `view source` command to see the full HTML that is being generated by the PHP and see whether the values are showing up correctly. If you can't get that far, have a look at the error log, (the one from right back at the start when we configured Apache). This file will list all of the PHP errors that occur. Have a look to ensure that your quotation marks are open and closed correctly and you haven't left any dangling. Check that you have a semicolon on the end of each of your statements (I always forget this one). Check that your parentheses, square brackets, and circle brackets are opened and closed correctly with no overlap.

Much of the stuff in PHP and web development is trial and error. You should check that each statement or block of code that you add functions and generates as expected. You can use a `print` statement to output variables as they are generated to see what exactly you are doing, which is a great way to diagnose any problems with your SQL statements. Finally, remember that you have access to the full system, go ahead and test inserting values into your SQL and delete them as needed; use the system to work out what it is doing and what causes your problems.

Where From Here?

We've done all this development and we have a functional to-do list that you can use to display, add, and delete entries with. We have used a variety of different programming tools to do things, but there are a lot of things we can do from here. These are some of the changes you could make to your to-do list:

- Add `isset()`; checks to each of the insert values to check that you are actually inserting values, not just blank spaces.
- Wrap each of the inserted values with the `mysqli_real_escape_string` function. This will increase the security of the to-do list because it prevents people writing "nasty" values into your application that will execute as separate SQL queries.

- Change the value option of each of the submit elements, so that they say what we are using them for: inserting or deleting.
- Create a table around the form elements for inserting and give each of the labels and insert fields their own row and own cell. Wrapping them in a table makes the display much more uniform.
- Start looking into CSS because it can make a big difference in how your to-do list looks. Twitters “bootstrap” is a great starting point.

Summary

So we've come a long way since we started. We've done a lot. We have installed and configured Apache, MySQL, and PHP. You have learned how to create simple SQL statements to create and delete databases and tables and then learned how to insert and delete data from these tables. Finally, you have learned how web content is displayed from an Apache web server. Then we combined all and went through some PHP and HTML to generate a whole application that provides access to a database to-do list table.

This has been a monumental undertaking because we have not only gone and installed three different applications and created an interconnected application stack but also gone over three languages to create an application on top of the application stack.

Congratulations—job well done!



WiPi: Wireless Computing

One of the biggest advantages of the Raspberry Pi is that it is tiny and can fit in just about any small nook or cranny you care to think of. This is why it is constantly touted as a great choice for people to use when designing—it can be placed just about anywhere and, with the right peripherals, can do almost any task. Additionally, its low power consumption means that there is less of a barrier to its usage in these situations as it doesn't need a large amount of electrical plumbing to "just work." Which is to say that Raspberry Pi is a super tiny machine that provides a monumental bang for your buck!

However, this is not without some strings attached (see Figure 9-1)! All that power is for naught without a small supply of external electricity. While 5 volts isn't much, it does need to somehow magically enter your Raspberry Pi to turn the lights on. In addition to requiring power, the Pi is almost useless on its own as there is little you can do with the Pi on 5 volts alone beyond making the power LED go on and off. You need to attach devices to make the Pi useful. To use the Pi as a simple home computer you need the following (at a minimum):

- HDMI or composite cable for video
- Ethernet cable for Internet and network access
- A keyboard to enter data with
- An SD card to provide storage



Figure 9-1. Constrained Pi

Suddenly, you see what I'm getting at; the Raspberry Pi, the tiny unattached super machine is actually quite tied down. Tied down with the needs of human usability and interactiveness, these constraints mean running a large number of cables to your Pi just to get some of the basic functionality out of it. As you can probably guess, all is not lost! There are a number of ways you can detach your Pi and make it a lot more mobile and give yourself a lot more flexibility in what you can do with it. That's the purpose of this chapter: to provide you with the tools and some background to remove the wires constraining your Raspberry Pi.

Preparing for WiPi

In addition to having your SD card loaded with a copy of Raspbian, you will want to ensure that your version of Raspbian has been upgraded because the September release included significant improvements in how WiFi access is managed. To get this system up, you first need to run a few quick commands:

```
$ sudo apt-get update  
$ sudo apt-get upgrade
```

These two commands will first tell your `apt-get` instance to update to know which versions of the software are available to you for usage. The second will tell `apt-get` to compare and then install every upgraded version of all software currently installed on the system. As you can imagine, running these commands on a regular basis is a very good idea because it allows you to keep your system running the latest version of the software.

Simple Constraint Removal

Before you begin chopping off cables willy-nilly, you need to understand what you want to achieve with your Pi. This will automatically place a number of constraints on what you want to do with it. I can't presuppose what you want to do with your Pi because there are genuinely so many options. So given this, what I will try and provide here are individual ways that you can cut the cord (not literally) and unconstrain your Pi. You may well be able to make use of your newfound cable-cutting skills with some of the content in later chapters to improve on the designs or to make them better suit your purpose.

Removing Human Input Devices

Start with the easiest and lowest hanging of fruits: the monitor, keyboard, and audio cables. I know it sounds funny to say that cutting out three peripherals at once is easiest, but you really are cutting something much different out of the equation: you. By far, the easiest set of connections to cut is that of the human interaction devices. Much of this book is, in fact, geared toward giving you the skills to work with your Pi without the aid of a monitor, a keyboard, or speakers. This is not to say you can't ever go without; you can always plug them back in as needed to diagnose any really critical issues, but you can start by removing them and managing your Pi remotely. The use of SSH is key here because you can use SSH to remotely connect to your Pi over your network connection. Once you have connected to your Pi, you have full command-line access to perform whatever tasks you need. There is also the added benefit that you can go and change the memory distribution within your Pi using `sudo raspi-config` and set the `memory_split` value to 240/16, which puts most of the memory into free RAM for your system to use and as little as possible into the Graphics display, which you won't normally be using.

Before you assume that this means no use of monitor and keyboard ever, stop. All the things you want to do with your Pi should be tested and configured in a working environment with keyboard, monitor, mouse, and so on, all available for usage. Once you have your system to a state where it works without your interference, you can power on and off without needing to jump back in and manage each time you know you should be able to remove all the human interface devices. This is a great start for anyone: being able to remove all the cables and interface with your system remotely. This is also means you can do other things such as manage your Raspberry Pi via your phone or from a

system at work across the Internet. Once you get the feeling for jumping across the Internet to manage your machine, I guarantee you will ask yourself how you did without it before.

Adding a Remote GUI

Alright, while I will readily admit that there are some true benefits from working within a console-only remote environment there are some people who just cannot get their head around it. And there are a number of cases where it is impractical for whatever reason. In these situations, it's best to use a tool called Virtual Network Computing (VNC) to give you the ability to interact with your Pi's desktop from another machine.

To do this, you will need to install a VNC server on your Pi and you will need a client on your machine. There are numerous different VNC solutions out there, and as long as they obey the VNC protocol, they should all be compatible. For your Pi, you will be using tightvnc because it aims to be efficient and "tight." You can download the clients for both Windows and Linux systems at www.tightvnc.com. But before you go jumping into clients, start by installing the server on your Pi. Run the following command:

```
$ sudo apt-get install tightvncserver
```

This command will install tightvnc with the package management tool; your output should be similar to the following:

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  x11-xserver-utils xfonts-base xfonts-encodings xfonts-utils
Suggested packages:
  tightvnc-java nickle cairo-5c xorg-docs-core
The following NEW packages will be installed:
  tightvncserver x11-xserver-utils xfonts-base xfonts-encodings xfonts-utils
0 upgraded, 5 newly installed, 0 to remove and 91 not upgraded.
Need to get 7,824 kB of archives.
After this operation, 11.7 MB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Once you have Tight VNC Server installed, you need to power it up and start it running so you can connect. So go ahead and execute this command:

```
$ tightvncserver
```

You will then be prompted for a pair of passwords: one is the actual VNC password that allows you to use the mouse and keyboard to manipulate the screen; the second is the optional "View Only" password that allows you to view what's on the screen but not interact. This is what you should expect to see after executing:

```
You will require a password to access your desktops.
Password:
Warning: password truncated to the length of 8.
Verify:
Would you like to enter a view-only password (y/n)? n
New 'X' desktop is raspberrypi:1
Creating default startup script /home/pi/.vnc/xstartup
Starting applications specified in /home/pi/.vnc/xstartup
Log file is /home/pi/.vnc/raspberrypi:1.log
```

Now all you need to do is go ahead and connect with the tightvnc client on the machine you wish to browse from. Go ahead and fire up the client. I am working from a Windows environment, and you should see a screen like in Figure 9-2.




Figure 9-2. TightVNC Client

In Figure 9-2, you can see that I have entered my Raspberry Pi's IP address and :1. Normally a single colon after an IP address designates a port number, but in this case it designates a screen number. If you go back over the output from the tight VNC startup, you can see that it says it created a new desktop of raspberrypi:1, which is what I have connected to. Now press Connect and boom, there is the lovely pink raspberry displayed on your screen and you are able to move around in the Raspberry Pi environment.

If you were to restart, however, you would be without a VNC server again, and running `tightvncserver` will ask you for credentials again. This is not the optimal way to do this, so instead you will run the following command if you wish to restart the server:

```
$ vncserver :1 -geometry 1920x1080 -depth 24
```

This command says to run a `vncserver` on screen :1 with the dimensions 1920 x 1080, which is the maximum resolution of your Pi (full 1080p high definition) and set the color depth to 24 bits. You can lower these values if need be to fit the dimensions of your client screen or to lower the number of resources used by your Pi in hosting a VNC connection. In addition to this, you may wish to run your VNC server on boot so you don't need to SSH in to set up a video connection. After all, that's just redundant! The simple way is to add the `vncserver` command to `/etc/rc.local` as everything in this file gets executed during the boot process.

By default, your `rc.local` file will look like this:

```
#!/bin/sh -e
#
# rc.local
#
# This script is executed at the end of each multiuser runlevel.
```

```

# Make sure that the script will "exit 0" on success or any other
# value on error.
#
# In order to enable or disable this script just change the execution
# bits.
#
# By default this script does nothing.

# Print the IP address
_IP=$(hostname -I) || true
if [ "$_IP" ]; then
    printf "My IP address is %s\n" "$_IP"
fi

exit 0

```

You can then add the following line to start tightvnc just before the exit 0:

```
vncserver :1 -geometry 1920x1080 -depth 24
```

If you want a proper start process, you can use something like the following script as the file /etc/init.d/tightvnc:

```

### BEGIN INIT INFO
# Provides: tightvnc
# Required-Start: $remote_fs $syslog
# Required-Stop: $remote_fs $syslog
# Default-Start: 2 3 4 5
# Default-Stop: 0 1 6
# Short-Description: Start Tight VNC Server at boot time
# Description: Start Tight VNC Server at boot time.
### END INIT INFO
#!/bin/sh
# /etc/init.d/tightvnc

USER=pi
HOME=/home/pi
export USER HOME

case "$1" in
start)
    echo "Starting Tight VNC Server"
    /usr/bin/vncserver :1 -geometry 1920x1080 -depth 24
;;
stop)
    echo "Stopping Tight VNC Server"
    /usr/bin/vncserver -kill :1
;;

```

```
*)
echo "Usage: /etc/init.d/tightvnc {start|stop}"
exit 1
;;
esac

exit 0
```

Once you have created this file, you need to execute the following command to allow the script to be executed:

```
$ sudo chmod +x /etc/init.d/tightvnc
```

Then you can start and stop the script with `/etc/init.d/tightvnc`. Be aware that the first time you run this script you will have to enter a VNC password again as this script is being run as root, rather than the Pi user. Finally, you can set the script to automatically boot with this:

```
$ sudo update-rc.d tightvnc defaults:
update-rc.d: using dependency based boot sequencing
```

Now that you have installed and configured VNC on your Pi and been able to remove the dependency on human input devices for working with the Pi, you can move into the next area.

True WiPi with WiFi

From your earlier work with VNC and SSH, you will have come to realize that in removing human input devices, you have increased the dependency on your network connection. This dependency comes in the form of management because you can't interact directly via keyboard and monitor anymore; you need to have another access method: the Ethernet network connection. Of course, this is still a wire, so you may as well keep the input devices, right? There is a very obvious solution here—WiFi—but to the chagrin of many people the Raspberry Pi doesn't come with WiFi on the board itself. You will need to go out and purchase a wireless adapter for your Pi if you wish to use one.

Okay, simple enough, you say, and off you pop down to the local computer hardware store to pick up a wireless adapter. You look around and urgh. There are so many wireless adapters available! So which of the dozens available is best for the Pi? Will all of them even work with the Pi? While most people would intuitively go and pick up the cheapest that will fit in with their WiFi setup at home, this is not the correct thing to do here as there are other factors at play.

First, all USB devices require power, and, as you are well aware, the Raspberry Pi runs on 5 volts of power. This means that you will need to get a wireless device that runs on a lower amount of power (if you place too much draw on your Pi's USB you can crash it), or you will need to invest in a powered USB device that connects to an external power source (not good for your purposes, but I will cover this later because there are other uses for them). So you need to be mindful of power; you also need to be mindful of compatibility as some of the stranger wireless adapters may not have support within Linux systems. I have included some data from the Raspberry Pi community on their experiences with various wireless adapters to supplement my own.

As you can see in Figure 9-3, I chose to use a D-Link DWA-131 wireless adapter purchased from my local tech store. The rest of this example will cover some of the basics of getting your Pi to communicate with your local wireless network using the DWA-131.



Figure 9-3. D-Link DWA-131 802.11n wireless N nano adapter

Note Your Pi may restart when you connect the wireless adapter, so save any work beforehand.

This doesn't mean you need to use this particular adapter alone; the setup of most wireless adapters to work with your Pi should be the same process once you have managed to get it to register correctly. So without further ado, start and plug in the wireless adapter and get to work. The current state of the board is shown in Figure 9-4.



Figure 9-4. Three down, two to go!

As you can see, I have still got my Ethernet cable attached for now. I will show you how to configure wireless on your Pi from both the command line and the GUI. Before you jump straight into configuration, you need to know exactly what you are working with, so jump into the shell and run the `lsusb` command, which is just like the `ls` command, but for USB devices.

```
$ lsusb  
Bus 001 Device 001: ID 1d6b:0002 Linux Foundation 2.0 root hub  
Bus 001 Device 002: ID 0424:9512 Standard Microsystems Corp.
```

```
Bus 001 Device 003: ID 0424:ec00 Standard Microsystems Corp.
Bus 001 Device 004: ID 07d1:3303 D-Link System DWA-131 802.11n Wireless N Nano Adapter(rev.A1)
[Realtek RTL8192SU]
```

Perfect! You can see from this command that my wireless adapter has been recognized by the system and should be able to function—fantastic. If not, don’t fear; you will simply need to install the driver for your particular network card. A small amount of Google research and you should be able to find the chipset of the WiFi card. The *chipset* refers to the small silicon chip inside the wireless adapter that “runs” the wireless adapter. The chipset is often manufactured by an entirely different company from the one whose badge appears on the device, and it is the chipset with which you need to interface and is thus what determines which drivers you need to find. Once you have found the chipset, a quick search on this should point out what driver packages you will need to install. Once you have installed your driver packages and you can see your wireless adapter in `lsusb`, you can move to the next step.

Since you have the device installed and registered, you should check that the device has registered correctly as a network device. When working with Linux, the most common tool for network management is `ifconfig`. So go ahead and run `ifconfig` and examine the output:

```
$ ifconfig
eth0      Link encap:Ethernet HWaddr b8:27:eb:8a:46:ba
          inet addr:10.0.0.20 Bcast:10.0.0.255 Mask:255.255.255.0
          UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
          RX packets:6558 errors:0 dropped:0 overruns:0 frame:0
          TX packets:268 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:374403 (365.6 KiB) TX bytes:28129 (27.4 KiB)

lo       Link encap:Local Loopback
          inet addr:127.0.0.1 Mask:255.0.0.0
          UP LOOPBACK RUNNING MTU:16436 Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)

wlan0     Link encap:Ethernet HWaddr 90:94:e4:51:81:7a
          UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:0 (0.0 B) TX bytes:0 (0.0 B)
```

The previous configuration actually shows information for three devices rather than just the wireless adapter (two of which you saw in Chapter 3):

- It shows first the `eth0` device that is the physical Ethernet device.
- Next is the `lo` adapter; recall that this is the system’s internal self-reference device that is used when you want to address traffic to within the Pi (or any system) from itself.
- Finally you have `wlan0`, which is the wireless device. You should be able to notice a few differences between the `eth0` and `lo` devices and the `wlan` device—it doesn’t have an IP address (denoted by `inet addr`) or a mask. These are things you will need to configure to get it talking with your wireless network—along with the details of your wireless networks SSID and security arrangements. So that will be your next step.

If your adapter didn't show up, but did register in the `lsusb`, reboot and perform the same troubleshooting steps. The system may be able to recognize your device, but be unable to interface with it correctly. Installing a driver package for your system will likely be the solution here.

To begin with, make this easy and do the configuration via the GUI, which ironically involves installing something via the command line. Go ahead and fire up `apt-get` to install `wpagui`. WPA GUI is an application that has recently been added to Raspbian to make the installation and configuration of wireless devices much, much simpler. It is so recent, in fact, that this chapter needed to be rewritten to include it because the version of Raspbian that was used during the first round of the writing did not include it! If you're unsure about whether you have the correct version, you should run this command because there is no harm in doing so. The command to install this application is the following:

```
$ sudo apt-get install wpagui
```

If you have the latest version of Raspbian, which includes `wpagui`, you will see the following output stating that the package is already installed:

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
wpagui is already the newest version.
0 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.
```

If you have not installed the package, your installation will look something like this:

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  libaudio2 liblcms1 libmng1 libqt4-svg libqtcrypt4 libqtgui4 menu
Suggested packages:
  nas liblcms-utils qt4-qtconfig menu-l10n
The following NEW packages will be installed:
  libaudio2 liblcms1 libmng1 libqt4-svg libqtcrypt4 libqtgui4 menu wpagui
0 upgraded, 8 newly installed, 0 to remove and 96 not upgraded.
Need to get 8,921 kB of archives.
After this operation, 22.6 MB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Fantastic! Now that you have installed `wpagui`, you should be able to make good use of it via the user interface. But before you jump in, go ahead and reboot your Pi because you have made a major change to how the Pi manages your network connections and you want to make sure it's bombproof before continuing. Booted okay? You can run `ifconfig` and the output looks the same, great.

Now you can begin configuring your wireless adapter. You'll start with the GUI and then see how to do it over the command line.

Note If you notice your Pi slowing or failing to connect to WiFi, it may be your keyboard drawing too much power. Try removing it and using VNC or SSH to configure your wireless.

GUI WiFi Configuration

Log in to the GUI and there should be a new icon called WiFi Config, as shown in Figure 9-5.



Figure 9-5. WiFi Config icon

Double-click this icon to open the application and you should be presented with the wpagui window, as shown in Figure 9-6.



Figure 9-6. wpa_gui window

From within the wpagui window, press the Scan button. It will take a few moments, but another window will pop up, as shown in Figure 9-7.




Figure 9-7. Scan results window

This window should look familiar to everyone. This is the window for selecting your wireless network. Find your particular network and then double-click it to open the configuration window (see Figure 9-8).




Figure 9-8. Network details

I am assuming that you should have your details available to set up your configuration. Most of this is intuitive, but there are a few things you need to be aware of:

- SSID is the “name” given to your wireless network, so you need to type this name.
- Next is Authentication, which specifies the type of authentication that you are using: WEP, WPA, WPA2, and so on.
- Next is Encryption, which selects the encryption mode that you will be using.
- The last is PSK, which stands for preshared key. This is your WiFi password.

Now change to the Wireless Security tab and set up whichever authentication method, encryption method, and password you are using for your WiFi. This information is normally inside your router’s configuration system, but can also be written on the router itself or within the router manual. Now click Add. Your wireless settings should be saved into the system and it will pick these up and begin attempting to connect to your adapter; after a few moments, you should expect to see something like Figure 9-9.



Figure 9-9. WiFi connected!

If you didn't see this, have a look at the mouse-over message it should give you an idea of what is going wrong with your connection. Check that your settings are correct and that you have picked the right security method. If in doubt, open your router configuration to double-check. Finally, you can go back into the console and issue `ifconfig` again to check out your wireless adapter settings!

```
wlan0      Link encap:Ethernet  HWaddr 90:94:e4:51:81:7a
          inet addr:10.0.0.59  Bcast:10.0.0.255  Mask:255.255.255.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:772 errors:0 dropped:10 overruns:0 frame:0
          TX packets:35 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:100749 (98.3 KiB)  TX bytes:4254 (4.1 KiB)
```

That looks much more like what you expect; it has an IP address, mask, and so on. It also shows that it has been sending and receiving packets (those are the RX and TX packet numbers) without error. The final test is to ping that IP address from a second machine, and that's it. You can now go ahead and remove that Ethernet cable. As a final test, go ahead and reboot; you should be able to ping your system on the `wlan0` IP address after the reboot, showing that you have successfully kept the wireless adapter details stored across a reboot. Now you can go ahead and remove that next cable and reboot as you want to ensure your system is using the wireless adapter by default.

Managing WiFi from the Command Line

For those of you who don't really want to use the GUI or who consider themselves purists to the true Linux tradition of using only the command line to manage your systems, this is the bit for you! To do this, you will be modifying the NetworkManager configuration. The basic way the NetworkManager system decides how it should manage your Ethernet devices is twofold:

- It checks to see whether the `managed` flag in `/etc/NetworkManager/NetworkManager.conf` is set.
- It takes over management of anything configured within `/etc/network/interfaces`; these settings say how this device should be managed.

You don't want the NetworkManager system to jump in and take over your settings from above, so you need to ensure that it's not going to. This will be an issue only if you installed the GUI earlier, or if NetworkManager came bundled with your Raspbian image, but it is best to be thorough. So go ahead and check `/etc/NetworkManager/NetworkManager.conf` and ensure that the `managed` flag is set to `false`. The file should look like:

```
[main]
plugins=ifupdown,keyfile

[ifupdown]
managed=false
```

Once this is done, go ahead and reboot your Pi; you want to ensure that the Pi comes up correctly and all existing Ethernet connections are working after this first reboot. Now you need to add the configuration into `/etc/network/interfaces` so that it knows how to manage the device and connect to your WiFi. So go ahead and open up `/etc/network/interfaces`. You will see that the file is already partially populated with the following lines:

```
auto lo

iface lo inet loopback
iface eth0 inet dhcp
```

You may also see the following lines in your file, depending on whether your system has already done some wireless configuration or not.

```
allow-hotplug wlan0
iface wlan0 inet manual
wpa-roam /etc/wpa_supplicant/wpa_supplicant.conf
iface default inet dhcp
```

These lines correspond to the loopback and eth0 devices, respectively, and provide some simple details about their management; in this case, the loopback is an automatically managed loopback device, and the eth0 device is given an IP address via DHCP. There may also be some existing configuration for your wlan0 device; if they exist, you should comment them out by adding a # to the front of them or delete them from the file. You now need to add a new configuration for the wlan0 wireless device, so go ahead and add the following to your config below the existing content:

```
auto wlan0
iface wlan0 inet dhcp
    wpa-ssid <Your WiFi SSID>
    wpa-psk <Your WiFi password>
```

This block says that you will have an automatically managed wlan0 device, which gets its connection information via DHCP. You also pass it the SSID and PSK. This configuration will work for both WPA and WPA2 network security implementations. Once you've finished, it's time to test your new device, so go ahead and run the following, which will start the interface. It may take a few moments, but your output should resemble mine:

```
$ sudo ifdown wlan0
$ sudo ifup wlan0
Internet Systems Consortium DHCP Client 4.2.2
Copyright 2004-2011 Internet Systems Consortium.
All rights reserved.
For info, please visit https://www.isc.org/software/dhcp/

Listening on LPF/wlan0/90:94:e4:51:81:7a
Sending on LPF/wlan0/90:94:e4:51:81:7a
Sending on Socket/fallback
DHCPDISCOVER on wlan0 to 255.255.255.255 port 67 interval 8
DHCPDISCOVER on wlan0 to 255.255.255.255 port 67 interval 15
DHCPDISCOVER on wlan0 to 255.255.255.255 port 67 interval 9
DHCPREQUEST on wlan0 to 255.255.255.255 port 67
DHCPOffer from 10.0.0.1
DHCPACK from 10.0.0.1
bound to 10.0.0.59 -- renewal in 444671 seconds.
```

All this shows that the device has connected to my wireless network and has been given an IP address via DHCP, which is exactly what you told it to do! Now suppose that you don't want to use DHCP, but would rather give your WiFi adapter a static IP address. You simply need to change the iface line to read static rather than dhcp and then add details for your IP address, subnetmask and gateway. Your config will look like this instead:

```
auto wlan0
iface wlan0 inet static
    address 10.0.0.57
    netmask 255.255.255.0
    gateway 10.0.0.1
    wpa-ssid <Your WiFi SSID>
    wpa-psk <Your WiFi password>
```

Now that you have successfully configured wireless on your system, you can go ahead and reboot your Pi to check that the wireless adapter automatically comes up. If not, go back and ensure that you set auto wlan0 and that you enabled the managed=false flag in the NetworkManager config file.

If that doesn't work, go over your configuration settings for SSID and WiFi password. You can confirm whether or not you can start your wireless adapter with ifup wlan0; if there are problems connecting and getting an IP address your system should show them here giving you a chance to work out what's going wrong. Finally, if you just can't get it to run, try removing any other attached USB devices; they can draw valuable power away from your WiFi adapter.

One last thing on the command line that may help you is the iw series of commands:

- You can run iwlist with the scan option to get a scan of information about nearby wireless hotspots.
- Or you can use iwconfig to get configuration details about your wireless adapter.

Down to One

Congratulations, you are on your way to a fully wireless Raspberry Pi! You have successfully removed the need for an Ethernet cable and you can now access your Pi wirelessly. It should now be down to just one cable! (See Figure 9-10.)



Figure 9-10. Last cable standing

Removing the Need for Power

There is very little you can do to shortcut the need to provide your Raspberry Pi with power; it needs it, and there are only a finite number of solutions to this problem. By far the most easy and convenient is to get a large USB battery. These are available just about everywhere these days for charging any of the many devices people commonly carry on the go. The simple solution here is to attach your Pi to one of these batteries for use because this kind of battery should guarantee the correct level of power output to your Pi—as they are designed to power USB devices. The important thing when looking at batteries to power your Pi is to be certain they can provide the full 1 amp of power, as some USB batteries are okay for charging but don't pack quite enough juice to power a Pi. Figure 9-11 shows the final WiPi.



Figure 9-11. WiPi at long last!

There are a few hacks out there with a combination of rechargeable batteries and solar panels, which will go much further for you, but those are well beyond the scope of this book.

Powered USB Adapters

I made mention earlier that if you place too much draw on your Pi's USB, you can crash it as it has only got a limited power supply of its own. The solution is to get your hands on a powered USB hub. These powered hubs normally have one connection out and a number of connections in and will also have a mains connection with which to supply power. The principle with these is that people rarely use the full data transmission volume of a USB but will often run out of power, so having a way to supplement the power of a USB device with a mains power source is a good way to give people more flexibility with their devices. If you find that when you attach your device your Pi switches itself off, it's a fair bet that you will need something like one of these to help spread out the power load (see Figure 9-12).



Figure 9-12. Powered USB hub in action

Summary

After all that work, you should have seen a large payoff. You should now know how to configure your Pi to make it wireless. You should be able to set up and connect to your Pi via a VNC client in addition to the normal SSH connections. You have learned how to configure your Raspberry Pi to take a wireless adapter and can attach it to your existing wireless network. Finally, you should be aware of some of the limitations of your Pi with regard to power consumption and how you can resolve these issues.



The Raspberry sPi

Everyone wants to be James Bond. Or at least I know I do. He seems to get into all sorts of situations; then whips out a handy gadget to save the day. The Raspberry Pi is, as we have shown, supremely versatile for such a tiny little piece of hardware. This makes it the perfect tool for a secret agent because with just a little bit of energy you can provide a fully-fledged technological solution at a fraction of the expense of employing John Cleese (admittedly the Pi is not as funny). So, with this in mind, we can now don our tuxedos and get down to secret agent business.

All good secret agents know that you need to be able to detect and be alerted to intruders in your secret hideout. So this makes it the perfect time to introduce you to the Raspberry sPi! The Raspberry sPi is a combination spy camera and alerting system, the perfect solution for spotting dastardly intruders into your inner sanctum! Or for monitoring your pets over the Internet. The basic concept here is that we will be configuring our Raspberry Pi. So without further ado, let's get started.

Materials Needed

Here are the materials needed for this project (also shown in Figure 10-1):



Figure 10-1. The Raspberry sPi kit

- 1x Raspberry Pi
- 1x Micro USB cable (for power)
- 1x Ethernet cable
- 1x USB webcam
- 1x SD card

At this point, you should note that not all webcams are compatible with the Raspberry Pi. This document was created with a Logitech C525 webcam. Thankfully, the Pi community maintains a list of compatible hardware, which is available at http://elinux.org/RPi_VerifiedPeripherals#USB_Webcams. If you are in doubt, please check there to see which peripherals will work with your Pi.

Presetup

Nothing special here; just attach your micro USB for power, attach the Ethernet for network access, attach the SD for storage, and make sure you leave the camera unattached—we will do that later. The rest of this process assumes you are familiar with the following:

- Installing the Raspbian OS
- Knowledge of how to work within a Raspbian Linux environment

If you aren't familiar with either of these two, that's perfectly okay. That's what this book is for after all, so please go back and read Chapter 1 for how to get started with Raspbian and Part 2 of this book to gain familiarity with working in a Raspbian Linux environment.

Getting Started

So, without further ado, let's get started. Go ahead and load up Raspbian and connect your Pi to your network. Go through the Raspbian setup procedures and make sure to enable SSH because you will need to remote in later. (We don't see James Bond carrying round a keyboard, HDMI cable, and a monitor, do we?) Once you are set up, go ahead and SSH into your Pi (covered in Chapter 3).

Okay, once you're logged in, go ahead and issue the command `dmesg` and have a look at the output.

dmesg

The `dmesg` command is a tool called a *driver message* that shows a logged output of all the messages to the kernel, which includes messages from devices and the drivers that are attached to your system. `dmesg` is infinitely useful in determining which devices are attached, and where and how you should interface with them. It also lists driver and kernel messages that are incredibly useful when diagnosing low-level problems. So, let's have a look at the relevant section of the output.

Note `less` is more. If you find that the output for a command is running over the top of the screen and you can't see it any more, reissue the command and then pipe it to `less` (i.e., `dmesg | less`). This should allow you to navigate the output of just about any command using the arrow keys and spacebar to page down.

```
[ 1.998581] Waiting for root device /dev/mmcblk0p2...
[ 2.071268] mmc0: new high speed SD card at address 7d37
[ 2.079917] mmcblk0: mmc0:7d37 SD02G 1.83 GiB
[ 2.088723] mmcblk0: p1 p2
[ 2.132278] EXT4-fs (mmcblk0p2): mounted filesystem with ordered data mode. Opts: (null)
[ 2.145275] VFS: Mounted root (ext4 filesystem) on device 179:2.
[ 2.155272] Freeing init memory: 200K
[ 2.187333] usb 1-1: new high speed USB device number 2 using dwc_otg
[ 2.418045] usb 1-1: New USB device found, idVendor=0424, idProduct=9512
[ 2.427750] usb 1-1: New USB device strings: Mfr=0, Product=0, SerialNumber=0
[ 2.438708] hub 1-1:1.0: USB hub found
[ 2.445041] hub 1-1:1.0: 3 ports detected
[ 2.727674] usb 1-1.1: new high speed USB device number 3 using dwc_otg
[ 2.837920] usb 1-1.1: New USB device found, idVendor=0424, idProduct=ec00
[ 2.857413] usb 1-1.1: New USB device strings: Mfr=0, Product=0, SerialNumber=0
[ 2.878968] smsc95xx v1.0.4
[ 2.942985] smsc95xx 1-1.1:1.0: eth0: register 'smsc95xx' at usb-bcm2708_usb-1.1, smsc95xx
USB 2.0 Ethernet, b8:27:eb:8a:46:ba
[ 11.006171] EXT4-fs (mmcblk0p2): re-mounted. Opts: (null)
[ 11.426960] ### snd_bcm2835_alsa_probe c05c88e0 ##### PROBING FOR bcm2835 ALSA
device (0):(1) #####
[ 11.442869] Creating card...
[ 11.448268] Creating device/chip ..
[ 11.454771] Adding controls ..
[ 11.460340] Registering card ....
[ 11.475463] bcm2835 ALSA CARD CREATED!
[ 11.487786] ### BCM2835 ALSA driver init OK ###
[ 18.493739] smsc95xx 1-1.1:1.0: eth0: link up, 100Mbps, full-duplex, lpa 0x45E1
[ 24.672322] Adding 102396k swap on /var/swap. Priority:-1 extents:1 across:102396k SS
```

As you can see, the output from `dmesg` is a little offputting and hard to understand. Much of the information relies on understanding a few basic tenets of how the `dmesg` system works. First is the numerical values within the square brackets [and]. This value is a timer value that registers the amount of time that has passed since system launch. This is a good way to follow the sequence of events in your system and to tell recent ones from those that occurred in the past.

Next is the content itself. Most of these lines are diagnostically useless alone, but together they form a very interesting picture of what exactly has gone on in my Pi. Let's start with the first line:

`waiting on root device /dev/mmcblk0p2`

The first few lines are obvious: the system is waiting for something, but what is `/dev/mmcblk0p2`? Well, the first giveaway is the `/dev`; this is where the Linux OS lists all its devices so we know `/dev/mmcblk0p2` is a device. Second is the `mmcblk0p2`, which looks like a string of garbage but is actually a series of abbreviations.

- `mmc` means multimediacard (the flash device standard)
- `blk` means a block storage device
- `0` means the first logical device of this type
- `p2` means partition two of the device

So we are waiting for the SD card. The second through fourth lines are what we are waiting for: the SD card to register as you can see it actually registers as each of the subdevices we mentioned, in order no less! After that, the

next two lines are the device mounting the filesystem on this device as readable. Continuing down, you can see the registration of the Pi's USB devices (lines 8–15), its Ethernet port (lines 16, 17, and 26) and finally an Advanced Linux Sound Architecture (ALSA) device, which is the audio port (lines 19–25).

Now that you understand why we didn't just attach the webcam to begin with, I wanted to view the `dmesg` output. Given that you know what to look for, now go ahead and attach your webcam and run `dmesg` again. Have a look at the new lines that have shown up. These should all be related to your newly attached webcam! Here is the output from me attaching mine:

```
[ 8168.793423] usb 1-1.2: new high speed USB device number 4 using dwc_otg
[ 8169.147691] usb 1-1.2: New USB device found, idVendor=046d, idProduct=081d
[ 8169.147736] usb 1-1.2: New USB device strings: Mfr=0, Product=0, SerialNumber=1
[ 8169.147758] usb 1-1.2: SerialNumber: 8627F4C0
[ 8169.314171] Linux video capture interface: v2.00
[ 8169.336482] uvcvideo: Found UVC 1.00 device <unnamed> (046d:081d)
[ 8169.402071] input: UVC Camera (046d:081d) as /devices/platform/bcm2708_usb/usb1/1-1/1-1.2/
1-1.2:1.2/input/input0
[ 8169.402280] usbcore: registered new interface driver uvcvideo
[ 8169.402299] USB Video Class driver (1.1.1)
[ 8169.460903] usbcore: registered new interface driver snd-usb-audio
```

So, what can you tell from these lines? For starters, you can see a large difference in timing for this particular entry in the `dmesg` log, which represents the delay from start up until I connected my device. Next, you will recognize that the first four lines represent the USB ports picking up a new attached device; in fact, the first line references the module used to do this `dwc_otg` (short for DesignWare Cores—On The Go). Line 5 shows that it's a video capture interface, and line 6 shows that the module `uvcvideo` found a USB Video Class (UVC) device that it identifies and registers. Finally, we can also see that the webcam's microphone has also been registered on the last line as a `snd-usb-audio` device. From this we can see that my webcam has been attached and has registered itself with a driver into the OS.

Now that the device has been registered, we need a way to reference it so it can be used by an application. You will notice that the `dmesg` output didn't actually specify which device we added, so let's have a look and see if we can't match it the other way. Let's start by listing the contents of `/dev` (results shown in Figure 10-2).

```
$ ls /dev/
```

```
pi@raspberrypi ~ $ ls /dev/
autofs          loop0      network_latency   ram4      tty0     tty23  tty38  tty52  ttyAMA0    vcs5
block           loop1      network_throughput ram5      tty1     tty24  tty39  tty53  ttyprintk  vcs6
btrfs-control  loop2      null                ram6      tty10    tty25  tty4   tty54  uinput     vcsa
bus             loop3      ppp                 ram7      tty11    tty26  tty40  tty55  urandom   vcsa1
cachefiles     loop4      ptmx                ram8      tty12    tty27  tty41  tty56  usbdev1.1 vcsa2
char            loop5      pts                 ram9      tty13    tty28  tty42  tty57  usbdev1.2 vcsa3
console         loop6      ram0                 ram0     random   tty14    tty29  tty43  tty58  usbdev1.3 vcsa4
cpu_dma_latency loop7      ram1                 ram1     raw      tty15    tty3   tty44  tty59  usbdev1.4 vcsa5
disk            loop-control ram10                ram10    root     tty16    tty30  tty45  tty6   v4l      vcsa6
fb0             MAKEDEV   ram11                ram11    shm      tty17    tty31  tty46  tty60  vchiq    video0
fd              mapper    ram12                ram12    snd      tty18    tty32  tty47  tty61  vc-mem   xconsole
full            mem       ram13                ram13    sndstat  tty19    tty33  tty48  tty62  vcs      zero
fuse            mmcblk0   ram14                ram14    stderr   tty2     tty34  tty49  tty63  vcs1
input           mmcblk0p1  ram15                ram15    stdin    tty20    tty35  tty5   tty7   vcs2
kmsg            mmcblk0p2  ram2                 ram2     stdout   tty21    tty36  tty50  tty8   vcs3
log             net       ram3                 ram3     tty     tty22    tty37  tty51  tty9   vcs4
pi@raspberrypi ~ $
```

Figure 10-2. Contents of `/dev` directory

That is a few devices. We can quickly eliminate a few right off the bat: all the `tty` devices because they are teletype devices, all the `ram` devices because that's the Pi's RAM, all the `loop` devices because they are loop back connectors, all the `memblk` devices because they relate to the SD card, and all the `vcs` devices because they are virtual consoles. There are still a few, but the best looking is `video0`. We are looking for a webcam, right? It registered as a video capture device didn't it? Okay, so we think that's right—but how can we tell?

`udev`, more specifically the `udev` administration function, is the device manager for the Linux kernel. It governs how and where devices register themselves so that they can be accessed as part of the OS by applications. In Debian (and thus Raspbian), the inner workings of `udev` can be accessed with the `udevadm` command, so let's go ahead and use it. Execute the following:

```
$ udevadm info -q all -n /dev/video0
P: /devices/platform/bcm2708_usb/usb1/1-1/1-1.2/1-1.2:1.2/video4linux/video0
N: video0
S: v4l/by-id/usb-046d_081d_8627F4C0-video-index0
S: v4l/by-path/platform-bcm2708_usb-usb-0:1.2:1.2-video-index0
E: DEVLINKS=/dev/v4l/by-id/usb-046d_081d_8627F4C0-video-index0 /dev/v4l/by-path/platform-bcm2708_usb-usb-0:1.2:1.2-video-index0
E: DEVNAME=/dev/video0
E: DEVPATH=/devices/platform/bcm2708_usb/usb1/1-1/1-1.2/1-1.2:1.2/video4linux/video0
E: ID_BUS=usb
E: ID_MODEL=081d
E: ID_MODEL_ENC=081d
E: ID_MODEL_ID=081d
E: ID_PATH=platform-bcm2708_usb-usb-0:1.2:1.2
E: ID_PATH_TAG=platform-bcm2708_usb-usb-0_1_2_1_2
E: ID_REVISION=0010
E: ID_SERIAL=046d_081d_8627F4C0
E: ID_SERIAL_SHORT=8627F4C0
E: ID_TYPE=video
E: ID_USB_DRIVER=uvcvideo
E: ID_USB_INTERFACES=:010100:010200:0e0100:0e0200:
E: ID_USB_INTERFACE_NUM=02
E: ID_V4L_CAPABILITIES=:capture:
E: ID_V4L_PRODUCT=UVC Camera (046d:081d)
E: ID_V4L_VERSION=2
E: ID_VENDOR=046d
E: ID_VENDOR_ENC=046d
E: ID_VENDOR_ID=046d
E: MAJOR=81
E: MINOR=0
E: SUBSYSTEM=video4linux
E: TAGS=:udev-acl:
E: UDEV_LOG=3
E: USEC_INITIALIZED=8168852755
```

Wow, okay, that's a big block of scary looking output. As I'm sure you've guessed by now, we just asked the `udev` system to give us an `info` query for all info relevant to the device with name `/dev/video0`, which it well and truly has. However, before hyperventilation sets in, actually have a look at the first few lines. The first line should spring out as familiar—it's the device identifier that we saw in `dmesg` from when the webcam was attached! In fact, most of the information in here is all data that was referenced somewhere in the `dmesg` block. Okay, so given that we can match the device from `dmesg` to the device here in our system, it's safe to say that the USB webcam we attached is available on the filesystem as `/dev/video0`.

Troubleshooting

If, unlike mine, your webcam had issues when registering with the kernel in `dmesg`, you may need to install a driver. You should check with the manufacturer and online to see whether there is a driver that's available for your webcam. If so, install the driver and then try reattaching your webcam to see if it can be registered. You may even get lucky, and your device driver will have the kernel list where the webcam is registered in `/dev` in `dmesg`—as it did with my SD card in the preceding example.

Additionally, if you don't have a device such as `/dev/video0` listed, you should check to see whether any of the other devices match up with output that relates to your device from `dmesg`. While Linux tries to be sensible in how the `/dev` filesystem is laid out, manufacturers and software engineers can do funny things to try and stand out—with the upshot that you sit there screaming about how you cannot find a device. So just be patient and have a look around and use the process of elimination to try and work out where your device has been attached. The goal here is to ensure that you know which `/dev` file references your webcam.

Motion Capture

Okay, system configured, check. Webcam attached, check. Now let's make all this actually do something! To do the webcam capture, we will be using the Linux motion capture ironically named *motion*. So let's go ahead and install it. Run the following command to direct apt-get to download and install the motion package:

```
$ sudo apt-get install motion
```

Your output should resemble the following:

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  ffmpeg libav-tools libavcodec53 libavdevice53 libavfilter2 libavformat53 libavutil51
libdc1394-22 libdirac-You
  libjack-jackd2-0 libmp3lame0 libpostproc52 libpq5 libraw1394-11 libschoedinger-1.0-0
libspeex1 libswscale2 libtheora0 libva1
  libvpx1 libx264-123 libxvidcore4
Suggested packages:
  jackd2 libraw1394-doc speex mysql-client postgresql-client
The following NEW packages will be installed:
  ffmpeg libav-tools libavcodec53 libavdevice53 libavfilter2 libavformat53 libavutil51
libdc1394-22 libdirac-encoder0 libgsm1
  libjack-jackd2-0 libmp3lame0 libpostproc52 libpq5 libraw1394-11 libschoedinger-1.0-0
libspeex1 libswscale2 libtheora0 libva1
  libvpx1 libx264-123 libxvidcore4 motion
0 upgraded, 24 newly installed, 0 to remove and 71 not upgraded.
Need to get 8,365 kB of archives.
After this operation, 17.6 MB of additional disk space will be used.
Do you want to continue [Y/n]? y
```

I don't know about you, but I'm way too excited to bother with all the setup and config; let's just fire up motion and see what it does! Start it up with this:

```
$ motion -s
```

The light on the camera goes on and action! Start waving and moving about, because you want it to capture you!

```
[1] Changes: 3374 - noise level: 15
[1] Changes: 3198 - noise level: 15
[1] Changes: 3011 - noise level: 15
[1] Changes: 2922 - noise level: 15
[1] Changes: 2555 - noise level: 15
[1] Changes: 2390 - noise level: 15
[1] Changes: 2491 - noise level: 15
[1] Changes: 2874 - noise level: 15
[1] Changes: 2817 - noise level: 15
[1] Changes: 3238 - noise level: 15
[1] Changes: 3093 - noise level: 15
[1] Motion detected - starting event 1
[1] File of type 1 saved to: ./01-20120910203217-05.jpg
[1] Changes: 2912 - noise level: 15
[1] File of type 1 saved to: ./01-20120910203217-06.jpg
[1] Changes: 2480 - noise level: 15
[1] File of type 1 saved to: ./01-20120910203217-07.jpg
```

There you go, you can see it capturing changes in the image and registering any changes in noise level as you move. And finally, it's capturing images! If it may take a moment to write the image file out, depending on your SD card speed—so be patient. Awesome, works like a charm, to cancel out of the application you can go ahead and press Ctrl+C. If you want, you can go ahead and log in with console to have a look at any images it's captured. Inspecting the output is generally a good idea because your webcam driver may need fine tuning or may not generate valid output. If you don't want to get out a cable and HDMI, you can copy these images off with tools such as `scp`. From a Windows machine, you can use a tool such as `winscp` to copy any files off. If you are copying to a Mac or Linux system, you just need to copy the files by issuing the following command:

```
$ scp <raspberry pi's ip>:/home/pi/*.jpg.
```

Note Please remember the period on the end of the `scp` command It's important because it is the destination identifier for the command.

This command says to copy from the server on the Raspberry Pi identified by IP in the directory `/home/pi` the home directory of the `pi` user and the first place you will be when you log in (I said I was excited, didn't I?) and match any file that ends in `.jpg`. The final period on the end is the destination: in Linux a period is a reference to the current directory, so that's where we are copying.

OK, now that we have had fun with the pictures, let's get down to business. The aim for the Raspberry sPi is to have it function with no intervention; this means we will need to have the motion application function without us telling it to. Raspbian has handled most of this for us already, but there are still a few little things we need to do.

Note You will need to edit these files as root, so remember to start your text editor with `sudo`.

First, open and edit the file `/etc/default/motion` and change the line `start_motion_daemon=no` to be `start_motion_daemon=yes`; this file governs whether the daemon will start when we issue the daemon's `start` command,

and having it run on command is exactly what we want. Now we tackle the config, so go ahead and open it. The file is `/etc/motion/motion.conf`, so open it up. First thing you will note is that daemon mode will be set to off, so go ahead and set that to on because we want motion to be started as a daemon, which means it will run in the background. Next thing to check is the videodevice line and ensure that the listed device matches the devices we confirmed earlier (for most of us, that should be `/dev/video0`).

Following this, we begin getting into the nuts and bolts of the situation. First we can modify the height, width, and framerate. The defaults are to capture two images per second at a resolution of 320 x 240. This creates a file about 12 Kb in size—you can check by examining the files that we generated earlier during our testing of the motion capture. You can increase this as you see fit; just remember not to exceed the maximum available resolution of your webcam and also that each file will increase in size as resolution increases.

Next is threshold. The threshold value is the amount of change that needs to occur so that an image will be captured. You can leave this value for now, but remember it later because you may want to make the capture more or less sensitive.

Next is `ffmpeg_cap_new`, which determines whether or not you want to capture a video file. I have turned this one off, but you can keep it on if you want to generate video files of your motion captures. These files are generated as .swf files now (think YouTube) but you can change them by editing the `ffmpeg_video_codec` variable.

Finally, you can change the `target_dir`, which says where you will output the image files when running as a daemon. This one is important because you will need to know where you are to copy your files from. Additionally, it's important to be aware that the `/tmp` directory is just that, a temporary space. This temporary space is emptied each time your OS is restarted.

This kind of functionality is perfect for what we intend because we don't want to have to deal with extra files that will stay on the disk until we clean it out; simply reboot your Raspberry sPi and off you go again. However, there are cases where we may want these files to be saved beyond a restart of the system—but more on them later.

Finally, save all your changes and exit so we can test the new config! My config is displayed here for your reference:

```
daemon on
process_id_file /var/run/motion/motion.pid
setup_mode off
videodevice /dev/video0
v4l2_palette 8
input 8
norm 0
frequency 0
rotate 0
width 320
height 240
framerate 2
minimum_frame_time 0
netcam_tolerant_check off
auto_brightness off
brightness 0
contrast 0
saturation 0
hue 0
roundrobin_frames 1
roundrobin_skip 1
switchfilter off
threshold 1500
threshold_tune off
noise_level 32
noise_tune on
despeckle EedDf
```

```
smart_mask_speed 0
lightswitch 0
minimum_motion_frames 1
pre_capture 0
post_capture 0
gap 60
max_mpeg_time 0
output_all off
output_normal on
output_motion off
quality 75
ppm off
ffmpeg_cap_new off
ffmpeg_cap_motion off
ffmpeg_timelapse 0
ffmpeg_timelapse_mode daily
ffmpeg_bps 500000
ffmpeg_variable_bitrate 0
ffmpeg_video_codec swf
ffmpeg_deinterlace off
snapshot_interval 0
locate off
text_right %Y-%m-%d\n%T-%q
text_changes off
text_event %Y%m%d%H%M%S
target_dir /tmp/motion
snapshot_filename %v-%Y%m%d%H%M%S-snapshot
jpeg_filename %v-%Y%m%d%H%M%S-%q
movie_filename %v-%Y%m%d%H%M%S
timelapse_filename %Y%m%d-timelapse
webcam_port 8081
webcam_quality 50
webcam_motion off
webcam_maxrate 1
webcam_localhost on
webcam_limit 0
control_port 8080
control_localhost on
control_html_output on
track_type 0
track_auto off
track_motorx 0
track_motory 0
track_maxx 0
track_maxy 0
track_iomojo_id 0
track_step_angle_x 10
track_step_angle_y 10
track_move_wait 10
track_speed 255
track_stepsize 40
```

This test will be just like the one before; we simply want to start motion and check that it will run and capture images. This time, however, we want to use the config file and the Linux start commands because it allows us to emulate a power on of our Raspberry sPi. So go ahead and execute `sudo /etc/init.d/motion start` this should start motion. You should see your camera light come on and if you move about, you should start seeing image captures appear almost immediately in `/tmp/motion` (or wherever you pointed `target_dir`). You can also check that the process is running by typing the following to check for a running motion process:

```
$ ps -ef | grep motion
```

Troubleshooting

Still not working? Go back over the section on connecting your webcam, unplug it, plug it back in, and check for the output from `dmesg`. Validate that the output you see in `dmesg` matches what you get from `udev` in examining the device you have said is your webcam. Did the motion package install correctly? examine the output from `apt-get` to see that it did; if not, try installing again or try this:

```
$ apt-get --reinstall install motion
```

Does motion run when you execute it on the command line with `motion -s`? If not, examine the messages on the screen, which will help you understand what has gone wrong.

If all that failed, go over your motion configuration file. Did you accidentally mis-set one of the variables? Did you forget to set the value in `/etc/default/motion` to yes? Is the motion daemon running, but you can't see images appear in your given output directory? If not, examine the file `/var/log/messages`, which is a repository of messages from daemons started by the system, including motion. The output that was onscreen before should now be in this file.

Alerting Yourself

Thus far, we have attached a webcam, configured the webcam, installed motion detection software and configured it to start automatically along with your Raspberry sPi. Now we have a fully fledged monitoring system that you can use to spy on your enemies and keep you notified about intruders entering your inner sanctum. Well, almost. You now need to create a system that will allow you to send messages to yourself so that you can be alerted when intruders enter.

To do this, we will be taking advantage of one of the oldest and most efficient messaging systems ever developed: e-mail. We won't, however, be going as far as to set up our own mail server and domain; that's a big task and well outside the scope of what we will be doing in this chapter. Instead, we will be aiming to set up a lightweight mail client to forward messages to a mail service in which it can then be sent on to you on whatever handy spy gadget you use to receive e-mail. Now that you understand the plan, let's begin.

Installing SSMTP

The application we will be using to send our mail to the mail server is called SSMTP. SSMTP is a highly simplified mail transfer agent (MTA) that is used when a system administrator needs to be able to send mail from a server, but doesn't need to be set up as a fully fledged mail domain. The advantage here should be obvious: a significantly less intensive setup process and much lower overhead—all of which should scream Raspberry Pi to you.

With all this in mind, let's begin. Start by issuing the command to download and install the SSMTP application and install it:

```
$ sudo apt-get install ssmtp
```

Your output should look like mine:

```
$ sudo apt-get install ssmtp
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  libgnutls-openssl27
The following NEW packages will be installed:
  libgnutls-openssl27 ssmtp
0 upgraded, 2 newly installed, 0 to remove and 71 not upgraded.
Need to get 272 kB of archives.
After this operation, 279 kB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Once we have installed SSMTP we need to configure it. The SSMTP configuration file is `/etc/ssmtp/ssmtp.conf`. Open it up with your favorite text editor and get started.

The first thing we need to configure is where the mail is destined to go. Most people who read this book will have some different and specific location they want their mail to go. Although all good spies want to be as efficient as possible, I can deal only with what's in front of me, so we will be configuring this for the most generic and free option available to everyone: Gmail.

So we need to configure where the mail is to be sent; in SSMTP, this is governed by the `mailhub` value. This value is the DNS name of the mail server for the domain we will use to send outbound mail. This name is traditionally `mail.<yourdomain>.<whatever>`, but for us it will be `smtp.gmail.com:587`, which is Google's SMTP server.

Having set the Gmail server, you may have noticed that we included a :587 on the end. The studious among you will recognize this as a port number. The even more studious would be aware that 587 is not the default port for SMTP (it is port 25). It is the port dedicated for Transport Layer Security (TLS) e-mail, (secure e-mail over SSL). This means we need to add another pair of options to ensure that we have encrypted communications:

```
UseTLS=YES
UseSTARTTLS=YES
```

The next option to configure is the `hostname` option, which is the hostname listed by the Raspberry sPi. Go ahead and list a hostname for the sPi if you have a domain. If not, just give it something like `raspberry.spi`. The next option should be `FromLineOverride=yes`, which allows us to set the "from" field on an e-mail if we so desire.

Finally, the last sets of options are related to authentication. There are three options you need to configure here. The first two are the most obvious: `username` and `password`. So go ahead and add `AuthUser=username@gmail.com` and `AuthPass=password`. The third, which is a little more circuitous, is the `AuthMethod` option, which is used to specify which authentication method will be used to allow users to register with the server. The method we will be using for Google is the `LOGIN` method. Although other mail servers may not require this to be specified, with Google we should set `AuthMethod=LOGIN`.

That's it; your config should now look something like this:

```
root=
Mailhub=smtp.gmail.com:587
UseTLS=YES
UseSTARTTLS=YES
Hostname=raspberry.spi
FromLineOverride=yes
AuthUser=username@gmail.com
AuthPass=password
AuthMethod=LOGIN
```

Note It should go without saying, but you should always have a secure root password.

One final change to make is to secure this data so that nobody except those that we intend can read the file. This is incredibly important because you have put your precious GMail password in the file! The solution is to change permissions so that nobody but the intended users can read it. Execute the following:

```
$ sudo chmod 640 /etc/ssmtp/ssmtp.conf
```

This will change the file so that the only user who can access it is root, and the only users who can are members of the mail group.

Once SSMTP is configured, all you need to do is invoke it to send the e-mail and get it to send the captured... Wait. There's something more we need; we need a way to invoke SSMTP so that we can tell it to send the file to our e-mail and alert us to the intruders. And we need a tool that can do all this from the command line. The best tool here is one called **mutt**.

Installing and Using mutt

The **mutt** tool is a text-based e-mail client with the fantastic motto of "All mail clients suck. This one just sucks less." (Isn't it obvious why we are using it?) The other benefit of **mutt** over other mail clients is that it is happy to send attachments from files on the filesystem—which is key to what we need to make the sPi send images of our intruders.

Now you know that it is **mutt** that we are installing, go ahead and install it by using the following:

```
$ sudo apt-get install mutt
```

Your output should look like this:

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
 libgpgme11 libpth20 libtokyocabinet9
Suggested packages:
 gpgsm gnupg2 urlview mixmaster
The following NEW packages will be installed:
 libgpgme11 libpth20 libtokyocabinet9 mutt
0 upgraded, 4 newly installed, 0 to remove and 71 not upgraded.
Need to get 1,985 kB of archives.
After this operation, 7,181 kB of additional disk space will be used.
```

Did **mutt** install okay? Good. As I mentioned earlier, **mutt** is designed to be small, simple, and easy to use and configure—lofty goals. With that in mind, the setup process for **mutt** is incredibly simple: we just need to tell **mutt** how to use our MTA (which is SSMTP). To do this, we need to create a file that will live hidden in the Pi users' home directory. The file is **.muttrc**, which will contain only one line:

```
set sendmail="/usr/sbin/ssmtp"
```

Instead of using a text editor we will use a "cheat" to do this. Run the following command:

```
$ echo "set sendmail=\"/usr/sbin/ssmtp\" > ~/.muttrc
```

That's it.

Note The backslashes before the quotation marks are escapes.


You can go and look in the root user's home directory for the file `.muttrc` and check the contents. Yep, it's there and exactly as expected. The `.muttrc` file is a config file used by `mutt` to load important configuration options automatically on startup. Many different applications use hidden files that end in `rc`, so if you are wondering how an application is storing certain configuration options, have a look in your home directory for files starting with a period and ending in `rc`.

Testing the Alerting System

Now that we have installed and configured SSMTP and `mutt`, we need to test them in combination. The way to use `mutt` is, strangely enough, to run the `mutt` command. The syntax is rather simple, so send yourself an e-mail from the command line. Execute the following command:

```
$ echo "the quick brown fox jumps over the lazy dog" | sudo mutt -s "[INTRUDER ALERT] Test of intruder system" <your email>@gmail.com
```

I know it's a long command, but don't worry because I will explain it in a moment. Go ahead and check your e-mail; you should have one there from the user root just as shown in Figure 10-3!



the quick brown fox jumps over the lazy
dog



Figure 10-3. E-mail from the Raspberry sPi

Now I'll explain the syntax and how we will use this to send alert messages from the sPi. The preceding message can be broken down into a few much smaller segments. The basic syntax that I use is this:

```
echo <mail content> | mutt -s <subject> <recipient>
```

You may be wondering about the use of echo. You need to generate the value that will be passed as output and then have the pipe direct it into mutt, rather than providing it as an argument.

Note You can include whole sentences with spaces in a shell command by enclosing the whole thing in quotation marks. You can even escape quotes within a command to create quotes within quotes!

We need to also be able to send attachments with our e-mail. mutt has us covered here because it provides the capability to attach files to its outgoing e-mail. The reason why I elected to use mutt to create the Raspberry sPi is because it is not only easy to configure but also provides the capability to send attachments.

The syntax to send an attachment with the e-mail we just sent is as follows:

```
echo <mail content> | mutt -s <subject> -a <filename> -- <recipient>
```

Yes, it's as easy as that! The -a signifies attachments, and the -- is used to separate attachments from recipients.

Troubleshooting

Did everything work as expected? If not, try configuring both SSMTP and mutt from scratch. When you send an e-mail, does it output an error regarding login details? If so, double-check your username and password. Is it an error regarding unsupported login type? If so check the TLS settings and the AuthMethod setting. Did you spell your mail server incorrectly? Is their trailing whitespace? That is spaces or tabs at the end of any of your entries. When you changed permissions did you move them too far so that now you can no longer read the file?

You may also encounter an error if you try and send an e-mail as the pi user, without sudo because the pi user is not a member of the mail group on your system. If you wish to allow the pi user or any other user for that matter, you need to add them using gpasswd, the syntax is this:

```
$ gpasswd -a <username> mail
```

Okay, so all sorted? You can capture the presence of intruders. You can alert yourself from the Raspberry sPi. You can send the captured images of your intruders to yourself from your sPi. So what's left? Automating it all.

Bringing It All Together

So far we have done the following:

- Attached a webcam to the Raspberry sPi
- Installed and configured the monitor application to use the webcam
- Used the webcam and monitor application to function as a motion sensor and capture images
- Installed and configured SSMTP as an MTA

- Installed and configured `mutt` mail application to use SSMTP
- Sent e-mail from the command line to ourselves using SSMTP and `mutt`

The remaining problem is that we have done most of this with human intervention. We need a solution that enables the capture of images to automatically trigger an e-mail being sent to us, with the image attached. Although I'm sure that someone has dreamed up a software solution to this problem, there is no reason why we can't write a solution ourselves!

Understanding the Problem

The first step in any software design problem, no matter how small a job, is to understand what the exact needs of the software solution are. So let's start by listing our goals:

- Capture of a new image should trigger the solution
- An e-mail should be sent with any new captured image(s)
- Should function automatically without needing to be started manually

Okay, that seems like the solution we need. But there are a few other things that having those two main requirements brings into play. These extra requirements are these:

- Only new images should be sent
- Only images captured by motion should be sent
- The application should always be running and not need to be manually started, just like motion

Okay, that seems a bit better; we are now pretty firm on what exactly our software solution should do. But what about how it should do it? Maybe we should be a little specific on how the application we are to write will achieve its goal.

Some things we should consider when developing the application include these:

- We should send the time that the image was captured
- We should check and send regularly, but given that e-mail has a delay, a semiregular check should be enough

Okay, now we're cooking. Those are looking like some good requirements to remember when building our application. We were specific in what we wanted and about how we should achieve it. Now that we know exactly what we want in our application, we can move on to the next step.

Making Decisions

The next step after working out what we want to achieve is to make some decisions about how it should be achieved. So, first decision: how should we write this? Well, much of what we want to achieve here is manipulation of shell commands. We want to list image files and then use those image files within another shell command to send us the e-mail alert. Additionally, you learned bash earlier in this book, so that makes it a perfect fit for what we want to develop.

Okay, we know what language we will be working in. Next, how will we start this one? I see two options here and both will inform which direction we take:

- One is to have the application executed from within the start script for `motion` because we want them to run together. This means that we can ensure that they are both running at the same time and that the script would need to be always running and then check periodically.
- The second option is to have the script automatically executed regularly. This means we wouldn't need to start it and keep the script running permanently. It also means that if the worst happens and the script dies, we don't need to deal with that and reboot it because it will be automatically started again. It does mean that we need to have something that will execute the script on a timer, but that can be handled by the cron daemon.

So, we have two options to decide from, but let's put off that decision for a moment because there is another thing we need to look at first.

What will we do with the images? We have two options again:

- First is to delete each image after we send it. This is good because it means we can save on space as we clean up after ourselves.
- The second option is to keep track of all the images that we have and send only the new ones. This means keeping a running track of which images we have in stock so we can determine which ones to send out.

So, now that you understand the options, it is time to make decisions. I know that it's hard to decide, so let me make it easier for you. Instead of four possible scripts, let's narrow it down to two:

- A script that will be run from `cron` and will search for images, send them, and then remove them
- A script that will start with the `motion` daemon and will monitor all the files, keep a running count and then recognize when there is a new image.

Okay, now that we have two working designs, let's head on to the next step.

Designing Solution

"Wait; you have two solutions and you didn't actually pick one. You giant flake." I hear you say that and it's true; I didn't. Because we had two very different ways to achieve our one goal we felt it would be best to cover both, give you the options, and share some guidance. So, let's start by looking at the commonalities between the two; first, they will both need to have a reference for the folder where `motion`'s captured files will be going. We could keep this in a variable and just edit it when we change the output folder, but that is a lot of effort to go to, changing two whole files. That's not the Linux way. So let's just load it from the `motion` config file and be done with it. It means that we need two variables: a variable that will contain the `motion` config file and the second that we will create dynamically to contain the file location. The variable that will contain the filename is simple enough with `MOTIONCONF=/etc/motion/motion.conf`.

Note All the following is being done as the root user (you can `sudo su` to get a root shell) because it is the root user who will spawn `motion`.

The next step is to pull the variable out, and the best way to do this is with a combination of two commands: one is grep, which will find any information in a file that contains a particular pattern, in this case any line that contains the line `target_dir`, but we also need to ensure that we avoid comments, so we need to grab lines that start with `target_dir`. This is achieved by simply adding one extra symbol so that grep can know that we are looking for a value that starts with our desired value. This gives us a command of `grep "^\target_dir" $MOTIONCONF`. If you execute this command you will see the following result:

```
root@raspberrypi ~ $ grep "^\target_dir" $MOTIONCONF
target_dir /tmp/motion
```

Now the next problem with our command rears its head: we need to get rid of the `target_dir` portion so we only have the location we are after. To remove the first part, we need to use the awk command, which will split the given data at whitespace values by default. This is perfect and then we just need to print the second value that results. So building this on to our command, we get this:

```
root@raspberrypi ~ $ grep "^\target_dir" $MOTIONCONF | awk '{ print $2}'
/tmp/motion
```

Now all that's left is to push the result into a variable, which is accomplished with the back-tick symbol, giving us a final result of this:

```
MOTIONDIR=' grep "^\target_dir" $MOTIONCONF | awk '{ print $2}''
```

Now, one thing that you may not have considered: if this directory is not created, motion will create it when it captures its first image, but that can take hours. We need it now, so let's go ahead and do a check to see whether it exists and if not, create it. The check to see whether a directory exists is `if [-d <directory name>]`. We also need to add a note to see if the directory doesn't exist, so go ahead and add an exclamation point before the `-d`, which means *not*. Finally, mix this with a `mkdir` and you should have the following `if` statement:

```
if [ ! -d $MOTIONDIR ]; then
    mkdir $MOTIONDIR
fi
```

Along with making the directory, you need to ensure that it is owned by the correct user. The `motion` daemon is run by root, but this ownership is passed to `motion`, so you need to add a `chown` to this command, which gives this value:

```
if [ ! -d $MOTIONDIR ]; then
    mkdir $MOTIONDIR
    chown motion $MOTIONDIR
fi
```

Okay, so now we have the directory where the files will be found, and we need to write some code to check it and count the number of files. Normally, you would think “Sweet, let's just use `ls` to list the files in the directory.” But this won't help in the event that there are no files because `ls` will give us an error. To do this, we need to use the `find` command. To find all the `.jpg` files within only that directory, we need to add a pair of arguments to `find` the following:

- `-maxdepth 1` says not to go within any subdirectories to search
- `-type f` to tell it to only search for files

I also wrapped the command in a pair of parentheses, which means that the file is meant to be treated as an array. This gives us this command:

```
LISTFILES=('find $MOTIONDIR -maxdepth 1 -type f')
```

Now, in case there are non-image files, we should also add in a grep to pull out only the .jpg files. A simple grep for anything ending in .jpg (using the \$ operator) should suffice. This will give us the following:

```
LISTFILES=('find $MOTIONDIR -maxdepth 1 -type f | grep jpg$')
```

Now we need to get the count of files in our list using the wc -l command. This will list all the files in a directory and if we pipe that to wc with the -l argument, we get the number of files within a given folder. This gives us our next line:

```
NUMFILES=' find $MOTIONDIR -maxdepth 1 -type f | grep jpg$ | wc -l'
```

There is one final common piece of code that will be common to both scripts that we can work up now: the mutt mailer line that will send the image file and a message. The original one we sent was good, but I've modified it slightly here to include a variable name for the attachment that is the variable IMAGEFILE. I've also added a new -F /root/.muttrc option, which forces mutt to use the root muttrc file we created. The new mail line is this:

```
echo -e "Warning,\nAn intruder was detected at `date`\nPlease see the image attached for details"\n| mutt -F /root/.muttrc -s "[INTURDER ALERT] Intruder Detected" -a $IMAGEFILE -- you@gmail.com
```

As you can see, it's a bit of a mouthful. I've added some nice body text that will come out as follows:

```
Warning,
An intruder was detected at Sat Sep 15 22:50:32 EST 2012
Please see the image attached for details
```

You may be asking yourself how I got those newlines in there, and how I was able to get a very nice date and time output in the e-mail. The date and time come from the shell date command, which I have encased in back-ticks to give us the output. I've also added some \ns, which are references to newlines, but only when you give echo the -e argument to tell it to interpret special escape values. Now that we have a block of code to work from, let's start with script number 1, the cron executed run once script.

Script 1

To begin, let's go over what we have so far; we can get the number of files in the correct working directory. The next step then is to check whether that number shows that there are indeed files to work with, which is to say > 0. This means we need the following if statement:

```
if [ $NUMFILES -gt 0 ]; then
```

With our if statement in place, we should now know whether there are any files we need to send urgently. Now we just need to iterate through the list and send each as an e-mail and then delete it. To iterate through anything, you need to use a loop; in this case, we will be using a for loop. This loop will split each image off from the list as its own variable, allowing us to take action on each one after another. The loop will look like the following, where IMAGEFILE is the current image from the list; this image changes to the next one in the list each time the for loop starts again. The for loop will finish when there are no more images to process from the list.

```
for IMAGEFILE in $LISTFILES
```

Note Remember to `chmod +x` your script so it can be executed!

Okay, that's it. We have all our code blocks, so let's assemble. The final script should look something like this:

```
#!/bin/bash
MOTIONCONF=/etc/motion/motion.conf
MOTIONDIR=' sudo grep "^target_dir" $MOTIONCONF | awk '{ print $2}''
if [ ! -d $MOTIONDIR ]; then
    mkdir $MOTIONDIR
    chown motion $MOTIONDIR
fi
LISTFILES=('find $MOTIONDIR -maxdepth 1 -type f | grep jpg$')
NUMFILES=' find $MOTIONDIR -maxdepth 1 -type f | grep jpg$ | wc -l'
if [ $NUMFILES -gt 0 ]; then
    for IMAGEFILE in $LISTFILES
    do
        echo -e "Warning, \nAn intruder was detected at 'date'\nPlease see the image attached
for details" \
            | mutt -s "[INTRUDER ALERT] Intruder Detected" \
            -F /root/.muttrc \
            -a $IMAGEFILE -- <your email>
        rm $IMAGEFILE
    done
fi
```

Script 2

Because we have finished discussing script 1, let's work on script 2. This one is slightly more complex. The first thing we need to do is add a counter that will tell us how many images we have dealt with already. Initially, we can set this to 0. So let's initialize this variable with `LASTCOUNT=$ NUMFILES`. Next is to create a simple loop that will keep the application running forever; this is, in fact, the easiest kind of loop and is simply `while true` (for more on while loops, see Chapter 7).

Now within the loop we need to update `LASTCOUNT` and `NUMFILES`, so have those commands run again to update their values. We need to compare to see if we have had anything new added (i.e., that `NUMFILES` is greater than `LASTCOUNT`). If it is, we need to read off the latest files and send them. Normally, this sounds like the place for an `if` statement, but in this case we will use a `while` loop because we want to do things while `LASTCOUNT` is greater than `NUMFILES`, so our second inner `while` loop is this:

```
while [ $LASTCOUNT -lt $NUMFILES ]
```

Now we simply need to update; first we need to get each new image a filename. Thankfully they are found in order because the motion package names them sequentially. This means that we simply need to pull off the top X to make up the difference. This means we will need an iterator value so we can count each of the top X as we go. So create an `ITERATOR` set to 0 outside the second loop but inside the first, so we reset it every round.

Now that we have this iterator, we can use it with the LISTFILES to pull off each of the numbers. Because we are treating LISTFILES as an array, we just need to use the ITERATOR to access that array element. We need to wrap the whole array output within a pair of curly braces because this will perform the array dereferencing and give us the output value:

```
IMAGEFILE=${LISTFILES[$ITERATOR]}
```

Finally, we need to increment LASTCOUNT and ITERATOR so that they can count off each image as we process it. This is done with the following:

```
LASTCOUNT='expr $LASTCOUNT + 1' and ITERATOR='expr $ITERATOR + 1'
```

The expr function treats given values as mathematical expressions and returns their result, which is perfect when we want to increment a value by 1.

So, now we need to assemble each of these items of code. When put together, they should look like this:

```
#!/bin/bash
MOTIONCONF=/etc/motion/motion.conf
MOTIONDIR=' sudo grep "target_dir" $MOTIONCONF | awk '{ print $2}''
if [ ! -d $MOTIONDIR ]; then
    mkdir $MOTIONDIR
    chown motion $MOTIONDIR
fi
LISTFILES='find $MOTIONDIR -maxdepth 1 -type f | grep jpg$'
NUMFILES='find $MOTIONDIR -maxdepth 1 -type f | grep jpg$ | wc -l'
LASTCOUNT=0
while true
do
    LISTFILES=('find $MOTIONDIR -maxdepth 1 -type f | grep jpg$')
    NUMFILES='find $MOTIONDIR -maxdepth 1 -type f | grep jpg$ | wc -l'
    ITERATOR=0
    while [ $LASTCOUNT -lt $NUMFILES ];
    do
        IMAGEFILE=${LISTFILES[$ITERATOR]}
        echo -e "Warning,\nAn intruder was detected at `date`\nPlease see the image
attached for details" \
        | mutt -F /root/.muttrc \
        -s "[INTRUDER ALERT] Intruder Detected" \
        -a $IMAGEFILE -- <your email>
LASTCOUNT='expr $LASTCOUNT + 1' ITERATOR='expr
$ITERATOR + 1' done sleep 1
done
```

Testing

So, let's take this super script for a spin. Ensure that motion is up and running and execute your script.

Note To cancel running the script, press Ctrl + c to stop execution. This is the only way to stop script 2.

You shouldn't expect to see any output, so just watch your e-mail inbox (see Figure 10-4).




Figure 10-4. One full inbox!

As you can see, it works. In fact, it works rather well; I have 18 unread e-mails from the intruder detection system. Perfect! Now that we have tested that the script will function in sending us e-mails of new images from motion, we need to configure the script so that it will start automatically—the one final requirement that remains unmet. Looking back, we decided that we would start each of the two scripts in different ways: script 1 would be added to cron and script 2 would be added to the motion start script, so that it would run with motion!

For script 1, things are fairly simple. cron will automatically evaluate if it should run something once a minute, so all we need to do is have a cron job that will run at all times. So as root (since we want to run this as root), open up the cron table with `crontab -e` and enter the following:

```
* * * * * /root/script1.sh
```

That's it; we have no added script 1 to start and run from a crontab entry every minute. So that every minute we will check to see if there are any files in `/tmp/motion`, send them to us and then delete them, which fulfills all our requirements.

The second script's launch method requires modification to the launch script for motion `/etc/init.d/motion`, so go ahead and open that up to be edited as root. The basic format of any of these start scripts is to have a large case statement that processes the given action it is to perform be that start, stop, restart, or any other number of required cases. So look for the line `start`, which is the action we want to enhance. Within this, we can see that there are a number of nested checks and preliminaries; we want our script only to be started after motion, which is invoked by the `start-stop-daemon` command.

So go ahead and add a link to your script after that; then add a `&` on the end. The ampersand says “run this script and then let it go in the background.” If we don’t do this, the start script will assume that its next function is to wait for that script to end, which it won’t! Your final updated start section should look like this:

```
start)
if check_daemon_enabled ; then
    if ! [ -d /var/run/motion ]; then
```

```

mkdir /var/run/motion
fi
chown motion:motion /var/run/motion

log_daemon_msg "Starting $DESC" "$NAME"
if start-stop-daemon --start --oknodo --exec $DAEMON -b --chuid motion ; then
    /root/script2.sh &
    log_end_msg 0
else
    log_end_msg 1
    RET=1
fi
fi
;;

```

Once you've finished editing, go and restart your Pi; then log back in and execute a check to see if the script is running! It is? Fantastic, you now have a fully automated Raspberry sPi of your very own!

```
$ ps -ef | grep script
root      2058      1  0 11:23 pts/0    00:00:00 /bin/bash /home/pi/script2.sh
```

Troubleshooting

As I'm sure you're aware by now, there are so many things that can go wrong when doing software development. So there is plenty of space for troubleshooting, and it makes my giving you advice about where to look a little difficult. That being said, there are a number of things you can do to make diagnosis easier.

There's a reason why I started off scripting by writing out each little section of code and then adding it. This is a version of what's called *isolation testing*, which involves doing as much as you can to isolate a particular section of code to execute it with controlled inputs and then evaluating the outputs to see that they work as you intend.

As you can imagine, doing this block by block means that you can test each little bit of an application as it grows.

The next thing is to take advantage of the echo command to output a working value or a location in the code. You can see what an application is working with and where within the code it is running, which lets you pick up on unexpected values and where your software is failing.

Finally, if all else fails, change the shell from `#!/bin/bash` to `#!/bin/bash -x`. Adding the `-x` will put your instance of bash into debug mode and it will output every variable, every manipulation, and every change it goes through. While there can be a lot of output this is the ideal way to see what's going on within a troublesome application.

Where to Go from Here

Wow. We've come a long way from where we started. We have:

- Attached a webcam to our Raspberry sPi
- Installed and configured the monitor application to use the webcam
- Used the webcam and monitor application to function as a motion sensor and capture images
- Installed and configured ssmtmp as an MTA
- Installed and configured mutt mail application to use SSMTP

- Send an e-mail from the command line to ourselves using SSMTMP and mutt
- Written two scripts to automatically pick up any new image files and then e-mail them to ourselves using all of the above

This is quite an accomplishment, but what now? As spies, you all know how adaptability is the key to success. What we have gone through here should be a guide on how you can take advantage of software like motion, SSMTMP and mutt. You should now also be familiar with using dmesg and how to locate newly loaded devices and integrate them into your system. Finally, you should see how you can use scripts written in bash to fill a gap in your application stack and solve complex problems. So what can you do from here?

- You can add a wireless adapter, so you don't need to be cabled in, giving you a wireless Raspberry sPi that you can hide in the most innocuous of places.
- You can add a USB battery pack to remove the need for access to a power point.
- You can change which directory the files are stored in out of /tmp. That way, you can have your images persisted to disk, where you can review them at a later date.
- Or you can go full Bond. I've taken my inspiration further afield, modified the software to capture full video, and disabled the image capture. I changed the storage location to be away from /tmp and I've attached a USB battery pack. I then added a little tape to hold this all together and mounted the webcam inside a hat (see Figure 10-5).



Figure 10-5. Maybe I should have invested in a stetson!

Summary

You should be able to understand the output of dmesg and know how udev works. You should also be able to configure the motion application, SSMTMP application, and mutt application. Finally, you should be able to send e-mail with different recipients, body text, subjects, and attachments.



Pi Media Center

The beautiful things about a Raspberry Pi are that it is so flexible a device and its potential uses are nearly limitless. You can do just about anything with Pi, and it has an array of peripheral input ports that can be used to great effect.

You should also be aware that while your Pi can display through a standard HDMI, the power of the device doesn't end there. It can display at a full 1920 x 1080 resolution, commonly called 1080P, which is the benchmark level for high definition. Your tiny little Raspberry Pi is capable of displaying video at full high-definition resolution, and thanks to the sound interleave of HDMI it can also play back audio on the same channel as video. That is the goal of this chapter; we are going to take advantage of the Pi's media playback capabilities to create two different varieties of Pi media center. One is a video playback center that you can use to play videos over your network. It can sit attached to your TV and allow you to tap into the Internet for video. The second is a wireless playback device that enables you to remotely stream your music collection and play back across a network.

Video with XBMC

As mentioned earlier, one big advantage of Raspberry Pi is that it can play back in full high definition, but we need a way to take advantage of this, which is where XBMC comes in.

XBMC (which is short for *XBox Media Center*) was originally designed as a media center for the original (Fat) XBox. People liked the project so much that it has subsequently been ported to just about every operating system under the sun, including but not limited to Windows, Android, OS X, and iOS. There is even a full-blown, stand-alone version called XMBCbuntu that is a port of Ubuntu that has XBMC installed, preconfigured, and ready to go out of the box. Given its wide use, open source nature, and easy-to-work-with nature, it has been ported and forked many times with its progeny, including MediaPortal, Plex, Voddlern, and Boxee. This goes to show how popular media center systems are; as a further example, Boxee makes a small device called the Boxee Box that costs upward of \$200. What we are creating is a tool that is just as powerful and only costs \$35—the price of the Pi. So now that you're all excited about the bright and shiny media center future, it's time to take advantage of it.

You've reached the end of the book, so it's time to crack out the big guns and do a proper Linux install from the application source. While most applications come prepackaged via `apt-get`, there are some that do not and they need to be compiled. Installing XBMC from source is a somewhat involved process and is very time-consuming. We have chosen to compile it from source because not only does it help you through one of the more complex ways of getting an application running in Linux but it also means that our installation of XBMC is optimized for our system because we have built it on our own system. Historically, compiling from source was considered to be an arduous process, but as the number of people using Linux has grown so has the need for simple-to-use installation tools and the rise of precompiled binary installation systems such as `apt-get`. But as with everything, there are still times when using a prebuild is never as good as doing it from scratch, and as time as gone on, tools for compilation have grown and become much more approachable. It's also easier to do when someone steps you through the process—which is what this book is all about.

So let's get started. First, I have stepped away from the primary build of Raspbian to use one of the slightly cut-down versions of Raspbian called Pisces that is available from Raspbian at www.raspbian.org/PiscesImages. I am using Pisces because it has been cut down and doesn't include as much overhead, which means more power for doing what we want: playing movies! Download and install this image on an SD card (you will need one with 4 GB of space) and then boot up and get the OS attached to the Internet so you can work with it. The default user is `raspbian` and the password is also `raspbian`, which is also the root password.

RASPMC

There is a much more simple way to get XBMC working on your Pi: using a prebuilt version of Raspbian called RaspMC. You can download a RaspMC installer and push the image onto an SD card just as you did with Raspbian. You can download this image from <http://www.raspbmc.com/download/>. This installation is very simple and will take between 15–25 minutes. You simply attach your Pi to your network, turn it on, and off the installer goes. It will download everything it needs from the Internet and then set itself up. This is by far the easiest way to get XBMC running on your Pi. For those of you who take the path of using RaspMC, you can skip to the "Starting and Using XBMC" section of this chapter.

Note If `ifconfig` is missing, install the `net-tools` package using `apt-get`. See Chapter 3 for more on `apt-get`.

Once you are inside the operating system, it's time to issue some commands and start getting XBMC installed. This is a very long process: the compiler is much slower than using `apt-get` because all the software installed by `apt-get` is precompiled, thus cutting out some steps. The compile in this case takes about 12 hours—yes hours, half a day. I find it best to leave the final compile to run overnight so I won't be tempted to sit there for hours on end and watch it.

The process takes such a long time because of the nature of the ARM processor and the speed at which it can work. This means that compiling takes longer, but we can do as much work on a lower level of power.

Setup to Build

Let's start with the build. The first step is to ensure that we can divert as much power and memory to the compile as possible. In the core Raspbian we did this with the `rapsi-config` command, but Pisces doesn't have that control script, so you have to change this manually by replacing the `start.elf` file with a different one that will change how the system's resources are allocated. Run this command:

```
$ sudo cp /boot/arm224_start.elf /boot/start.elf
```

This will copy the version that allocates 224 MB to system RAM and 32 MB to graphics into the system to be used. Now we need to make this configuration change registered within the system and then reboot so that the changes can take effect by issuing the following two commands:

```
$ sudo rpi-update  
$ sudo reboot
```

Once your system has rebooted, and you are back up and running, it is best if you upgrade the OS to ensure it is running the latest versions of software. To do this, we will leverage `apt-get` by telling it to update itself with the latest

version of its configuration from the Internet and then upgrade all available packages on your system. To do this, issue the following commands:

```
$ sudo apt-get update
$ sudo apt-get upgrade
```

This execution will take a little bit of time because it needs to download, unpack, and replace a large number of system components. So get this up and running and then go have a coffee; it should be done when you return, and you can move on to the next step.

Note Some of these packages might already be installed on your system. We include them as they might not be installed for everyone.

The next step is to install packages to give us the ability to compile; they include the essential build tools, automatic configuration tools, zip and unzip tools, and so on. To install these tools, run the following command:

```
$ sudo apt-get install build-essential autoconf ccache gawk gperf mesa-utils zip unzip
Reading package lists... Done
Building dependency tree
Reading state information... Done
build-essential is already the newest version.
The following packages were automatically installed and are no longer required:
  libcdio-cddao libcdio-paranoia0 libcdio10 libcelto-0 libdb4.8 librpmio2
Use 'apt-get autoremove' to remove them.
The following extra packages will be installed:
  automake autotools-dev libglew1.7 libsigsegv2
Suggested packages:
  autoconf2.13 autoconf-archive gnu-standards autoconf-doc libtool gettext distcc gawk-doc
  glew-utils
The following NEW packages will be installed:
  autoconf automake autotools-dev ccache gawk gperf libglew1.7 libsigsegv2 mesa-utils
The following packages will be upgraded:
  unzip zip
2 upgraded, 9 newly installed, 0 to remove and 187 not upgraded.
Need to get 2,647 kB/3,174 kB of archives.
After this operation, 7,143 kB of additional disk space will be used.
Do you want to continue [Y/n]?
```

Once you have installed these basic packages, move on to the far bigger install job. We now need to install a number of packages that will install video and decoder libraries and specialized graphics libraries. In addition to all these, you will need to install a few other packages to enable you to connect remotely to the XBMC. This command is very long, so I recommend you download the source code from the Apress website. There is a script within that package that will execute these commands for you. The command to install this giant number of packages is as follows:

```
$ sudo apt-get install autotools-dev comerr-dev dpkg-dev libalsaplayer-dev \
libapt-pkg-dev libasound2-dev libass-dev libatk1.0-dev \
libavahi-client-dev libavahi-common-dev libavcodec-dev libavformat-dev \
libavutil-dev libbison-dev libbluray-dev libboost1.49-dev \
```

```

libbz2-dev libc-dev-bin libc6-dev libcaca-dev libcairo2-dev \
libcdio-dev libclalsadrv-dev libcrypto++-dev libcurl3-gnutls-dev \
libdbus-1-dev libdbus-glib-1-dev libdirectfb-dev libdrm-dev libegl1-mesa-dev \
libelf-dev libenca-dev libept-dev libevent-dev libexpat1-dev libflac-dev \
libfontconfig1-dev libfreetype6-dev libfribidi-dev libgconf2-dev \
libgcrypt11-dev libgdk-pixbuf2.0-dev libgl1-mesa-dev libgles2-mesa-dev \
libglew-dev libglewmx-dev libgl2.0-dev libglu1-mesa-dev \
libgnome-keyring-dev libgnutls-dev libgpg-error-dev libgtk2.0-dev libhal-dev \
lib hunspell-dev libice-dev libicu-dev libidn11-dev libiso9660-dev \
libjasper-dev libjbig-dev libjconv-dev libjpeg8-dev libkrb5-dev \
libldap2-dev libltdl-dev liblz0-dev libmado-dev libmicrohttpd-dev \
libmodplug-dev libmp3lame-dev libmpeg2-4-dev libmysqlclient-dev \
libncurses5-dev libnspr4-dev libnss3-dev libogg-dev libopenal-dev \
libp11-kit-dev libpam0g-dev libpango1.0-dev libpcre++-dev libpcre3-dev \
libpixman-1-dev libpng12-dev libprotobuf-dev libpthread-stubs0-dev \
libpulse-dev librtmp-dev libsamplerate0-dev \
libsdl-image1.2-dev libsdl1.2-dev libslang2-dev \
libsm-dev libsmclient-dev libspeex-dev \
libsqLite3-dev libssh-dev libssh2-1-dev libssl-dev libstdc++6-4.6-dev \
libtagcoll2-dev libtasn1-3-dev libtiff4-dev libtinfo-dev libtinyxml-dev \
libts-dev libudev-dev libv8-dev libva-dev libvdpa-dev \
libvorbis-dev libvpx-dev libwebp-dev libwibble-dev \
libx11-dev libx11-xcb-dev libxapian-dev libxau-dev \
libxcb-glx0-dev libxcb-render0-dev libxcb-shm0-dev \
libxcb1-dev libcomposite-dev libcursor-dev libxdamage-dev \
libxdmcp-dev libxext-dev libfixes-dev libxft-dev libxi-dev \
libxinerama-dev libxml2-dev libxmu-dev libxrandr-dev \
libxrender-dev libxslt1-dev libxss-dev libxt-dev \
libxtst-dev libxf86vm-dev libyajl-dev libzip-dev linux-libc-dev \
lzma-dev mesa-common-dev python-dev python2.7-dev x11proto-composite-dev \
x11proto-core-dev x11proto-damage-dev x11proto-dri2-dev x11proto-fixes-dev \
x11proto-gl-dev x11proto-input-dev x11proto-kb-dev x11proto-randr-dev \
x11proto-record-dev x11proto-render-dev x11proto-scrnsaver-dev \
x11proto-xext-dev x11proto-xf86vidmode-dev x11proto-xinerama-dev xtrans-dev \
zlib1g-dev

```

This list of packages has a number of slashes on the end; they are line breaks that say that the current command should be executed and that there will be another line of command following. Line breaks are very useful when you need to break up long commands over a large number of lines, as I have done here. This command will take a while to run because there are a large number of packages to be downloaded and installed. So execute this command, go cook dinner, and then come back.

Once you have installed all these packages, you need to copy some special files into place and create other ones that are linked into the correct location so they can be used. Again, they will be included within the script that we have created to help ease the execution of these commands.

The first of these is to copy the VideCoreIV include files from the Raspberry Pi firmware into /usr/include. This directory is a special directory from which the compiler will search for libraries and files that are to be included within a build. To copy all these files into the correct place, execute the following commands:

```

$ sudo cp -R /opt/vc/include/* /usr/include
$ sudo cp /opt/vc/include/interface/vcos/pthreads/* /usr/include/interface/vcos

```

Now in addition to copying these files into place, there are a number we need to link correctly into the right location so they can be read. While in the last step we copied include files that are predominantly source code, this time we need to link in precompiled portions of the firmware into the correct spot to be read by the compiler at compile time. Instead of copying, we are linking because we are happy to use the files in their current location and save on some space. Run the following:

```
$ sudo ln -fs /opt/vc/lib/libEGL.so /usr/lib/libEGL.so
$ sudo ln -fs /opt/vc/lib/libEGL.so /usr/lib/arm-linux-gnueabihf/libEGL.so
$ sudo ln -fs /opt/vc/lib/libEGL.so /usr/lib/arm-linux-gnueabihf/libEGL.so.1
$ sudo ln -fs /opt/vc/lib/libEGL_static.a /usr/lib/libEGL_static.a
$ sudo ln -fs /opt/vc/lib/libEGL_static.a /usr/lib/arm-linux-gnueabihf/libEGL_static.a
$ sudo ln -fs /opt/vc/lib/libGLESv2.so /usr/lib/libGLESv2.so
$ sudo ln -fs /opt/vc/lib/libGLESv2.so /usr/lib/arm-linux-gnueabihf/libGLESv2.so
$ sudo ln -fs /opt/vc/lib/libGLESv2.so /usr/lib/arm-linux-gnueabihf/libGLESv2.so.2
$ sudo ln -fs /opt/vc/lib/libGLESv2_static.a /usr/lib/libGLESv2_static.a
$ sudo ln -fs /opt/vc/lib/libGLESv2_static.a /usr/lib/arm-linux-gnueabihf/libGLESv2_static.a
$ sudo ln -fs /opt/vc/lib/libbcm_host.so /usr/lib/libbcm_host.so
$ sudo ln -fs /opt/vc/lib/libbcm_host.so /usr/lib/arm-linux-gnueabihf/libbcm_host.so
$ sudo ln -fs /opt/vc/lib/libvchiq_arm.a /usr/lib/libvchiq_arm.a
$ sudo ln -fs /opt/vc/lib/libvchiq_arm.a /usr/lib/arm-linux-gnueabihf/libvchiq_arm.a
$ sudo ln -fs /opt/vc/lib/libvchiq_arm.so /usr/lib/libvchiq_arm.so
$ sudo ln -fs /opt/vc/lib/libvchiq_arm.so /usr/lib/arm-linux-gnueabihf/libvchiq_arm.so
$ sudo ln -fs /opt/vc/lib/libvcos.a /usr/lib/libvcos.a
$ sudo ln -fs /opt/vc/lib/libvcos.a /usr/lib/arm-linux-gnueabihf/libvcos.a
$ sudo ln -fs /opt/vc/lib/libvcos.so /usr/lib/libvcos.so
$ sudo ln -fs /opt/vc/lib/libvcos.so /usr/lib/arm-linux-gnueabihf/libvcos.so
```

Note We found that these links can be removed by other processes and sometimes by rebooting; if they do, you need to re-create them.

Now that you have linked all the files into the correct place, go ahead and get the source code. Change the directory into your home directory with this:

```
$ cd ~
```

If you can't type a tilde (~), you can try reconfiguring the Pi keyboard layout with `dpkg-reconfigure keyboard-configuration` and then rebooting the Pi. You can also just use `cd` with no arguments because it will also take you to your home directory. The tilde is just a symbol to reference the current user's home directory.

Now that you have changed into your home directory, it is time to download the source code so you can work with it and correctly compile everything. Thankfully, a modern tool has appeared that makes doing source code downloads much easier: Git, which is named for Linus Torvalds of Linux Kernel fame. Git is used for source code management and version control; all users can download the current source version of a set of source code and submit modifications back to the central repository.

Note This means that we always download the latest version, but we might get a less stable build if something is wrong with the current version in Git. Most developers strive to keep their Git repositories in good working order, but it is something to be aware of.

All we want to do at this point is download the source code into our home directory to work with it. To do this, use the `git` command and tell it to clone the repository of source code onto the system. We have also added the `--depth 1` to the command to say that we only want to clone the most recent 1 revision of the source code (the absolute latest version of the code) and avoid getting a lot of extra and unwanted historical code.

Start your clone by running the following:

```
$ git clone --depth 1 git://github.com/xbmc/xbmc-rbp.git
Cloning into 'xbmc-rbp'...
remote: Counting objects: 35172, done.
remote: Compressing objects: 100% (22895/22895), done.
remote: Total 35172 (delta 15265), reused 27885 (delta 10698)
Receiving objects: 100% (35172/35172), 158.26 MiB | 385 KiB/s, done.
Resolving deltas: 100% (15265/15265), done.
```

Now that we have cloned the source code of XBMC, we need to start building the application. This is the part that will take the longest by far. Change the directory into the newly created `xbmc-rbp` directory because there are a few last steps we need to take.

Changing Files with Sed and Regular Expressions

Before we go ahead and kick off the build of our XBMC application, we need to make a few minor adjustments to the source code so that it compiles in the exact manner we want it to. Instead of making these changes by hand, we will use a tool called *Sed* (short for *Stream Editor*). It does just as its name implies: it takes a stream of text and then performs edits on that text. The tricky part of *Sed* is that it makes use of a highly specialized language called *regular expressions*, which are used to perform the equivalent of a find-and-replace function.

Regular expressions evolved out of the need to validate a given set of data to see if it conforms to a given standard and to edit when the given example is found. While a simple find and replace is enough for most people when editing text, how do you perform a find-and-replace when you want to edit only words that start with *T* as the third letter in a given sentence? Probably one of the easiest cases for a regular expression is e-mail. You want to confirm that a given string is a valid e-mail: an e-mail will contain a username (who it's from), a domain (where it's from, e.g., Hotmail or Gmail), and there will be an *at* (@) symbol in between. Okay, we can do that with just a wildcard like *. So something that looks like *@* will be an e-mail, right? Well the * that we use to denote anything will pick up anything. Even spaces, numbers, strange symbols, and so on. So we need a way to pick up something that is made up of a string of valid e-mail characters (letters, numbers, period, underscores) and then @, and finally a valid domain. This will be a collection of letters, followed by a period and then perhaps another set (or a few sets) of letters. Describing this is exhausting, which is why we have a specialized language for performing them.

Now that you understand what we are trying to achieve, run the following two commands to alter the file `tools/rbp/setup-sdk.sh`. These sed commands make use of the `-i` option, which says to edit the given file:

```
$ sed -i 's/USE_BUILDROOT=1/USE_BUILDROOT=0/' tools/rbp/setup-sdk.sh
$ sed -i 's/TOOLCHAIN=\usr\local\bcm-gcc/TOOLCHAIN=\usr/' tools/rbp/setup-sdk.sh
```

These commands change the `USE_BUILDROOT` variable from a value of 1 to a value of 0 wherever it is found within that file and to replace any instance of the phrase `TOOLCHAIN=\usr\local\bcm-gcc` with `TOOLCHAIN=\usr`. Once you have edited the `tools/rbp/setup-sdk.sh` file, you should execute it with the following command:

```
$ sudo sh tools/rbp/setup-sdk.sh
```

The command will generate a makefile, but won't display anything onscreen. Once it is finished it's time to edit the newly created tools/rbp/depends/xbmc/Makefile. All you need to do then is run one last sed on the file to add a # to the start of any instance of cd \$(SOURCE); \$(CONFIGURE):

```
$ sed -i 's/cd $(SOURCE); $(CONFIGURE)/#cd $(SOURCE); $(CONFIGURE)/*' tools/rbp/depends/xbmc/Makefile
```

What we did is change the way our system will compile the XBMC software and make it use some of our system libraries rather than use the ones that can be found within the XBMC source code.

Now that the Makefile is created, it is time to compile!

Compiling Source Code

Most compilations use a very simple logic: you run the command ./configure to generate a configuration file that knows about the makeup of your system. This generates a file that is called a *makefile*. This file is in a slightly specialized language that describes the way in which a compile for a collection should be performed. It is not a script that executes code in and of itself; it is a collection of directions that will be used by the make application.

Unfortunately, this compile is slightly more advanced. We need to make our configuration tools to match the setup we are trying to do before we run the configuration and the compile. This will take a little while to execute, so kick off and go have another cuppa.

The command to execute is this:

```
$ make -C tools/rbp/depends/xbmc/
```

This command will run the make within the directory tools/rbp/depends/xbmc/, which will generate the config file. The output for this is very long, but when it's done successfully, you should see something like this on your console window:

```
examples/Makefile.am: installing './depcomp'
Makefile.am: installing './INSTALL'
autoreconf: Leaving directory 'lib/libdvd/libdvdnav'
Please (re)run configure...
#cd ../../..; ./configure --prefix=/opt/xbmc-bcm/xbmc-bin --build=i686-linux --host=arm-
bcm2708-linux-gnueabi --enable-gles --disable-sdl --disable-x11 --disable-xrandr --disable-openmax
--disable-optical-drive --disable-dvcss --disable-joystick --disable-debug --disable-crystalhd
--disable-vtbdecoder --disable-vaapi --disable-vdpau --disable-pulse --disable-projectm
--with-platform=raspberry-pi --disable-optimizations --enable-rpi-cec-api
#cd ../../..; make -j 1
#cd ../../..; make install
make: Leaving directory '/home/raspbian/xbmc-rbp/tools/rbp/depends/xbmc'
```

Now that the compile of the configuration utility has gone through its time to perform the configure. This configure command is very long and again, you will find code that will perform this for you in the Apress source code repository, the configure command is this:

```
$ ./configure --prefix=/usr --build=arm-linux-gnueabihf --host=arm-linux-gnueabihf \
--localstatedir=/var/lib --with-platform=raspberry-pi --disable-gl --enable-gles \
--disable-x11 --disable-sdl --enable-ccache --enable-optimizations \
--enable-external-libraries --disable-goom --disable-hal --disable-pulse \
--disable-vaapi --disable-vdpau --disable-xrandr --disable-airplay \
```

```
--disable-alsa --enable-avahi --disable-libbluray --disable-dvdcss \
--disable-debug --disable-joystick --enable-mid --disable-nfs --disable-profiling \
--disable-projectm --enable-rsxs --enable-rtmp --disable-vaaapi \
--disable-vdadercoder --disable-external-ffmpeg --disable-optical-drive
```

Specify a number of arguments to the `configure` command that specifies where this installation should go (`--prefix=/usr`), which architecture to build (`--build=arm-linux-gnueabihf`), where to keep its running files (`--localstatedir=/var/lib`), and which platform (`--with-platform=raspberry-pi`), along with a whole host of disable and e-mail options to remove features such as CD drives and joysticks and enable things such as optimizations and Avahi automated detection of peripherals. Ultimately, when the command has finished executing, the output on your screen should look like this:

XBM C Configuration:

```
Debugging:      No
Profiling:     No
Optimization:   Yes
Crosscomp.:    No
target ARCH:   arm
target CPU:    arm1176jzf-s
OpenGL ES:     Yes
ALSA:          No
DBUS:          Yes
VDPAU:         No
VAAPI:         No
CrystalHD:    No
VDADecoder:   No
VTBDecoder:   No
OpenMax:       No
Joystick:     No
XRandR:        No
GOOM:          No
RSXS:          Yes
ProjectM:     No
Skin Touched: No
X11:           No
Bluray:        No
TexturePacker: Yes
MID Support:   Yes
ccache:        Yes
ALSA Support:  No
PulseAudio:    No
HAL Support:   No
DVDCSS:        No
Avahi:          Yes
Non-free:      Yes
ASAP Codec:   No
MySQL:          Yes
Webserver:     Yes
libRTMP support: Yes
libsmbclient support: Yes
```

```

libnfs client support:No
libafpclient support: No
AirPlay support:      No
AirTunes support:     No
Optical drive:        No
libudev support:      Yes
libusb support:        No
libcec support:        No
libmp3lame support:   Yes
libvorbisenc support: Yes
libcap support:        No
External FFmpeg:      No
prefix:               /usr

```

This configuration specifies all the flags we asked for in our installation, which is designed to get the best from the Raspberry Pi. So, now we have done our configure onto the actual compilation. The command itself *highly* anticlimactic for what it will do: tell the system to spend the next 12 hours generating code for our system. All these other jobs taking a little time is nothing compared to how long this compile will run for. So saddle up, and execute this simple four-letter command `make`:

```
$ make
```

Don't be alarmed if you see lines such as the following:

```
/tmp/ccGvUe1g.s:507: Warning: swp{b} use is deprecated for this architecture
```

The only thing that will stop the compile is a critical error or success. These warnings are related to some of the functions being used within the XBMC compile being a bit old for our Pi. When the compile has succeeded, you should expect to see this:

XBMC built successfully

This shows that we have successfully build XBMC and can do the last step in this process, which is to install our newly compiled software. The make system will actually take care of installation, too. To install, you simply run `make install`, which copies all the freshly compiled binaries into the correct spot in your OS. You will need to prefix this with `sudo` because the location we are copying too (`/usr`) is a system location, and only root is given write access. Your output should look like this:

```

$ sudo make install
Copying XBMC binary to /usr/lib/xbmc xbmc.bin
You can run XBMC with the command 'xbmc'
Copying support and legal files...
Done!
Copying system files to /usr/share/xbmc

```

And we are done! While it was long and required a little bit of work to begin with, the actual process of compiling is relatively pain-free, you know, beyond the pain of waiting. So now that the XBMC is installed, how do we use it?

Troubleshooting

Before we jump ahead and run XBMC, we should have a look at some of the issues you could have run into. Most of these steps should take care of themselves because this setup is generic and meant to work on any Pi because they

all share the same hardware. If you are ever in doubt, remove the `xbmc-rbp` directory and start over. Ensure that you execute the copies and the linking in full. I found that when doing the git clone, things could sometimes get stuck and not move.

When this is the case, it is best to stop the current job with `Ctrl + C`, remove the directory, reboot, and start again. It was also better to remove peripherals from the system during the compile and install stages because having too many peripheral devices caused load issues on my system. Additionally, you should ensure that your output for commands such as `configure` match the ones given as these flags can play a much bigger role when compiling and save you a nasty headache and another 12-hour compile job. Finally, ensure that all your `apt-gets` run and complete successfully because these packages are required to provide the libraries that your Pi will use when building XBMC.

Starting and Using XBMC

Now that we have gone through that mega compile process, it is time to start up XBMC and get going! First, though, we need to reset the options we set earlier, directing a larger amount of memory into system RAM. We will need to go to the graphics device because we want to do some serious display work with XBMC. Execute the following command to change over the Pi boot system:

```
$ sudo cp /boot/arm128_start.elf /boot/start.elf
```

Go back and check that the linked files mentioned earlier are correct. Now we can also make some changes to the `/boot/config.txt` file because it governs how the Pi's low-level hardware runs. The most simple change to make is to add the line `disable_overscan=1` as it will remove the need to do overscan processing because overscan processing will make the amount of video we play back larger, thus consuming more resources. Once you have made these changes, issue the following commands to reconfigure and reboot your Pi one last time:

```
$ sudo rpi-update  
$ sudo reboot
```

Now that you have changed the settings back to 50/50 resource allocation, start XBMC by logging in as root and issuing the following command:

```
$ /usr/lib/xbmc/xbmc.bin
```

Your screen will freeze for a moment, and Figure 11-1 will greet you.



Figure 11-1. XBMC boot screen

You should get to the XBMC menu, which looks like Figure 11-2.



Figure 11-2. XBMC main menu

From within the main menu, there are options to select different types of playback, and XBMC is a fully graphical system that will present you with options and locations for getting files from when you are within it. As long as your media are accessible on a network, XBMC should be able to access it and play it back! Now on the topic of playback, you will probably be sitting there thinking, “Will I be stuck with this keyboard for all my playback?” The answer is a resounding *no!* There are two solutions:

- Purchase a USB remote and attach it to your Pi.
- The second is much sneakier: use your smartphone!

Most smartphones have the capability to download an XBMC remote control application that simply needs to be paired with your XBMC system to function. To do this, we need to enable two options within XBMC to allow it to receive remote control. Go into the settings menu and then into services, and change the remote control settings to “Allow programs on other systems to control XBMC.” Once you have allowed remote control, you need to go into Webserver and set “Allow control of XBMC via HTTP”; you should also go ahead and set a username and password so that only people with the password can control your XMBC instance, not just anyone on your network with an XBMC app!

Once you have enabled these remote control options, it is simply a matter of downloading the application from your favorite app store and then setting the options within the application. You should only need to give the application the local address of your Pi, which you can see via the System ▶ System info menu, set the port to 8080, and furnish an instance name for the XBMC and the username/password, as shown in Figure 11-3.




Figure 11-3. XBMC remote application

Once you have given all the details, your phone should turn into a remote for your XBMC instance! Complete with a whole host of directional buttons and playback controls, as shown in Figure 11-4.



Figure 11-4. XBMC remote in action!

XBM C on Boot

Although we have started XBMC by hand thus far, most people will want it to occur automatically because what is the point of a media center that you need to plug a keyboard into each time you want to start it up?

The following is a simple start script that will allow you to execute XBMC using a start script. Create this in the file /etc/init.d/x xbmc (like all other long files, it can be found in the Apress repository for this book):

```
#!/bin/sh
### BEGIN INIT INFO
# Provides:          xbmc
# Required-Start:    $all
# Required-Stop:     $all
# Default-Start:    2 3 4 5
# Default-Stop:     0 1 6
# Short-Description: Start XBMC
# Description:       Start XBMC
### END INIT INFO
DAEMON=/usr/bin/xinit
DAEMON_OPTS="/usr/lib/xbmc/xbmc.bin"
NAME=xbmc
DESC=XBMC
RUN_AS=root
PID_FILE=/var/run/xbmc.pid
test -x $DAEMON || exit 0
set -e
case "$1" in
  start)
    echo "Starting $DESC"
    start-stop-daemon --start -c $RUN_AS --background --pidfile $PID_FILE --make-pidfile
--exec $DAEMON -- $DAEMON_OPTS
    ;;
  stop)
    echo "Stopping $DESC"
    start-stop-daemon --stop --pidfile $PID_FILE
    ;;
  restart|force-reload)
    echo "Restarting $DESC"
    start-stop-daemon --stop --pidfile $PID_FILE
    sleep 5
    start-stop-daemon --start -c $RUN_AS --background --pidfile $PID_FILE --make-pidfile
--exec $DAEMON -- $DAEMON_OPTS
    ;;
  *)
    echo "Usage: /etc/init.d/$NAME{start|stop|restart|force-reload}" >&2
    exit 1
    ;;
esac
exit 0
```

Once this file is created, run the following commands to make it executable and have it loaded into the Pi's boot process:

```
$ sudo chmod +x /etc/init.d/xbmc
$ sudo update-rc.d xbmc defaults
```

Now go ahead and reboot your system; it will bring you automatically into XBMC! Congratulations; at this point, you should have a fully functional XBMC instance complete with working smartphone remote!

Troubleshooting

Because you compiled from source successfully, there shouldn't be any issues with getting your XBMC instance up and going. The first thing to check is that the compile did indeed complete fully and so did the install. They are the two most likely culprits. If you start your XBMC instance, you might see warnings like the following:

```
libEGL warning: DRI2: xcb_connect failed
```

If you do, you have to re-run the previous copy and linking commands. This issue relates to being unable to load those library files.

Stream Music with Airplay

People often want to use the Raspberry Pi as a streaming music system, which means that it is connected to a set of speakers and can be remotely controlled and has music streamed to it from a central controlling system. There are already a large number of systems out there for doing this kind of playback, but for us the plan is to take advantage of the Apple AirPlay protocol. You can continue working in this on Pisces or you can revert to the original Raspbian. This install will work the same way on both devices.

This setup is much easier to do than the video update; we still need to do some compiling to get things to work, but we won't need to run a `./configure` because there are far fewer moving parts. Additionally, we want to have the ability to play audio out to a given device, so you need to ensure you have speakers attached.

The first thing to do is to install the following precursor packages:

```
$ sudo apt-get install build-essential libssl-dev libcrypt-openssl-rsa-perl libao-dev
libio-socket-inet6-perl libwww-perl avahi-utils pkg-config alsa-utils
Reading package lists... Done
Building dependency tree
Reading state information... Done
build-essential is already the newest version.
libio-socket-inet6-perl is already the newest version.
libio-socket-inet6-perl set to manually installed.
libssl-dev is already the newest version.
libwww-perl is already the newest version.
libwww-perl set to manually installed.
pkg-config is already the newest version.
pkg-config set to manually installed.

The following packages were automatically installed and are no longer required:
 libcdio-cddao libcdio-paranoia libcdio10 libcelt0-0 libdb4.8 librpmio2
Use 'apt-get autoremove' to remove them.
```

The following extra packages will be installed:

```
alsa-base avahi-daemon libao-common libao4 libavahi-core7 libcrypt-openssl-bignum-perl
libdaemon0 libnss-mdns
```

Suggested packages:

```
alsa-oss oss-compat avahi-autoipd libaudio2 libesd0 libesd-alsa0
```

The following NEW packages will be installed:

```
alsa-base alsound avahi-daemon avahi-utils libao-common libao-dev libao4 libavahi-core7
libcrypt-openssl-bignum-perl libcrypt-openssl-rsa-perl libdaemon0
libnss-mdns
```

0 upgraded, 12 newly installed, 0 to remove and 139 not upgraded.

Need to get 1,699 kB of archives.

After this operation, 3,702 kB of additional disk space will be used.

Do you want to continue [Y/n]?

Once you have installed all these packages, load the sound module into the system to take advantage of it. To load a module, use the modprobe command. In this instance, load the snd_bcm2835 module, which is the Broadcom sound adapter in the Pi. Issue this command:

```
$ sudo modprobe snd_bcm2835
```

The command will run, but it doesn't generate any output to tell us the status of the execution or just about anything else. We need to check and see whether the module has been loaded. Just like the ls command for directories and lsusb for listing USB devices, there is an lsmod command for listing the modules that are currently installed within your system. Issue the lsmod command and check that you can see the snd_bcm2835 module loaded within your system (like mine):

```
$ lsmod
Module           Size  Used by
snd_bcm2835      21485  0
snd_pcm          82208  1 snd_bcm2835
snd_page_alloc    5383   1 snd_pcm
snd_seq           59808  0
snd_seq_device    6920   1 snd_seq
snd_timer         21905  2 snd_seq,snd_pcm
snd               57668  5 snd_timer,snd_seq_device,snd_seq,snd_pcm,snd_bcm2835
ipv6              290227 34
r8712u            182646  0
spi_bcm2708       4815   0
i2c_bcm2708       3818   0
```

Because you are loading the sound module for the first time, you also need to select which audio device is being used (because none was used before). To change the device, use the amixer command as follows:

```
$ sudo amixer cset numid=3 1
numid=3,iface=MIXER,name='PCM Playback Route'
; type=INTEGER,access=rw-----
; values=1,min=0,max=3,step=0
: values=1
```

We changed the numid=3 value, which is the PCM Playback Route, to be equal to 1, which is the value of the 3.5 jack. You can change the 1 on the end to a 2, which signifies that you should use the HDMI's inbuilt audio channel to send the sound.

Now that the hardware is sorted out, it is time to take care of the software by downloading the shairport software from its repository on the Internet. By far the easiest way to do this is by using the wget (short for *Web GET*) command to download all the available contents on a given URL. In this manner, we can download a whole software package from the Internet with a simple URL. Execute the following:

```
$ wget https://github.com/albertz/shairport/zipball/master
--2012-09-30 18:13:36-- https://github.com/albertz/shairport/zipball/master
Resolving github.com (github.com)... 207.97.227.239
Connecting to github.com (github.com)|207.97.227.239|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://nodeload.github.com/albertz/shairport/zipball/master [following]
--2012-09-30 18:13:42-- https://nodeload.github.com/albertz/shairport/zipball/master
Resolving nodeload.github.com (nodeload.github.com)... 207.97.227.252
Connecting to nodeload.github.com (nodeload.github.com)|207.97.227.252|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 46413 (45K) [application/zip]
Saving to: 'master'

100%[=====] 46,413      60.1K/s   in 0.8s
2012-09-30 18:13:50 (60.1 KB/s) - 'master' saved [46413/46413]
```

Now that we have downloaded the shairport software we need to work with it. The file is called `master`, so let's neaten it up and name it correctly. Just use the `mv` command to change the filename. Once you have moved the file, unzip it with the `unzip` command and `cd` into the newly created `albertz-shairport-b58f156` directory:

```
$ mv master albertz-shairport-b58f156.zip
$ unzip albertz-shairport-b58f156.zip
$ cd albertz-shairport-b58f156
```

Now you need to build the application, so run the `make` command. Thankfully this application is far simpler than the XBMC one and thus doesn't need to be preconfigured.

```
$ make
cc -O2 -Wall -c alac.c -o alac.o
cc -O2 -Wall -DHAIRTUNES_STANDALONE hairtunes.c alac.o -o hairtunes -lm -lpthread -lssl
-lcrypto -lao
cc -O2 -Wall -c socketlib.c -o socketlib.o
cc -O2 -Wall -c shairport.c -o shairport.o
cc -O2 -Wall -c hairtunes.c -o hairtunes.o
cc -O2 -Wall socketlib.o shairport.o alac.o hairtunes.o -o shairport -lm -lpthread -lssl
-lcrypto -lao
```

Once the compile is finished, it is time to test; shairport is run by executing the `shairport.pl` file. So execute it as such:

```
$ ./shairport.pl
Can't locate Net/SDP.pm in @INC (@INC contains: /usr/lib/perl5/site_perl /usr/share/perl5/site_perl
/usr/lib/perl5/vendor_perl /usr/share/perl5/vendor_perl /usr/lib/perl5/core_perl /usr/share/perl5/
core_perl .) at ./shairport.pl line 45.
```

Whoops! There's an error related to a particular module of code not being available to the Perl system that it needs to perform the install. In this case, it is the Net/SDP module that we need to install.

Unfortunately, apt-get cannot fix this one, so you have to get it working by downloading the module from the Perl CPAN library (which is much like a giant apt-get repository for Perl software that developers can borrow and use) and then install it. You will again use wget to download the file from the Internet. Issue the following command:

```
$ wget http://search.cpan.org/CPAN/authors/id/N/NJ/NJH/Net-SDP-0.07.tar.gz
--2012-09-30 19:01:11-- http://search.cpan.org/CPAN/authors/id/N/NJ/NJH/Net-SDP-0.07.tar.gz
Resolving search.cpan.org (search.cpan.org)... 199.15.176.161
Connecting to search.cpan.org (search.cpan.org)|199.15.176.161|:80... connected.
HTTP request sent, awaiting response... 302 Found
Location: http://mirror.westfield.com.au/cpan/authors/id/N/NJ/NJH/Net-SDP-0.07.tar.gz [following]
--2012-09-30 19:01:11-- http://mirror.westfield.com.au/cpan/authors/id/N/NJ/NJH/Net-SDP-0.07.tar.gz
Resolving mirror.westfield.com.au (mirror.westfield.com.au)... 203.42.62.21
Connecting to mirror.westfield.com.au (mirror.westfield.com.au)|203.42.62.21|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 20679 (20K) [application/x-gzip]
Saving to: 'Net-SDP-0.07.tar.gz'

100%[=====] 20,679
--.-K/s in 0.05s
2012-09-30 19:01:12 (382 KB/s) - 'Net-SDP-0.07.tar.gz' saved [20679/20679]
```

Once the file is downloaded, you need to extract it with the tar command. Issue a tar command to unpack the file with the arguments -zvxf to unzip and extract this tarball. Then change into the newly created Net-SDP-0.07 directory:

```
$ tar -zvxf Net-SDP-0.07.tar.gz
$ cd Net-SDP-0.07
```

Now that the package is here, you just need to issue commands to the Perl build system here. The first is to execute the Build.PL script to generate the build file:

```
$ perl Build.PL
Created MYMETA.yml and MYMETA.json
Creating new 'Build' script for 'Net-SDP' version '0.07'
```

Once the build file is created, you need to run it with this:

```
$ ./Build
Building Net-SDP
```

And once the build is finished, run the inbuilt tests to ensure there are no problems by issuing the test command to the build script:

```
$ ./Build test
t/00use.t ..... ok
t/10generate.t .. ok
t/10parse.t ..... ok
t/20repeat.t .... ok
t/30asstring.t .. ok
All tests successful.
Files=5, Tests=69, 5 wallclock secs ( 0.97 usr 0.05 sys + 3.40 cusr 0.29 csys = 4.71 CPU)
Result: PASS
```

Finally, you need to install the modules, so issue the `install` command to the `build` script (you might need to press Enter after the command to get the install to finish and display):

```
$ sudo ./Build install  
[sudo] password for raspbian:  
Building Net-SDP  
Installing /usr/local/man/man1/sdp2rat.1p  
Installing /usr/local/share/perl/5.14.2/Net/SDP.pm  
Installing /usr/local/share/perl/5.14.2/Net/SDP/Time.pm  
Installing /usr/local/share/perl/5.14.2/Net/SDP/Media.pm  
Installing /usr/local/man/man3/Net::SDP::Media.3pm  
Installing /usr/local/man/man3/Net::SDP::Time.3pm  
Installing /usr/local/man/man3/Net::SDP.3pm  
Installing /usr/local/bin/sdp2rat
```

Once the build is done, it is a simple matter to run the `shairport.pl` script again, and this time it should work! The output should look like this:

```
$ cd..  
$ ./shairport.pl  
Established under name '891BEA3BF8A1@ShairPort 2113 on pisces'
```

Now you need to establish connectivity with the application and test that you can send music to it. Open up iTunes and then change the output destination in the bottom-right corner of iTunes (see Figure 11-5). You should see a shairport name there. Select it and press Play on your favorite song and presto: streaming music!




Figure 11-5. ShairPort enabled

Troubleshooting

As with XBMC, this process is fairly painless because the Raspberry Pi has a fixed hardware platform, and without all these variables there is not as much to worry about—things should just work.

If you find that no sound is coming out of your speakers when they are attached, check that you correctly issued both the `modprobe` and the `amixer` commands because both are needed to make the audio function.

Other than this issue, most of the commands should work. If not, you should delete any downloaded content and start over.

Airport on Boot

As with using the XBMC instance, most people won't want to have shairport started with their system without needing to manually log in to your Pi and start the application. This is easy for us because the shairport install has all the files needed to do this bundled with it! To perform this, first install the shairport system: run `sudo make install`, and your output should look like this:

```
$ sudo make install
install -D -m 0755 haitunes /usr/local/bin/haitunes
install -D -m 0755 shairport.pl /usr/local/bin/shairport.pl
install -D -m 0755 shairport /usr/local/bin/shairport
```

Once the install is finished, you can copy the provided `init` script into the correct location:

```
$ sudo cp shairport.init.sample /etc/init.d/shairport
```

Now that the file is copied, you need to make one small modification: add the `modprobe` command from earlier—`modprobe snd_bcm2835`—into the `/etc/init.d/shairport` file just before you start the shairport application. The `start` section should look like this:

```
start() {
    echo -n "Starting shairport: "
    modprobe snd_bcm2835
    start-stop-daemon --start --quiet --pidfile "$PIDFILE" \
                      --exec "$DAEMON" -b --oknodo -- $DAEMON_ARGS
    log_end_msg $?
}
```

The last commands we need to run are these:

```
$ sudo chmod +x /etc./init.d/shairport
$ sudo update-rc.d shairport defaults
```

These commands make the new `start` script writeable and update the boot sequence to include the new shairport scripts! Reboot your Pi, and shairport should be up and running for you on boot! Congratulations.

Summary

In this comprehensive chapter, you installed a whole course of precursor applications and learned how to download source code, configure it, and compile from it. This is quite a lot of work and the XBMC compile took in excess of 12 hours!

You also created and installed start scripts to get these new media playback systems to run on the Pi at boot, avoiding the need to manually intervene.

Finally, you connected the XBMC instance to the remote control application for smartphones. This has been an incredible amount of work for one chapter, so well done!

Index

A

Advanced Linux Sound Architecture (ALSA) device, 202

Airplay

- amixer command, 237
- boot, 241
- central controlling system, 236
- build script, 239–240
- lsmod command, 237
- Net-SDP-0.07 directory, 239
- execute, 238
- precursor packages, 236
- ShairPort enable, 240
- troubleshooting, 241
- wget, 238–239

Apache

- apt-get tool, 150
- browser window, 152
- configuration, 152–155
- output, 151
- running, 151
- stories, 150
- troubleshooting, 155
- web server, 150

Linux system, 142

run files, 143

XBMC, 141–142

interpreted versus compiled, 132

logical operation

- if, 134
- infinite loop, 139
- loop, 139–140

output, 132–133

recap, 133

security and user management

new user adding, 147

rules-Raspbian security, 146

test based arithmetics

- advantage, 136
- arithmetic operation, 137
- assignment mode, 135
- dereferencing mode, 136
- modulus, 135
- output, 137
- testing operators, 136
- variable, 136

troubleshooting, 137, 139, 141

variables, 133–134

B

BASH

basic coding

- FizzBuzz test, 131
- goals, 130
- meaning, 131
- text editor, 131

init script

- blocks, 142
- case statement, 143
- creation, 144–145
- forking, 143

C

Combine commands

buffer, 102

environment, 104

inverse grep, 103

PATH, 105–106

ps and grep, 102

results, 103

xargs, 104

Commands. *See also* System resources monitoring

bit confusing, 90

filesystem, 89

Commands (*cont.*)

Linux distributions, 91

RTFM, 91

sudo command, 89

system/administrative function, 89

Common Gateway Interface (CGI), 165

Compiled languages, 132

Computer terminals. *See also* Three terminals

command line, 66–67

shells, 68

TeleVideo 925, 53

till-duh, 67

Ye Olde computer, 54–58

always available, 57

convinced, 57

dumb terminal, 54

fast, 57

GUIs, 55

key problem, 54

lightweight, 57

modern terminals, 55–56

power, 57

simple, 56

TUI, 55

virtual terminal, 58

D**dmesg command**

abbreviations, 201

/dev, 202

logged output, 200–201

output, 202

troubleshooting, 204

udev administration function, 203

Driver message. *See* dmesg command

Dynamic Host Configuration Protocol (DHCP), 60

E**Editing text files**

arbitrary format, 110

binary data format, 109

Chinese text, 110

contenders, 110–111

foreign languages, 110

nano command

changes, 114

disk file, 115

edited file, 115

file saving, 113–114

line number, 117

Mary, 116

mean, 112–113

move around, 115–116

new file, 112

shortcut bar, 113

style, 114

text editor, 111

wrapping up, 117

WriteOut command, 114

overview, 109

semantics, 109

statement, 109

structure, 109

vim

charm, 120

copy and paste style, 127

delete command, 124–125

error message, 119

escape (ESC) key, 119

features, 117

getting out, 120

indenting text, 126

insert mode, 118

install, 117

mass pasting, 128

miscellaneous little commands, 126

move around, 123

navigation shortcuts, 123

regular expressions, 122

save changes, 119

search, 121–122

specific line number, 122–123

straightforward, 117–118

tildes, 118

visual mode, 126

F**File commands. *See also* Commands**

combine commands

buffer, 102

environment, 104

inverse grep, 103

PATH, 105–106

ps and grep, 102

results, 103

xargs, 104

command-line arguments, 101

cron Command, 107

generalized regular expression parser, 102

kind, 101

usage, 101

File-paths

command, 72

disk set, 72

file permissions

chmod command, 84

execute bit, 84

Linux, 83
 ls command, 82
 pi, 83
 Raspberry file, 82–83
 roles, 82
 set, 84–87

filingsystem
 all together, 71
 definition, 69
 directory, 70
 filesystem, 70
 libraries, 69
 Mac different, 71
 paper-based system, 70
 separate roots, 70–71
 unified filesystem, 71

hard disk, 72

layouts
 /bin and /sbin, 74
 /boot, 74
 /dev, 74
 /etc., 73
 /home, 74
 /lib, 75
 /lost+found, 75
 /media, 75
 /mnt, 75
 /opt, 75
 /proc, 73
 /root, 73
 / root directory, 73
 /srv, 75
 /sys, 76
 /tmp, 76
 /usr, 75
 /var, 74
 wrapper, 76

Linux, 73
 SCSI-based disks, 72
 shortcuts and links, 87–88
 user and groups, 81–82
 working
 directories, 78
 cp command, 79–80
 file/directory, 81
 mv command, 79
 rm command, 80
 mkdir, 78
 overview, 76
 pwd, 76
 touch, 77–78
 use of, 77

Filing system
 all together, 71
 definition, 69

directory, 70
 filesystem, 70
 libraries, 69
 Mac different, 71
 separate roots, 70–71
 unified filesystem, 71

■ G

Graphical user interface (GUI), 55

■ H

Hyper Text Markup Language (HTML). *See also MySQL tags, 168–169 websites, 168–169*

■ I, J, K

Integrated development environments (IDEs), 48
 Interpreted languages, 132
 init script
 blocks, 142
 case statement, 143
 creation, 144–145
 forking, 143
 Linux system, 142
 run files, 143
 XBMC, 141–142

■ L

LAMP
 A stand for Apache
 apt-get tool, 150
 browser window, 152
 configuration, 152
 output, 151
 running, 151
 stories, 150
 troubleshooting, 155
 web server, 150

L stand for Linux, 150
 P stand for PHP, 165
 CGIs, 165
 command, 165
 dynamic functions, 165
 functions, 166
 info page, 166
 MySQL config info, 167
 php5-mysql package, 167
 points, 166

web application
 basic concept, 169
 connection and database, 170

LAMP (*cont.*)
 create statement, 170
 database content, 171–172
 DB-remove data, 178–179
 HTML, 168–169
 index.php, 170
 insert case, 175–176
 overview, 168
 PHP code, 170
 remove entries, 176–178
 to-do list, 179
 troubleshooting, 179
 website data insertion, 173–174
 website content, 149
Layouts, File-paths
 /bin and /sbin, 74
 /boot, 74
 /dev, 74
 /etc., 73
 /home, 74
 /lib, 75
 /lost+found, 75
 /media, 75
 /mnt, 75
 /opt, 75
 /proc, 73
 /root, 73
 / root directory, 73
 /srv, 75
 /sys, 76
 /tmp, 76
 /usr, 75
 /var, 74
 wrapper, 76
Lightweight Desktop Environment (LXDE), 59
Lightweight X11 Desktop Environment (LXDE)
 meaning, 32
Pi, 32
right side-task bar, 33–34
shutdown/rebooting, 34
start menu concept, 34–45
 accessories, 36–43
 and bottom left, 33
 Education, 43–45
 Internet, 46–47
 operating system, 34
 preferences, 50–52
 programming menu, 47–49
 run command, 35
 system tools, 49–50
 window, 35
X-server, 31
Linux. *See also* LAMP
 definition, 9
 features, 9

histroy, 9
open source, 9
Raspbian, 10

■ M

Mail transfer agent (MTA), 208
Media center. *See also* Airplay
 XBox Media Center
 boot screen, 232, 235–236
 main menu, 233
 operating system, 224
 Pi boot system, 232
 Pisces, 224
 program running, 234
 RaspMC, 224
 regular expressions and sed,
 228–229
 remote application, 234
 setup-build, 224–228
 solutions, 233
 source code, 229–231
 troubleshooting, 232, 236
 XBMCbuntu, 223
MySQL
 HTML, 155
 installation, 156
 Structured Query Language
 (SQL)
 DELETE, 164
 insert data, 160–161
 pi, 157
 SELECT, 161–162
 syntax, 157
 table creation, 157–158, 160
 types, 157
 UPDATE, 163–164

■ N, O, P, Q

Nano command
 changes, 114
 disk file, 115
 edited file, 115
 file saving, 113–114
 line number, 117
 Mary, 116
 mean, 112–113
 move around, 115–116
 new file, 112
 shortcut bar, 113
 style, 114
 text editor, 111
 wrapping up, 117
WriteOut command, 114

P

Paper-based systems. *See* Filing system
PHP

- CGIs, 165
- command, 165
- dynamic functions, 165
- php5-mysql package, 167
- functions, 166
- info page, 166
- MySQL config info, 167
- points, 166
- Pi. *See also* Media center
 - boot-it, 17–18
 - filesystem, 19
 - first boot, 18–19
 - keyboard, 19–23
 - configuration, 21–22
 - customization, 21, 23
 - Generic 105 key (Intl) PC, 19
 - layout, 20–21
 - X server, 23
 - LED, 18
 - locale
 - choose, 24
 - default, 24–25
 - symbols, 24
 - memory allocation
 - memory split, 28
 - memory_split option, 27
 - mini desktop computer, 27
 - reboot process
 - desktop, 29–30
 - output window, 28
 - remote access
 - command line, 129
 - DNS, 130
 - Linux environment, 129
 - networking, 130
 - SSH, 130
 - SD card, 18
 - set/change password, 23–24
 - time zone
 - country/city, 26–27
 - regions, 26
 - TV/monitor, 18

R

Raspberry Pi. *See also* Pi, Raspbian, Linux

- average-sized mouse, 3
- envelope, 1–2
- ingredients
 - BIOS, 8
 - capable display-HDMI, 5

- HDMI lead, 5
- keyboard and mouse, 7
- lists, 3
- loaded Pi, 8
- micro USB lead, 4
- multicard reader, 7
- reader-SD card, 6–7
- SD card, 5–6
- USB power adapter, 4
- picture, 2
- Raspberry Pi
 - altering system
 - advantage, 208
 - mutt tool, 210–211
 - SSMTP, 208, 210
 - testing, 211–212
 - troubleshooting, 212
 - designing solution, 214–215
- dmesg command
 - abbreviations, 201
 - /dev, 202
 - logged output, 200–201
 - output, 202
 - troubleshooting, 204
 - udev administration function, 203
- making decisions, 213–214
- materials needed, 199
- motion, 204–206, 208
 - camera, 205
 - command, 205
 - ffmpeg_cap_new, 206
 - open and edit file, 205
 - output, 204
 - package, 204
 - process, 208
 - reference, 206, 208
 - target_dir, 206
 - threshold, 206
 - troubleshooting, 208
- presetup, 200
- problems
 - goals, 213
 - requirements, 213
- Script 1, 216
- Script 2, 217
- so far, 212–213
- stack and solve complex problems, 221
- start up, 220
- stetson, 221
- testing
 - cron, 219
 - inbox-output, 218
 - script launch method, 219
 - updates, 219
- troubleshooting, 220

Raspbian

Image Writer, 11–13
action shot, 13
binary version, 11
disk image, 11–12
error message, 12

Mac
image writing, 14–17
terminals, 14

meaning, 10

Read The Flaming Manual (RTFM), 91

Regular expressions, 228–229

S

Scripting languages. *See* Interpreted languages

Secure Shell (SSH)

client-sort out, 61–66

Mac, 65–66

putty, 61–64

de facto standard, 60

server to Pi

ifconfig, 60–61

interface configuration, 60

IP address, 60

Send our mail to the mail server (SSMTP),

208–210

Source code. *See* Compiled languages

Start menu concept, LXDE

accessories

Debian reference, 36–37

File Manager, 37–38

Image Viewer, 38–39

Leafpad, 40

LXTerminal, 41

menu, 36

Root Terminal, 41–43

XArchiver, 43

Education

menu, 43–44

Scratch, 44

Squeak, 45–46

Internet menu, 46–47

operating system, 34

preferences

customize look and feel, 51–52

desktop, 50–51

programming menu

IDEs, 48

IDLE, 49

Python, 48

window, 47–48

run command, 35–36

system tools

menu, 49

RAM usage, 50

task manager list, 49–50

window, 35

Stream Editor (sed), 228

Structured Query Language (SQL)

DELETE, 164

insert data, 160–161

pi, 157

SELECT, 161–162

syntax, 157

table creation, 157–158, 160

types, 157

UPDATE, 163–164

System resources monitoring

CPU utilization percentage, 93

disk usage, 95–97

free to view memory allocation, 95

managing processes, 97–100

current processes, 97

in /proc, 100

kill processes, 99

ps –ef, 98

ps means, 99

memory usage, 94

output, 91–92

process table, 94

tasks, 93

uptime and load average, 93

T, U

Terminals, 53. *See also* Computer terminals

Test based arithmetics

advantage, 136

arithmetic operation, 137

assignment mode, 135

dereferencing mode, 136

modulus, 135

output, 137

testing operators, 136

variable, 136

Text-based user interface (TUI), 55

Three terminals

console, 58

LXDE, 59

SSH

client, 61–66

de facto standard, 60

equation, 59

server to Pi, 60–61

virtual terminal, 59

V

- Video, 223
- Vim
 - charm, 120
 - copy and paste style, 127
 - delete command, 124–125
 - error message, 119
 - escape (ESC) key, 119
 - features, 117
 - getting out, 120
 - indenting text, 126
 - insert mode, 118
 - install, 117
 - mass pasting, 128
 - miscellaneous little commands, 126
 - move around, 123
 - navigation shortcuts, 123
 - regular expressions, 122
 - save changes, 119
 - search, 121–122
 - specific line number, 122–123
 - straightforward, 117–118
 - tildes, 118
 - visual mode, 126
- Virtual Network Computing (VNC)
 - execution, 183
 - install, 183
 - output, 183
 - TightVNC Client, 184
 - vncserver, 184, 186
- Virtual terminal, 58

W

- WiFi
 - command line, 194–195
 - configuration
 - aware of, 192
 - connect, 192
 - icon, 190
 - ifconfig, 193
 - network details, 191–192
 - scan results window, 191
 - wpagui window, 190
 - D-Link DWA-131, 186
 - Ethernet cable, 196
 - installation, 189
 - lsusb command, 187
 - output, 188

overview, 186

setup, 187

USB devices, 186

wireless adapter, 188

wpa-gui, 189

Wireless Computing (WiPi). *See also WiFi*

commands, 182

Pi usage, 181–182

power output, 197

simple constraint removal, 182–184, 186

human input devices, 182

remote GUI, 183–184, 186

USB adapters, 197

X

XBox Media Center (XBMC)

boot screen, 232, 235–236

build-setup, 224–228

main menu, 233

operating system, 224

Pi boot system, 232

Pisces, 224

program running, 234

RaspMC, 224

remote application, 234

sed and regular expressions, 228–229

solutions, 233

source code, 229–231

troubleshooting, 232, 236

XMCBuntu, 223

X-server, 31

Y, Z

Ye Olde computer

always available, 57

convinced, 57

dumb terminal, 54

fast, 57

GUIs, 55

key problem, 54

lightweight, 57

modern terminals, 55–56

pipes, 57

power, 57

simple, 56

TUI, 55

virtual terminal, 58

Learn Raspberry Pi with Linux



Peter Membrey
David Hows

Apress®

Learn Raspberry Pi with Linux

Copyright © 2013 by Peter Membrey and David Hows

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed. Exempted from this legal reservation are brief excerpts in connection with reviews or scholarly analysis or material supplied specifically for the purpose of being entered and executed on a computer system, for exclusive use by the purchaser of the work. Duplication of this publication or parts thereof is permitted only under the provisions of the Copyright Law of the Publisher's location, in its current version, and permission for use must always be obtained from Springer. Permissions for use may be obtained through RightsLink at the Copyright Clearance Center. Violations are liable to prosecution under the respective Copyright Law.

ISBN-13 (pbk): 978-1-4302-4821-7

ISBN-13 (electronic): 978-1-4302-4822-4

Trademarked names, logos, and images may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, logo, or image we use the names, logos, and images only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

The use in this publication of trade names, trademarks, service marks, and similar terms, even if they are not identified as such, is not to be taken as an expression of opinion as to whether or not they are subject to proprietary rights.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

President and Publisher: Paul Manning

Lead Editor: Michelle Lowman

Developmental Editor: Matthew Moodie

Technical Reviewer: Stewart Watkiss

Editorial Board: Steve Anglin, Mark Beckner, Ewan Buckingham, Gary Cornell, Louise Corrigan, Morgan Ertel, Jonathan Gennick, Jonathan Hassell, Robert Hutchinson, Michelle Lowman, James Markham, Matthew Moodie, Jeff Olson, Jeffrey Pepper, Douglas Pundick, Ben Renow-Clarke, Dominic Shakeshaft, Gwenan Spearing, Matt Wade, Tom Welsh

Coordinating Editor: Jill Balzano

Copy Editor: Nancy Sixsmith

Compositor: SPi Global

Indexer: SPi Global

Artist: SPi Global

Cover Designer: Anna Ishchenko

Distributed to the book trade worldwide by Springer Science+Business Media New York, 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax (201) 348-4505, e-mail orders-ny@springer-sbm.com, or visit www.springeronline.com. Apress Media, LLC is a California LLC and the sole member (owner) is Springer Science + Business Media Finance Inc (SSBM Finance Inc). SSBM Finance Inc is a Delaware corporation.

For information on translations, please e-mail rights@apress.com, or visit www.apress.com.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales–eBook Licensing web page at www.apress.com/bulk-sales.

Any source code or other supplementary materials referenced by the author in this text is available to readers at www.apress.com. For detailed information about how to locate your book's source code, go to www.apress.com/source-code/.

To my parents, Peter and Bev, for all their support.

— David Hows

*For my wife Sarah, son Kaydyn, and my close friends Jacqui Koo and Mark Stephens,
who together make my world go 'round.*

— Peter Membrey

Contents

About the Authors	xv
About the Technical Reviewer	xvii
Acknowledgments	xix
Introduction	xxi
■ Chapter 1: Your First Bite of Raspberry Pi.....	1
Your Freshly Baked Pi Arrives.....	1
List of Ingredients	3
Micro USB lead	4
USB Power Adapter	4
HDMI Lead	5
HDMI Capable Display	5
SD Card.....	5
SD Card Reader	6
USB Keyboard and Mouse	7
Whew, We're Done!.....	7
Don't Panic!.....	8
Linux.....	8
What Is Linux?	9
Downloading Raspbian.....	10
Getting Raspbian onto Your SD Card	10
Using Image Writer on Windows.....	11
Using dd on the Mac.....	14
First Boot.....	17

Configuring Your Pi	18
Expanding the Filesystem.....	19
Configuring the Keyboard	20
Changing the Password.....	23
Configuring the Locale.....	24
Changing the Time Zone	26
Allocating Memory.....	27
At Last! It's Configured!	28
Summary.....	30
■ Chapter 2: Surveying the Landscape.....	31
Welcome to LXDE	31
What Do We Have Here?	32
The Start Menu and Bottom Left of the Task Bar.....	33
And on the Right Side.....	33
The Start Menu.....	34
Accessories	36
Education.....	43
Internet.....	46
Programming.....	47
System Tools.....	49
Preferences	50
Summary.....	52
■ Chapter 3: Getting Comfortable	53
Ye Olde Computer.....	54
Say Hello to the Dumb Terminal	54
Modern Terminals.....	55
Why Do We Still Care About These Things?	56
Three Terminals.....	58
Console	58
Opening a Virtual Terminal in the GUI	59
Connecting via SSH	59

Welcome to the Command Line	66
Different Shells.....	68
Summary.....	68
Chapter 4: The File-Paths to Success	69
What Is a Filing System?	69
More than One Filesystem.....	70
Separate Roots	70
Unified Filesystem	71
The Mac has to be Different	71
Bring it All Together	71
Everything as a File	72
Filesystem Layout	73
/ (Root Directory)	73
/root	73
/etc	73
/proc	73
/var	74
/boot	74
/bin and /sbin	74
/dev.....	74
/home	74
/lib	75
/lost+found.....	75
/mnt.....	75
/media	75
/usr	75
/opt	75
/srv	75
/sys	76
/tmp	76
Wrapping it Up	76

Putting it to Work.....	76
Where Are We? Using pwd.....	76
What's in Here with Us? Using ls.....	77
Creating Files to Play with: Using Touch.....	77
Somewhere to Store our Files: Using mkdir	78
Making Use of a New Directory: Using the mv Command	79
Time for Some Cloning: How to Use the cp Command	79
The Power to Destroy: Using the rm Command.....	80
Fully Qualified and Relative Paths	81
Users and Groups	81
File Permissions	82
Setting File Permissions	84
Shortcuts and Links	87
Summary.....	88
■Chapter 5: Essential Commands.....	89
Become the Boss	89
RTFM	91
System Resource Monitoring	91
Uptime and Load Average	92
Tasks.....	93
CPU Utilization Percentage	93
Memory Usage.....	94
Process Table.....	94
Using free to View Memory Allocation	95
Disk Usage.....	95
Managing Processes	97
File Commands.....	100
Combining Commands.....	102
User Environment	104
The cron Command	107
Summary.....	107

■ Chapter 6: Editing Files on the Command Line.....	109
What is a Text File?	109
The Contenders	110
Starting out with nano.....	111
So What Does This All Mean?	112
Saving Your File	113
Moving Around in nano.....	115
Wrapping Up nano	117
Getting Started with vim	117
The vim Modes	118
Saving Your Changes	119
Getting Out of vim.....	120
Searching in vim.....	121
Moving to a Specific Line Number.....	122
Moving About in vim	123
Deleting in vim.....	123
Miscellaneous Little Commands.....	126
Visual Mode	126
Indenting Text	126
Summary.....	128
■ Chapter 7: Managing Your Pi.....	129
Remotely Accessing the Pi	129
Networking	129
DNS.....	130
SSH.....	130
BASH: Basic Coding.....	130
What Is BASH?	131
Starting in BASH	131
Interpreted versus Compiled	132
Output in BASH	132

Recap	133
Variables	133
Logical Operation: if	134
Test Based Arithmetic	135
Troubleshooting	137
Logical Operation: Loop	139
Troubleshooting	141
Practical BASH: An Init Script	141
Pick and Match with the case Statement	143
Application within Application: Forking	143
Update the Run Files	143
Creating Your Own init Script	144
Security and User Management	146
The Rules of Raspbian Security	146
Adding a New User	147
Summary	148
■ Chapter 8: A LAMP of Your Own	149
First Steps	149
<i>L</i> is for Linux	150
<i>A</i> is for Apache	150
Apache Configuration	152
Troubleshooting	155
<i>M</i> is for MySQL	155
Installing MySQL	156
Structured Query Language	156
<i>P</i> is for PHP	165
Simple Web App	168
HTML	168
Starting Our Page	169
Display the Database Content	171

Website Data Insertion.....	173
Inserting to the Database	175
Remove Entries.....	176
Remove Data from DB	178
Troubleshooting	179
Where From Here?.....	179
Summary.....	180
■ Chapter 9: WiPi: Wireless Computing	181
Preparing for WiPi	182
Simple Constraint Removal	182
Removing Human Input Devices.....	182
Adding a Remote GUI.....	183
True WiPi with WiFi	186
GUI WiFi Configuration	190
Managing WiFi from the Command Line.....	194
Down to One	196
Removing the Need for Power.....	196
Powered USB Adapters	197
Summary.....	197
■ Chapter 10: The Raspberry sPi	199
Materials Needed	199
Presetup	200
Getting Started	200
dmesg.....	200
Troubleshooting	204
Motion Capture.....	204
Troubleshooting	208
Alerting Yourself	208
Installing SSMTP	208
Installing and Using mutt.....	210

CONTENTS

Testing the Alerting System.....	211
Troubleshooting	212
Bringing It All Together	212
Understanding the Problem.....	213
Making Decisions	213
Designing Solution.....	214
Script 1	216
Script 2	217
Testing	218
Troubleshooting	220
Where to Go from Here.....	220
Summary.....	221
Chapter 11: Pi Media Center.....	223
Video with XBMC	223
Setup to Build	224
Changing Files with Sed and Regular Expressions.....	228
Compiling Source Code	229
Troubleshooting	231
Starting and Using XBMC	232
XBMC on Boot.....	235
Troubleshooting	236
Stream Music with Airplay	236
Troubleshooting	241
Airport on Boot.....	241
Summary.....	242
Index.....	243

About the Authors



Peter Membrey is a Chartered IT Professional with nearly 15 years of experience using Linux and open source solutions to solve problems in the real world. An RHCE since the age of 17, he has also had the honor of working for Red Hat and writing several books covering open source solutions. He holds a masters degree in IT (Information Security) from the University of Liverpool and is currently a PhD candidate at the Hong Kong Polytechnic University where his research interests include cloud computing, big data science and security. He lives in Hong Kong with his wonderful wife Sarah and son Kaydyn. His Cantonese continues to regress.



David Hows is an Honors graduate from the University of Woolongong in NSW Australia. He got his start in computing trying to drive more performance out of his family PC without spending a fortune. This lead to a career in IT where David has worked as a Systems Administrator, Performance Engineer, Software Developer, Solutions Architect and Database Engineer. David has tried in vain for many years to play soccer well and his coffee mug reads “Grumble Bum.”

About the Technical Reviewer



Stewart Watkiss graduated from The University of Hull, UK, with a Master's degree in electronic engineering. He has been a fan of Linux since first installing it on a home computer during the late 1990s, back when Linux was installed using floppy disks. While working as a Linux system administrator, he was awarded Advanced Linux Certification (LPIC 2) in 2006, and created the PenguinTutor website to help others working toward Linux certification (<http://www.penguintutor.com>).

After ordering his first Raspberry Pi on the launch day, he has since gotten three more, which he uses as a family friendly proxy, a media center, for working on hobby electronic projects, and teaching his two children about computers.

Acknowledgments

Without a doubt, the true heroes of this book are the Raspberry Pi Foundation and the Raspberry Pi community because without either of them, we wouldn't be here. The Raspberry Pi Foundation, the group who bought us the Pi, is a nonprofit charity with the stated goal with the Pi was to reinvigorate the development spirit of the '80s and to provide an affordable platform for generations of youngsters to get into the wonderful world of computing in a truly meaningful way. It has done so much to bridge the digital divide and bring computing to the common man, and for this it deserve our thanks.

If the Raspberry Pi Foundation is the heart, the Raspberry Pi community has to be the soul. These people have sought to provide the tools and support people need to get started with the Pi regardless of skill level. They have invested countless hours (and dollars) in writing, fixing, porting, hacking, and testing things for the Pi. Please feel free to pay a visit to them at <http://www.raspberrypi.org/> and see what we mean.